

MathOptInterface

The JuMP core developers and contributors

May 16, 2022

Contents

Contents	ii
I Introduction	1
1 Introduction	2
1.1 What is MathOptInterface?	2
1.2 How the documentation is structured	2
1.3 Citing MathOptInterface	3
2 Motivation	4
II Tutorials	5
3 Solving a problem using MathOptInterface	6
3.1 Required packages	6
3.2 Define the data	6
3.3 Add the variables	7
3.4 Set the objective	7
3.5 Add the constraints	7
3.6 Optimize the model	7
3.7 Understand why the solver stopped	7
3.8 Understand what solution was returned	8
3.9 Query the objective	8
3.10 Query the primal solution	8
4 Implementing a solver interface	9
4.1 A note on the API	9
4.2 Preliminaries	9
4.3 Structuring the package	11
4.4 Initial code	12
4.5 The big decision: copy-to or incremental modifications?	15
4.6 Names	18
4.7 Solutions	19
4.8 Other tips	20
5 Transitioning from MathProgBase	22
5.1 Transitioning a solver interface	22
5.2 Transitioning the high-level functions	22
6 Implementing a constraint bridge	24
6.1 Preliminaries	24
6.2 Four mandatory parts in a constraint bridge	24
6.3 Bridge registration	27
6.4 Bridge improvements	27
7 Manipulating expressions	30
7.1 Creating functions	30

7.2	Canonicalizing functions	31
7.3	Exploring functions	32
8	Latency	33
8.1	Background	33
8.2	Causes	33
8.3	Resolutions	36
III	Manual	39
9	Standard form problem	40
9.1	Functions	40
9.2	One-dimensional sets	41
9.3	Vector cones	41
9.4	Matrix cones	42
9.5	Multi-dimensional sets with combinatorial structure	42
10	Models	43
10.1	Attributes	43
10.2	ModelLike API	44
10.3	AbstractOptimizer API	44
11	Variables	46
11.1	Add a variable	46
11.2	Delete a variable	46
11.3	Variable attributes	47
12	Constraints	48
12.1	Add a constraint	48
12.2	Delete a constraint	49
12.3	Constraint attributes	49
12.4	Constraints by function-set pairs	49
12.5	JuMP mapping	50
13	Solutions	52
13.1	Solving and retrieving the results	52
13.2	Why did the solver stop?	52
13.3	Primal solutions	52
13.4	Dual solutions	53
13.5	Common status situations	53
13.6	Querying solution attributes	54
14	Problem modification	56
14.1	Modify the set of a constraint	56
14.2	Modify the function of a constraint	57
14.3	Modify constant term in a scalar function	58
14.4	Modify constant terms in a vector function	58
14.5	Modify affine coefficients in a scalar function	59
14.6	Modify affine coefficients in a vector function	59
IV	Background	61
15	Duality	62
15.1	Duality and scalar product	63
15.2	Dual for problems with quadratic functions	64
16	Infeasibility certificates	66
16.1	Conic duality	66
16.2	Unbounded problems	67
16.3	Infeasible problems	68

17	Naming conventions	69
17.1	Sets	69
V	API Reference	70
18	Standard form	71
18.1	Functions	71
18.2	Sets	74
18.3	Scalar sets	76
18.4	Vector sets	78
18.5	Matrix sets	82
19	Models	87
19.1	Attribute interface	87
19.2	Model interface	91
19.3	Model attributes	93
19.4	Optimizer interface	95
19.5	Optimizer attributes	96
20	Variables	104
20.1	Functions	104
20.2	Attributes	107
21	Constraints	109
21.1	Types	109
21.2	Functions	109
21.3	Attributes	111
22	Modifications	115
23	Nonlinear programming	118
23.1	Types	118
23.2	Attributes	119
23.3	Functions	119
24	Callbacks	122
24.1	Attributes	122
24.2	Lazy constraints	123
24.3	User cuts	124
24.4	Heuristic solutions	125
25	Errors	127
VI	Submodules	132
26	Benchmarks	133
26.1	Overview	133
26.2	API Reference	134
27	Bridges	136
27.1	Overview	136
27.2	Implementation	139
27.3	API Reference	143
28	FileFormats	163
28.1	Overview	163
28.2	API Reference	167
29	Nonlinear	171
29.1	Overview	171
29.2	API Reference	181
30	Utilities	191
30.1	Overview	191

	30.2	API Reference	199
31	Test		225
	31.1	Overview	225
	31.2	API Reference	230

Part I

Introduction

Chapter 1

Introduction

Welcome to the documentation for MathOptInterface.

Note

This documentation is also available in PDF format: [MathOptInterface.pdf](#).

1.1 What is MathOptInterface?

[MathOptInterface.jl](#) (MOI) is an abstraction layer designed to provide a unified interface to mathematical optimization solvers so that users do not need to understand multiple solver-specific APIs.

Tip

This documentation is aimed at developers writing software interfaces to solvers and modeling languages using the MathOptInterface API. If you are a user interested in solving optimization problems, we encourage you instead to use MOI through a higher-level modeling interface like [JuMP](#) or [Convex.jl](#).

1.2 How the documentation is structured

Having a high-level overview of how this documentation is structured will help you know where to look for certain things.

- The **Tutorials** section contains articles on how to use and implement the MathOptInterface API. Look here if you want to write a model in MOI, or write an interface to a new solver.
- The **Manual** contains short code-snippets that explain how to use the MOI API. Look here for more details on particular areas of MOI.
- The **Background** section contains articles on the theory behind MathOptInterface. Look here if you want to understand why, rather than how.
- The **API Reference** contains a complete list of functions and types that comprise the MOI API. Look here if you want to know how to use (or implement) a particular function.
- The **Submodules** section contains stand-alone documentation for each of the submodules within MOI. These submodules are not required to interface a solver with MOI, but they make the job much easier.

1.3 Citing MathOptInterface

A [paper describing the design and features of MathOptInterface](#) is available on [arXiv](#).

If you find MathOptInterface useful in your work, we kindly request that you cite the following paper:

```
@article{legat2021mathoptinterface,  
  title={MathOptInterface: a data structure for mathematical optimization problems},  
  author={Legat, Beno{\^i}t and Dowson, Oscar and Garcia, Joaquim Dias and Lubin, Miles},  
  journal={INFORMS Journal on Computing},  
  year={2021},  
  doi={10.1287/ijoc.2021.1067},  
  publisher={INFORMS}  
}
```


Chapter 2

Motivation

MathOptInterface (MOI) is a replacement for [MathProgBase](#), the first-generation abstraction layer for mathematical optimization previously used by [JuMP](#) and [Convex.jl](#).

To address a number of limitations of MathProgBase, MOI is designed to:

- Be simple and extensible
 - unifying linear, quadratic, and conic optimization,
 - seamlessly facilitating extensions to essentially arbitrary constraints and functions (e.g., indicator constraints, complementarity constraints, and piecewise-linear functions)
- Be fast
 - by allowing access to a solver's in-memory representation of a problem without writing intermediate files (when possible)
 - by using multiple dispatch and avoiding requiring containers of nonconcrete types
- Allow a solver to return multiple results (e.g., a pool of solutions)
- Allow a solver to return extra arbitrary information via attributes (e.g., variable- and constraint-wise membership in an irreducible inconsistent subset for infeasibility analysis)
- Provide a greatly expanded set of status codes explaining what happened during the optimization procedure
- Enable a solver to more precisely specify which problem classes it supports
- Enable both primal and dual warm starts
- Enable adding and removing both variables and constraints by indices that are not required to be consecutive
- Enable any modification that the solver supports to an existing model
- Avoid requiring the solver wrapper to store an additional copy of the problem data

Part II

Tutorials

Chapter 3

Solving a problem using MathOptInterface

In this tutorial we demonstrate how to use MathOptInterface to solve the binary-constrained knapsack problem:

$$\begin{aligned} \max \quad & c^\top x \\ \text{s.t.} \quad & w^\top x \leq C \\ & x_i \in \{0, 1\}, \quad \forall i = 1, \dots, n \end{aligned}$$

3.1 Required packages

Load the MathOptInterface module and define the shorthand MOI:

```
using MathOptInterface
const MOI = MathOptInterface
```

As an optimizer, we choose GLPK:

```
using GLPK
optimizer = GLPK.Optimizer()
```

3.2 Define the data

We first define the constants of the problem:

```
julia> c = [1.0, 2.0, 3.0]
3-element Vector{Float64}:
 1.0
 2.0
 3.0

julia> w = [0.3, 0.5, 1.0]
3-element Vector{Float64}:
 0.3
 0.5
 1.0

julia> C = 3.2
3.2
```

3.3 Add the variables

```
| julia> x = MOI.add_variables(optimizer, length(c));
```

3.4 Set the objective

```
| julia> MOI.set(
    optimizer,
    MOI.ObjectiveFunction{MOI.ScalarAffineFunction{Float64}}{()},
    MOI.ScalarAffineFunction(MOI.ScalarAffineTerm.(c, x), 0.0),
    );
| julia> MOI.set(optimizer, MOI.ObjectiveSense(), MOI.MAX_SENSE)
```

Tip

`MOI.ScalarAffineTerm.(c, x)` is a shortcut for `[MOI.ScalarAffineTerm(c[i], x[i]) for i = 1:3]`. This is Julia's broadcast syntax in action, and is used quite often throughout MOI.

3.5 Add the constraints

We add the knapsack constraint and integrality constraints:

```
| julia> MOI.add_constraint(
    optimizer,
    MOI.ScalarAffineFunction(MOI.ScalarAffineTerm.(w, x), 0.0),
    MOI.LessThan(C),
    );
```

Add integrality constraints:

```
| julia> for x_i in x
    MOI.add_constraint(optimizer, x_i, MOI.ZeroOne())
end
```

3.6 Optimize the model

```
| julia> MOI.optimize!(optimizer)
```

3.7 Understand why the solver stopped

The first thing to check after optimization is why the solver stopped, e.g., did it stop because of a time limit or did it stop because it found the optimal solution?

```
| julia> MOI.get(optimizer, MOI.TerminationStatus())
OPTIMAL::TerminationStatusCode = 1
```

Looks like we found an optimal solution!

3.8 Understand what solution was returned

```
julia> MOI.get(optimizer, MOI.ResultCount())
1

julia> MOI.get(optimizer, MOI.PrimalStatus())
FEASIBLE_POINT::ResultStatusCode = 1

julia> MOI.get(optimizer, MOI.DualStatus())
NO_SOLUTION::ResultStatusCode = 0
```

3.9 Query the objective

What is its objective value?

```
julia> MOI.get(optimizer, MOI.ObjectiveValue())
6.0
```

3.10 Query the primal solution

And what is the value of the variables x?

```
julia> MOI.get(optimizer, MOI.VariablePrimal(), x)
3-element Vector{Float64}:
 1.0
 1.0
 1.0
```

Chapter 4

Implementing a solver interface

This guide outlines the basic steps to implement an interface to MathOptInterface for a new solver.

Danger

Implementing an interface to MathOptInterface for a new solver is a lot of work. Before starting, we recommend that you join the [Developer chatroom](#) and explain a little bit about the solver you are wrapping. If you have questions that are not answered by this guide, please ask them in the [Developer chatroom](#) so we can improve this guide!

4.1 A note on the API

The API of MathOptInterface is large and varied. In order to support the diversity of solvers and use-cases, we make heavy use of [duck-typing](#). That is, solvers are not expected to implement the full API, nor is there a well-defined minimal subset of what must be implemented. Instead, you should implement the API as necessary in order to make the solver function as you require.

The main reason for using duck-typing is that solvers work in different ways and target different use-cases.

For example:

- Some solvers support incremental problem construction, support modification after a solve, and have native support for things like variable names.
- Other solvers are "one-shot" solvers that require all of the problem data to construct and solve the problem in a single function call. They do not support modification or things like variable names.
- Other "solvers" are not solvers at all, but things like file readers. These may only support functions like [read_from_file](#), and may not even support the ability to add variables or constraints directly!
- Finally, some "solvers" are layers which take a problem as input, transform it according to some rules, and pass the transformed problem to an inner solver.

4.2 Preliminaries

Before starting on your wrapper, you should do some background research and make the solver accessible via Julia.

Decide if MathOptInterface is right for you

The first step in writing a wrapper is to decide whether implementing an interface is the right thing to do.

MathOptInterface is an abstraction layer for unifying constrained mathematical optimization solvers. If your solver doesn't fit in the category, i.e., it implements a derivative-free algorithm for unconstrained objective functions, MathOptInterface may not be the right tool for the job.

Tip

If you're not sure whether you should write an interface, ask in the [Developer chatroom](#).

Find a similar solver already wrapped

The next step is to find (if possible) a similar solver that is already wrapped. Although not strictly necessary, this will be a good place to look for inspiration when implementing your wrapper.

The [JuMP documentation](#) has a good list of solvers, along with the problem classes they support.

Tip

If you're not sure which solver is most similar, ask in the [Developer chatroom](#).

Create a low-level interface

Before writing a MathOptInterface wrapper, you first need to be able to call the solver from Julia.

Wrapping solvers written in Julia

If your solver is written in Julia, there's nothing to do here! Go to the next section.

Wrapping solvers written in C

Julia is well suited to wrapping solvers written in C.

Info

This is not true for C++. If you have a solver written in C++, first write a C interface, then wrap the C interface.

Before writing a MathOptInterface wrapper, there are a few extra steps.

Create a JLL If the C code is publicly available under an open-source license, create a JLL package via [Yggdrasil](#). The easiest way to do this is to copy an existing solver. Good examples to follow are the [COIN-OR solvers](#).

Warning

Building the solver via Yggdrasil is non-trivial. Please ask the [Developer chatroom](#) for help.

If the code is commercial or not publicly available, the user will need to manually install the solver. See [Gurobi.jl](#) or [CPLEX.jl](#) for examples of how to structure this.

Use Clang.jl to wrap the C API The next step is to use [Clang.jl](#) to automatically wrap the C API. The easiest way to do this is to follow an example. Good examples to follow are [Cbc.jl](#) and [HiGHS.jl](#).

Sometimes, you will need to make manual modifications to the resulting files.

Solvers written in other languages

Ask the [Developer chatroom](#) for advice. You may be able to use one of the [JuliaInterop](#) packages to call out to the solver.

For example, [SeDuMi.jl](#) uses [MATLAB.jl](#) to call the SeDuMi solver written in MATLAB.

4.3 Structuring the package

Structure your wrapper as a Julia package. Consult the [Julia documentation](#) if you haven't done this before.

MOI solver interfaces may be in the same package as the solver itself (either the C wrapper if the solver is accessible through C, or the Julia code if the solver is written in Julia, for example), or in a separate package which depends on the solver package.

Note

The JuMP [core contributors](#) request that you do not use "JuMP" in the name of your package without prior consent.

Your package should have the following structure:

```
/.github
  /workflows
    ci.yml
    format_check.yml
    TagBot.yml
/gen
  gen.jl # Code to wrap the C API
/src
  NewSolver.jl
  /gen
    libnewsolver_api.jl
    libnewsolver_common.jl
  /MOI_wrapper
    MOI_wrapper.jl
    other_files.jl
/test
  runtests.jl
  /MOI_wrapper
    MOI_wrapper.jl
.gitignore
.JuliaFormatter.toml
README.md
LICENSE.md
Project.toml
```

- The `/.github` folder contains the scripts for GitHub actions. The easiest way to write these is to copy the ones from an existing solver.
- The `/gen` and `/src/gen` folders are only needed if you are wrapping a [solver written in C](#).

- The `/src/MOI_wrapper` folder contains the Julia code for the MOI wrapper.
- The `/test` folder contains code for testing your package. See [Setup tests](#) for more information.
- The `.JuliaFormatter.toml` and `.github/workflows/format_check.yml` enforce code formatting using [JuliaFormatter.jl](#). Check existing solvers or `JuMP.jl` for details.

Documentation

Your package must include documentation explaining how to use the package. The easiest approach is to include documentation in your `README.md`. A more involved option is to use [Documenter.jl](#).

Examples of packages with README-based documentation include:

- [Cbc.jl](#)
- [HiGHS.jl](#)
- [SCS.jl](#)

Examples of packages with Documenter-based documentation include:

- [Alpine.jl](#)
- [COSMO.jl](#)
- [Juniper.jl](#)

Setup tests

The best way to implement an interface to `MathOptInterface` is via [test-driven development](#).

The `MOI.Test` submodule contains a large test suite to help check that you have implemented things correctly.

Follow the guide [How to test a solver](#) to set up the tests for your package.

Tip

Run the tests frequently when developing. However, at the start there is going to be a lot of errors! Start by excluding large classes of tests (e.g., `exclude = ["test_basic_", "test_model_"]`), implement any missing methods until the tests pass, then remove an exclusion and repeat.

4.4 Initial code

By this point, you should have a package setup with tests, formatting, and access to the underlying solver. Now it's time to start writing the wrapper.

The Optimizer object

The first object to create is a subtype of [AbstractOptimizer](#). This type is going to store everything related to the problem.

By convention, these optimizers should not be exported and should be named `PackageName.Optimizer`.

```

import MathOptInterface
const MOI = MathOptInterface

struct Optimizer <: MOI.AbstractOptimizer
    # Fields go here
end

```

Optimizer objects for C solvers

Warning

This section is important if you wrap a solver written in C.

Wrapping a solver written in C will require the use of pointers, and for you to manually free the solver's memory when the `Optimizer` is garbage collected by Julia.

Never pass a pointer directly to a Julia `ccall` function.

Instead, store the pointer as a field in your `Optimizer`, and implement `Base.cconvert` and `Base.unsafe_convert`. Then you can pass `Optimizer` to any `ccall` function that expects the pointer.

In addition, make sure you implement a finalizer for each model you create.

If `newsolver_createProblem()` is the low-level function that creates the problem pointer in C, and `newsolver_freeProblem(::Ptr{Cvoid})` is the low-level function that frees memory associated with the pointer, your `Optimizer()` function should look like this:

```

struct Optimizer <: MOI.AbstractOptimizer
    ptr::Ptr{Cvoid}

    function Optimizer()
        ptr = newsolver_createProblem()
        model = Optimizer(ptr)
        finalizer(model) do m
            newsolver_freeProblem(m)
        end
        return model
    end
end

Base.cconvert(::Type{Ptr{Cvoid}}, model::Optimizer) = model.ptr
Base.unsafe_convert(::Type{Ptr{Cvoid}}, model::Optimizer) = model.ptr

```

Implement methods for `Optimizer`

All `Optimizers` must implement the following methods:

- `empty!`
- `is_empty`
- `optimize!`

Other methods, detailed below, are optional or depend on how you implement the interface.

Tip

For this and all future methods, read the docstrings to understand what each method does, what it expects as input, and what it produces as output. If it isn't clear, let us know and we will improve the docstrings! It is also very helpful to look at an existing wrapper for a similar solver.

You should also implement `Base.show(::IO, ::Optimizer)` to print a nice string when someone prints your model. For example

```
function Base.show(io::IO, model::Optimizer)
    return print(io, "NewSolver with the pointer $(model.ptr)")
end
```

Implement attributes

MathOptInterface uses attributes to manage different aspects of the problem.

For each attribute

- `get` gets the current value of the attribute
- `set` sets a new value of the attribute. Not all attributes can be set. For example, the user can't modify the `SolverName`.
- `supports` returns a `Bool` indicating whether the solver supports the attribute.

Info

Use `attribute_value_type` to check the value expected by a given attribute. You should make sure that your `get` function correctly infers to this type (or a subtype of it).

Each column in the table indicates whether you need to implement the particular method for each attribute.

Attribute	<code>get</code>	<code>set</code>	<code>supports</code>
<code>SolverName</code>	Yes	No	No
<code>SolverVersion</code>	Yes	No	No
<code>RawSolver</code>	Yes	No	No
<code>Name</code>	Yes	Yes	Yes
<code>Silent</code>	Yes	Yes	Yes
<code>TimeLimitSec</code>	Yes	Yes	Yes
<code>RawOptimizerAttribute</code>	Yes	Yes	Yes
<code>NumberOfThreads</code>	Yes	Yes	Yes

For example:

```
function MOI.get(model::Optimizer, ::MOI.Silent)
    return # true if MOI.Silent is set
end

function MOI.set(model::Optimizer, ::MOI.Silent, v::Bool)
    if v
        # Set a parameter to turn off printing
    else
    end
end
```

```

        # Restore the default printing
    end
    return
end

MOI.supports(::Optimizer, ::MOI.Silent) = true

```

Define `supports_constraint`

The next step is to define which constraints and objective functions you plan to support.

For each function-set constraint pair, define `supports_constraint`:

```

function MOI.supports_constraint(
    ::Optimizer,
    ::Type{MOI.VariableIndex},
    ::Type{MOI.ZeroOne},
)
    return true
end

```

To make this easier, you may want to use Unions:

```

function MOI.supports_constraint(
    ::Optimizer,
    ::Type{MOI.VariableIndex},
    ::Type{<:Union{MOI.LessThan, MOI.GreaterThan, MOI.EqualTo}},
)
    return true
end

```

Tip

Only support a constraint if your solver has native support for it.

4.5 The big decision: copy-to or incremental modifications?

Now you need to decide whether to support incremental modification or not.

Incremental modification means that the user can add variables and constraints one-by-one without needing to rebuild the entire problem, and they can modify the problem data after an `optimize!` call. Supporting incremental modification means implementing functions like `add_variable` and `add_constraint`.

The alternative is to accept the problem data in a single `copy_to` function call, after which it cannot be modified. Because `copy_to` sees all of the data at once, it can typically call a more efficient function to load data into the underlying solver.

Good examples of solvers supporting incremental modification are MILP solvers like `GLPK.jl` and `Gurobi.jl`. Examples of `copy_to` solvers are `AmplNLWriter.jl` and `SCS.jl`.

It is possible to implement both approaches, but you should probably start with one for simplicity.

Tip

Only support incremental modification if your solver has native support for it.

In general, supporting incremental modification is more work, and it usually requires some extra book-keeping. However, it provides a more efficient interface to the solver if the problem is going to be resolved multiple times with small modifications. Moreover, once you've implemented incremental modification, it's usually not much extra work to add a `copy_to` interface. The converse is not true.

Tip

If this is your first time writing an interface, start with `copy_to`.

The `copy_to` interface

To implement the `copy_to` interface, implement the following function:

- `copy_to`

The incremental interface

Warning

Writing this interface is a lot of work. The easiest way is to consult the source code of a similar solver!

To implement the incremental interface, implement the following functions:

- `add_variable`
- `add_variables`
- `add_constraint`
- `add_constraints`
- `is_valid`
- `delete`

Info

Solvers do not have to support `AbstractScalarFunction` in `GreaterThan`, `LessThan`, `EqualTo`, or `Interval` with a nonzero constant in the function. Throw `ScalarFunctionConstantNotZero` if the function constant is not zero.

In addition, you should implement the following model attributes:

Attribute	<code>get</code>	<code>set</code>	<code>supports</code>
<code>ListOfModelAttributesSet</code>	Yes	No	No
<code>ObjectiveFunctionType</code>	Yes	No	No
<code>ObjectiveFunction</code>	Yes	Yes	Yes
<code>ObjectiveSense</code>	Yes	Yes	Yes
<code>Name</code>	Yes	Yes	Yes

Variable-related attributes:

Constraint-related attributes:

Attribute	get	set	supports
ListOfVariableAttributesSet	Yes	No	No
NumberOfVariables	Yes	No	No
ListOfVariableIndices	Yes	No	No

Attribute	get	set	supports
ListOfConstraintAttributesSet	Yes	No	No
NumberOfConstraints	Yes	No	No
ListOfConstraintTypesPresent	Yes	No	No
ConstraintFunction	Yes	Yes	No
ConstraintSet	Yes	Yes	No

Modifications

If your solver supports modifying data in-place, implement `modify` for the following AbstractModifications:

- `ScalarConstantChange`
- `ScalarCoefficientChange`
- `VectorConstantChange`
- `MultirowChange`

Variables constrained on creation

Some solvers require variables be associated with a set when they are created. This conflicts with the incremental modification approach, since you cannot first add a free variable and then constrain it to the set.

If this is the case, implement:

- `add_constrained_variable`
- `add_constrained_variables`
- `supports_add_constrained_variables`

By default, MathOptInterface assumes solvers support free variables. If your solver does not support free variables, define:

```
MOI.supports_add_constrained_variables(::Optimizer, ::Type{Reals}) = false
```

Incremental and copy_to

If you implement the incremental interface, you have the option of also implementing `copy_to`.

If you don't want to implement `copy_to`, e.g., because the solver has no API for building the problem in a single function call, define the following fallback:

```
MOI.supports_incremental_interface(::Optimizer) = true

function MOI.copy_to(dest::Optimizer, src::MOI.ModelLike)
    return MOI.Utilities.default_copy_to(dest, src)
end
```

4.6 Names

Regardless of which interface you implement, you have the option of implementing the `Name` attribute for variables and constraints:

Attribute	get	set	supports
<code>VariableName</code>	Yes	Yes	Yes
<code>ConstraintName</code>	Yes	Yes	Yes

If you implement names, you must also implement the following three methods:

```
function MOI.get(model::Optimizer, ::Type{MOI.VariableIndex}, name::String)
    return # The variable named `name`.
end

function MOI.get(model::Optimizer, ::Type{MOI.ConstraintIndex}, name::String)
    return # The constraint any type named `name`.
end

function MOI.get(
    model::Optimizer,
    ::Type{MOI.ConstraintIndex{F,S}},
    name::String,
) where {F,S}
    return # The constraint of type F-in-S named `name`.
end
```

These methods have the following rules:

- If there is no variable or constraint with the name, return nothing
- If there is a single variable or constraint with that name, return the variable or constraint
- If there are multiple variables or constraints with the name, throw an error.

Warning

You should not implement `ConstraintName` for `VariableIndex` constraints. If you implement `ConstraintName` for other constraints, you can add the following two methods to disable `ConstraintName` for `VariableIndex` constraints.

```
function MOI.supports(
    ::Optimizer,
    ::MOI.ConstraintName,
    ::Type{<:MOI.ConstraintIndex{MOI.VariableIndex,<:MOI.AbstractScalarSet}},
)
    return throw(MOI.VariableIndexConstraintNameError())
end

function MOI.set(
    ::Optimizer,
    ::MOI.ConstraintName,
    ::MOI.ConstraintIndex{MOI.VariableIndex,<:MOI.AbstractScalarSet},
    ::String,
)
    return throw(MOI.VariableIndexConstraintNameError())
end
```

4.7 Solutions

Implement `optimize!` to solve the model:

- `optimize!`

All Optimizers must implement the following attributes:

- `DualStatus`
- `PrimalStatus`
- `RawStatusString`
- `ResultCount`
- `TerminationStatus`

Info

You only need to implement `get` for solution attributes. Don't implement `set` or `supports`.

Note

Solver wrappers should document how the low-level statuses map to the MOI statuses. Statuses like `NEARLY_FEASIBLE_POINT` and `INFEASIBLE_POINT`, are designed to be used when the solver explicitly indicates that relaxed tolerances are satisfied or the returned point is infeasible, respectively.

You should also implement the following attributes:

- `ObjectiveValue`
- `SolveTimeSec`
- `VariablePrimal`

Tip

Attributes like `VariablePrimal` and `ObjectiveValue` are indexed by the result count. Use `MOI.check_result_index_bound(attr)` to throw an error if the attribute is not available.

If your solver returns dual solutions, implement:

- `ConstraintDual`
- `DualObjectiveValue`

For integer solvers, implement:

- `ObjectiveBound`
- `RelativeGap`

If applicable, implement:

- [SimplexIterations](#)
- [BarrierIterations](#)
- [NodeCount](#)

If your solver uses the Simplex method, implement:

- [ConstraintBasisStatus](#)

If your solver accepts primal or dual warm-starts, implement:

- [VariablePrimalStart](#)
- [ConstraintDualStart](#)

4.8 Other tips

Here are some other points to be aware of when writing your wrapper.

Unsupported constraints at runtime

In some cases, your solver may support a particular type of constraint (e.g., quadratic constraints), but only if the data meets some condition (e.g., it is convex).

In this case, declare that you support the constraint, and throw [AddConstraintNotAllowed](#).

Dealing with multiple variable bounds

MathOptInterface uses [VariableIndex](#) constraints to represent variable bounds. Defining multiple variable bounds on a single variable is not allowed.

Throw [LowerBoundAlreadySet](#) or [UpperBoundAlreadySet](#) if the user adds a constraint that results in multiple bounds.

Only throw if the constraints conflict. It is okay to add [VariableIndex-in-GreaterThan](#) and then [VariableIndex-in-LessThan](#), but not [VariableIndex-in-Interval](#) and then [VariableIndex-in-LessThan](#),

Expect duplicate coefficients

Solvers must expect that functions such as [ScalarAffineFunction](#) and [VectorQuadraticFunction](#) may contain duplicate coefficients.

For example, `ScalarAffineFunction([ScalarAffineTerm(x, 1), ScalarAffineTerm(x, 1)], 0.0)`.

Use [Utilities.canonical](#) to return a new function with the duplicate coefficients aggregated together.

Don't modify user-data

All data passed to the solver must be copied immediately to internal data structures. Solvers may not modify any input vectors and must assume that input vectors will not be modified by users in the future.

This applies, for example, to the terms vector in [ScalarAffineFunction](#). Vectors returned to the user, e.g., via [ObjectiveFunction](#) or [ConstraintFunction](#) attributes, must not be modified by the solver afterwards. The in-place version of [get!](#) can be used by users to avoid extra copies in this case.

Column Generation

There is no special interface for column generation. If the solver has a special API for setting coefficients in existing constraints when adding a new variable, it is possible to queue modifications and new variables and then call the solver's API once all of the new coefficients are known.

Solver-specific attributes

You don't need to restrict yourself to the attributes defined in the MathOptInterface.jl package.

Solver-specific attributes should be specified by creating an appropriate subtype of [AbstractModelAttribute](#), [AbstractOptimizerAttribute](#), [AbstractVariableAttribute](#), or [AbstractConstraintAttribute](#).

For example, Gurobi.jl adds attributes for multiobjective optimization by [defining](#):

```
struct NumberOfObjectives <: MOI.AbstractModelAttribute end

function MOI.set(model::Optimizer, ::NumberOfObjectives, n::Integer)
    # Code to set NumberOfObjectives
    return
end

function MOI.get(model::Optimizer, ::NumberOfObjectives)
    n = # Code to get NumberOfObjectives
    return n
end
```

Then, the user can write:

```
model = Gurobi.Optimizer()
MOI.set(model, Gurobi.NumberofObjectives(), 3)
```

Chapter 5

Transitioning from MathProgBase

MathOptInterface is a replacement for [MathProgBase.jl](#). However, it is not a direct replacement.

5.1 Transitioning a solver interface

MathOptInterface is more extensive than MathProgBase which may make its implementation seem daunting at first. There are however numerous utilities in MathOptInterface that simplify the implementation process.

For more information, read [Implementing a solver interface](#).

5.2 Transitioning the high-level functions

MathOptInterface doesn't provide replacements for the high-level interfaces in MathProgBase. We recommend you use [JuMP](#) as a modeling interface instead.

Tip

If you haven't used JuMP before, start with the tutorial [Getting started with JuMP](#)

linprog

Here is one way of transitioning from linprog:

```
using JuMP

function linprog(c, A, sense, b, l, u, solver)
    N = length(c)
    model = Model(solver)
    @variable(model, l[i] <= x[i=1:N] <= u[i])
    @objective(model, Min, c' * x)
    eq_rows, ge_rows, le_rows = sense .== '=', sense .== '>', sense .== '<'
    @constraint(model, A[eq_rows, :] * x .== b[eq_rows])
    @constraint(model, A[ge_rows, :] * x .>= b[ge_rows])
    @constraint(model, A[le_rows, :] * x .<= b[le_rows])
    optimize!(model)
    return (
        status = termination_status(model),
        objval = objective_value(model),
        sol = value.(x)
    )
end
```

mixintprog

Here is one way of transitioning from mixintprog:

```
using JuMP

function mixintprog(c, A, rowlb, rowub, vartypes, lb, ub, solver)
    N = length(c)
    model = Model(solver)
    @variable(model, lb[i] <= x[i=1:N] <= ub[i])
    for i in 1:N
        if vartypes[i] == :Bin
            set_binary(x[i])
        elseif vartypes[i] == :Int
            set_integer(x[i])
        end
    end
    @objective(model, Min, c' * x)
    @constraint(model, rowlb .<= A * x .<= rowub)
    optimize!(model)
    return (
        status = termination_status(model),
        objval = objective_value(model),
        sol = value.(x)
    )
end
```

quadprog

Here is one way of transitioning from quadprog:

```
using JuMP

function quadprog(c, Q, A, rowlb, rowub, lb, ub, solver)
    N = length(c)
    model = Model(solver)
    @variable(model, lb[i] <= x[i=1:N] <= ub[i])
    @objective(model, Min, c' * x + 0.5 * x' * Q * x)
    @constraint(model, rowlb .<= A * x .<= rowub)
    optimize!(model)
    return (
        status = termination_status(model),
        objval = objective_value(model),
        sol = value.(x)
    )
end
```

Chapter 6

Implementing a constraint bridge

This guide outlines the basic steps to create a new bridge from a constraint expressed in the formalism Function-in-Set.

6.1 Preliminaries

First, decide on the set you want to bridge. Then, study its properties: the most important one is whether the set is scalar or vector, which impacts the dimensionality of the functions that can be used with the set.

- A scalar function only has one dimension. MOI defines three types of scalar functions: a variable ([VariableIndex](#)), an affine function ([ScalarAffineFunction](#)), or a quadratic function ([ScalarQuadraticFunction](#)).
- A vector function has several dimensions (at least one). MOI defines three types of vector functions: several variables ([VectorOfVariables](#)), an affine function ([VectorAffineFunction](#)), or a quadratic function ([VectorQuadraticFunction](#)). The main difference with scalar functions is that the order of dimensions can be very important: for instance, in an indicator constraint ([Indicator](#)), the first dimension indicates whether the constraint about the second dimension is active.

To explain how to implement a bridge, we present the example of [Bridges.Constraint.FlipSignBridge](#). This bridge maps \leq ([LessThan](#)) constraints to \geq ([GreaterThan](#)) constraints. This corresponds to reversing the sign of the inequality. We focus on scalar affine functions (we disregard the cases of a single variable or of quadratic functions). This example is a simplified version of the code included in MOI.

6.2 Four mandatory parts in a constraint bridge

The first part of a constraint bridge is a new concrete subtype of [Bridges.Constraint.AbstractBridge](#). This type must have fields to store all the new variables and constraints that the bridge will add. Typically, these types are parametrized by the type of the coefficients in the model.

Then, three sets of functions must be defined:

1. [Bridges.Constraint.bridge_constraint](#): this function implements the bridge and creates the required variables and constraints.
2. [supports_constraint](#): these functions must return true when the combination of function and set is supported by the bridge. By default, the base implementation always returns false and the bridge does not have to provide this implementation.

3. `Bridges.added_constrained_variable_types` and `Bridges.added_constraint_types`: these functions return the types of variables and constraints that this bridge adds. They are used to compute the set of other bridges that are required to use the one you are defining, if need be.

More functions can be implemented, for instance to retrieve properties from the bridge or deleting a bridged constraint.

1. Structure for the bridge

A typical struct behind a bridge depends on the type of the coefficients that are used for the model (typically `Float64`, but coefficients might also be integers or complex numbers).

This structure must hold a reference to all the variables and the constraints that are created as part of the bridge.

The type of this structure is used throughout MOI as an identifier for the bridge. It is passed as argument to most functions related to bridges.

The best practice is to have the name of this type end with `Bridge`.

In our example, the bridge maps any `ScalarAffineFunction{T}-in-LessThan{T}` constraint to a single `ScalarAffineFunction{T}-in-GreaterThan{T}` constraint. The affine function has coefficients of type `T`. The bridge is parametrized with `T`, so that the constraint that the bridge creates also has coefficients of type `T`.

```
struct SignBridge{T<:Number} <: Bridges.Constraint.AbstractBridge
    constraint::ConstraintIndex{ScalarAffineFunction{T}, GreaterThan{T}}
end
```

2. Bridge creation

The function `Bridges.Constraint.bridge_constraint` is called whenever the bridge is instantiated for a specific model, with the given function and set. The arguments to `bridge_constraint` are similar to `add_constraint`, with the exception of the first argument: it is the Type of the struct defined in the first step (for our example, `Type{SignBridge{T}}`).

`bridge_constraint` returns an instance of the struct defined in the first step. the first step.

In our example, the bridge constraint could be defined as:

```
function Bridges.Constraint.bridge_constraint(
    ::Type{SignBridge{T}}, # Bridge to use.
    model::ModelLike, # Model to which the constraint is being added.
    f::ScalarAffineFunction{T}, # Function to rewrite.
    s::LessThan{T}, # Set to rewrite.
) where {T}
    # Create the variables and constraints required for the bridge.
    con = add_constraint(model, -f, GreaterThan(-s.upper))

    # Return an instance of the bridge type with a reference to all the
    # variables and constraints that were created in this function.
    return SignBridge(con)
end
```

3. Supported constraint types

The function `supports_constraint` determines whether the bridge type supports a given combination of function and set.

This function must closely match `bridge_constraint`, because it will not be called if `supports_constraint` returns false.

```
function supports_constraint(
  ::Type{SignBridge{T}}, # Bridge to use.
  ::Type{ScalarAffineFunction{T}}, # Function to rewrite.
  ::Type{LessThan{T}}, # Set to rewrite.
) where {T}
  # Do some computation to ensure that the constraint is supported.
  # Typically, you can directly return true.
  return true
end
```

4. Metadata about the bridge

To determine whether a bridge can be used, MOI uses a shortest-path algorithm that uses the variable types and the constraints that the bridge can create. This information is communicated from the bridge to MOI using the functions `Bridges.added_constrained_variable_types` and `Bridges.added_constraint_types`. Both return lists of tuples: either a list of 1-tuples containing the variable types (typically, `ZeroOne` or `Integer`) or a list of 2-tuples containing the functions and sets (like `ScalarAffineFunction{T}-GreaterThan`).

For our example, the bridge does not create any constrained variables, and only `ScalarAffineFunction{T}-in-GreaterThan{T}` constraints:

```
function Bridges.added_constrained_variable_types(::Type{SignBridge{T}}) where {T}
  # The bridge does not create variables, return an empty list of tuples:
  return Tuple{Type}[]
end

function Bridges.added_constraint_types(::Type{SignBridge{T}}) where {T}
  return Tuple{Type,Type}[
    # One element per F-in-S the bridge creates.
    (ScalarAffineFunction{T}, GreaterThan{T}),
  ]
end
```

A bridge that creates binary variables would rather have this definition of `added_constrained_variable_types`:

```
function Bridges.added_constrained_variable_types(::Type{SomeBridge{T}}) where {T}
  # The bridge only creates binary variables:
  return Tuple{Type}[(ZeroOne,)]
end
```

Warning

If you declare the creation of constrained variables in `added_constrained_variable_types`, the corresponding constraint type `VariableIndex` must not be indicated in `added_constraint_types`. This would restrict the use of the bridge to solvers that can add such a constraint after the variable is created.

More concretely, if you declare in `added_constrained_variable_types` that your bridge creates binary variables (`ZeroOne`), and if you never add such a constraint afterward (you do not call `add_constraint(model, var, ZeroOne())`), then you must not list (`VariableIndex, ZeroOne`) in `added_constraint_types`.

Typically, the function `Bridges.Constraint.concrete_bridge_type` does not have to be defined for most bridges.

6.3 Bridge registration

For a bridge to be used by MOI, it must be known by MOI.

SingleBridgeOptimizer

The first way to do so is to create a single-bridge optimizer. This type of optimizer wraps another optimizer and adds the possibility to use only one bridge. It is especially useful when unit testing bridges.

It is common practice to use the same name as the type defined for the bridge (`SignBridge`, in our example) without the suffix `Bridge`.

```
const Sign{T,OT<: ModelLike} =
    SingleBridgeOptimizer{SignBridge{T}, OT}
```

In the context of unit tests, this bridge is used in conjunction with a `Utilities.MockOptimizer`:

```
mock = Utilities.MockOptimizer(
    Utilities.UniversalFallback(Utilities.Model{Float64}()),
)
bridged_mock = Sign{Float64}(mock)
```

New bridge for a LazyBridgeOptimizer

Typical user-facing models for MOI are based on `Bridges.LazyBridgeOptimizer`. For instance, this type of model is returned by `Bridges.full_bridge_optimizer`. These models can be added more bridges by using `Bridges.add_bridge`:

```
inner_optimizer = Utilities.Model{Float64}()
optimizer = Bridges.full_bridge_optimizer(inner_optimizer, Float64)
Bridges.add_bridge(optimizer, SignBridge{Float64})
```

6.4 Bridge improvements

Attribute retrieval

Like models, bridges have attributes that can be retrieved using `get` and `set`. The most important ones are the number of variables and constraints, but also the lists of variables and constraints.

In our example, we only have one constraint and only have to implement the `NumberOfConstraints` and `ListOfConstraintIndices` attributes:


```

function get(
  ::SignBridge{T},
  ::NumberOfConstraints{
    ScalarAffineFunction{T},
    GreaterThan{T},
  },
) where {T}
  return 1
end

function get(
  bridge::SignBridge{T},
  ::ListOfConstraintIndices{
    ScalarAffineFunction{T},
    GreaterThan{T},
  },
) where {T}
  return [bridge.constraint]
end

```

You must implement one such pair of functions for each type of constraint the bridge adds to the model.

Warning

Avoid returning a list from the bridge object without copying it. Users must be able to change the contents of the returned list without altering the bridge object.

For variables, the situation is simpler. If your bridge creates new variables, you must implement the [NumberOfVariables](#) and [ListOfVariableIndices](#) attributes. However, these attributes do not have parameters, unlike their constraint counterparts. Only two functions suffice:

```

function get(
  ::SignBridge{T},
  ::NumberOfVariables,
) where {T}
  return 0
end

function get(
  ::SignBridge{T},
  ::ListOfVariableIndices,
) where {T}
  return VariableIndex[]
end

```

Model modifications

To avoid copying the model when the user request to change a constraint, MOI provides [modify](#). Bridges can also implement this API to allow certain changes, such as coefficient changes.

In our case, a modification of a coefficient in the original constraint (i.e. replacing the value of the coefficient of a variable in the affine function) must be transmitted to the constraint created by the bridge, but with a sign change.

```
function modify(  
  model::ModelLike,  
  bridge::SignBridge,  
  change::ScalarCoefficientChange,  
)  
  modify(  
    model,  
    bridge.constraint,  
    ScalarCoefficientChange(change.variable, -change.new_coefficient),  
  )  
  return  
end
```

Bridge deletion

When a bridge is deleted, the constraints it added must be deleted too.

```
function delete(model::ModelLike, bridge::SignBridge)  
  delete(model, bridge.constraint)  
  return  
end
```

Chapter 7

Manipulating expressions

This guide highlights a syntactically appealing way to build expressions at the MOI level, but also to look at their contents. It may be especially useful when writing models or bridge code.

7.1 Creating functions

This section details the ways to create functions with MathOptInterface.

Creating scalar affine functions

The simplest scalar function is simply a variable:

```
julia> x = MOI.add_variable(model) # Create the variable x
MathOptInterface.VariableIndex(1)
```

This type of function is extremely simple; to express more complex functions, other types must be used. For instance, a [ScalarAffineFunction](#) is a sum of linear terms (a factor times a variable) and a constant. Such an object can be built using the standard constructor:

```
julia> f = MOI.ScalarAffineFunction([MOI.ScalarAffineTerm{Int64}(1, x)], 2) # x + 2
MathOptInterface.ScalarAffineFunction{Int64}(MathOptInterface.ScalarAffineTerm{Int64}[MathOptInterface.ScalarAffineTerm{Int64}(1, MathOptInterface.VariableIndex(1))], 2)
```

However, you can also use operators to build the same scalar function:

```
julia> f = x + 2
MathOptInterface.ScalarAffineFunction{Int64}(MathOptInterface.ScalarAffineTerm{Int64}[MathOptInterface.ScalarAffineTerm{Int64}(1, MathOptInterface.VariableIndex(1))], 2)
```

Creating scalar quadratic functions

Scalar quadratic functions are stored in [ScalarQuadraticFunction](#) objects, in a way that is highly similar to scalar affine functions. You can obtain a quadratic function as a product of affine functions:

```
julia> 1 * x * x
MathOptInterface.ScalarQuadraticFunction{Int64}(MathOptInterface.ScalarQuadraticTerm{Int64}[MathOptInterface.ScalarQuadraticTerm{Int64}(1, MathOptInterface.VariableIndex(1), MathOptInterface.VariableIndex(1))],
MathOptInterface.ScalarAffineTerm{Int64}[], 0)
```

```
julia> f * f # (x + 2)^2
MathOptInterface.ScalarQuadraticFunction{Int64}(MathOptInterface.ScalarQuadraticTerm{Int64}[MathOptInterface.ScalarQuadraticTerm{Int64}(MathOptInterface.VariableIndex(1), MathOptInterface.VariableIndex(1))],
↳ MathOptInterface.ScalarAffineTerm{Int64}[MathOptInterface.ScalarAffineTerm{Int64}(2,
↳ MathOptInterface.VariableIndex(1)), MathOptInterface.ScalarAffineTerm{Int64}(2,
↳ MathOptInterface.VariableIndex(1))], 4)

julia> f^2 # (x + 2)^2 too
MathOptInterface.ScalarQuadraticFunction{Int64}(MathOptInterface.ScalarQuadraticTerm{Int64}[MathOptInterface.ScalarQuadraticTerm{Int64}(MathOptInterface.VariableIndex(1), MathOptInterface.VariableIndex(1))],
↳ MathOptInterface.ScalarAffineTerm{Int64}[MathOptInterface.ScalarAffineTerm{Int64}(2,
↳ MathOptInterface.VariableIndex(1)), MathOptInterface.ScalarAffineTerm{Int64}(2,
↳ MathOptInterface.VariableIndex(1))], 4)
```

Creating vector functions

A vector function is a function with several values, irrespective of the number of input variables. Similarly to scalar functions, there are three main types of vector functions: [VectorOfVariables](#), [VectorAffineFunction](#), and [VectorQuadraticFunction](#).

The easiest way to create a vector function is to stack several scalar functions using [Utilities.vectorize](#). It takes a vector as input, and the generated vector function (of the most appropriate type) has each dimension corresponding to a dimension of the vector.

```
julia> g = MOI.Utilities.vectorize([f, 2 * f])
MathOptInterface.VectorAffineFunction{Int64}(MathOptInterface.VectorAffineTerm{Int64}[MathOptInterface.VectorAffineTerm{Int64}(MathOptInterface.ScalarAffineTerm{Int64}(1, MathOptInterface.VariableIndex(1))),
↳ MathOptInterface.VectorAffineTerm{Int64}(2, MathOptInterface.ScalarAffineTerm{Int64}(2,
↳ MathOptInterface.VariableIndex(1)))], [2, 4])
```

Warning

[Utilities.vectorize](#) only takes a vector of similar scalar functions: you cannot mix [VariableIndex](#) and [ScalarAffineFunction](#), for instance. In practice, it means that `Utilities.vectorize([x, f])` does not work; you should rather use `Utilities.vectorize([1 * x, f])` instead to only have [ScalarAffineFunction](#) objects.

7.2 Canonicalizing functions

In more advanced use cases, you might need to ensure that a function is "canonical". Functions are stored as an array of terms, but there is no check that these terms are redundant: a [ScalarAffineFunction](#) object might have two terms with the same variable, like $x + x + 1$. These terms could be merged without changing the semantics of the function: $2x + 1$.

Working with these objects might be cumbersome. Canonicalization helps maintain redundancy to zero.

[Utilities.is_canonical](#) checks whether a function is already in its canonical form:

```
julia> MOI.Utilities.is_canonical(f + f) # (x + 2) + (x + 2) is stored as x + x + 4
false
```

[Utilities.canonical](#) returns the equivalent canonical version of the function:

```
julia> MOI.Utilities.canonical(f + f) # Returns 2x + 4
MathOptInterface.ScalarAffineFunction{Int64}(MathOptInterface.ScalarAffineTerm{Int64}[MathOptInterface.ScalarAffineTerm{Int64}(2, MathOptInterface.VariableIndex(1))], 4)
```

7.3 Exploring functions

At some point, you might need to dig into a function, for instance to map it into solver constructs.

Vector functions

`Utilities.scalarize` returns a vector of scalar functions from a vector function:

```
julia> MOI.Utilities.scalarize(g) # Returns a vector [f, 2 * f].
2-element Vector{MathOptInterface.ScalarAffineFunction{Int64}}:
↪ MathOptInterface.ScalarAffineFunction{Int64}(MathOptInterface.ScalarAffineTerm{Int64}[MathOptInterface.ScalarAffin
↪ MathOptInterface.VariableIndex(1)]], 2)
↪ MathOptInterface.ScalarAffineFunction{Int64}(MathOptInterface.ScalarAffineTerm{Int64}[MathOptInterface.ScalarAffin
↪ MathOptInterface.VariableIndex(1)]], 4)
```

Note

`Utilities.eachscalar` returns an iterator on the dimensions, which serves the same purpose as `Utilities.scalarize`.

`output_dimension` returns the number of dimensions of the output of a function:

```
julia> MOI.output_dimension(g)
2
```

Chapter 8

Latency

MathOptInterface suffers the "time-to-first-solve" problem of start-up latency.

This hurts both the user- and developer-experience of MathOptInterface. In the first case, because simple models have a multi-second delay before solving, and in the latter, because our tests take so long to run!

This page contains some advice on profiling and fixing latency-related problems in the MathOptInterface.jl repository.

8.1 Background

Before reading this part of the documentation, you should familiarize yourself with the reasons for latency in Julia and how to fix them.

- Read the blogposts on julialang.org on [precompilation](#) and [SnoopCompile](#)
- Read the [SnoopCompile](#) documentation.
- Watch Tim Holy's [talk at JuliaCon 2021](#)
- Watch the [package development workshop at JuliaCon 2021](#)

8.2 Causes

There are three main causes of latency in MathOptInterface:

1. A large number of types
2. Lack of method ownership
3. Type-instability in the bridge layer

A large number of types

Julia is very good at specializing method calls based on the input type. Each specialization has a compilation cost, but the benefit of faster run-time performance.

The best-case scenario is for a method to be called a large number of times with a single set of argument types. The worst-case scenario is for a method to be called a single time for a large set of argument types.

Because of MathOptInterface's function-in-set formulation, we fall into the worst-case situation.

This is a fundamental limitation of Julia, so there isn't much we can do about it. However, if we can precompile MathOptInterface, much of the cost can be shifted from start-up latency to the time it takes to precompile a package on installation.

However, there are two things which make MathOptInterface hard to precompile...

Lack of method ownership

Lack of method ownership happens when a call is made using a mix of structs and methods from different modules. Because of this, no single module "owns" the method that is being dispatched, and so it cannot be precompiled.

Tip

This is a slightly simplified explanation. Read the [precompilation tutorial](#) for a more in-depth discussion on back-edges.

Unfortunately, the design of MOI means that this is a frequent occurrence! We have a bunch of types in MOI.Utilities that wrap types defined in external packages (i.e., the Optimizers), which implement methods of functions defined in MOI (e.g., add_variable, add_constraint).

Here's a simple example of method-ownership in practice:

```
module MyMOI
  struct Wrapper{T}
    inner::T
  end
  optimize!(x::Wrapper) = optimize!(x.inner)
end # MyMOI

module MyOptimizer
  using ..MyMOI
  struct Optimizer end
  MyMOI.optimize!(x::Optimizer) = 1
end # MyOptimizer

using SnoopCompile
model = MyMOI.Wrapper(MyOptimizer.Optimizer())

julia> tinf = @snoopi_deep MyMOI.optimize!(model)
InferenceTimingNode: 0.008256/0.008543 on InferenceFrameInfo for Core.Compiler.Timings.ROOT() with 1
↳ direct children
```

The result is that there was one method that required type inference. If we visualize tinf:

```
using ProfileView
ProfileView.view(flamegraph(tinf))
```

we see a flamegraph with a large red-bar indicating that the method `MyMOI.optimize(MyMOI.Wrapper{MyOptimizer.Optimizer})` cannot be precompiled.

To fix this, we need to designate a module to "own" that method (i.e., create a back-edge). The easiest way to do this is for MyOptimizer to call `MyMOI.optimize(MyMOI.Wrapper{MyOptimizer.Optimizer})` during using MyOptimizer. Let's see that in practice:

```

module MyMOI
struct Wrapper{T}
    inner::T
end
optimize(x::Wrapper) = optimize(x.inner)
end # MyMOI

module MyOptimizer
using ..MyMOI
struct Optimizer end
MyMOI.optimize(x::Optimizer) = 1
# The syntax of this let-while loop is very particular:
# * `let ... end` keeps everything local to avoid polluting the MyOptimizer
#   namespace
# * `while true ... break end` runs the code once, and forces Julia to compile
#   the inner loop, rather than interpret it.
let
    while true
        model = MyMOI.Wrapper(Optimizer())
        MyMOI.optimize(model)
        break
    end
end
end # MyOptimizer

using SnoopCompile
model = MyMOI.Wrapper(MyOptimizer.Optimizer())

julia> tinf = @snoopi_deep MyMOI.optimize(model)
InferenceTimingNode: 0.006822/0.006822 on InferenceFrameInfo for Core.Compiler.Timings.ROOT() with 0
↳ direct children

```

There are now 0 direct children that required type inference because the method was already stored in MyOptimizer!

Unfortunately, this trick only works if the call-chain is fully inferrable. If there are breaks (due to type instability), then the benefit of doing this is reduced. And unfortunately for us, the design of MathOptInterface has a lot of type instabilities...

Type instability in the bridge layer

Most of MathOptInterface is pretty good at ensuring type-stability. However, a key component is not type stable, and that is the bridging layer.

In particular, the bridging layer defines [Bridges.LazyBridgeOptimizer](#), which has fields like:

```

struct LazyBridgeOptimizer
    constraint_bridge_types::Vector{Any}
    constraint_node::Dict{Tuple{Type, Type}, ConstraintNode}
    constraint_types::Vector{Tuple{Type, Type}}
end

```

This is because the LazyBridgeOptimizer needs to be able to deal with any function-in-set type passed to it, and we also allow users to pass additional bridges that they defined in external packages.

So to recap, MathOptInterface suffers package latency because:

1. there are a large number of types and functions...
2. and these are split between multiple modules, including external packages...
3. and there are type-instabilities like those in the bridging layer.

8.3 Resolutions

There are no magic solutions to reduce latency. [Issue #1313](#) tracks progress on reducing latency in MathOptInterface.

A useful script is the following (replace GLPK as needed):

```
using MathOptInterface, GLPK
const MOI = MathOptInterface

function example_diet(optimizer, bridge)
    category_data = [
        1800.0 2200.0;
        91.0    Inf;
        0.0    65.0;
        0.0    1779.0
    ]
    cost = [2.49, 2.89, 1.50, 1.89, 2.09, 1.99, 2.49, 0.89, 1.59]
    food_data = [
        410 24 26 730;
        420 32 10 1190;
        560 20 32 1800;
        380 4 19 270;
        320 12 10 930;
        320 15 12 820;
        320 31 12 1230;
        100 8 2.5 125;
        330 8 10 180
    ]
    bridge_model = if bridge
        MOI.instantiate(optimizer; with_bridge_type=Float64)
    else
        MOI.instantiate(optimizer)
    end
    model = MOI.Utilities.CachingOptimizer(
        MOI.Utilities.UniversalFallback(MOI.Utilities.Model{Float64}()),
        MOI.Utilities.AUTOMATIC,
    )
    MOI.Utilities.reset_optimizer(model, bridge_model)
    MOI.set(model, MOI.Silent(), true)
    nutrition = MOI.add_variables(model, size(category_data, 1))
    for (i, v) in enumerate(nutrition)
        MOI.add_constraint(model, v, MOI.GreaterThan(category_data[i, 1]))
        MOI.add_constraint(model, v, MOI.LessThan(category_data[i, 2]))
    end
    buy = MOI.add_variables(model, size(food_data, 1))
    MOI.add_constraint.(model, buy, MOI.GreaterThan(0.0))
    MOI.set(model, MOI.ObjectiveSense(), MOI.MIN_SENSE)
    f = MOI.ScalarAffineFunction(MOI.ScalarAffineTerm.(cost, buy), 0.0)
    MOI.set(model, MOI.ObjectiveFunction{typeof(f)}(), f)
```

```

    for (j, n) in enumerate(nutrition)
        f = MOI.ScalarAffineFunction(
            MOI.ScalarAffineTerm.(food_data[:, j], buy),
            0.0,
        )
        push!(f.terms, MOI.ScalarAffineTerm(-1.0, n))
        MOI.add_constraint(model, f, MOI.EqualTo(0.0))
    end
    MOI.optimize!(model)
    term_status = MOI.get(model, MOI.TerminationStatus())
    @assert term_status == MOI.OPTIMAL
    MOI.add_constraint(
        model,
        MOI.ScalarAffineFunction(
            MOI.ScalarAffineTerm.(1.0, [buy[end-1], buy[end]]),
            0.0,
        ),
        MOI.LessThan(6.0),
    )
    MOI.optimize!(model)
    @assert MOI.get(model, MOI.TerminationStatus()) == MOI.INFEASIBLE
    return
end

if length(ARGS) > 0
    bridge = get(ARGS, 2, "") != "--no-bridge"
    println("Running: $(ARGS[1]) $(get(ARGS, 2, ""))")
    @time example_diet(GLPK.Optimizer, bridge)
    @time example_diet(GLPK.Optimizer, bridge)
    exit(0)
end

```

You can create a flame-graph via

```

using SnoopComile
tinf = @snoopi_deep example_diet(GLPK.Optimizer, true)
using ProfileView
ProfileView.view(flamegraph(tinf))

```

Here's how things looked in mid-August 2021:

There are a few opportunities for improvement (non-red flames, particularly on the right). But the main problem is a large red (non-precompilable due to method ownership) flame.




Figure 8.1: flamegraph

Part III

Manual

Chapter 9

Standard form problem

MathOptInterface represents optimization problems in the standard form:

$$\min_{x \in \mathbb{R}^n} f_0(x) \tag{9.1}$$

$$\text{s.t.} \quad f_i(x) \in \mathcal{S}_i \quad i = 1 \dots m \tag{9.2}$$

where:

- the functions f_0, f_1, \dots, f_m are specified by [AbstractFunction](#) objects
- the sets $\mathcal{S}_1, \dots, \mathcal{S}_m$ are specified by [AbstractSet](#) objects

Tip

For more information on this standard form, read [our paper](#).

MOI defines some commonly used functions and sets, but the interface is extensible to other sets recognized by the solver.

9.1 Functions

The function types implemented in MathOptInterface.jl are:

- [VariableIndex](#): x_j , i.e., projection onto a single coordinate defined by a variable index j .
- [VectorOfVariables](#): projection onto multiple coordinates (i.e., extracting a subvector).
- [ScalarAffineFunction](#): $a^T x + b$, where a is a vector and b scalar.
- [VectorAffineFunction](#): $Ax + b$, where A is a matrix and b is a vector.
- [ScalarQuadraticFunction](#): $\frac{1}{2}x^T Q x + a^T x + b$, where Q is a symmetric matrix, a is a vector, and b is a constant.
- [VectorQuadraticFunction](#): a vector of scalar-valued quadratic functions.

Extensions for nonlinear programming are present but not yet well documented.

9.2 One-dimensional sets

The one-dimensional set types implemented in MathOptInterface.jl are:

- `LessThan(upper)`: $\{x \in \mathbb{R} : x \leq \text{upper}\}$
- `GreaterThan(lower)`: $\{x \in \mathbb{R} : x \geq \text{lower}\}$
- `EqualTo(value)`: $\{x \in \mathbb{R} : x = \text{value}\}$
- `Interval(lower, upper)`: $\{x \in \mathbb{R} : x \in [\text{lower}, \text{upper}]\}$
- `Integer()`: \mathbb{Z}
- `ZeroOne()`: $\{0, 1\}$
- `Semicontinuous(lower, upper)`: $\{0\} \cup [\text{lower}, \text{upper}]$
- `Semiinteger(lower, upper)`: $\{0\} \cup \{\text{lower}, \text{lower} + 1, \dots, \text{upper} - 1, \text{upper}\}$

9.3 Vector cones

The vector-valued set types implemented in MathOptInterface.jl are:

- `Reals(dimension)`: $\mathbb{R}^{\text{dimension}}$
- `Zeros(dimension)`: $\{0\}^{\text{dimension}}$
- `Nonnegatives(dimension)`: $\{x \in \mathbb{R}^{\text{dimension}} : x \geq 0\}$
- `Nonpositives(dimension)`: $\{x \in \mathbb{R}^{\text{dimension}} : x \leq 0\}$
- `SecondOrderCone(dimension)`: $\{(t, x) \in \mathbb{R}^{\text{dimension}} : t \geq \|x\|_2\}$
- `RotatedSecondOrderCone(dimension)`: $\{(t, u, x) \in \mathbb{R}^{\text{dimension}} : 2tu \geq \|x\|_2^2, t, u \geq 0\}$
- `ExponentialCone()`: $\{(x, y, z) \in \mathbb{R}^3 : y \exp(x/y) \leq z, y > 0\}$
- `DualExponentialCone()`: $\{(u, v, w) \in \mathbb{R}^3 : -u \exp(v/u) \leq \exp(1)w, u < 0\}$
- `GeometricMeanCone(dimension)`: $\{(t, x) \in \mathbb{R}^{n+1} : x \geq 0, t \leq \sqrt[n]{x_1 x_2 \cdots x_n}\}$ where n is `dimension - 1`
- `PowerCone(exponent)`: $\{(x, y, z) \in \mathbb{R}^3 : x^{\text{exponent}} y^{1-\text{exponent}} \geq |z|, x, y \geq 0\}$
- `DualPowerCone(exponent)`: $\{(u, v, w) \in \mathbb{R}^3 : \frac{u}{\text{exponent}}^{\text{exponent}} \frac{v}{1-\text{exponent}}^{1-\text{exponent}} \geq |w|, u, v \geq 0\}$
- `NormOneCone(dimension)`: $\{(t, x) \in \mathbb{R}^{\text{dimension}} : t \geq \|x\|_1\}$ where $\|x\|_1 = \sum_i |x_i|$
- `NormInfinityCone(dimension)`: $\{(t, x) \in \mathbb{R}^{\text{dimension}} : t \geq \|x\|_\infty\}$ where $\|x\|_\infty = \max_i |x_i|$.
- `RelativeEntropyCone(dimension)`: $\{(u, v, w) \in \mathbb{R}^{\text{dimension}} : u \geq \sum_i w_i \log(\frac{w_i}{v_i}), v_i \geq 0, w_i \geq 0\}$

9.4 Matrix cones

The matrix-valued set types implemented in `MathOptInterface.jl` are:

- `RootDetConeTriangle(dimension)`: $\{(t, X) \in \mathbb{R}^{1+\text{dimension}(\text{dimension}+1)/2} : t \leq \det(X)^{1/\text{dimension}}, X \text{ is the upper triangular part of a PSD matrix}\}$
- `RootDetConeSquare(dimension)`: $\{(t, X) \in \mathbb{R}^{1+\text{dimension}^2} : t \leq \det(X)^{1/\text{dimension}}, X \text{ is a PSD matrix}\}$
- `PositiveSemidefiniteConeTriangle(dimension)`: $\{X \in \mathbb{R}^{\text{dimension}(\text{dimension}+1)/2} : X \text{ is the upper triangular part of a PSD matrix}\}$
- `PositiveSemidefiniteConeSquare(dimension)`: $\{X \in \mathbb{R}^{\text{dimension}^2} : X \text{ is a PSD matrix}\}$
- `LogDetConeTriangle(dimension)`: $\{(t, u, X) \in \mathbb{R}^{2+\text{dimension}(\text{dimension}+1)/2} : t \leq u \log(\det(X/u)), X \text{ is the upper triangular part of a PSD matrix}, u > 0\}$
- `LogDetConeSquare(dimension)`: $\{(t, u, X) \in \mathbb{R}^{2+\text{dimension}^2} : t \leq u \log(\det(X/u)), X \text{ is a PSD matrix}, u > 0\}$
- `NormSpectralCone(row_dim, column_dim)`: $\{(t, X) \in \mathbb{R}^{1+\text{row_dim} \times \text{column_dim}} : t \geq \sigma_1(X), X \text{ is a matrix with row_dim rows and column_dim columns}\}$
- `NormNuclearCone(row_dim, column_dim)`: $\{(t, X) \in \mathbb{R}^{1+\text{row_dim} \times \text{column_dim}} : t \geq \sum_i \sigma_i(X), X \text{ is a matrix with row_dim rows and column_dim columns}\}$

Some of these cones can take two forms: `XXXConeTriangle` and `XXXConeSquare`.

In `XXXConeTriangle` sets, the matrix is assumed to be symmetric, and the elements are provided by a vector, in which the entries of the upper-right triangular part of the matrix are given column by column (or equivalently, the entries of the lower-left triangular part are given row by row).

In `XXXConeSquare` sets, the entries of the matrix are given column by column (or equivalently, row by row), and the matrix is constrained to be symmetric. As an example, given a 2-by-2 matrix of variables X and a one-dimensional variable t , we can specify a root-det constraint as $[t, X_{11}, X_{12}, X_{22}] \in \text{RootDetConeTriangle}$ or $[t, X_{11}, X_{12}, X_{21}, X_{22}] \in \text{RootDetConeSquare}$.

We provide both forms to enable flexibility for solvers who may natively support one or the other. Transformations between `XXXConeTriangle` and `XXXConeSquare` are handled by bridges, which removes the chance of conversion mistakes by users or solver developers.

9.5 Multi-dimensional sets with combinatorial structure

- `SOS1(weights)`: A special ordered set of Type I.
- `SOS2(weights)`: A special ordered set of Type II.
- `Indicator(set)`: A set to specify indicator constraints.
- `Complements(dimension)`: A set for mixed complementarity constraints.

Chapter 10

Models

The most significant part of MOI is the definition of the **model API** that is used to specify an instance of an optimization problem (e.g., by adding variables and constraints). Objects that implement the model API must inherit from the [ModelLike](#) abstract type.

Notably missing from the model API is the method to solve an optimization problem. [ModelLike](#) objects may store an instance (e.g., in memory or backed by a file format) without being linked to a particular solver. In addition to the model API, MOI defines [AbstractOptimizer](#) and provides methods to solve the model and interact with solutions. See the [Solutions](#) section for more details.

Info

Throughout the rest of the manual, `model` is used as a generic [ModelLike](#), and `optimizer` is used as a generic [AbstractOptimizer](#).

Tip

MOI does not export functions, but for brevity we often omit qualifying names with the MOI module. Best practice is to have

```
using MathOptInterface
const MOI = MathOptInterface
```

and prefix all MOI methods with `MOI.` in user code. If a name is also available in base Julia, we always explicitly use the module prefix, for example, with `MOI.get`.

10.1 Attributes

Attributes are properties of the model that can be queried and modified. These include constants such as the number of variables in a model ([NumberOfVariables](#)), and properties of variables and constraints such as the name of a variable ([VariableName](#)).

There are four types of attributes:

- Model attributes (subtypes of [AbstractModelAttribute](#)) refer to properties of a model.
- Optimizer attributes (subtypes of [AbstractOptimizerAttribute](#)) refer to properties of an optimizer.
- Constraint attributes (subtypes of [AbstractConstraintAttribute](#)) refer to properties of an individual constraint.

- Variable attributes (subtypes of `AbstractVariableAttribute`) refer to properties of an individual variable.

Some attributes are values that can be queried by the user but not modified, while other attributes can be modified by the user.

All interactions with attributes occur through the `get` and `set` functions.

Consult the docstrings of each attribute for information on what it represents.

10.2 ModelLike API

The following attributes are available:

- `ListOfConstraintAttributesSet`
- `ListOfConstraintIndices`
- `ListOfConstraintTypesPresent`
- `ListOfModelAttributesSet`
- `ListOfVariableAttributesSet`
- `ListOfVariableIndices`
- `NumberOfConstraints`
- `NumberOfVariables`
- `Name`
- `ObjectiveFunction`
- `ObjectiveFunctionType`
- `ObjectiveSense`

10.3 AbstractOptimizer API

The following attributes are available:

- `DualStatus`
- `PrimalStatus`
- `RawStatusString`
- `ResultCount`
- `TerminationStatus`
- `BarrierIterations`
- `DualObjectiveValue`
- `NodeCount`

- `NumberOfThreads`
- `ObjectiveBound`
- `ObjectiveValue`
- `RelativeGap`
- `RawOptimizerAttribute`
- `RawSolver`
- `Silent`
- `SimplexIterations`
- `SolverName`
- `SolverVersion`
- `SolveTimeSec`
- `TimeLimitSec`

Chapter 11

Variables

11.1 Add a variable

Use `add_variable` to add a single variable.

```
julia> x = MOI.add_variable(model)
MathOptInterface.VariableIndex(1)
```

`add_variable` returns a `VariableIndex` type, which is used to refer to the added variable in other calls.

Check if a `VariableIndex` is valid using `is_valid`.

```
julia> MOI.is_valid(model, x)
true
```

Use `add_variables` to add a number of variables.

```
julia> y = MOI.add_variables(model, 2)
2-element Vector{MathOptInterface.VariableIndex}:
 MathOptInterface.VariableIndex(2)
 MathOptInterface.VariableIndex(3)
```

Warning

The integer does not necessarily correspond to the column inside an optimizer!

11.2 Delete a variable

Delete a variable using `delete`.

```
julia> MOI.delete(model, x)

julia> MOI.is_valid(model, x)
false
```

Warning

Not all `ModelLike` models support deleting variables. A `DeleteNotAllowed` error is thrown if this is not supported.

11.3 Variable attributes

The following attributes are available for variables:

- `VariableName`
- `VariablePrimalStart`
- `VariablePrimal`

Get and set these attributes using `get` and `set`.

```
julia> MOI.set(model, MOI.VariableName(), x, "var_x")  
  
julia> MOI.get(model, MOI.VariableName(), x)  
"var_x"
```

Chapter 12

Constraints

12.1 Add a constraint

Use `add_constraint` to add a single constraint.

```
julia> c = MOI.add_constraint(model, MOI.VectorOfVariables(x), MOI.Nonnegatives(2))
MathOptInterface.ConstraintIndex{MathOptInterface.VectorOfVariables,
↳ MathOptInterface.Nonnegatives}(1)
```

`add_constraint` returns a `ConstraintIndex` type, which is used to refer to the added constraint in other calls.

Check if a `ConstraintIndex` is valid using `is_valid`.

```
julia> MOI.is_valid(model, c)
true
```

Use `add_constraints` to add a number of constraints of the same type.

```
julia> c = MOI.add_constraints(
    model,
    [x[1], x[2]],
    [MOI.GreaterThan(0.0), MOI.GreaterThan(1.0)]
)
2-element Vector{MathOptInterface.ConstraintIndex{MathOptInterface.VariableIndex,
↳ MathOptInterface.GreaterThan{Float64}}}:
 MathOptInterface.ConstraintIndex{MathOptInterface.VariableIndex,
↳ MathOptInterface.GreaterThan{Float64}}(1)
 MathOptInterface.ConstraintIndex{MathOptInterface.VariableIndex,
↳ MathOptInterface.GreaterThan{Float64}}(2)
```

This time, a vector of `ConstraintIndex` are returned.

Use `supports_constraint` to check if the model supports adding a constraint type.

```
julia> MOI.supports_constraint(
    model,
    MOI.VariableIndex,
    MOI.GreaterThan{Float64},
)
true
```

12.2 Delete a constraint

Use `delete` to delete a constraint.

```
julia> MOI.delete(model, c)

julia> MOI.is_valid(model, c)
false
```

12.3 Constraint attributes

The following attributes are available for constraints:

- `ConstraintName`
- `ConstraintPrimalStart`
- `ConstraintDualStart`
- `ConstraintPrimal`
- `ConstraintDual`
- `ConstraintBasisStatus`
- `ConstraintFunction`
- `CanonicalConstraintFunction`
- `ConstraintSet`

Get and set these attributes using `get` and `set`.

```
julia> MOI.set(model, MOI.ConstraintName(), c, "con_c")

julia> MOI.get(model, MOI.ConstraintName(), c)
"con_c"
```

12.4 Constraints by function-set pairs

Below is a list of common constraint types and how they are represented as function-set pairs in MOI. In the notation below, x is a vector of decision variables, x_i is a scalar decision variable, α, β are scalar constants, a, b are constant vectors, A is a constant matrix and \mathbb{R}_+ (resp. \mathbb{R}_-) is the set of nonnegative (resp. nonpositive) real numbers.

Linear constraints

By convention, solvers are not expected to support nonzero constant terms in the `ScalarAffineFunctions` the first four rows above, because they are redundant with the parameters of the sets. For example, encode $2x + 1 \leq 2$ as $2x \leq 1$.

Constraints with `VariableIndex` in `LessThan`, `GreaterThan`, `EqualTo`, or `Interval` sets have a natural interpretation as variable bounds. As such, it is typically not natural to impose multiple lower- or upper-bounds on the same variable, and the solver interfaces will throw respectively `LowerBoundAlreadySet` or `UpperBoundAlreadySet`.

Mathematical Constraint	MOI Function	MOI Set
$a^T x \leq \beta$	ScalarAffineFunction	LessThan
$a^T x \geq \alpha$	ScalarAffineFunction	GreaterThan
$a^T x = \beta$	ScalarAffineFunction	EqualTo
$\alpha \leq a^T x \leq \beta$	ScalarAffineFunction	Interval
$x_i \leq \beta$	VariableIndex	LessThan
$x_i \geq \alpha$	VariableIndex	GreaterThan
$x_i = \beta$	VariableIndex	EqualTo
$\alpha \leq x_i \leq \beta$	VariableIndex	Interval
$Ax + b \in \mathbb{R}_+^n$	VectorAffineFunction	Nonnegatives
$Ax + b \in \mathbb{R}_-^n$	VectorAffineFunction	Nonpositives
$Ax + b = 0$	VectorAffineFunction	Zeros

Moreover, adding two [VariableIndex](#) constraints on the same variable with the same set is impossible because they share the same index as it is the index of the variable, see [ConstraintIndex](#).

It is natural, however, to impose upper- and lower-bounds separately as two different constraints on a single variable. The difference between imposing bounds by using a single [Interval](#) constraint and by using separate [LessThan](#) and [GreaterThan](#) constraints is that the latter will allow the solver to return separate dual multipliers for the two bounds, while the former will allow the solver to return only a single dual for the interval constraint.

Conic constraints

Mathematical Constraint	MOI Function	MOI Set
$\ Ax + b\ _2 \leq c^T x + d$	VectorAffineFunction	SecondOrderCone
$y \geq \ x\ _2$	VectorOfVariables	SecondOrderCone
$2yz \geq \ x\ _2^2, y, z \geq 0$	VectorOfVariables	RotatedSecondOrderCone
$(a_1^T x + b_1, a_2^T x + b_2, a_3^T x + b_3) \in \mathcal{E}$	VectorAffineFunction	ExponentialCone
$A(x) \in \mathcal{S}_+$	VectorAffineFunction	PositiveSemidefiniteConeTriangle
$B(x) \in \mathcal{S}_+$	VectorAffineFunction	PositiveSemidefiniteConeSquare
$x \in \mathcal{S}_+$	VectorOfVariables	PositiveSemidefiniteConeTriangle
$x \in \mathcal{S}_+$	VectorOfVariables	PositiveSemidefiniteConeSquare

where \mathcal{E} is the exponential cone (see [ExponentialCone](#)), \mathcal{S}_+ is the set of positive semidefinite symmetric matrices, A is an affine map that outputs symmetric matrices and B is an affine map that outputs square matrices.

Quadratic constraints

Mathematical Constraint	MOI Function	MOI Set
$\frac{1}{2}x^T Qx + a^T x + b \geq 0$	ScalarQuadraticFunction	GreaterThan
$\frac{1}{2}x^T Qx + a^T x + b \leq 0$	ScalarQuadraticFunction	LessThan
$\frac{1}{2}x^T Qx + a^T x + b = 0$	ScalarQuadraticFunction	EqualTo
Bilinear matrix inequality	VectorQuadraticFunction	PositiveSemidefiniteCone...

Discrete and logical constraints

12.5 JuMP mapping

The following bullet points show examples of how JuMP constraints are translated into MOI function-set pairs:

Mathematical Constraint	MOI Function	MOI Set
$x_i \in \mathbb{Z}$	VariableIndex	Integer
$x_i \in \{0, 1\}$	VariableIndex	ZeroOne
$x_i \in \{0\} \cup [l, u]$	VariableIndex	Semicontinuous
$x_i \in \{0\} \cup \{l, l+1, \dots, u-1, u\}$	VariableIndex	Semiinteger
At most one component of x can be nonzero	VectorOfVariables	SOS1
At most two components of x can be nonzero, and if so they must be adjacent components	VectorOfVariables	SOS2
$y = 1 \implies a^T x \in S$	VectorAffineFunctionIndicator	

- `@constraint(m, 2x + y <= 10)` becomes `ScalarAffineFunction-in-LessThan`
- `@constraint(m, 2x + y >= 10)` becomes `ScalarAffineFunction-in-GreaterThan`
- `@constraint(m, 2x + y == 10)` becomes `ScalarAffineFunction-in-EqualTo`
- `@constraint(m, 0 <= 2x + y <= 10)` becomes `ScalarAffineFunction-in-Interval`
- `@constraint(m, 2x + y in ArbitrarySet())` becomes `ScalarAffineFunction-in-ArbitrarySet`.

Variable bounds are handled in a similar fashion:

- `@variable(m, x <= 1)` becomes `VariableIndex-in-LessThan`
- `@variable(m, x >= 1)` becomes `VariableIndex-in-GreaterThan`

One notable difference is that a variable with an upper and lower bound is translated into two constraints, rather than an interval. i.e.:

- `@variable(m, 0 <= x <= 1)` becomes `VariableIndex-in-LessThan` and `VariableIndex-in-GreaterThan`.

Chapter 13

Solutions

13.1 Solving and retrieving the results

Once an optimizer is loaded with the objective function and all of the constraints, we can ask the solver to solve the model by calling `optimize!`.

```
| MOI.optimize!(optimizer)
```

13.2 Why did the solver stop?

The optimization procedure may terminate for a number of reasons. The `TerminationStatus` attribute of the optimizer returns a `TerminationStatusCode` object which explains why the solver stopped.

The termination statuses distinguish between proofs of optimality, infeasibility, local convergence, limits, and termination because of something unexpected like invalid problem data or failure to converge.

A typical usage of the `TerminationStatus` attribute is as follows:

```
| status = MOI.get(optimizer, TerminationStatus())  
| if status == MOI.OPTIMAL  
|     # Ok, we solved the problem!  
| else  
|     # Handle other cases.  
| end
```

After checking the `TerminationStatus`, check `ResultCount`. This attribute returns the number of results that the solver has available to return. A result is defined as a primal-dual pair, but either the primal or the dual may be missing from the result. While the `OPTIMAL` termination status normally implies that at least one result is available, other statuses do not. For example, in the case of infeasibility, a solver may return no result or a proof of infeasibility. The `ResultCount` attribute distinguishes between these two cases.

13.3 Primal solutions

Use the `PrimalStatus` optimizer attribute to return a `ResultStatusCode` describing the status of the primal solution.

Common returns are described below in the `Common status situations` section.

Query the primal solution using the `VariablePrimal` and `ConstraintPrimal` attributes.

Query the objective function value using the `ObjectiveValue` attribute.

13.4 Dual solutions

Warning

See [Duality](#) for a discussion of the MOI conventions for primal-dual pairs and certificates.

Use the [DualStatus](#) optimizer attribute to return a [ResultStatusCode](#) describing the status of the dual solution.

Query the dual solution using the [ConstraintDual](#) attribute.

Query the dual objective function value using the [DualObjectiveValue](#) attribute.

13.5 Common status situations

The sections below describe how to interpret typical or interesting status cases for three common classes of solvers. The example cases are illustrative, not comprehensive. Solver wrappers may provide additional information on how the solver's statuses map to MOI statuses.

Info

* in the tables indicate that multiple different values are possible.

Primal-dual convex solver

Linear programming and conic optimization solvers fall into this category.

What happened?	TerminationStatus	ResultCount	PrimalStatus	DualStatus
Proved optimality	OPTIMAL	1	FEASIBLE_POINT	FEASIBLE_POINT
Proved infeasible	INFEASIBLE	1	NO_SOLUTION	INFEASIBILITY_CERTIFICATE
Optimal within relaxed tolerances	ALMOST_OPTIMAL	1	FEASIBLE_POINT	FEASIBLE_POINT
Optimal within relaxed tolerances	ALMOST_OPTIMAL	1	ALMOST_FEASIBLE_POINT	ALMOST_FEASIBLE_POINT
Detected an unbounded ray of the primal	DUAL_INFEASIBLE	1	INFEASIBILITY_CERTIFICATE	NO_SOLUTION
Stall	SLOW_PROGRESS	1	*	*

Global branch-and-bound solvers

Mixed-integer programming solvers fall into this category.

What happened?	TerminationStatus	ResultCount	PrimalStatus	DualStatus
Proved optimality	OPTIMAL	1	FEASIBLE_POINT	NO_SOLUTION
Presolve detected infeasibility or unboundedness	INFEASIBLE_OR_UNBOUNDED	0	NO_SOLUTION	NO_SOLUTION
Proved infeasibility	INFEASIBLE	0	NO_SOLUTION	NO_SOLUTION
Timed out (no solution)	TIME_LIMIT	0	NO_SOLUTION	NO_SOLUTION
Timed out (with a solution)	TIME_LIMIT	1	FEASIBLE_POINT	NO_SOLUTION
CPXMIP_OPTIMAL_INFEAS	ALMOST_OPTIMAL	1	INFEASIBLE_POINT	NO_SOLUTION

Info

`CPXMIP_OPTIMAL_INFEAS` is a CPLEX status that indicates that a preprocessed problem was solved to optimality, but the solver was unable to recover a feasible solution to the original problem. Handling this status was one of the motivating drivers behind the design of MOI.

Local search solvers

Nonlinear programming solvers fall into this category. It also includes non-global tree search solvers like `Juniper`.

What happened?	TerminationStatus	ResultCount	PrimalStatus	DualStatus
Converged to a stationary point	LOCALLY_SOLVED	1	FEASIBLE_POINT	FEASIBLE_POINT
Completed a non-global tree search (with a solution)	LOCALLY_SOLVED	1	FEASIBLE_POINT	FEASIBLE_POINT
Converged to an infeasible point	LOCALLY_INFEASIBLE	1	INFEASIBLE_POINT	*
Completed a non-global tree search (no solution found)	LOCALLY_INFEASIBLE	0	NO_SOLUTION	NO_SOLUTION
Iteration limit	ITERATION_LIMIT	1	*	*
Diverging iterates	NORM_LIMIT or OBJECTIVE_LIMIT	1	*	*

13.6 Querying solution attributes

Some solvers will not implement every solution attribute. Therefore, a call like `MOI.get(model, MOI.SolveTimeSec())` may throw an `UnsupportedAttribute` error.

If you need to write code that is agnostic to the solver (for example, you are writing a library that an end-user passes their choice of solver to), you can work-around this problem using a try-catch:

```
function get_solve_time(model)
    try
        return MOI.get(model, MOI.SolveTimeSec())
    catch err
        if err isa MOI.UnsupportedAttribute
            return NaN # Solver doesn't support. Return a placeholder value.
        end
        rethrow(err) # Something else went wrong. Rethrow the error
    end
end
```

If, after careful profiling, you find that the try-catch is taking a significant portion of your runtime, you can improve performance by caching the result of the try-catch:

```
mutable struct CachedSolveTime{M}
    model::M
    supports_solve_time::Bool
    CachedSolveTime(model::M) where {M} = new(model, true)
end

function get_solve_time(model::CachedSolveTime)
    if !model.supports_solve_time
        return NaN
    end
end
```

```
    try
      return MOI.get(model, MOI.SolveTimeSec())
    catch err
      if err isa MOI.UnsupportedAttribute
        model.supports_solve_time = false
        return NaN
      end
      rethrow(err) # Something else went wrong. Rethrow the error
    end
  end
end
```

Chapter 14

Problem modification

In addition to adding and deleting constraints and variables, MathOptInterface supports modifying, in-place, coefficients in the constraints and the objective function of a model.

These modifications can be grouped into two categories:

- modifications which replace the set of function of a constraint with a new set or function
- modifications which change, in-place, a component of a function

Warning

Solve ModelLike objects do not support problem modification.

14.1 Modify the set of a constraint

Use `set` and `ConstraintSet` to modify the set of a constraint by replacing it with a new instance of the same type.

```
julia> c = MOI.add_constraint(
    model,
    MOI.ScalarAffineFunction([MOI.ScalarAffineTerm(1.0, x)], 0.0),
    MOI.EqualTo(1.0),
)
MathOptInterface.ConstraintIndex{MathOptInterface.ScalarAffineFunction{Float64},
↪ MathOptInterface.EqualTo{Float64}}{1}

julia> MOI.set(model, MOI.ConstraintSet(), c, MOI.EqualTo(2.0));

julia> MOI.get(model, MOI.ConstraintSet(), c) == MOI.EqualTo(2.0)
true
```

However, the following will fail as the new set is of a different type to the original set:

```
julia> MOI.set(model, MOI.ConstraintSet(), c, MOI.GreaterThan(2.0))
ERROR: [...]
```

Special cases: set transforms

If our constraint is an affine inequality, then this corresponds to modifying the right-hand side of a constraint in linear programming.

In some special cases, solvers may support efficiently changing the set of a constraint (for example, from [LessThan](#) to [GreaterThan](#)). For these cases, `MathOptInterface` provides the `transform` method.

The `transform` function returns a new constraint index, and the old constraint index (i.e., `c`) is no longer valid.

```
julia> c = MOI.add_constraint(
    model,
    MOI.ScalarAffineFunction([MOI.ScalarAffineTerm(1.0, x)], 0.0),
    MOI.LessThan(1.0),
)
MathOptInterface.ConstraintIndex{MathOptInterface.ScalarAffineFunction{Float64},
↪ MathOptInterface.LessThan{Float64}}(1)

julia> new_c = MOI.transform(model, c, MOI.GreaterThan(2.0));

julia> MOI.is_valid(model, c)
false

julia> MOI.is_valid(model, new_c)
true
```

Note

`transform` cannot be called with a set of the same type. Use `set` instead.

14.2 Modify the function of a constraint

Use `set` and `ConstraintFunction` to modify the function of a constraint by replacing it with a new instance of the same type.

```
julia> c = MOI.add_constraint(
    model,
    MOI.ScalarAffineFunction([MOI.ScalarAffineTerm(1.0, x)], 0.0),
    MOI.EqualTo(1.0),
)
MathOptInterface.ConstraintIndex{MathOptInterface.ScalarAffineFunction{Float64},
↪ MathOptInterface.EqualTo{Float64}}(1)

julia> new_f = MOI.ScalarAffineFunction([MOI.ScalarAffineTerm(2.0, x)], 1.0);

julia> MOI.set(model, MOI.ConstraintFunction(), c, new_f);

julia> MOI.get(model, MOI.ConstraintFunction(), c) ≈ new_f
true
```

However, the following will fail as the new function is of a different type to the original function:

```
julia> MOI.set(model, MOI.ConstraintFunction(), c, x)
ERROR: [...]
```

14.3 Modify constant term in a scalar function

Use `modify` and `ScalarConstantChange` to modify the constant term in a `ScalarAffineFunction` or `ScalarQuadraticFunction`.

```
julia> c = MOI.add_constraint(
    model,
    MOI.ScalarAffineFunction([MOI.ScalarAffineTerm(1.0, x)], 0.0),
    MOI.EqualTo(1.0),
)
MathOptInterface.ConstraintIndex{MathOptInterface.ScalarAffineFunction{Float64},
↳ MathOptInterface.EqualTo{Float64}}(1)

julia> MOI.modify(model, c, MOI.ScalarConstantChange(1.0));

julia> new_f = MOI.ScalarAffineFunction([MOI.ScalarAffineTerm(1.0, x)], 1.0);

julia> MOI.get(model, MOI.ConstraintFunction(), c) ≈ new_f
true
```

Tip

`ScalarConstantChange` can also be used to modify the objective function by passing an instance of `ObjectiveFunction` instead of the constraint index `c` as we saw above.

```
julia> MOI.set(
    model,
    MOI.ObjectiveFunction{MOI.ScalarAffineFunction{Float64}}(),
    new_f,
);

julia> MOI.modify(
    model,
    MOI.ObjectiveFunction{MOI.ScalarAffineFunction{Float64}}(),
    MOI.ScalarConstantChange(-1.0)
);

julia> MOI.get(
    model,
    MOI.ObjectiveFunction{MOI.ScalarAffineFunction{Float64}}(),
) ≈ MOI.ScalarAffineFunction([MOI.ScalarAffineTerm(1.0, x)], -1.0)
true
```

14.4 Modify constant terms in a vector function

Use `modify` and `VectorConstantChange` to modify the constant vector in a `VectorAffineFunction` or `VectorQuadraticFunction`.

```
julia> c = MOI.add_constraint(
    model,
    MOI.VectorAffineFunction([
        MOI.VectorAffineTerm(1, MOI.ScalarAffineTerm(1.0, x)),
        MOI.VectorAffineTerm(2, MOI.ScalarAffineTerm(2.0, x)),
    ],
    [0.0, 0.0],
),
```

```

        MOI.Nonnegatives(2),
    )
    MathOptInterface.ConstraintIndex{MathOptInterface.VectorAffineFunction{Float64}},
    ↪ MathOptInterface.Nonnegatives}(1)

julia> MOI.modify(model, c, MOI.VectorConstantChange([3.0, 4.0]));

julia> new_f = MOI.VectorAffineFunction(
    [
        MOI.VectorAffineTerm(1, MOI.ScalarAffineTerm(1.0, x)),
        MOI.VectorAffineTerm(2, MOI.ScalarAffineTerm(2.0, x)),
    ],
    [3.0, 4.0],
);

julia> MOI.get(model, MOI.ConstraintFunction(), c) ≈ new_f
true

```

14.5 Modify affine coefficients in a scalar function

Use `modify` and `ScalarCoefficientChange` to modify the affine coefficient of a `ScalarAffineFunction` or `ScalarQuadraticFunction`.

```

julia> c = MOI.add_constraint(
    model,
    MOI.ScalarAffineFunction([MOI.ScalarAffineTerm(1.0, x)], 0.0),
    MOI.EqualTo(1.0),
)
MathOptInterface.ConstraintIndex{MathOptInterface.ScalarAffineFunction{Float64}},
↪ MathOptInterface.EqualTo{Float64}}(1)

julia> MOI.modify(model, c, MOI.ScalarCoefficientChange(x, 2.0));

julia> new_f = MOI.ScalarAffineFunction([MOI.ScalarAffineTerm(2.0, x)], 0.0);

julia> MOI.get(model, MOI.ConstraintFunction(), c) ≈ new_f
true

```

Tip

`ScalarCoefficientChange` can also be used to modify the objective function by passing an instance of `ObjectiveFunction` instead of the constraint index `c` as we saw above.

14.6 Modify affine coefficients in a vector function

Use `modify` and `MultirowChange` to modify a vector of affine coefficients in a `VectorAffineFunction` or a `VectorQuadraticFunction`.

```

julia> c = MOI.add_constraint(
    model,
    MOI.VectorAffineFunction([
        MOI.VectorAffineTerm(1, MOI.ScalarAffineTerm(1.0, x)),
        MOI.VectorAffineTerm(2, MOI.ScalarAffineTerm(2.0, x)),
    ],
    ],
)

```



```

        [0.0, 0.0],
    ),
    MOI.Nonnegatives(2),
)
MathOptInterface.ConstraintIndex{MathOptInterface.VectorAffineFunction{Float64},
↪ MathOptInterface.Nonnegatives}(1)

julia> MOI.modify(model, c, MOI.MultirowChange(x, [(1, 3.0), (2, 4.0)]));

julia> new_f = MOI.VectorAffineFunction(
    [
        MOI.VectorAffineTerm(1, MOI.ScalarAffineTerm(3.0, x)),
        MOI.VectorAffineTerm(2, MOI.ScalarAffineTerm(4.0, x)),
    ],
    [0.0, 0.0],
);

julia> MOI.get(model, MOI.ConstraintFunction(), c) ≈ new_f
true

```

Part IV

Background

Chapter 15

Duality

Conic duality is the starting point for MOI's duality conventions. When all functions are affine (or coordinate projections), and all constraint sets are closed convex cones, the model may be called a conic optimization problem.

For a minimization problem in geometric conic form, the primal is:

$$\min_{x \in \mathbb{R}^n} \quad a_0^T x + b_0 \quad (15.1)$$

$$\text{s.t.} \quad A_i x + b_i \in \mathcal{C}_i \quad i = 1 \dots m \quad (15.2)$$

and the dual is a maximization problem in standard conic form:

$$\max_{y_1, \dots, y_m} \quad - \sum_{i=1}^m b_i^T y_i + b_0 \quad (15.3)$$

$$\text{s.t.} \quad a_0 - \sum_{i=1}^m A_i^T y_i = 0 \quad (15.4)$$

$$y_i \in \mathcal{C}_i^* \quad i = 1 \dots m \quad (15.5)$$

where each \mathcal{C}_i is a closed convex cone and \mathcal{C}_i^* is its dual cone.

For a maximization problem in geometric conic form, the primal is:

$$\max_{x \in \mathbb{R}^n} \quad a_0^T x + b_0 \quad (15.6)$$

$$\text{s.t.} \quad A_i x + b_i \in \mathcal{C}_i \quad i = 1 \dots m \quad (15.7)$$

and the dual is a minimization problem in standard conic form:

$$\min_{y_1, \dots, y_m} \quad \sum_{i=1}^m b_i^T y_i + b_0 \quad (15.8)$$

$$\text{s.t.} \quad a_0 + \sum_{i=1}^m A_i^T y_i = 0 \quad (15.9)$$

$$y_i \in \mathcal{C}_i^* \quad i = 1 \dots m \quad (15.10)$$

A linear inequality constraint $a^T x + b \geq c$ is equivalent to $a^T x + b - c \in \mathbb{R}_+$, and $a^T x + b \leq c$ is equivalent to $a^T x + b - c \in \mathbb{R}_-$. Variable-wise constraints are affine constraints with the appropriate identity mapping in place of A_i .

For the special case of minimization LPs, the MOI primal form can be stated as:

$$\min_{x \in \mathbb{R}^n} \quad a_0^T x + b_0 \quad (15.11)$$

$$\text{s.t.} \quad A_1 x \geq b_1 \quad (15.12)$$

$$A_2 x \leq b_2 \quad (15.13)$$

$$A_3 x = b_3 \quad (15.14)$$

By applying the stated transformations to conic form, taking the dual, and transforming back into linear inequality form, one obtains the following dual:

$$\max_{y_1, y_2, y_3} \quad b_1^T y_1 + b_2^T y_2 + b_3^T y_3 + b_0 \quad (15.15)$$

$$\text{s.t.} \quad A_1^T y_1 + A_2^T y_2 + A_3^T y_3 = a_0 \quad (15.16)$$

$$y_1 \geq 0 \quad (15.17)$$

$$y_2 \leq 0 \quad (15.18)$$

For maximization LPs, the MOI primal form can be stated as:

$$\max_{x \in \mathbb{R}^n} \quad a_0^T x + b_0 \quad (15.19)$$

$$\text{s.t.} \quad A_1 x \geq b_1 \quad (15.20)$$

$$A_2 x \leq b_2 \quad (15.21)$$

$$A_3 x = b_3 \quad (15.22)$$

and similarly, the dual is:

$$\min_{y_1, y_2, y_3} \quad -b_1^T y_1 - b_2^T y_2 - b_3^T y_3 + b_0 \quad (15.23)$$

$$\text{s.t.} \quad A_1^T y_1 + A_2^T y_2 + A_3^T y_3 = -a_0 \quad (15.24)$$

$$y_1 \geq 0 \quad (15.25)$$

$$y_2 \leq 0 \quad (15.26)$$

Warning

For the LP case, the signs of the feasible dual variables depend only on the sense of the corresponding primal inequality and not on the objective sense.

15.1 Duality and scalar product

The scalar product is different from the canonical one for the sets [PositiveSemidefiniteConeTriangle](#), [LogDetConeTriangle](#), [RootDetConeTriangle](#).

If the set C_i of the section [Duality](#) is one of these three cones, then the rows of the matrix A_i corresponding to off-diagonal entries are twice the value of the coefficients field in the [VectorAffineFunction](#) for the corresponding rows. See [PositiveSemidefiniteConeTriangle](#) for details.

15.2 Dual for problems with quadratic functions

Quadratic Programs (QPs)

For quadratic programs with only affine conic constraints,

$$\begin{aligned} \min_{x \in \mathbb{R}^n} \quad & \frac{1}{2} x^T Q_0 x + a_0^T x + b_0 \\ \text{s.t.} \quad & A_i x + b_i \in \mathcal{C}_i \quad i = 1 \dots m. \end{aligned}$$

with cones $\mathcal{C}_i \subseteq \mathbb{R}^{m_i}$ for $i = 1, \dots, m$, consider the Lagrangian function

$$L(x, y) = \frac{1}{2} x^T Q_0 x + a_0^T x + b_0 - \sum_{i=1}^m y_i^T (A_i x + b_i).$$

Let $z(y)$ denote $\sum_{i=1}^m A_i^T y_i - a_0$, the Lagrangian can be rewritten as

$$L(x, y) = \frac{1}{2} x^T Q_0 x - z(y)^T x + b_0 - \sum_{i=1}^m y_i^T b_i.$$

The condition $\nabla_x L(x, y) = 0$ gives

$$0 = \nabla_x L(x, y) = Q_0 x + a_0 - \sum_{i=1}^m y_i^T b_i$$

which gives $Q_0 x = z(y)$. This allows to obtain that

$$\min_{x \in \mathbb{R}^n} L(x, y) = -\frac{1}{2} z(y)^T Q_0^{-1} z(y) + b_0 - \sum_{i=1}^m y_i^T b_i$$

so the dual problem is

$$\max_{y_i \in \mathcal{C}_i^*} \min_{x \in \mathbb{R}^n} -\frac{1}{2} z(y)^T Q_0^{-1} z(y) + b_0 - \sum_{i=1}^m y_i^T b_i.$$

If Q_0 is invertible, we have $x = Q_0^{-1} z(y)$ hence

$$\min_{x \in \mathbb{R}^n} L(x, y) = -\frac{1}{2} z(y)^T Q_0^{-1} z(y) + b_0 - \sum_{i=1}^m y_i^T b_i$$

so the dual problem is

$$\max_{y_i \in \mathcal{C}_i^*} -\frac{1}{2} z(y)^T Q_0^{-1} z(y) + b_0 - \sum_{i=1}^m y_i^T b_i.$$

Quadratically Constrained Quadratic Programs (QCQPs)

Given a problem with both quadratic function and quadratic objectives:

$$\begin{array}{ll} \min_{x \in \mathbb{R}^n} & \frac{1}{2}x^T Q_0 x + a_0^T x + b_0 \\ \text{s.t.} & \frac{1}{2}x^T Q_i x + a_i^T x + b_i \in \mathcal{C}_i \quad i = 1 \dots m. \end{array}$$

with cones $\mathcal{C}_i \subseteq \mathbb{R}$ for $i = 1 \dots m$, consider the Lagrangian function

$$L(x, y) = \frac{1}{2}x^T Q_0 x + a_0^T x + b_0 - \sum_{i=1}^m y_i \left(\frac{1}{2}x^T Q_i x + a_i^T x + b_i \right)$$

A pair of primal-dual variables (x^*, y^*) is optimal if

- x^* is a minimizer of

$$\min_{x \in \mathbb{R}^n} L(x, y^*).$$

That is,

$$0 = \nabla_x L(x, y^*) = Q_0 x + a_0 - \sum_{i=1}^m y_i^* (Q_i x + a_i).$$

- and y^* is a maximizer of

$$\max_{y_i \in \mathcal{C}_i^*} L(x^*, y).$$

That is, for all $i = 1, \dots, m$, $\frac{1}{2}x^T Q_i x + a_i^T x + b_i$ is either zero or in the **normal cone** of \mathcal{C}_i^* at y^* . For instance, if \mathcal{C}_i is $\{z \in \mathbb{R} : z \leq 0\}$, this means that if $\frac{1}{2}x^T Q_i x + a_i^T x + b_i$ is nonzero at x^* then $y_i^* = 0$. This is the classical complementary slackness condition.

If \mathcal{C}_i is a vector set, the discussion remains valid with $y_i(\frac{1}{2}x^T Q_i x + a_i^T x + b_i)$ replaced with the scalar product between y_i and the vector of scalar-valued quadratic functions.

Chapter 16

Infeasibility certificates

When given a conic problem that is infeasible or unbounded, some solvers can produce a certificate of infeasibility. This page explains what a certificate of infeasibility is, and the related conventions that MathOptInterface adopts.

16.1 Conic duality

MathOptInterface uses conic duality to define infeasibility certificates. A full explanation is given in the section [Duality](#), but here is a brief overview.

Minimization problems

For a minimization problem in geometric conic form, the primal is:

$$\min_{x \in \mathbb{R}^n} \quad a_0^\top x + b_0 \quad (16.1)$$

$$\text{s.t.} \quad A_i x + b_i \in \mathcal{C}_i \quad i = 1 \dots m, \quad (16.2)$$

and the dual is a maximization problem in standard conic form:

$$\max_{y_1, \dots, y_m} \quad - \sum_{i=1}^m b_i^\top y_i + b_0 \quad (16.3)$$

$$\text{s.t.} \quad a_0 - \sum_{i=1}^m A_i^\top y_i = 0 \quad (16.4)$$

$$y_i \in \mathcal{C}_i^* \quad i = 1 \dots m, \quad (16.5)$$

where each \mathcal{C}_i is a closed convex cone and \mathcal{C}_i^* is its dual cone.

Maximization problems

For a maximization problem in geometric conic form, the primal is:

$$\max_{x \in \mathbb{R}^n} \quad a_0^\top x + b_0 \quad (16.6)$$

$$\text{s.t.} \quad A_i x + b_i \in \mathcal{C}_i \quad i = 1 \dots m, \quad (16.7)$$

and the dual is a minimization problem in standard conic form:

$$\min_{y_1, \dots, y_m} \sum_{i=1}^m b_i^\top y_i + b_0 \quad (16.8)$$

$$\text{s.t.} \quad a_0 + \sum_{i=1}^m A_i^\top y_i = 0 \quad (16.9)$$

$$y_i \in \mathcal{C}_i^* \quad i = 1 \dots m. \quad (16.10)$$

16.2 Unbounded problems

A problem is unbounded if and only if:

1. there exists a feasible primal solution
2. the dual is infeasible.

A feasible primal solution—if one exists—can be obtained by setting `ObjectiveSense` to `FEASIBILITY_SENSE` before optimizing. Therefore, most solvers terminate after they prove the dual is infeasible via a certificate of dual infeasibility, but before they have found a feasible primal solution. This is also the reason that MathOptInterface defines the `DUAL_INFEASIBLE` status instead of `UNBOUNDED`.

A certificate of dual infeasibility is an improving ray of the primal problem. That is, there exists some vector d such that for all $\eta > 0$:

$$A_i(x + \eta d) + b_i \in \mathcal{C}_i, \quad i = 1 \dots m,$$

and (for minimization problems):

$$a_0^\top(x + \eta d) + b_0 < a_0^\top x + b_0,$$

for any feasible point x . The latter simplifies to $a_0^\top d < 0$. For maximization problems, the inequality is reversed, so that $a_0^\top d > 0$.

If the solver has found a certificate of dual infeasibility:

- `TerminationStatus` must be `DUAL_INFEASIBLE`
- `PrimalStatus` must be `INFEASIBILITY_CERTIFICATE`
- `VariablePrimal` must be the corresponding value of d
- `ConstraintPrimal` must be the corresponding value of $A_i d$
- `ObjectiveValue` must be the value $a_0^\top d$. Note that this is the value of the objective function at d , ignoring the constant b_0 .

Note

The choice of whether to scale the ray d to have magnitude 1 is left to the solver.

16.3 Infeasible problems

A certificate of primal infeasibility is an improving ray of the dual problem. However, because infeasibility is independent of the objective function, we first homogenize the primal problem by removing its objective.

For a minimization problem, a dual improving ray is some vector d such that for all $\eta > 0$:

$$-\sum_{i=1}^m A_i^\top (y_i + \eta d_i) = 0 \quad (16.11)$$

$$(y_i + \eta d_i) \in \mathcal{C}_i^* \quad i = 1 \dots m, \quad (16.12)$$

and:

$$-\sum_{i=1}^m b_i^\top (y_i + \eta d_i) > -\sum_{i=1}^m b_i^\top y_i,$$

for any feasible dual solution y . The latter simplifies to $-\sum_{i=1}^m b_i^\top d_i > 0$. For a maximization problem, the inequality is $\sum_{i=1}^m b_i^\top d_i < 0$. (Note that these are the same inequality, modulo a - sign.)

If the solver has found a certificate of primal infeasibility:

- `TerminationStatus` must be `INFEASIBLE`
- `DualStatus` must be `INFEASIBILITY_CERTIFICATE`
- `ConstraintDual` must be the corresponding value of d
- `DualObjectiveValue` must be the value $-\sum_{i=1}^m b_i^\top d_i$ for minimization problems and $\sum_{i=1}^m b_i^\top d_i$ for maximization problems.

Note

The choice of whether to scale the ray d to have magnitude 1 is left to the solver.

Infeasibility certificates of variable bounds

Many linear solvers (e.g., Gurobi) do not provide explicit access to the primal infeasibility certificate of a variable bound. However, given a set of linear constraints:

$$l_A \leq Ax \leq u_A \quad (16.13)$$

$$l_x \leq x \leq u_x, \quad (16.14)$$

the primal certificate of the variable bounds can be computed using the primal certificate associated with the affine constraints, d . (Note that d will have one element for each row of the A matrix, and that some or all of the elements in the vectors l_A and u_A may be $\pm\infty$. If both l_A and u_A are finite for some row, the corresponding element in d must be 0.)

Given d , compute $\bar{d} = d^\top A$. If the bound is finite, a certificate for the lower variable bound of x_i is $\max\{\bar{d}_i, 0\}$, and a certificate for the upper variable bound is $\min\{\bar{d}_i, 0\}$.

Chapter 17

Naming conventions

MOI follows several conventions for naming functions and structures. These should also be followed by packages extending MOI.

17.1 Sets

Sets encode the structure of constraints. Their names should follow the following conventions:

- Abstract types in the set hierarchy should begin with `Abstract` and end in `Set`, e.g., `AbstractScalarSet`, `AbstractVectorSet`.
- Vector-valued conic sets should end with `Cone`, e.g., `NormInfinityCone`, `SecondOrderCone`.
- Vector-valued Cartesian products should be plural and not end in `Cone`, e.g., `Nonnegatives`, not `NonnegativeCone`.
- Matrix-valued conic sets should provide two representations: `ConeSquare` and `ConeTriangle`, e.g., `RootDetConeTriangle` and `RootDetConeSquare`. See [Matrix cones](#) for more details.
- Scalar sets should be singular, not plural, e.g., `Integer`, not `Integers`.
- As much as possible, the names should follow established conventions in the domain where this set is used: for instance, convex sets should have names close to those of `CVX`, and constraint-programming sets should follow `MiniZinc`'s constraints.

Part V

API Reference

Chapter 18

Standard form

18.1 Functions

[MathOptInterface.AbstractFunction](#) – Type.

| AbstractFunction

Abstract supertype for function objects.

[source](#)

[MathOptInterface.AbstractScalarFunction](#) – Type.

| AbstractScalarFunction

Abstract supertype for scalar-valued function objects.

[source](#)

[MathOptInterface.AbstractVectorFunction](#) – Type.

| AbstractVectorFunction

Abstract supertype for vector-valued function objects.

[source](#)

[MathOptInterface.VariableIndex](#) – Type.

| VariableIndex

A type-safe wrapper for Int64 for use in referencing variables in a model. To allow for deletion, indices need not be consecutive.

[source](#)

[MathOptInterface.VectorOfVariables](#) – Type.

| VectorOfVariables(variables)

The function that extracts the vector of variables referenced by `variables`, a `Vector{VariableIndex}`. This function is naturally be used for constraints that apply to groups of variables, such as an "all different" constraint, an indicator constraint, or a complementarity constraint.

[source](#)

[MathOptInterface.ScalarAffineTerm](#) – Type.

```
struct ScalarAffineTerm{T}
    coefficient::T
    variable::VariableIndex
end
```

Represents cx_i where c is coefficient and x_i is the variable identified by variable.

[source](#)

[MathOptInterface.ScalarAffineFunction](#) – Type.

```
ScalarAffineFunction{T}(terms, constant)
```

The scalar-valued affine function $a^T x + b$, where:

- a is a sparse vector specified by a list of [ScalarAffineTerm](#) structs.
- b is a scalar specified by `constant::T`

Duplicate variable indices in terms are accepted, and the corresponding coefficients are summed together.

[source](#)

[MathOptInterface.VectorAffineTerm](#) – Type.

```
struct VectorAffineTerm{T}
    output_index::Int64
    scalar_term::ScalarAffineTerm{T}
end
```

A [ScalarAffineTerm](#) plus its index of the output component of a [VectorAffineFunction](#) or [VectorQuadraticFunction](#). `output_index` can also be interpreted as a row index into a sparse matrix, where the `scalar_term` contains the column index and coefficient.

[source](#)

[MathOptInterface.VectorAffineFunction](#) – Type.

```
VectorAffineFunction{T}(terms, constants)
```

The vector-valued affine function $Ax + b$, where:

- A is a sparse matrix specified by a list of [VectorAffineTerm](#) objects.
- b is a vector specified by constants

Duplicate indices in the A are accepted, and the corresponding coefficients are summed together.

[source](#)

[MathOptInterface.ScalarQuadraticTerm](#) – Type.

```
struct ScalarQuadraticTerm{T}
    coefficient::T
    variable_1::VariableIndex
    variable_2::VariableIndex
end
```

Represents cx_ix_j where c is coefficient, x_i is the variable identified by `variable_1` and x_j is the variable identified by `variable_2`.

[source](#)

`MathOptInterface.ScalarQuadraticFunction` - Type.

```
| ScalarQuadraticFunction{T}(quadratic_terms, affine_terms, constant)
```

The scalar-valued quadratic function $\frac{1}{2}x^T Qx + a^T x + b$, where:

- a is a sparse vector specified by a list of `ScalarAffineTerm` structs.
- b is a scalar specified by `constant`.
- Q is a symmetric matrix specified by a list of `ScalarQuadraticTerm` structs.

Duplicate indices in a or Q are accepted, and the corresponding coefficients are summed together. "Mirrored" indices (q, r) and (r, q) (where r and q are `VariableIndexes`) are considered duplicates; only one need be specified.

For example, for two scalar variables y, z , the quadratic expression $yz + y^2$ is represented by the terms `ScalarQuadraticTerm.([1.0, 2.0], [y, y], [z, y])`.

[source](#)

`MathOptInterface.VectorQuadraticTerm` - Type.

```
| struct VectorQuadraticTerm{T}
|     output_index::Int64
|     scalar_term::ScalarQuadraticTerm{T}
| end
```

A `ScalarQuadraticTerm` plus its index of the output component of a `VectorQuadraticFunction`. Each output component corresponds to a distinct sparse matrix Q_i .

[source](#)

`MathOptInterface.VectorQuadraticFunction` - Type.

```
| VectorQuadraticFunction{T}(quadratic_terms, affine_terms, constants)
```

The vector-valued quadratic function with i th component ("output index") defined as $\frac{1}{2}x^T Q_i x + a_i^T x + b_i$, where:

- a_i is a sparse vector specified by the `VectorAffineTerms` with `output_index == i`.
- b_i is a scalar specified by `constants[i]`
- Q_i is a symmetric matrix specified by the `VectorQuadraticTerm` with `output_index == i`.

Duplicate indices in a_i or Q_i are accepted, and the corresponding coefficients are summed together. "Mirrored" indices (q, r) and (r, q) (where r and q are `VariableIndexes`) are considered duplicates; only one need be specified.

[source](#)

Utilities

[MathOptInterface.output_dimension](#) – Function.

```
| output_dimension(f::AbstractFunction)
```

Return 1 if f has a scalar output and the number of output components if f has a vector output.

[source](#)

[MathOptInterface.constant](#) – Method.

```
| constant(f::Union{ScalarAffineFunction, ScalarQuadraticFunction})
```

Returns the constant term of the scalar function

[source](#)

[MathOptInterface.constant](#) – Method.

```
| constant(f::Union{VectorAffineFunction, VectorQuadraticFunction})
```

Returns the vector of constant terms of the vector function

[source](#)

[MathOptInterface.constant](#) – Method.

```
| constant(f::VariableIndex, ::Type{T}) where {T}
```

The constant term of a VariableIndex function is the zero value of the specified type T.

[source](#)

[MathOptInterface.constant](#) – Method.

```
| constant(f::VectorOfVariables, ::Type{T}) where {T}
```

The constant term of a VectorOfVariables function is a vector of zero values of the specified type T.

[source](#)

18.2 Sets

[MathOptInterface.AbstractSet](#) – Type.

```
| AbstractSet
```

Abstract supertype for set objects used to encode constraints. A set object should not contain any [VariableIndex](#) or [ConstraintIndex](#) as the set is passed unmodified during [copy_to](#).

[source](#)

[MathOptInterface.AbstractScalarSet](#) – Type.

```
| AbstractScalarSet
```

Abstract supertype for subsets of \mathbb{R} .

[source](#)

[MathOptInterface.AbstractVectorSet](#) – Type.

```
| AbstractVectorSet
```

Abstract supertype for subsets of \mathbb{R}^n for some n .

[source](#)

Utilities

[MathOptInterface.dimension](#) – Function.

```
| dimension(s::AbstractSet)
```

Return the [output_dimension](#) that an [AbstractFunction](#) should have to be used with the set s .

Examples

```
| julia> dimension(Reals(4))
4
| julia> dimension(LessThan(3.0))
1
| julia> dimension(PositiveSemidefiniteConeTriangle(2))
3
```

[source](#)

[MathOptInterface.dual_set](#) – Function.

```
| dual_set(s::AbstractSet)
```

Return the dual set of s , that is the dual cone of the set. This follows the definition of duality discussed in [Duality](#).

See [Dual cone](#) for more information.

If the dual cone is not defined it returns an error.

Examples

```
| julia> dual_set(Reals(4))
Zeros(4)
| julia> dual_set(SecondOrderCone(5))
SecondOrderCone(5)
| julia> dual_set(ExponentialCone())
DualExponentialCone()
```

[source](#)

[MathOptInterface.dual_set_type](#) – Function.

```
| dual_set_type(S::Type{<:AbstractSet})
```

Return the type of dual set of sets of type S , as returned by [dual_set](#). If the dual cone is not defined it returns an error.

Examples


```
julia> dual_set_type(Reals)
Zeros

julia> dual_set_type(SecondOrderCone)
SecondOrderCone

julia> dual_set_type(ExponentialCone)
DualExponentialCone
```

[source](#)

MathOptInterface.constant – Method.

```
constant(s::Union{EqualTo, GreaterThan, LessThan})
```

Returns the constant of the set.

[source](#)

MathOptInterface.supports_dimension_update – Function.

```
supports_dimension_update(S::Type{<:MOI.AbstractVectorSet})
```

Return a Bool indicating whether the elimination of any dimension of n-dimensional sets of type S give an n-1-dimensional set S. By default, this function returns false so it should only be implemented for sets that supports dimension update.

For instance, `supports_dimension_update(MOI.Nonnegatives)` is true because the elimination of any dimension of the n-dimensional nonnegative orthant gives the n-1-dimensional nonnegative orthant. However `supports_dimension_update(MOI.ExponentialCone)` is false.

[source](#)

MathOptInterface.update_dimension – Function.

```
update_dimension(s::AbstractVectorSet, new_dim)
```

Returns a set with the dimension modified to new_dim.

[source](#)

18.3 Scalar sets

List of recognized scalar sets.

MathOptInterface.GreaterThan – Type.

```
GreaterThan{T <: Real}(lower::T)
```

The set $[lower, \infty) \subseteq \mathbb{R}$.

[source](#)

MathOptInterface.LessThan – Type.

```
LessThan{T <: Real}(upper::T)
```

The set $(-\infty, upper] \subseteq \mathbb{R}$.

[source](#)

[MathOptInterface.EqualTo](#) – Type.

```
| EqualTo{T <: Number}(value::T)
```

The set containing the single point $x \in \mathbb{R}$ where x is given by value.

[source](#)

[MathOptInterface.Interval](#) – Type.

```
| Interval{T <: Real}(lower::T,upper::T)
```

The interval $[lower, upper] \subseteq \mathbb{R}$. If lower or upper is -Inf or Inf, respectively, the set is interpreted as a one-sided interval.

```
| Interval(s::GreaterThan{<:AbstractFloat})
```

Construct a (right-unbounded) Interval equivalent to the given [GreaterThan](#) set.

```
| Interval(s::LessThan{<:AbstractFloat})
```

Construct a (left-unbounded) Interval equivalent to the given [LessThan](#) set.

```
| Interval(s::EqualTo{<:Real})
```

Construct a (degenerate) Interval equivalent to the given [EqualTo](#) set.

[source](#)

[MathOptInterface.Integer](#) – Type.

```
| Integer()
```

The set of integers \mathbb{Z} .

[source](#)

[MathOptInterface.ZeroOne](#) – Type.

```
| ZeroOne()
```

The set $\{0, 1\}$.

[source](#)

[MathOptInterface.Semicontinuous](#) – Type.

```
| Semicontinuous{T <: Real}(lower::T,upper::T)
```

The set $\{0\} \cup [lower, upper]$.

[source](#)

[MathOptInterface.Semiinteger](#) – Type.

```
| Semiinteger{T <: Real}(lower::T,upper::T)
```

The set $\{0\} \cup \{lower, lower + 1, \dots, upper - 1, upper\}$.

[source](#)

18.4 Vector sets

List of recognized vector sets.

[MathOptInterface.Reals](#) - Type.

| `Reals(dimension)`

The set $\mathbb{R}^{dimension}$ (containing all points) of dimension dimension.

[source](#)

[MathOptInterface.Zeros](#) - Type.

| `Zeros(dimension)`

The set $\{0\}^{dimension}$ (containing only the origin) of dimension dimension.

[source](#)

[MathOptInterface.Nonnegatives](#) - Type.

| `Nonnegatives(dimension)`

The nonnegative orthant $\{x \in \mathbb{R}^{dimension} : x \geq 0\}$ of dimension dimension.

[source](#)

[MathOptInterface.Nonpositives](#) - Type.

| `Nonpositives(dimension)`

The nonpositive orthant $\{x \in \mathbb{R}^{dimension} : x \leq 0\}$ of dimension dimension.

[source](#)

[MathOptInterface.NormInfinityCone](#) - Type.

| `NormInfinityCone(dimension)`

The ℓ_∞ -norm cone $\{(t, x) \in \mathbb{R}^{dimension} : t \geq \|x\|_\infty = \max_i |x_i|\}$ of dimension dimension.

[source](#)

[MathOptInterface.NormOneCone](#) - Type.

| `NormOneCone(dimension)`

The ℓ_1 -norm cone $\{(t, x) \in \mathbb{R}^{dimension} : t \geq \|x\|_1 = \sum_i |x_i|\}$ of dimension dimension.

[source](#)

[MathOptInterface.SecondOrderCone](#) - Type.

| `SecondOrderCone(dimension)`

The second-order cone (or Lorenz cone or ℓ_2 -norm cone) $\{(t, x) \in \mathbb{R}^{dimension} : t \geq \|x\|_2\}$ of dimension dimension.

[source](#)

[MathOptInterface.RotatedSecondOrderCone](#) – Type.

| RotatedSecondOrderCone(dimension)

The rotated second-order cone $\{(t, u, x) \in \mathbb{R}^{\text{dimension}} : 2tu \geq \|x\|_2^2, t, u \geq 0\}$ of dimension dimension.

[source](#)

[MathOptInterface.GeometricMeanCone](#) – Type.

| GeometricMeanCone(dimension)

The geometric mean cone $\{(t, x) \in \mathbb{R}^{n+1} : x \geq 0, t \leq \sqrt[n]{x_1 x_2 \cdots x_n}\}$, where dimension = n + 1 >= 2.

Duality note

The dual of the geometric mean cone is $\{(u, v) \in \mathbb{R}^{n+1} : u \leq 0, v \geq 0, -u \leq n \sqrt[n]{\prod_i v_i}\}$, where dimension = n + 1 >= 2.

[source](#)

[MathOptInterface.ExponentialCone](#) – Type.

| ExponentialCone()

The 3-dimensional exponential cone $\{(x, y, z) \in \mathbb{R}^3 : y \exp(x/y) \leq z, y > 0\}$.

[source](#)

[MathOptInterface.DualExponentialCone](#) – Type.

| DualExponentialCone()

The 3-dimensional dual exponential cone $\{(u, v, w) \in \mathbb{R}^3 : -u \exp(v/u) \leq \exp(1)w, u < 0\}$.

[source](#)

[MathOptInterface.PowerCone](#) – Type.

| PowerCone{T <: Real}(exponent::T)

The 3-dimensional power cone $\{(x, y, z) \in \mathbb{R}^3 : x^{\text{exponent}} y^{1-\text{exponent}} \geq |z|, x \geq 0, y \geq 0\}$ with parameter exponent.

[source](#)

[MathOptInterface.DualPowerCone](#) – Type.

| DualPowerCone{T <: Real}(exponent::T)

The 3-dimensional power cone $\{(u, v, w) \in \mathbb{R}^3 : (\frac{u}{\text{exponent}})^{\text{exponent}} (\frac{v}{1-\text{exponent}})^{1-\text{exponent}} \geq |w|, u \geq 0, v \geq 0\}$ with parameter exponent.

[source](#)

[MathOptInterface.RelativeEntropyCone](#) – Type.

| RelativeEntropyCone(dimension)

The relative entropy cone $\{(u, v, w) \in \mathbb{R}^{1+2n} : u \geq \sum_{i=1}^n w_i \log(\frac{w_i}{v_i}), v_i \geq 0, w_i \geq 0\}$, where dimension = $2n + 1 \geq 1$.

Duality note

The dual of the relative entropy cone is $\{(u, v, w) \in \mathbb{R}^{1+2n} : \forall i, w_i \geq u(\log(\frac{u}{v_i}) - 1), v_i \geq 0, u > 0\}$ of dimension dimension = $2n + 1$.

[source](#)

`MathOptInterface.NormSpectralCone` - Type.

```
| NormSpectralCone(row_dim, column_dim)
```

The epigraph of the matrix spectral norm (maximum singular value function) $\{(t, X) \in \mathbb{R}^{1+row_dim \times column_dim} : t \geq \sigma_1(X)\}$, where σ_i is the i th singular value of the matrix X of row dimension row_dim and column dimension column_dim.

The matrix X is vectorized by stacking the columns, matching the behavior of Julia's `vec` function.

[source](#)

`MathOptInterface.NormNuclearCone` - Type.

```
| NormNuclearCone(row_dim, column_dim)
```

The epigraph of the matrix nuclear norm (sum of singular values function) $\{(t, X) \in \mathbb{R}^{1+row_dim \times column_dim} : t \geq \sum_i \sigma_i(X)\}$, where σ_i is the i th singular value of the matrix X of row dimension row_dim and column dimension column_dim.

The matrix X is vectorized by stacking the columns, matching the behavior of Julia's `vec` function.

[source](#)

`MathOptInterface.SOS1` - Type.

```
| SOS1{T <: Real}(weights::Vector{T})
```

The set corresponding to the special ordered set (SOS) constraint of type 1. Of the variables in the set, at most one can be nonzero. The weights induce an ordering of the variables; as such, they should be unique values. The k th element in the set corresponds to the k th weight in weights. See [here](#) for a description of SOS constraints and their potential uses.

[source](#)

`MathOptInterface.SOS2` - Type.

```
| SOS2{T <: Real}(weights::Vector{T})
```

The set corresponding to the special ordered set (SOS) constraint of type 2. Of the variables in the set, at most two can be nonzero, and if two are nonzero, they must be adjacent in the ordering of the set. The weights induce an ordering of the variables; as such, they should be unique values. The k th element in the set corresponds to the k th weight in weights. See [here](#) for a description of SOS constraints and their potential uses.

[source](#)

`MathOptInterface.Indicator` - Type.

```
| Indicator{A<:ActivationCondition,S<:AbstractScalarSet}(set::S)
```

The set corresponding to an indicator constraint.

When A is `ACTIVATE_ON_ZERO`, this means: $\{(y, x) \in \{0, 1\} \times \mathbb{R}^n : y = 0 \implies x \in \text{set}\}$

When A is `ACTIVATE_ON_ONE`, this means: $\{(y, x) \in \{0, 1\} \times \mathbb{R}^n : y = 1 \implies x \in \text{set}\}$

Notes

Most solvers expect that the first row of the function is interpretable as a variable index `x_i` (e.g., `1.0 * x + 0.0`). An error will be thrown if this is not the case.

Example

The constraint $\{(y, x) \in \{0, 1\} \times \mathbb{R}^2 : y = 1 \implies x_1 + x_2 \leq 9\}$ is defined as

```
f = MOI.VectorAffineFunction(
  [
    MOI.VectorAffineTerm(1, MOI.ScalarAffineTerm(1.0, y)),
    MOI.VectorAffineTerm(2, MOI.ScalarAffineTerm(1.0, x1)),
    MOI.VectorAffineTerm(2, MOI.ScalarAffineTerm(1.0, x2)),
  ],
  [0.0, 0.0],
)
s = MOI.Indicator{MOI.ACTIVATE_ON_ONE}(MOI.LessThan(9.0))
MOI.add_constraint(model, f, s)
```

[source](#)

`MathOptInterface.Complements` – Type.

```
| Complements(dimension::Base.Integer)
```

The set corresponding to a mixed complementarity constraint.

Complementarity constraints should be specified with an `AbstractVectorFunction`-in-`Complements(dimension)` constraint.

The dimension of the vector-valued function `F` must be `dimension`. This defines a complementarity constraint between the scalar function `F[i]` and the variable in `F[i + dimension/2]`. Thus, `F[i + dimension/2]` must be interpretable as a single variable `x_i` (e.g., `1.0 * x + 0.0`), and `dimension` must be even.

The mixed complementarity problem consists of finding `x_i` in the interval `[lb, ub]` (i.e., in the set `Interval(lb, ub)`), such that the following holds:

1. $F_i(x) == 0$ if $lb_i < x_i < ub_i$
2. $F_i(x) >= 0$ if $lb_i == x_i$
3. $F_i(x) <= 0$ if $x_i == ub_i$

Classically, the bounding set for `x_i` is `Interval(0, Inf)`, which recovers: $0 \leq F_i(x) \perp x_i \leq 0$, where the \perp operator implies $F_i(x) * x_i = 0$.

Examples

The problem:

```
| x -in- Interval(-1, 1)
| [-4 * x - 3, x] -in- Complements(2)
```

defines the mixed complementarity problem where the following holds:

1. $-4 * x - 3 == 0$ if $-1 < x < 1$
2. $-4 * x - 3 >= 0$ if $x == -1$
3. $-4 * x - 3 <= 0$ if $x == 1$

There are three solutions:

1. $x = -3/4$ with $F(x) = 0$
2. $x = -1$ with $F(x) = 1$
3. $x = 1$ with $F(x) = -7$

The function F can also be defined in terms of single variables. For example, the problem:

```
[x_3, x_4] -in- Nonnegatives(2)
[x_1, x_2, x_3, x_4] -in- Complements(4)
```

defines the complementarity problem where $0 \leq x_1 \perp x_3 \leq 0$ and $0 \leq x_2 \perp x_4 \leq 0$.

[source](#)

18.5 Matrix sets

Matrix sets are vectorized in order to be subtypes of [AbstractVectorSet](#).

For sets of symmetric matrices, storing both the (i, j) and (j, i) elements is redundant. Use the [AbstractSymmetricMatrixSet](#) set to represent only the vectorization of the upper triangular part of the matrix.

When the matrix of expressions constrained to be in the set is not symmetric, and hence additional constraints are needed to force the equality of the (i, j) and (j, i) elements, use the [AbstractSymmetricMatrixSetSquare](#) set.

The [Bridges.Constraint.SquareBridge](#) can transform a set from the square form to the [triangular_form](#) by adding appropriate constraints if the (i, j) and (j, i) expressions are different.

[MathOptInterface.AbstractSymmetricMatrixSetTriangle](#) – Type.

```
abstract type AbstractSymmetricMatrixSetTriangle <: AbstractVectorSet end
```

Abstract supertype for subsets of the (vectorized) cone of symmetric matrices, with [side_dimension](#) rows and columns. The entries of the upper-right triangular part of the matrix are given column by column (or equivalently, the entries of the lower-left triangular part are given row by row). A vectorized cone of [dimension](#) n corresponds to a square matrix with side dimension $\sqrt{1/4 + 2n} - 1/2$. (Because a $d \times d$ matrix has $d(d + 1)/2$ elements in the upper or lower triangle.)

Examples

The matrix

$$\begin{bmatrix} 1 & 2 & 4 \\ 2 & 3 & 5 \\ 4 & 5 & 6 \end{bmatrix}$$

has [side_dimension](#) 3 and vectorization $(1, 2, 3, 4, 5, 6)$.

Note

Two packed storage formats exist for symmetric matrices, the respective orders of the entries are:

- upper triangular column by column (or lower triangular row by row);
- lower triangular column by column (or upper triangular row by row).

The advantage of the first format is the mapping between the (i, j) matrix indices and the k index of the vectorized form. It is simpler and does not depend on the side dimension of the matrix. Indeed,

- the entry of matrix indices (i, j) has vectorized index $k = \text{div}((j - 1) * j, 2) + i$ if $i \leq j$ and $k = \text{div}((i - 1) * i, 2) + j$ if $j \leq i$;
- and the entry with vectorized index k has matrix indices $i = \text{div}(1 + \text{isqrt}(8k - 7), 2)$ and $j = k - \text{div}((i - 1) * i, 2)$ or $j = \text{div}(1 + \text{isqrt}(8k - 7), 2)$ and $i = k - \text{div}((j - 1) * j, 2)$.

Duality note

The scalar product for the symmetric matrix in its vectorized form is the sum of the pairwise product of the diagonal entries plus twice the sum of the pairwise product of the upper diagonal entries; see [p. 634, 1]. This has important consequence for duality.

Consider for example the following problem ([PositiveSemidefiniteConeTriangle](#) is a subtype of [AbstractSymmetricMatrix](#)).

$$\begin{array}{ll} \max_{x \in \mathbb{R}} & x \\ \text{s.t.} & (1, -x, 1) \in \text{PositiveSemidefiniteConeTriangle}(2). \end{array}$$

The dual is the following problem

$$\begin{array}{ll} \min_{y \in \mathbb{R}^3} & y_1 + y_3 \\ \text{s.t.} & 2y_2 = 1 \\ & y \in \text{PositiveSemidefiniteConeTriangle}(2). \end{array}$$

Why do we use $2y_2$ in the dual constraint instead of y_2 ? The reason is that $2y_2$ is the scalar product between y and the symmetric matrix whose vectorized form is $(0, 1, 0)$. Indeed, with our modified scalar products we have

$$\langle (0, 1, 0), (y_1, y_2, y_3) \rangle = \text{trace} \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix} \begin{pmatrix} y_1 & y_2 \\ y_2 & y_3 \end{pmatrix} = 2y_2.$$

References

[1] Boyd, S. and Vandenberghe, L.. Convex optimization. Cambridge university press, 2004.

[source](#)

[MathOptInterface.AbstractSymmetricMatrixSetSquare](#) - Type.

```
| abstract type AbstractSymmetricMatrixSetSquare <: AbstractVectorSet end
```

Abstract supertype for subsets of the (vectorized) cone of symmetric matrices, with [side_dimension](#) rows and columns. The entries of the matrix are given column by column (or equivalently, row by row). The matrix is both constrained to be symmetric and to have its [triangular_form](#) belong to the corresponding

set. That is, if the functions in entries (i, j) and (j, i) are different, then a constraint will be added to make sure that the entries are equal.

Examples

[PositiveSemidefiniteConeSquare](#) is a subtype of [AbstractSymmetricMatrixSetSquare](#) and constraining the matrix

$$\begin{bmatrix} 1 & -y \\ -z & 0 \end{bmatrix}$$

to be symmetric positive semidefinite can be achieved by constraining the vector $(1, -z, -y, 0)$ (or $(1, -y, -z, 0)$) to belong to the [PositiveSemidefiniteConeSquare\(2\)](#). It both constrains $y = z$ and $(1, -y, 0)$ (or $(1, -z, 0)$) to be in [PositiveSemidefiniteConeTriangle\(2\)](#), since `triangular_form(PositiveSemidefiniteConeSquare)` is [PositiveSemidefiniteConeTriangle](#).

[source](#)

[MathOptInterface.side_dimension](#) – Function.

```
side_dimension(set::Union{AbstractSymmetricMatrixSetTriangle,
                          AbstractSymmetricMatrixSetSquare})
```

Side dimension of the matrices in set. By convention, it should be stored in the `side_dimension` field but if it is not the case for a subtype of [AbstractSymmetricMatrixSetTriangle](#), the method should be implemented for this subtype.

[source](#)

[MathOptInterface.triangular_form](#) – Function.

```
triangular_form(S::Type{<:AbstractSymmetricMatrixSetSquare})
triangular_form(set::AbstractSymmetricMatrixSetSquare)
```

Return the [AbstractSymmetricMatrixSetTriangle](#) corresponding to the vectorization of the upper triangular part of matrices in the [AbstractSymmetricMatrixSetSquare](#) set.

[source](#)

List of recognized matrix sets.

[MathOptInterface.PositiveSemidefiniteConeTriangle](#) – Type.

```
PositiveSemidefiniteConeTriangle(side_dimension) <: AbstractSymmetricMatrixSetTriangle
```

The (vectorized) cone of symmetric positive semidefinite matrices, with `side_dimension` rows and columns.

See [AbstractSymmetricMatrixSetTriangle](#) for more details on the vectorized form.

[source](#)

[MathOptInterface.PositiveSemidefiniteConeSquare](#) – Type.

```
PositiveSemidefiniteConeSquare(side_dimension) <: AbstractSymmetricMatrixSetSquare
```

The cone of symmetric positive semidefinite matrices, with side length `side_dimension`.

See [AbstractSymmetricMatrixSetSquare](#) for more details on the vectorized form.

The entries of the matrix are given column by column (or equivalently, row by row).

The matrix is both constrained to be symmetric and to be positive semidefinite. That is, if the functions in entries (i, j) and (j, i) are different, then a constraint will be added to make sure that the entries are equal.

Examples

Constraining the matrix

$$\begin{bmatrix} 1 & -y \\ -z & 0 \end{bmatrix}$$

to be symmetric positive semidefinite can be achieved by constraining the vector $(1, -z, -y, 0)$ (or $(1, -y, -z, 0)$) to belong to the `PositiveSemidefiniteConeSquare(2)`.

It both constrains $y = z$ and $(1, -y, 0)$ (or $(1, -z, 0)$) to be in `PositiveSemidefiniteConeTriangle(2)`.

[source](#)

`MathOptInterface.LogDetConeTriangle` - Type.

```
| LogDetConeTriangle(side_dimension)
```

The log-determinant cone $\{(t, u, X) \in \mathbb{R}^{2+d(d+1)/2} : t \leq u \log(\det(X/u)), u > 0\}$, where the matrix X is represented in the same symmetric packed format as in the `PositiveSemidefiniteConeTriangle`.

The argument `side_dimension` is the side dimension of the matrix X , i.e., its number of rows or columns.

[source](#)

`MathOptInterface.LogDetConeSquare` - Type.

```
| LogDetConeSquare(side_dimension)
```

The log-determinant cone $\{(t, u, X) \in \mathbb{R}^{2+d^2} : t \leq u \log(\det(X/u)), X \text{ symmetric}, u > 0\}$, where the matrix X is represented in the same format as in the `PositiveSemidefiniteConeSquare`.

Similarly to `PositiveSemidefiniteConeSquare`, constraints are added to ensure that X is symmetric.

The argument `side_dimension` is the side dimension of the matrix X , i.e., its number of rows or columns.

[source](#)

`MathOptInterface.RootDetConeTriangle` - Type.

```
| RootDetConeTriangle(side_dimension)
```

The root-determinant cone $\{(t, X) \in \mathbb{R}^{1+d(d+1)/2} : t \leq \det(X)^{1/d}\}$, where the matrix X is represented in the same symmetric packed format as in the `PositiveSemidefiniteConeTriangle`.

The argument `side_dimension` is the side dimension of the matrix X , i.e., its number of rows or columns.

[source](#)

`MathOptInterface.RootDetConeSquare` - Type.

```
| RootDetConeSquare(side_dimension)
```

The root-determinant cone $\{(t, X) \in \mathbb{R}^{1+d^2} : t \leq \det(X)^{1/d}, X \text{ symmetric}\}$, where the matrix X is represented in the same format as [PositiveSemidefiniteConeSquare](#).

Similarly to [PositiveSemidefiniteConeSquare](#), constraints are added to ensure that X is symmetric.

The argument `side_dimension` is the side dimension of the matrix X , i.e., its number of rows or columns.

[source](#)

Chapter 19

Models

19.1 Attribute interface

`MathOptInterface.is_set_by_optimize` – Function.

```
| is_set_by_optimize(::AnyAttribute)
```

Return a Bool indicating whether the value of the attribute is modified during an `optimize!` call, that is, the attribute is used to query the result of the optimization.

Important note when defining new attributes

This function returns false by default so it should be implemented for attributes that are modified by `optimize!`.

[source](#)

`MathOptInterface.is_copyable` – Function.

```
| is_copyable(::AnyAttribute)
```

Return a Bool indicating whether the value of the attribute may be copied during `copy_to` using `set`.

Important note when defining new attributes

By default `is_copyable(attr)` returns `!is_set_by_optimize(attr)`. A specific method should be defined for attributes which are copied indirectly during `copy_to`. For instance, both `is_copyable` and `is_set_by_optimize` return false for the following attributes:

- `ListOfOptimizerAttributesSet`, `ListOfModelAttributesSet`, `ListOfConstraintAttributesSet` and `ListOfVariableAttributesSet`.
- `SolverName` and `RawSolver`: these attributes cannot be set.
- `NumberOfVariables` and `ListOfVariableIndices`: these attributes are set indirectly by `add_variable` and `add_variables`.
- `ObjectiveFunctionType`: this attribute is set indirectly when setting the `ObjectiveFunction` attribute.
- `NumberOfConstraints`, `ListOfConstraintIndices`, `ListOfConstraintTypesPresent`, `CanonicalConstraintFunction`, `ConstraintFunction` and `ConstraintSet`: these attributes are set indirectly by `add_constraint` and `add_constraints`.

[source](#)

[MathOptInterface.get](#) – Function.

```
| get(optimizer::AbstractOptimizer, attr::AbstractOptimizerAttribute)
```

Return an attribute `attr` of the optimizer `optimizer`.

```
| get(model::ModelLike, attr::AbstractModelAttribute)
```

Return an attribute `attr` of the model `model`.

```
| get(model::ModelLike, attr::AbstractVariableAttribute, v::VariableIndex)
```

If the attribute `attr` is set for the variable `v` in the model `model`, return its value, return nothing otherwise. If the attribute `attr` is not supported by model then an error should be thrown instead of returning nothing.

```
| get(model::ModelLike, attr::AbstractVariableAttribute, v::Vector{VariableIndex})
```

Return a vector of attributes corresponding to each variable in the collection `v` in the model `model`.

```
| get(model::ModelLike, attr::AbstractConstraintAttribute, c::ConstraintIndex)
```

If the attribute `attr` is set for the constraint `c` in the model `model`, return its value, return nothing otherwise. If the attribute `attr` is not supported by model then an error should be thrown instead of returning nothing.

```
| get(model::ModelLike, attr::AbstractConstraintAttribute, c::Vector{ConstraintIndex{F,S}})
```

Return a vector of attributes corresponding to each constraint in the collection `c` in the model `model`.

```
| get(model::ModelLike, ::Type{VariableIndex}, name::String)
```

If a variable with name `name` exists in the model `model`, return the corresponding index, otherwise return nothing. Errors if two variables have the same name.

```
| get(model::ModelLike, ::Type{ConstraintIndex{F,S}}, name::String) where {F<:AbstractFunction,S<:AbstractSet}
```

If an F-in-S constraint with name `name` exists in the model `model`, return the corresponding index, otherwise return nothing. Errors if two constraints have the same name.

```
| get(model::ModelLike, ::Type{ConstraintIndex}, name::String)
```

If any constraint with name `name` exists in the model `model`, return the corresponding index, otherwise return nothing. This version is available for convenience but may incur a performance penalty because it is not type stable. Errors if two constraints have the same name.

Examples

```
| get(model, ObjectiveValue())
| get(model, VariablePrimal(), ref)
| get(model, VariablePrimal(5), [ref1, ref2])
| get(model, OtherAttribute("something specific to cplex"))
| get(model, VariableIndex, "var1")
| get(model, ConstraintIndex{ScalarAffineFunction{Float64},LessThan{Float64}}, "con1")
| get(model, ConstraintIndex, "con1")
```

[source](#)

[MathOptInterface.get!](#) – Function.

```
| get!(output, model::ModelLike, args...)
```

An in-place version of [get](#).

The signature matches that of [get](#) except that the the result is placed in the vector output.

[source](#)

[MathOptInterface.set](#) – Function.

```
| set(optimizer::AbstractOptimizer, attr::AbstractOptimizerAttribute, value)
```

Assign value to the attribute attr of the optimizer optimizer.

```
| set(model::ModelLike, attr::AbstractModelAttribute, value)
```

Assign value to the attribute attr of the model model.

```
| set(model::ModelLike, attr::AbstractVariableAttribute, v::VariableIndex, value)
```

Assign value to the attribute attr of variable v in model model.

```
| set(model::ModelLike, attr::AbstractVariableAttribute, v::Vector{VariableIndex}, vector_of_values
    )
```

Assign a value respectively to the attribute attr of each variable in the collection v in model model.

```
| set(model::ModelLike, attr::AbstractConstraintAttribute, c::ConstraintIndex, value)
```

Assign a value to the attribute attr of constraint c in model model.

```
| set(model::ModelLike, attr::AbstractConstraintAttribute, c::Vector{ConstraintIndex{F,S}},
    vector_of_values)
```

Assign a value respectively to the attribute attr of each constraint in the collection c in model model.

An [UnsupportedAttribute](#) error is thrown if model does not support the attribute attr (see [supports](#)) and a [SetAttributeNotAllowed](#) error is thrown if it supports the attribute attr but it cannot be set.

Replace set in a constraint

```
| set(model::ModelLike, ::ConstraintSet, c::ConstraintIndex{F,S}, set::S)
```

Change the set of constraint c to the new set set which should be of the same type as the original set.

Examples

If c is a `ConstraintIndex{F,Interval}`

```
| set(model, ConstraintSet(), c, Interval(0, 5))
| set(model, ConstraintSet(), c, GreaterThan(0.0)) # Error
```

Replace function in a constraint

```
| set(model::ModelLike, ::ConstraintFunction, c::ConstraintIndex{F,S}, func::F)
```

Replace the function in constraint c with func. F must match the original function type used to define the constraint.

Note

Setting the constraint function is not allowed if F is `VariableIndex`, it throws a `SettingVariableIndexNotAllowed` error. Indeed, it would require changing the index c as the index of `VariableIndex` constraints should be the same as the index of the variable.

Examples

If c is a `ConstraintIndex{ScalarAffineFunction,S}` and $v1$ and $v2$ are `VariableIndex` objects,

```
set(model, ConstraintFunction(), c,
  ScalarAffineFunction(ScalarAffineTerm.([1.0, 2.0], [v1, v2]), 5.0))
set(model, ConstraintFunction(), c, v1) # Error
```

[source](#)

`MathOptInterface.supports` – Function.

```
supports(model::ModelLike, sub::AbstractSubmittable)::Bool
```

Return a `Bool` indicating whether `model` supports the submittable `sub`.

```
supports(model::ModelLike, attr::AbstractOptimizerAttribute)::Bool
```

Return a `Bool` indicating whether `model` supports the optimizer attribute `attr`. That is, it returns false if `copy_to(model, src)` shows a warning in case `attr` is in the `ListOfOptimizerAttributesSet` of `src`; see [copy_to](#) for more details on how unsupported optimizer attributes are handled in copy.

```
supports(model::ModelLike, attr::AbstractModelAttribute)::Bool
```

Return a `Bool` indicating whether `model` supports the model attribute `attr`. That is, it returns false if `copy_to(model, src)` cannot be performed in case `attr` is in the `ListOfModelAttributeSet` of `src`.

```
supports(model::ModelLike, attr::AbstractVariableAttribute, ::Type{VariableIndex})::Bool
```

Return a `Bool` indicating whether `model` supports the variable attribute `attr`. That is, it returns false if `copy_to(model, src)` cannot be performed in case `attr` is in the `ListOfVariableAttributesSet` of `src`.

```
supports(model::ModelLike, attr::AbstractConstraintAttribute, ::Type{ConstraintIndex{F,S}})::Bool
  where {F,S}
```

Return a `Bool` indicating whether `model` supports the constraint attribute `attr` applied to an F -in- S constraint. That is, it returns false if `copy_to(model, src)` cannot be performed in case `attr` is in the `ListOfConstraintAttributesSet` of `src`.

For all five methods, if the attribute is only not supported in specific circumstances, it should still return true.

Note that `supports` is only defined for attributes for which `is_copyable` returns true as other attributes do not appear in the list of attributes set obtained by `ListOf...AttributesSet`.

[source](#)

`MathOptInterface.attribute_value_type` – Function.

```
attribute_value_type(attr::AnyAttribute)
```

Given an attribute `attr`, return the type of value expected by `get`, or returned by `set`.

Notes

- Only implement this if it make sense to do so. If un-implemented, the default is `Any`.

[source](#)

19.2 Model interface

`MathOptInterface.ModelLike` - Type.

```
| ModelLike
```

Abstract supertype for objects that implement the "Model" interface for defining an optimization problem.

[source](#)

`MathOptInterface.is_empty` - Function.

```
| is_empty(model::ModelLike)
```

Returns false if the model has any model attribute set or has any variables or constraints.

Note that an empty model can have optimizer attributes set.

[source](#)

`MathOptInterface.empty!` - Function.

```
| empty!(model::ModelLike)
```

Empty the model, that is, remove all variables, constraints and model attributes but not optimizer attributes.

[source](#)

`MathOptInterface.write_to_file` - Function.

```
| write_to_file(model::ModelLike, filename::String)
```

Writes the current model data to the given file. Supported file types depend on the model type.

[source](#)

`MathOptInterface.read_from_file` - Function.

```
| read_from_file(model::ModelLike, filename::String)
```

Read the file filename into the model model. If model is non-empty, this may throw an error.

Supported file types depend on the model type.

Note

Once the contents of the file are loaded into the model, users can query the variables via `get(model, ListOfVariableIndices())`. However, some filetypes, such as LP files, do not maintain an explicit ordering of the variables. Therefore, the returned list may be in an arbitrary order. To avoid depending on the order of the indices, users should look up each variable index by name: `get(model, VariableIndex, "name")`.

[source](#)

`MathOptInterface.supports_incremental_interface` - Function.

```
| supports_incremental_interface(model::ModelLike)
```


Return a `Bool` indicating whether `model` supports building incrementally via `add_variable` and `add_constraint`.

The main purpose of this function is to determine whether a model can be loaded into `model` incrementally or whether it should be cached and copied at once instead.

[source](#)

`MathOptInterface.copy_to` – Function.

```
| copy_to(dest::ModelLike, src::ModelLike)::IndexMap
```

Copy the model from `src` into `dest`.

The target `dest` is emptied, and all previous indices to variables and constraints in `dest` are invalidated.

Returns an `IndexMap` object that translates variable and constraint indices from the `src` model to the corresponding indices in the `dest` model.

Notes

- If a constraint that in `src` is not supported by `dest`, then an `UnsupportedConstraint` error is thrown.
- If an `AbstractModelAttribute`, `AbstractVariableAttribute`, or `AbstractConstraintAttribute` is set in `src` but not supported by `dest`, then an `UnsupportedAttribute` error is thrown.

`AbstractOptimizerAttributes` are not copied to the `dest` model.

IndexMap

Implementations of `copy_to` must return an `IndexMap`. For technical reasons, this type is defined in the `Utilities` submodule as `MOI.Utilities.IndexMap`. However, since it is an integral part of the MOI API, we provide `MOI.IndexMap` as an alias.

Example

```
| # Given empty `ModelLike` objects `src` and `dest`.
|
| x = add_variable(src)
|
| is_valid(src, x) # true
| is_valid(dest, x) # false (`dest` has no variables)
|
| index_map = copy_to(dest, src)
| is_valid(dest, x) # false (unless index_map[x] == x)
| is_valid(dest, index_map[x]) # true
```

[source](#)

`MathOptInterface.IndexMap` – Type.

```
| IndexMap()
```

The dictionary-like object returned by `copy_to`.

IndexMap

Implementations of `copy_to` must return an `IndexMap`. For technical reasons, the `IndexMap` type is defined in the `Utilities` submodule as `MOI.Utilities.IndexMap`. However, since it is an integral part of the MOI API, we provide this `MOI.IndexMap` as an alias.

[source](#)

19.3 Model attributes

[MathOptInterface.AbstractModelAttribute](#) - Type.

```
| AbstractModelAttribute
```

Abstract supertype for attribute objects that can be used to set or get attributes (properties) of the model.

[source](#)

[MathOptInterface.Name](#) - Type.

```
| Name()
```

A model attribute for the string identifying the model. It has a default value of "" if not set'.

[source](#)

[MathOptInterface.ObjectiveFunction](#) - Type.

```
| ObjectiveFunction{F<:AbstractScalarFunction}()
```

A model attribute for the objective function which has a type `F<:AbstractScalarFunction`. `F` should be guaranteed to be equivalent but not necessarily identical to the function type provided by the user. Throws an `InexactError` if the objective function cannot be converted to `F`, e.g. the objective function is quadratic and `F` is `ScalarAffineFunction{Float64}` or it has non-integer coefficient and `F` is `ScalarAffineFunction{Int}`.

[source](#)

[MathOptInterface.ObjectiveFunctionType](#) - Type.

```
| ObjectiveFunctionType()
```

A model attribute for the type `F` of the objective function set using the `ObjectiveFunction{F}` attribute.

Examples

In the following code, `attr` should be equal to `MOI.VariableIndex`:

```
| x = MOI.add_variable(model)
| MOI.set(model, MOI.ObjectiveFunction{MOI.VariableIndex}(),
|           x)
| attr = MOI.get(model, MOI.ObjectiveFunctionType())
```

[source](#)

[MathOptInterface.ObjectiveSense](#) - Type.

```
| ObjectiveSense()
```

A model attribute for the objective sense of the objective function, which must be an `OptimizationSense`: `MIN_SENSE`, `MAX_SENSE`, or `FEASIBILITY_SENSE`. The default is `FEASIBILITY_SENSE`.

When setting the sense to `FEASIBILITY_SENSE`, the [ObjectiveFunction](#) attribute must be erased.

[source](#)

[MathOptInterface.NumberOfVariables](#) - Type.

```
| NumberOfVariables()
```

A model attribute for the number of variables in the model.

[source](#)

[MathOptInterface.ListOfVariableIndices](#) - Type.

```
| ListOfVariableIndices()
```

A model attribute for the `Vector{VariableIndex}` of all variable indices present in the model (i.e., of length equal to the value of `NumberOfVariables()`) in the order in which they were added.

[source](#)

[MathOptInterface.ListOfConstraintTypesPresent](#) - Type.

```
| ListOfConstraintTypesPresent()
```

A model attribute for the list of tuples of the form (F, S) , where F is a function type and S is a set type indicating that the attribute `NumberOfConstraints{F,S}()` has value greater than zero.

[source](#)

[MathOptInterface.NumberOfConstraints](#) - Type.

```
| NumberOfConstraints{F,S}()
```

A model attribute for the number of constraints of the type F -in- S present in the model.

[source](#)

[MathOptInterface.ListOfConstraintIndices](#) - Type.

```
| ListOfConstraintIndices{F,S}()
```

A model attribute for the `Vector{ConstraintIndex{F,S}}` of all constraint indices of type F -in- S in the model (i.e., of length equal to the value of `NumberOfConstraints{F,S}()`) in the order in which they were added.

[source](#)

[MathOptInterface.ListOfOptimizerAttributesSet](#) - Type.

```
| ListOfOptimizerAttributesSet()
```

An optimizer attribute for the `Vector{AbstractOptimizerAttribute}` of all optimizer attributes that were set.

[source](#)

[MathOptInterface.ListOfModelAttributeSet](#) - Type.

```
| ListOfModelAttributeSet()
```

A model attribute for the `Vector{AbstractModelAttribute}` of all model attributes `attr` such that 1) `is_copyable(attr)` returns true and 2) the attribute was set to the model.

[source](#)

[MathOptInterface.ListOfVariableAttributesSet](#) - Type.

```
| ListOfVariableAttributesSet()
```

A model attribute for the `Vector{AbstractVariableAttribute}` of all variable attributes `attr` such that 1) `is_copyable(attr)` returns true and 2) the attribute was set to variables.

[source](#)

[MathOptInterface.ListOfConstraintAttributesSet](#) - Type.

```
| ListOfConstraintAttributesSet{F, S}()
```

A model attribute for the `Vector{AbstractConstraintAttribute}` of all constraint attributes `attr` such that 1) `is_copyable(attr)` returns true and

2. the attribute was set to F-in-S constraints.

Note

The attributes [ConstraintFunction](#) and [ConstraintSet](#) should not be included in the list even if then have been set with [set](#).

[source](#)

19.4 Optimizer interface

[MathOptInterface.AbstractOptimizer](#) - Type.

```
| AbstractOptimizer <: ModelLike
```

Abstract supertype for objects representing an instance of an optimization problem tied to a particular solver. This is typically a solver's in-memory representation. In addition to `ModelLike`, `AbstractOptimizer` objects let you solve the model and query the solution.

[source](#)

[MathOptInterface.OptimizerWithAttributes](#) - Type.

```
| struct OptimizerWithAttributes
|     optimizer_constructor
|     params::Vector{Pair{AbstractOptimizerAttribute,<:Any}}
| end
```

Object grouping an optimizer constructor and a list of optimizer attributes. Instances are created with [instantiate](#).

[source](#)

[MathOptInterface.optimize!](#) - Function.

```
| optimize!(optimizer::AbstractOptimizer)
```

Optimize the problem contained in `optimizer`.

Before calling `optimize!`, the problem should first be constructed using the incremental interface (see [supports_incremental_interface](#)) or [copy_to](#).

[source](#)

[MathOptInterface.instantiate](#) – Function.

```

instantiate(
  optimizer_constructor,
  with_bridge_type: Union{Nothing, Type} = nothing,
)

```

Creates an instance of optimizer by either:

- calling `optimizer_constructor.optimizer_constructor()` and setting the parameters in `optimizer_constructor.p` if `optimizer_constructor` is a [OptimizerWithAttributes](#)
- calling `optimizer_constructor()` if `optimizer_constructor` is callable.

If `with_bridge_type` is not `nothing`, it enables all the bridges defined in the `MathOptInterface.Bridges` submodule with coefficient type `with_bridge_type`.

If the optimizer created by `optimizer_constructor` does not support loading the problem incrementally (see [supports_incremental_interface](#)), then a [Utilities.CachingOptimizer](#) is added to store a cache of the bridged model.

[source](#)

[MathOptInterface.default_cache](#) – Function.

```

default_cache(optimizer::ModelLike, ::Type{T}) where {T}

```

Return a new instance of the default model type to be used as cache for optimizer in a [Utilities.CachingOptimizer](#) for holding constraints of coefficient type `T`. By default, this returns `Utilities.UniversalFallback(Utilities.Model{T}())`. If copying from a instance of a given model type is faster for optimizer then a new method returning an instance of this model type should be defined.

[source](#)

19.5 Optimizer attributes

[MathOptInterface.AbstractOptimizerAttribute](#) – Type.

```

AbstractOptimizerAttribute

```

Abstract supertype for attribute objects that can be used to set or get attributes (properties) of the optimizer.

Note

The difference between `AbstractOptimizerAttribute` and `AbstractModelAttribute` lies in the behavior of `is_empty`, `empty!` and `copy_to`. Typically optimizer attributes only affect how the model is solved.

[source](#)

[MathOptInterface.SolverName](#) – Type.

```

SolverName()

```

An optimizer attribute for the string identifying the solver/optimizer.

[source](#)

[MathOptInterface.SolverVersion](#) – Type.

```
| SolverVersion()
```

An optimizer attribute for the string identifying the version of the solver.

Note

For solvers supporting [semantic versioning](#), the `SolverVersion` should be a string of the form "vMAJOR.MINOR.PATCH", so that it can be converted to a Julia `VersionNumber` (e.g., `VersionNumber("v1.2.3")`).

We do not require Semantic Versioning because some solvers use alternate versioning systems. For example, CPLEX uses Calendar Versioning, so `SolverVersion` will return a string like "202001".

[source](#)

`MathOptInterface.Silent` – Type.

```
| Silent()
```

An optimizer attribute for silencing the output of an optimizer. When set to `true`, it takes precedence over any other attribute controlling verbosity and requires the solver to produce no output. The default value is `false` which has no effect. In this case the verbosity is controlled by other attributes.

Note

Every optimizer should have verbosity on by default. For instance, if a solver has a solver-specific log level attribute, the MOI implementation should set it to 1 by default. If the user sets `Silent` to `true`, then the log level should be set to 0, even if the user specifically sets a value of log level. If the value of `Silent` is `false` then the log level set to the solver is the value given by the user for this solver-specific parameter or 1 if none is given.

[source](#)

`MathOptInterface.TimeLimitSec` – Type.

```
| TimeLimitSec()
```

An optimizer attribute for setting a time limit for an optimization. When set to `nothing`, it deactivates the solver time limit. The default value is `nothing`. The time limit is in seconds.

[source](#)

`MathOptInterface.RawOptimizerAttribute` – Type.

```
| RawOptimizerAttribute(name::String)
```

An optimizer attribute for the solver-specific parameter identified by `name`.

[source](#)

`MathOptInterface.NumberOfThreads` – Type.

```
| NumberOfThreads()
```

An optimizer attribute for setting the number of threads used for an optimization. When set to `nothing` uses solver default. Values are positive integers. The default value is `nothing`.

[source](#)

`MathOptInterface.RawSolver` – Type.

```
| RawSolver()
```

A model attribute for the object that may be used to access a solver-specific API for this optimizer.

[source](#)

List of attributes useful for optimizers

`MathOptInterface.TerminationStatus` – Type.

```
| TerminationStatus()
```

A model attribute for the `TerminationStatusCode` explaining why the optimizer stopped.

[source](#)

`MathOptInterface.TerminationStatusCode` – Type.

```
| TerminationStatusCode
```

An Enum of possible values for the `TerminationStatus` attribute. This attribute is meant to explain the reason why the optimizer stopped executing in the most recent call to `optimize!`.

If no call has been made to `optimize!`, then the `TerminationStatus` is:

- `OPTIMIZE_NOT_CALLED`: The algorithm has not started.

OK

These are generally OK statuses, i.e., the algorithm ran to completion normally.

- `OPTIMAL`: The algorithm found a globally optimal solution.
- `INFEASIBLE`: The algorithm concluded that no feasible solution exists.
- `DUAL_INFEASIBLE`: The algorithm concluded that no dual bound exists for the problem. If, additionally, a feasible (primal) solution is known to exist, this status typically implies that the problem is unbounded, with some technical exceptions.
- `LOCALLY_SOLVED`: The algorithm converged to a stationary point, local optimal solution, could not find directions for improvement, or otherwise completed its search without global guarantees.
- `LOCALLY_INFEASIBLE`: The algorithm converged to an infeasible point or otherwise completed its search without finding a feasible solution, without guarantees that no feasible solution exists.
- `INFEASIBLE_OR_UNBOUNDED`: The algorithm stopped because it decided that the problem is infeasible or unbounded; this occasionally happens during MIP presolve.

Solved to relaxed tolerances

- `ALMOST_OPTIMAL`: The algorithm found a globally optimal solution to relaxed tolerances.
- `ALMOST_INFEASIBLE`: The algorithm concluded that no feasible solution exists within relaxed tolerances.
- `ALMOST_DUAL_INFEASIBLE`: The algorithm concluded that no dual bound exists for the problem within relaxed tolerances.
- `ALMOST_LOCALLY_SOLVED`: The algorithm converged to a stationary point, local optimal solution, or could not find directions for improvement within relaxed tolerances.

Limits

The optimizer stopped because of some user-defined limit.

- `ITERATION_LIMIT`: An iterative algorithm stopped after conducting the maximum number of iterations.
- `TIME_LIMIT`: The algorithm stopped after a user-specified computation time.
- `NODE_LIMIT`: A branch-and-bound algorithm stopped because it explored a maximum number of nodes in the branch-and-bound tree.
- `SOLUTION_LIMIT`: The algorithm stopped because it found the required number of solutions. This is often used in MIPs to get the solver to return the first feasible solution it encounters.
- `MEMORY_LIMIT`: The algorithm stopped because it ran out of memory.
- `OBJECTIVE_LIMIT`: The algorithm stopped because it found a solution better than a minimum limit set by the user.
- `NORM_LIMIT`: The algorithm stopped because the norm of an iterate became too large.
- `OTHER_LIMIT`: The algorithm stopped due to a limit not covered by one of the above.

Problematic

This group of statuses means that something unexpected or problematic happened.

- `SLOW_PROGRESS`: The algorithm stopped because it was unable to continue making progress towards the solution.
- `NUMERICAL_ERROR`: The algorithm stopped because it encountered unrecoverable numerical error.
- `INVALID_MODEL`: The algorithm stopped because the model is invalid.
- `INVALID_OPTION`: The algorithm stopped because it was provided an invalid option.
- `INTERRUPTED`: The algorithm stopped because of an interrupt signal.
- `OTHER_ERROR`: The algorithm stopped because of an error not covered by one of the statuses defined above.

[source](#)

`MathOptInterface.PrimalStatus` – Type.

```
| PrimalStatus(result_index::Int = 1)
```

A model attribute for the `ResultStatusCode` of the primal result `result_index`. If `result_index` is omitted, it defaults to 1.

See `ResultCount` for information on how the results are ordered.

If `result_index` is larger than the value of `ResultCount` then `NO_SOLUTION` is returned.

[source](#)

`MathOptInterface.DualStatus` – Type.

```
| DualStatus(result_index::Int = 1)
```

A model attribute for the `ResultStatusCode` of the dual result `result_index`. If `result_index` is omitted, it defaults to 1.

See `ResultCount` for information on how the results are ordered.

If `result_index` is larger than the value of `ResultCount` then `NO_SOLUTION` is returned.

[source](#)

`MathOptInterface.ResultStatusCode` – Type.

`| ResultStatusCode`

An Enum of possible values for the `PrimalStatus` and `DualStatus` attributes. The values indicate how to interpret the result vector.

- `NO_SOLUTION`: the result vector is empty.
- `FEASIBLE_POINT`: the result vector is a feasible point.
- `NEARLY_FEASIBLE_POINT`: the result vector is feasible if some constraint tolerances are relaxed.
- `INFEASIBLE_POINT`: the result vector is an infeasible point.
- `INFEASIBILITY_CERTIFICATE`: the result vector is an infeasibility certificate. If the `PrimalStatus` is `INFEASIBILITY_CERTIFICATE`, then the primal result vector is a certificate of dual infeasibility. If the `DualStatus` is `INFEASIBILITY_CERTIFICATE`, then the dual result vector is a proof of primal infeasibility.
- `NEARLY_INFEASIBILITY_CERTIFICATE`: the result satisfies a relaxed criterion for a certificate of infeasibility.
- `REDUCTION_CERTIFICATE`: the result vector is an ill-posed certificate; see [this article](#) for details. If the `PrimalStatus` is `REDUCTION_CERTIFICATE`, then the primal result vector is a proof that the dual problem is ill-posed. If the `DualStatus` is `REDUCTION_CERTIFICATE`, then the dual result vector is a proof that the primal is ill-posed.
- `NEARLY_REDUCTION_CERTIFICATE`: the result satisfies a relaxed criterion for an ill-posed certificate.
- `UNKNOWN_RESULT_STATUS`: the result vector contains a solution with an unknown interpretation.
- `OTHER_RESULT_STATUS`: the result vector contains a solution with an interpretation not covered by one of the statuses defined above.

[source](#)

`MathOptInterface.RawStatusString` – Type.

`| RawStatusString()`

A model attribute for a solver specific string explaining why the optimizer stopped.

[source](#)

`MathOptInterface.ResultCount` – Type.

`| ResultCount()`

A model attribute for the number of results available.

Order of solutions

A number of attributes contain an index, `result_index`, which is used to refer to one of the available results. Thus, `result_index` must be an integer between 1 and the number of available results.

As a general rule, the first result (`result_index=1`) is the most important result (e.g., an optimal solution or an infeasibility certificate). Other results will typically be alternate solutions that the solver found during the search for the first result.

If a (local) optimal solution is available, i.e., `TerminationStatus` is `OPTIMAL` or `LOCALLY_SOLVED`, the first result must correspond to the (locally) optimal solution. Other results may be alternative optimal solutions, or they may be other suboptimal solutions; use `ObjectiveValue` to distinguish between them.

If a primal or dual infeasibility certificate is available, i.e., [TerminationStatus](#) is INFEASIBLE or DUAL_INFEASIBLE and the corresponding [PrimalStatus](#) or [DualStatus](#) is INFEASIBILITY_CERTIFICATE, then the first result must be a certificate. Other results may be alternate certificates, or infeasible points.

[source](#)

[MathOptInterface.ObjectiveValue](#) – Type.

```
| ObjectiveValue(result_index::Int = 1)
```

A model attribute for the objective value of the primal solution `result_index`.

If the solver does not have a primal value for the objective because the `result_index` is beyond the available solutions (whose number is indicated by the [ResultCount](#) attribute), getting this attribute must throw a [ResultIndexBoundsError](#). Otherwise, if the result is unavailable for another reason (for instance, only a dual solution is available), the result is undefined. Users should first check [PrimalStatus](#) before accessing the `ObjectiveValue` attribute.

See [ResultCount](#) for information on how the results are ordered.

[source](#)

[MathOptInterface.DualObjectiveValue](#) – Type.

```
| DualObjectiveValue(result_index::Int = 1)
```

A model attribute for the value of the objective function of the dual problem for the `result_index`th dual result.

If the solver does not have a dual value for the objective because the `result_index` is beyond the available solutions (whose number is indicated by the [ResultCount](#) attribute), getting this attribute must throw a [ResultIndexBoundsError](#). Otherwise, if the result is unavailable for another reason (for instance, only a primal solution is available), the result is undefined. Users should first check [DualStatus](#) before accessing the `DualObjectiveValue` attribute.

See [ResultCount](#) for information on how the results are ordered.

[source](#)

[MathOptInterface.ObjectiveBound](#) – Type.

```
| ObjectiveBound()
```

A model attribute for the best known bound on the optimal objective value.

[source](#)

[MathOptInterface.RelativeGap](#) – Type.

```
| RelativeGap()
```

A model attribute for the final relative optimality gap.

Warning

The definition of this gap is solver-dependent. However, most solvers implementing this attribute define the relative gap as some variation of $\frac{|b-f|}{|f|}$, where b is the best bound and f is the best feasible objective value.

[source](#)

[MathOptInterface.SolveTimeSec](#) - Type.

```
| SolveTimeSec()
```

A model attribute for the total elapsed solution time (in seconds) as reported by the optimizer.

[source](#)

[MathOptInterface.SimplexIterations](#) - Type.

```
| SimplexIterations()
```

A model attribute for the cumulative number of simplex iterations during the optimization process. In particular, for a mixed-integer program (MIP), the total simplex iterations for all nodes.

[source](#)

[MathOptInterface.BarrierIterations](#) - Type.

```
| BarrierIterations()
```

A model attribute for the cumulative number of barrier iterations while solving a problem.

[source](#)

[MathOptInterface.NodeCount](#) - Type.

```
| NodeCount()
```

A model attribute for the total number of branch-and-bound nodes explored while solving a mixed-integer program (MIP).

[source](#)

Conflict Status

[MathOptInterface.compute_conflict!](#) - Function.

```
| compute_conflict!(optimizer::AbstractOptimizer)
```

Computes a minimal subset of constraints such that the model with the other constraint removed is still infeasible.

Some solvers call a set of conflicting constraints an Irreducible Inconsistent Subsystem (IIS).

See also [ConflictStatus](#) and [ConstraintConflictStatus](#).

Note

If the model is modified after a call to `compute_conflict!`, the implementor is not obliged to purge the conflict. Any calls to the above attributes may return values for the original conflict without a warning. Similarly, when modifying the model, the conflict can be discarded.

[source](#)

[MathOptInterface.ConflictStatus](#) - Type.

```
| ConflictStatus()
```

A model attribute for the `ConflictStatusCode` explaining why the conflict refiner stopped when computing the conflict.

[source](#)

`MathOptInterface.ConflictStatusCode` – Type.

```
| ConflictStatusCode
```

An Enum of possible values for the `ConflictStatus` attribute. This attribute is meant to explain the reason why the conflict finder stopped executing in the most recent call to `compute_conflict!`.

Possible values are:

- `COMPUTE_CONFLICT_NOT_CALLED`: the function `compute_conflict!` has not yet been called
- `NO_CONFLICT_EXISTS`: there is no conflict because the problem is feasible
- `NO_CONFLICT_FOUND`: the solver could not find a conflict
- `CONFLICT_FOUND`: at least one conflict could be found

[source](#)

`MathOptInterface.ConstraintConflictStatus` – Type.

```
| ConstraintConflictStatus()
```

A constraint attribute indicating whether the constraint participates in the conflict. Its type is `ConflictParticipationStatusCode`.

[source](#)

`MathOptInterface.ConflictParticipationStatusCode` – Type.

```
| ConflictParticipationStatusCode
```

An Enum of possible values for the `ConstraintConflictStatus` attribute. This attribute is meant to indicate whether a given constraint participates or not in the last computed conflict.

Possible values are:

- `NOT_IN_CONFLICT`: the constraint does not participate in the conflict
- `IN_CONFLICT`: the constraint participates in the conflict
- `MAYBE_IN_CONFLICT`: the constraint may participate in the conflict, the solver was not able to prove that the constraint can be excluded from the conflict

[source](#)

Chapter 20

Variables

20.1 Functions

[MathOptInterface.add_variable](#) – Function.

```
| add_variable(model::ModelLike)::VariableIndex
```

Add a scalar variable to the model, returning a variable index.

A [AddVariableNotAllowed](#) error is thrown if adding variables cannot be done in the current state of the model `model`.

[source](#)

[MathOptInterface.add_variables](#) – Function.

```
| add_variables(model::ModelLike, n::Int)::Vector{VariableIndex}
```

Add `n` scalar variables to the model, returning a vector of variable indices.

A [AddVariableNotAllowed](#) error is thrown if adding variables cannot be done in the current state of the model `model`.

[source](#)

[MathOptInterface.add_constrained_variable](#) – Function.

```
| add_constrained_variable(  
|     model::ModelLike,  
|     set::AbstractScalarSet  
| )::Tuple{MOI.VariableIndex,  
|         MOI.ConstraintIndex{MOI.VariableIndex, typeof(set)}}
```

Add to model a scalar variable constrained to belong to `set`, returning the index of the variable created and the index of the constraint constraining the variable to belong to `set`.

By default, this function falls back to creating a free variable with [add_variable](#) and then constraining it to belong to `set` with [add_constraint](#).

[source](#)

[MathOptInterface.add_constrained_variables](#) – Function.

```

add_constrained_variables(
    model::ModelLike,
    sets::AbstractVector{<:AbstractScalarSet}
)::Tuple{
    Vector{MOI.VariableIndex},
    Vector{MOI.ConstraintIndex{MOI.VariableIndex,eltype(sets)}},
}

```

Add to model scalar variables constrained to belong to sets, returning the indices of the variables created and the indices of the constraints constraining the variables to belong to each set in sets. That is, if it returns variables and constraints, constraints[i] is the index of the constraint constraining variable[i] to belong to sets[i].

By default, this function falls back to calling `add_constrained_variable` on each set.

source

```

add_constrained_variables(
    model::ModelLike,
    set::AbstractVectorSet,
)::Tuple{
    Vector{MOI.VariableIndex},
    MOI.ConstraintIndex{MOI.VectorOfVariables,typeof(set)},
}

```

Add to model a vector of variables constrained to belong to set, returning the indices of the variables created and the index of the constraint constraining the vector of variables to belong to set.

By default, this function falls back to creating free variables with `add_variables` and then constraining it to belong to set with `add_constraint`.

source

`MathOptInterface.supports_add_constrained_variable` – Function.

```

supports_add_constrained_variable(
    model::ModelLike,
    S::Type{<:AbstractScalarSet}
)::Bool

```

Return a Bool indicating whether model supports constraining a variable to belong to a set of type S either on creation of the variable with `add_constrained_variable` or after the variable is created with `add_constraint`.

By default, this function falls back to `supports_add_constrained_variables(model, Reals) && supports_constraint(model, MOI.VariableIndex, S)` which is the correct definition for most models.

Example

Suppose that a solver supports only two kind of variables: binary variables and continuous variables with a lower bound. If the solver decides not to support `VariableIndex-in-Binary` and `VariableIndex-in-GreaterThan` constraints, it only has to implement `add_constrained_variable` for these two sets which prevents the user to add both a binary constraint and a lower bound on the same variable. Moreover, if the user adds a `VariableIndex-in-GreaterThan` constraint, implementing this interface (i.e., `supports_add_constrained_variable`) enables the constraint to be transparently bridged into a supported constraint.

source

`MathOptInterface.supports_add_constrained_variables` – Function.

```
supports_add_constrained_variables(
    model::ModelLike,
    S::Type{<:AbstractVectorSet}
)::Bool
```

Return a Bool indicating whether model supports constraining a vector of variables to belong to a set of type S either on creation of the vector of variables with `add_constrained_variables` or after the variable is created with `add_constraint`.

By default, if S is Reals then this function returns true and otherwise, it falls back to `supports_add_constrained_variables(Reals)` && `supports_constraint(model, MOI.VectorOfVariables, S)` which is the correct definition for most models.

Example

In the standard conic form (see [Duality](#)), the variables are grouped into several cones and the constraints are affine equality constraints. If Reals is not one of the cones supported by the solvers then it needs to implement `supports_add_constrained_variables(::Optimizer, ::Type{Reals}) = false` as free variables are not supported. The solvers should then implement `supports_add_constrained_variables(::Optimizer, ::Type{<:SupportedCones}) = true` where SupportedCones is the union of all cone types that are supported; it does not have to implement the method `supports_constraint(::Type{VectorOfVariables}, Type{<:SupportedCones})` as it should return false and it's the default. This prevents the user to constrain the same variable in two different cones. When a VectorOfVariables-in-S is added, the variables of the vector have already been created so they already belong to given cones. If bridges are enabled, the constraint will therefore be bridged by adding slack variables in S and equality constraints ensuring that the slack variables are equal to the corresponding variables of the given constraint function.

Note that there may also be sets for which `!supports_add_constrained_variables(model, S)` and `supports_constraint(model, MOI.VectorOfVariables, S)`. For instance, suppose a solver supports positive semidefinite variable constraints and two types of variables: binary variables and nonnegative variables. Then the solver should support adding VectorOfVariables-in-PositiveSemidefiniteConeTriangle constraints, but it should not support creating variables constrained to belong to the PositiveSemidefiniteConeTriangle because the variables in PositiveSemidefiniteConeTriangle should first be created as either binary or non-negative.

[source](#)

`MathOptInterface.is_valid` – Method.

```
is_valid(model::ModelLike, index::Index)::Bool
```

Return a Bool indicating whether this index refers to a valid object in the model model.

[source](#)

`MathOptInterface.delete` – Method.

```
delete(model::ModelLike, index::Index)
```

Delete the referenced object from the model. Throw `DeleteNotAllowed` if index cannot be deleted.

The following modifications also take effect if Index is `VariableIndex`:

- If index used in the objective function, it is removed from the function, i.e., it is substituted for zero.
- For each func-in-set constraint of the model:
 - If func isa VariableIndex and func == index then the constraint is deleted.

- If `func` is a `VectorOfVariables` and `index` in `func.variables` then
 - * if `length(func.variables) == 1` is one, the constraint is deleted;
 - * if `length(func.variables) > 1` and `supports_dimension_update(set)` then then the variable is removed from `func` and `set` is replaced by `update_dimension(set, MOI.dimension(set) - 1)`.
 - * Otherwise, a `DeleteNotAllowed` error is thrown.
- Otherwise, the variable is removed from `func`, i.e., it is substituted for zero.

source

`MathOptInterface.delete` – Method.

```
| delete(model::ModelLike, indices::Vector{R<:Index}) where {R}
```

Delete the referenced objects in the vector `indices` from the model. It may be assumed that `R` is a concrete type. The default fallback sequentially deletes the individual items in `indices`, although specialized implementations may be more efficient.

source

20.2 Attributes

`MathOptInterface.AbstractVariableAttribute` – Type.

```
| AbstractVariableAttribute
```

Abstract supertype for attribute objects that can be used to set or get attributes (properties) of variables in the model.

source

`MathOptInterface.VariableName` – Type.

```
| VariableName()
```

A variable attribute for a string identifying the variable. It is valid for two variables to have the same name; however, variables with duplicate names cannot be looked up using `get`. It has a default value of "" if not set'.

source

`MathOptInterface.VariablePrimalStart` – Type.

```
| VariablePrimalStart()
```

A variable attribute for the initial assignment to some primal variable's value that the optimizer may use to warm-start the solve. May be a number or nothing (unset).

source

`MathOptInterface.VariablePrimal` – Type.

```
| VariablePrimal(result_index::Int = 1)
```


A variable attribute for the assignment to some primal variable's value in result `result_index`. If `result_index` is omitted, it is 1 by default.

If the solver does not have a primal value for the variable because the `result_index` is beyond the available solutions (whose number is indicated by the [ResultCount](#) attribute), getting this attribute must throw a [ResultIndexBoundsError](#). Otherwise, if the result is unavailable for another reason (for instance, only a dual solution is available), the result is undefined. Users should first check [PrimalStatus](#) before accessing the `VariablePrimal` attribute.

See [ResultCount](#) for information on how the results are ordered.

[source](#)

[MathOptInterface.VariableBasisStatus](#) - Type.

```
| VariableBasisStatus(result_index::Int = 1)
```

A variable attribute for the `BasisStatusCode` of a variable in result `result_index`, with respect to an available optimal solution basis.

If the solver does not have a basis statue for the variable because the `result_index` is beyond the available solutions (whose number is indicated by the [ResultCount](#) attribute), getting this attribute must throw a [ResultIndexBoundsError](#). Otherwise, if the result is unavailable for another reason (for instance, only a dual solution is available), the result is undefined. Users should first check [PrimalStatus](#) before accessing the `VariableBasisStatus` attribute.

See [ResultCount](#) for information on how the results are ordered.

[source](#)

Chapter 21

Constraints

21.1 Types

[MathOptInterface.ConstraintIndex](#) – Type.

```
| ConstraintIndex{F, S}
```

A type-safe wrapper for Int64 for use in referencing F-in-S constraints in a model. The parameter F is the type of the function in the constraint, and the parameter S is the type of set in the constraint. To allow for deletion, indices need not be consecutive. Indices within a constraint type (i.e. F-in-S) must be unique, but non-unique indices across different constraint types are allowed. If F is [VariableIndex](#) then the index is equal to the index of the variable. That is for an `index::ConstraintIndex{VariableIndex}`, we always have

```
| index.value == MOI.get(model, MOI.ConstraintFunction(), index).value
```

[source](#)

21.2 Functions

[MathOptInterface.is_valid](#) – Method.

```
| is_valid(model::ModelLike, index::Index)::Bool
```

Return a Bool indicating whether this index refers to a valid object in the model model.

[source](#)

[MathOptInterface.add_constraint](#) – Function.

```
| add_constraint(model::ModelLike, func::F, set::S)::ConstraintIndex{F,S} where {F,S}
```

Add the constraint $f(x) \in \mathcal{S}$ where f is defined by func, and \mathcal{S} is defined by set.

```
| add_constraint(model::ModelLike, v::VariableIndex, set::S)::ConstraintIndex{VariableIndex,S}
  | where {S}
| add_constraint(model::ModelLike, vec::Vector{VariableIndex}, set::S)::ConstraintIndex{
  | VectorOfVariables,S} where {S}
```

Add the constraint $v \in \mathcal{S}$ where v is the variable (or vector of variables) referenced by v and \mathcal{S} is defined by set.

- An `UnsupportedConstraint` error is thrown if model does not support F-in-S constraints,
- a `AddConstraintNotAllowed` error is thrown if it supports F-in-S constraints but it cannot add the constraint(s) in its current state and
- a `ScalarFunctionConstantNotZero` error may be thrown if func is an `AbstractScalarFunction` with nonzero constant and set is `EqualTo`, `GreaterThan`, `LessThan` or `Interval`.
- a `LowerBoundAlreadySet` error is thrown if F is a `VariableIndex` and a constraint was already added to this variable that sets a lower bound.
- a `UpperBoundAlreadySet` error is thrown if F is a `VariableIndex` and a constraint was already added to this variable that sets an upper bound.

source

`MathOptInterface.add_constraints` – Function.

```
add_constraints(model::ModelLike, funcs::Vector{F},
  ↪ sets::Vector{S})::Vector{ConstraintIndex{F,S}} where {F,S}
```

Add the set of constraints specified by each function-set pair in `funcs` and `sets`. `F` and `S` should be concrete types. This call is equivalent to `add_constraint.(model, funcs, sets)` but may be more efficient.

source

`MathOptInterface.transform` – Function.

Transform Constraint Set

```
transform(model::ModelLike, c::ConstraintIndex{F,S1}, newset::S2)::ConstraintIndex{F,S2}
```

Replace the set in constraint `c` with `newset`. The constraint index `c` will no longer be valid, and the function returns a new constraint index with the correct type.

Solvers may only support a subset of constraint transforms that they perform efficiently (for example, changing from a `LessThan` to `GreaterThan` set). In addition, set modification (where `S1 = S2`) should be performed via the `modify` function.

Typically, the user should delete the constraint and add a new one.

Examples

If `c` is a `ConstraintIndex{ScalarAffineFunction{Float64},LessThan{Float64}}`,

```
c2 = transform(model, c, GreaterThan(0.0))
transform(model, c, LessThan(0.0)) # errors
```

source

`MathOptInterface.supports_constraint` – Function.

```
MOI.supports_constraint(
  BT::Type{<:AbstractBridge},
  F::Type{<:MOI.AbstractFunction},
  S::Type{<:MOI.AbstractSet},
)::Bool
```

Return a `Bool` indicating whether the bridges of type `BT` support bridging F-in-S constraints.

source

```

| supports_constraint(
|   model::ModelLike,
|   ::Type{F},
|   ::Type{S},
| )::Bool where {F<:AbstractFunction,S<:AbstractSet}

```

Return a Bool indicating whether model supports F-in-S constraints, that is, `copy_to(model, src)` does not throw `UnsupportedConstraint` when `src` contains F-in-S constraints. If F-in-S constraints are only not supported in specific circumstances, e.g. F-in-S constraints cannot be combined with another type of constraint, it should still return true.

source

21.3 Attributes

`MathOptInterface.AbstractConstraintAttribute` – Type.

```

| AbstractConstraintAttribute

```

Abstract supertype for attribute objects that can be used to set or get attributes (properties) of constraints in the model.

source

`MathOptInterface.ConstraintName` – Type.

```

| ConstraintName()

```

A constraint attribute for a string identifying the constraint.

It is valid for constraints variables to have the same name; however, constraints with duplicate names cannot be looked up using `get`, regardless of whether they have the same F-in-S type.

`ConstraintName` has a default value of "" if not set.

Notes

You should not implement `ConstraintName` for `VariableIndex` constraints.

source

`MathOptInterface.ConstraintPrimalStart` – Type.

```

| ConstraintPrimalStart()

```

A constraint attribute for the initial assignment to some constraint's `ConstraintPrimal` that the optimizer may use to warm-start the solve.

May be nothing (unset), a number for `AbstractScalarFunction`, or a vector for `AbstractVectorFunction`.

source

`MathOptInterface.ConstraintDualStart` – Type.

```

| ConstraintDualStart()

```

A constraint attribute for the initial assignment to some constraint's `ConstraintDual` that the optimizer may use to warm-start the solve.

May be nothing (unset), a number for `AbstractScalarFunction`, or a vector for `AbstractVectorFunction`.

source

[MathOptInterface.ConstraintPrimal](#) - Type.

```
| ConstraintPrimal(result_index::Int = 1)
```

A constraint attribute for the assignment to some constraint's primal value(s) in result `result_index`.

If the constraint is $f(x)$ in S , then in most cases the `ConstraintPrimal` is the value of f , evaluated at the corresponding [VariablePrimal](#) solution.

However, some conic solvers reformulate $b - Ax$ in S to $s = b - Ax$, s in S . These solvers may return the value of s for `ConstraintPrimal`, rather than $b - Ax$. (Although these are constrained by an equality constraint, due to numerical tolerances they may not be identical.)

If the solver does not have a primal value for the constraint because the `result_index` is beyond the available solutions (whose number is indicated by the [ResultCount](#) attribute), getting this attribute must throw a [ResultIndexBoundsError](#). Otherwise, if the result is unavailable for another reason (for instance, only a dual solution is available), the result is undefined. Users should first check [PrimalStatus](#) before accessing the `ConstraintPrimal` attribute.

If `result_index` is omitted, it is 1 by default. See [ResultCount](#) for information on how the results are ordered.

[source](#)

[MathOptInterface.ConstraintDual](#) - Type.

```
| ConstraintDual(result_index::Int = 1)
```

A constraint attribute for the assignment to some constraint's dual value(s) in result `result_index`. If `result_index` is omitted, it is 1 by default.

If the solver does not have a dual value for the variable because the `result_index` is beyond the available solutions (whose number is indicated by the [ResultCount](#) attribute), getting this attribute must throw a [ResultIndexBoundsError](#). Otherwise, if the result is unavailable for another reason (for instance, only a primal solution is available), the result is undefined. Users should first check [DualStatus](#) before accessing the `ConstraintDual` attribute.

See [ResultCount](#) for information on how the results are ordered.

[source](#)

[MathOptInterface.ConstraintBasisStatus](#) - Type.

```
| ConstraintBasisStatus(result_index::Int = 1)
```

A constraint attribute for the `BasisStatusCode` of some constraint in result `result_index`, with respect to an available optimal solution basis. If `result_index` is omitted, it is 1 by default.

If the solver does not have a basis statue for the constraint because the `result_index` is beyond the available solutions (whose number is indicated by the [ResultCount](#) attribute), getting this attribute must throw a [ResultIndexBoundsError](#). Otherwise, if the result is unavailable for another reason (for instance, only a dual solution is available), the result is undefined. Users should first check [PrimalStatus](#) before accessing the `ConstraintBasisStatus` attribute.

See [ResultCount](#) for information on how the results are ordered.

Notes

For the basis status of a variable, query [VariableBasisStatus](#).

`ConstraintBasisStatus` does not apply to `VariableIndex` constraints. You can infer the basis status of a `VariableIndex` constraint by looking at the result of `VariableBasisStatus`.

[source](#)

`MathOptInterface.BasisStatusCode` – Type.

| `BasisStatusCode`

An Enum of possible values for the `ConstraintBasisStatus` and `VariableBasisStatus` attributes, explaining the status of a given element with respect to an optimal solution basis.

Possible values are:

- `BASIC`: element is in the basis
- `NONBASIC`: element is not in the basis
- `NONBASIC_AT_LOWER`: element is not in the basis and is at its lower bound
- `NONBASIC_AT_UPPER`: element is not in the basis and is at its upper bound
- `SUPER_BASIC`: element is not in the basis but is also not at one of its bounds

Notes

- `NONBASIC_AT_LOWER` and `NONBASIC_AT_UPPER` should be used only for constraints with the `Interval` set. In this case, they are necessary to distinguish which side of the constraint is active. One-sided constraints (e.g., `LessThan` and `GreaterThan`) should use `NONBASIC` instead of the `NONBASIC_AT_*` values. This restriction does not apply to `VariableBasisStatus`, which should return `NONBASIC_AT_*` regardless of whether the alternative bound exists.
- In linear programs, `SUPER_BASIC` occurs when a variable with no bounds is not in the basis.

[source](#)

`MathOptInterface.ConstraintFunction` – Type.

| `ConstraintFunction()`

A constraint attribute for the `AbstractFunction` object used to define the constraint. It is guaranteed to be equivalent but not necessarily identical to the function provided by the user.

[source](#)

`MathOptInterface.CanonicalConstraintFunction` – Type.

| `CanonicalConstraintFunction()`

A constraint attribute for a canonical representation of the `AbstractFunction` object used to define the constraint. Getting this attribute is guaranteed to return a function that is equivalent but not necessarily identical to the function provided by the user.

By default, `MOI.get(model, MOI.CanonicalConstraintFunction(), ci)` fallbacks to `MOI.Utilities.canonical(MOI.get(model, MOI.ConstraintFunction(), ci))`. However, if `model` knows that the constraint function is canonical then it can implement a specialized method that directly return the function without calling `Utilities.canonical`. Therefore, the value returned **cannot** be assumed to be a copy of the function stored in `model`. Moreover, `Utilities.Model` checks with `Utilities.is_canonical` whether the function stored internally is already canonical and if it's the case, then it returns the function stored internally instead of a copy.

[source](#)

`MathOptInterface.ConstraintSet` - Type.

| `ConstraintSet()`

A constraint attribute for the `AbstractSet` object used to define the constraint.

[source](#)

Chapter 22

Modifications

`MathOptInterface.modify` – Function.

Constraint Function

```
modify(model::ModelLike, ci::ConstraintIndex, change::AbstractFunctionModification)
```

Apply the modification specified by `change` to the function of constraint `ci`.

An `ModifyConstraintNotAllowed` error is thrown if modifying constraints is not supported by the model.

Examples

```
modify(model, ci, ScalarConstantChange(10.0))
```

Objective Function

```
modify(model::ModelLike, ::ObjectiveFunction, change::AbstractFunctionModification)
```

Apply the modification specified by `change` to the objective function of `model`. To change the function completely, call `set` instead.

An `ModifyObjectiveNotAllowed` error is thrown if modifying objectives is not supported by the model.

Examples

```
modify(model, ObjectiveFunction{ScalarAffineFunction{Float64}}(), ScalarConstantChange(10.0))
```

Multiple modifications in Constraint Functions

```
modify(
  model::ModelLike,
  cis::AbstractVector{<:ConstraintIndex},
  changes::AbstractVector{<:AbstractFunctionModification},
)
```

Apply multiple modifications specified by `changes` to the functions of constraints `cis`.

A `ModifyConstraintNotAllowed` error is thrown if modifying constraints is not supported by `model`.

Examples


```

modify(
  model,
  [ci, ci],
  [
    ScalarCoefficientChange{Float64}(VariableIndex(1), 1.0),
    ScalarCoefficientChange{Float64}(VariableIndex(2), 0.5),
  ],
)

```

Multiple modifications in the Objective Function

```

modify(
  model::ModelLike,
  attr::ObjectiveFunction,
  changes::AbstractVector{<:AbstractFunctionModification},
)

```

Apply multiple modifications specified by `changes` to the functions of constraints `cis`.

A `ModifyObjectiveNotAllowed` error is thrown if modifying objective coefficients is not supported by `model`.

Examples

```

modify(
  model,
  ObjectiveFunction{ScalarAffineFunction{Float64}}(),
  [
    ScalarCoefficientChange{Float64}(VariableIndex(1), 1.0),
    ScalarCoefficientChange{Float64}(VariableIndex(2), 0.5),
  ],
)

```

[source](#)

`MathOptInterface.AbstractFunctionModification` – Type.

```

| AbstractFunctionModification

```

An abstract supertype for structs which specify partial modifications to functions, to be used for making small modifications instead of replacing the functions entirely.

[source](#)

`MathOptInterface.ScalarConstantChange` – Type.

```

| ScalarConstantChange{T}(new_constant::T)

```

A struct used to request a change in the constant term of a scalar-valued function. Applicable to `ScalarAffineFunction` and `ScalarQuadraticFunction`.

[source](#)

`MathOptInterface.VectorConstantChange` – Type.

```

| VectorConstantChange{T}(new_constant::Vector{T})

```

A struct used to request a change in the constant vector of a vector-valued function. Applicable to `VectorAffineFunction` and `VectorQuadraticFunction`.

[source](#)

`MathOptInterface.ScalarCoefficientChange` - Type.

```
| ScalarCoefficientChange{T}(variable::VariableIndex, new_coefficient::T)
```

A struct used to request a change in the linear coefficient of a single variable in a scalar-valued function. Applicable to `ScalarAffineFunction` and `ScalarQuadraticFunction`.

[source](#)

`MathOptInterface.MultirowChange` - Type.

```
| MultirowChange{T}(variable::VariableIndex, new_coefficients::Vector{Tuple{Int64, T}})
```

A struct used to request a change in the linear coefficients of a single variable in a vector-valued function. New coefficients are specified by `(output_index, coefficient)` tuples. Applicable to `VectorAffineFunction` and `VectorQuadraticFunction`.

[source](#)

Chapter 23

Nonlinear programming

23.1 Types

[MathOptInterface.AbstractNLPEvaluator](#) – Type.

```
| AbstractNLPEvaluator
```

Abstract supertype for the callback object that is used to query function values, derivatives, and expression graphs. It is used in `NLPBlock`.

[source](#)

[MathOptInterface.NLPBoundsPair](#) – Type.

```
| NLPBoundsPair(lower, upper)
```

A struct holding a pair of lower and upper bounds. `-Inf` and `Inf` can be used to indicate no lower or upper bound, respectively.

[source](#)

[MathOptInterface.NLPBlockData](#) – Type.

```
| struct NLPBlockData
|   constraint_bounds::Vector{NLPBoundsPair}
|   evaluator::AbstractNLPEvaluator
|   has_objective::Bool
| end
```

A struct encoding a set of nonlinear constraints of the form $lb \leq g(x) \leq ub$ and, if `has_objective == true`, a nonlinear objective function $f(x)$. `constraint_bounds` holds the pairs of `lb` and `ub` elements. Nonlinear objectives override any objective set by using the `ObjectiveFunction` attribute. The evaluator is a callback object that is used to query function values, derivatives, and expression graphs. If `has_objective == false`, then it is an error to query properties of the objective function, and in Hessian-of-the-Lagrangian queries, `σ` must be set to zero.

Note

Throughout the evaluator, all variables are ordered according to [ListOfVariableIndices](#). Hence, MOI copies of nonlinear problems should be done with attention.

[source](#)

23.2 Attributes

`MathOptInterface.NLPBlock` – Type.

```
| NLPBlock()
```

Holds the `NLPBlockData` that represents a set of nonlinear constraints, and optionally a nonlinear objective.

[source](#)

`MathOptInterface.NLPBlockDual` – Type.

```
| NLPBlockDual(result_index::Int)
| NLPBlockDual()
```

The Lagrange multipliers on the constraints from the `NLPBlock` in result `result_index`. If `result_index` is omitted, it is 1 by default.

[source](#)

`MathOptInterface.NLPBlockDualStart` – Type.

```
| NLPBlockDualStart()
```

An initial assignment of the Lagrange multipliers on the constraints from the `NLPBlock` that the solver may use to warm-start the solve.

[source](#)

23.3 Functions

`MathOptInterface.initialize` – Function.

```
| initialize(d::AbstractNLPEvaluator, requested_features::Vector{Symbol})
```

Must be called before any other methods. The vector `requested_features` lists features requested by the solver. These may include `:Grad` for gradients of the objective, `f`, `:Jac` for explicit Jacobians of constraints, `g`, `:JacVec` for Jacobian-vector products, `:HessVec` for Hessian-vector and Hessian-of-Lagrangian-vector products, `:Hess` for explicit Hessians and Hessian-of-Lagrangians, and `:ExprGraph` for expression graphs.

[source](#)

`MathOptInterface.features_available` – Function.

```
| features_available(d::AbstractNLPEvaluator)
```

Returns the subset of features available for this problem instance, as a vector of symbols in the same format as in `initialize`.

[source](#)

`MathOptInterface.eval_objective` – Function.

```
| eval_objective(d::AbstractNLPEvaluator, x)
```

Evaluate the objective $f(x)$, returning a scalar value.

[source](#)

`MathOptInterface.eval_constraint` – Function.

```
| eval_constraint(d::AbstractNLP evaluator, g, x)
```

Evaluate the constraint function $g(x)$, storing the result in the vector g which must be of the appropriate size.

[source](#)

`MathOptInterface.eval_objective_gradient` – Function.

```
| eval_objective_gradient(d::AbstractNLP evaluator, df, x)
```

Evaluate $\nabla f(x)$ as a dense vector, storing the result in the vector df which must be of the appropriate size.

[source](#)

`MathOptInterface.jacobian_structure` – Function.

```
| jacobian_structure(d::AbstractNLP evaluator)::Vector{Tuple{Int64,Int64}}
```

Returns the sparsity structure of the Jacobian matrix $J_g(x) = \begin{bmatrix} \nabla g_1(x) \\ \nabla g_2(x) \\ \vdots \\ \nabla g_m(x) \end{bmatrix}$ where g_i is the i th component of g . The sparsity structure is assumed to be independent of the point x . Returns a vector of tuples, (row, column), where each indicates the position of a structurally nonzero element. These indices are not required to be sorted and can contain duplicates, in which case the solver should combine the corresponding elements by adding them together.

[source](#)

`MathOptInterface.hessian_lagrangian_structure` – Function.

```
| hessian_lagrangian_structure(d::AbstractNLP evaluator)::Vector{Tuple{Int64,Int64}}
```

Returns the sparsity structure of the Hessian-of-the-Lagrangian matrix $\nabla^2 f + \sum_{i=1}^m \nabla^2 g_i$ as a vector of tuples, where each indicates the position of a structurally nonzero element. These indices are not required to be sorted and can contain duplicates, in which case the solver should combine the corresponding elements by adding them together. Any mix of lower and upper-triangular indices is valid. Elements (i, j) and (j, i) , if both present, should be treated as duplicates.

[source](#)

`MathOptInterface.eval_constraint_jacobian` – Function.

```
| eval_constraint_jacobian(d::AbstractNLP evaluator, J, x)
```

Evaluates the sparse Jacobian matrix $J_g(x) = \begin{bmatrix} \nabla g_1(x) \\ \nabla g_2(x) \\ \vdots \\ \nabla g_m(x) \end{bmatrix}$. The result is stored in the vector J in the same order as the indices returned by `jacobian_structure`.

[source](#)

`MathOptInterface.eval_constraint_jacobian_product` - Function.

```
| eval_constraint_jacobian_product(d::AbstractNLPEvaluator, y, x, w)
```

Computes the Jacobian-vector product $J_g(x)w$, storing the result in the vector y .

[source](#)

`MathOptInterface.eval_constraint_jacobian_transpose_product` - Function.

```
| eval_constraint_jacobian_transpose_product(d::AbstractNLPEvaluator, y, x, w)
```

Computes the Jacobian-transpose-vector product $J_g(x)^T w$, storing the result in the vector y .

[source](#)

`MathOptInterface.eval_hessian_lagrangian` - Function.

```
| eval_hessian_lagrangian(d::AbstractNLPEvaluator, H, x, σ, μ)
```

Given scalar weight σ and vector of constraint weights μ , computes the sparse Hessian-of-the-Lagrangian matrix $\sigma \nabla^2 f(x) + \sum_{i=1}^m \mu_i \nabla^2 g_i(x)$, storing the result in the vector H in the same order as the indices returned by `hessian_lagrangian_structure`.

[source](#)

`MathOptInterface.eval_hessian_lagrangian_product` - Function.

```
| eval_hessian_lagrangian_product(d::AbstractNLPEvaluator, h, x, v, σ, μ)
```

Given scalar weight σ and vector of constraint weights μ , computes the Hessian-of-the-Lagrangian-vector product $(\sigma \nabla^2 f(x) + \sum_{i=1}^m \mu_i \nabla^2 g_i(x)) v$, storing the result in the vector h .

[source](#)

`MathOptInterface.objective_expr` - Function.

```
| objective_expr(d::AbstractNLPEvaluator)
```

Returns an expression graph for the objective function as a standard Julia Expr object. All sums and products are flattened out as simple `Expr{:+, ...}` and `Expr{:*, ...}` objects. The symbol `x` is used as a placeholder for the vector of decision variables. No other undefined symbols are permitted; coefficients are embedded as explicit values. For example, the expression $x_1 + \sin(x_2 / \exp(x_3))$ would be represented as the Julia object `:(x[1] + sin(x[2]/exp(x[3])))`. Each integer index is wrapped in a `VariableIndex`. See the [Julia manual](#) for more information on the structure of Expr objects. There are currently no restrictions on recognized functions; typically these will be built-in Julia functions like `^`, `exp`, `log`, `cos`, `tan`, `sqrt`, etc., but modeling interfaces may choose to extend these basic functions.

[source](#)

`MathOptInterface.constraint_expr` - Function.

```
| constraint_expr(d::AbstractNLPEvaluator, i)
```

Returns an expression graph for the i th constraint in the same format as described above, with an additional comparison operator indicating the sense of and bounds on the constraint. The right-hand side of the comparison must be a constant; that is, `:(x[1]^3 <= 1)` is allowed, while `:(1 <= x[1]^3)` is not valid. Double-sided constraints are allowed, in which case both the lower bound and upper bounds should be constants; for example, `:(-1 <= cos(x[1]) + sin(x[2]) <= 1)` is valid.

[source](#)

Chapter 24

Callbacks

[MathOptInterface.AbstractCallback](#) – Type.

```
| abstract type AbstractCallback <: AbstractModelAttribute end
```

Abstract type for a model attribute representing a callback function. The value set to subtypes of `AbstractCallback` is a function that may be called during `optimize!`. As `optimize!` is in progress, the result attributes (i.e, the attributes `attr` such that `is_set_by_optimize(attr)`) may not be accessible from the callback, hence trying to get result attributes might throw a `OptimizeInProgress` error.

At most one callback of each type can be registered. If an optimizer already has a function for a callback type, and the user registers a new function, then the old one is replaced.

The value of the attribute should be a function taking only one argument, commonly called `callback_data`, that can be used for instance in [LazyConstraintCallback](#), [HeuristicCallback](#) and [UserCutCallback](#).

[source](#)

[MathOptInterface.AbstractSubmittable](#) – Type.

```
| AbstractSubmittable
```

Abstract supertype for objects that can be submitted to the model.

[source](#)

[MathOptInterface.submit](#) – Function.

```
| submit(optimizer::AbstractOptimizer, sub::AbstractSubmittable,  
|         values...)::Nothing
```

Submit values to the submittable `sub` of the optimizer `optimizer`.

An `UnsupportedSubmittable` error is thrown if model does not support the attribute `attr` (see [supports](#)) and a `SubmitNotAllowed` error is thrown if it supports the submittable `sub` but it cannot be submitted.

[source](#)

24.1 Attributes

[MathOptInterface.CallbackNodeStatus](#) – Type.

```
| CallbackNodeStatus(callback_data)
```

An optimizer attribute describing the (in)feasibility of the primal solution available from `CallbackVariablePrimal` during a callback identified by `callback_data`.

Returns a `CallbackNodeStatusCode` Enum.

[source](#)

`MathOptInterface.CallbackNodeStatusCode` – Type.

```
| CallbackNodeStatusCode
```

An Enum of possible return values from calling `get` with `CallbackNodeStatus`.

Possible values are:

- `CALLBACK_NODE_STATUS_INTEGER`: the primal solution available from `CallbackVariablePrimal` is integer feasible.
- `CALLBACK_NODE_STATUS_FRACTIONAL`: the primal solution available from `CallbackVariablePrimal` is integer infeasible.
- `CALLBACK_NODE_STATUS_UNKNOWN`: the primal solution available from `CallbackVariablePrimal` might be integer feasible or infeasible.

[source](#)

`MathOptInterface.CallbackVariablePrimal` – Type.

```
| CallbackVariablePrimal(callback_data)
```

A variable attribute for the assignment to some primal variable's value during the callback identified by `callback_data`.

[source](#)

24.2 Lazy constraints

`MathOptInterface.LazyConstraintCallback` – Type.

```
| LazyConstraintCallback() <: AbstractCallback
```

The callback can be used to reduce the feasible set given the current primal solution by submitting a `LazyConstraint`. For instance, it may be called at an incumbent of a mixed-integer problem. Note that there is no guarantee that the callback is called at every feasible primal solution.

The current primal solution is accessed through `CallbackVariablePrimal`. Trying to access other result attributes will throw `OptimizeInProgress` as discussed in `AbstractCallback`.

Examples

```
x = MOI.add_variables(optimizer, 8)
MOI.set(optimizer, MOI.LazyConstraintCallback(), callback_data -> begin
    sol = MOI.get(optimizer, MOI.CallbackVariablePrimal(callback_data), x)
    if # should add a lazy constraint
        func = # computes function
        set = # computes set
        MOI.submit(optimizer, MOI.LazyConstraint(callback_data), func, set)
    end
end)
```


source

`MathOptInterface.LazyConstraint` - Type.

```
| LazyConstraint(callback_data)
```

Lazy constraint func-in-set submitted as func, set. The optimal solution returned by `VariablePrimal` will satisfy all lazy constraints that have been submitted.

This can be submitted only from the `LazyConstraintCallback`. The field `callback_data` is a solver-specific callback type that is passed as the argument to the feasible solution callback.

Examples

Suppose `x` and `y` are `VariableIndex`s of optimizer. To add a `LazyConstraint` for $2x + 3y \leq 1$, write

```
| func = 2.0x + 3.0y
| set = MOI.LessThan(1.0)
| MOI.submit(optimizer, MOI.LazyConstraint(callback_data), func, set)
```

inside a `LazyConstraintCallback` of data `callback_data`.

source

24.3 User cuts

`MathOptInterface.UserCutCallback` - Type.

```
| UserCutCallback() <: AbstractCallback
```

The callback can be used to submit `UserCut` given the current primal solution. For instance, it may be called at fractional (i.e., non-integer) nodes in the branch and bound tree of a mixed-integer problem. Note that there is not guarantee that the callback is called everytime the solver has an infeasible solution.

The infeasible solution is accessed through `CallbackVariablePrimal`. Trying to access other result attributes will throw `OptimizeInProgress` as discussed in `AbstractCallback`.

Examples

```
| x = MOI.add_variables(optimizer, 8)
| MOI.set(optimizer, MOI.UserCutCallback(), callback_data -> begin
|   sol = MOI.get(optimizer, MOI.CallbackVariablePrimal(callback_data), x)
|   if # can find a user cut
|     func = # computes function
|     set = # computes set
|     MOI.submit(optimizer, MOI.UserCut(callback_data), func, set)
|   end
| end
```

source

`MathOptInterface.UserCut` - Type.

```
| UserCut(callback_data)
```

Constraint func-to-set suggested to help the solver detect the solution given by `CallbackVariablePrimal` as infeasible. The cut is submitted as func, set. Typically `CallbackVariablePrimal` will violate integrality constraints, and a cut would be of the form `ScalarAffineFunction-in-LessThan` or `ScalarAffineFunction-in-GreaterThan`. Note that, as opposed to `LazyConstraint`, the provided constraint cannot modify the feasible set, the constraint should be redundant, e.g., it may be a consequence of affine and integrality constraints.

This can be submitted only from the `UserCutCallback`. The field `callback_data` is a solver-specific callback type that is passed as the argument to the infeasible solution callback.

Note that the solver may silently ignore the provided constraint.

[source](#)

24.4 Heuristic solutions

`MathOptInterface.HeuristicCallback` – Type.

```
| HeuristicCallback() <: AbstractCallback
```

The callback can be used to submit `HeuristicSolution` given the current primal solution. For instance, it may be called at fractional (i.e., non-integer) nodes in the branch and bound tree of a mixed-integer problem. Note that there is not guarantee that the callback is called everytime the solver has an infeasible solution.

The current primal solution is accessed through `CallbackVariablePrimal`. Trying to access other result attributes will throw `OptimizeInProgress` as discussed in `AbstractCallback`.

Examples

```
x = MOI.add_variables(optimizer, 8)
MOI.set(optimizer, MOI.HeuristicCallback(), callback_data -> begin
    sol = MOI.get(optimizer, MOI.CallbackVariablePrimal(callback_data), x)
    if # can find a heuristic solution
        values = # computes heuristic solution
        MOI.submit(optimizer, MOI.HeuristicSolution(callback_data), x,
                    values)
    end
end
```

[source](#)

`MathOptInterface.HeuristicSolutionStatus` – Type.

```
| HeuristicSolutionStatus
```

An Enum of possible return values for `submit` with `HeuristicSolution`. This informs whether the heuristic solution was accepted or rejected. Possible values are:

- `HEURISTIC_SOLUTION_ACCEPTED`: The heuristic solution was accepted.
- `HEURISTIC_SOLUTION_REJECTED`: The heuristic solution was rejected.
- `HEURISTIC_SOLUTION_UNKNOWN`: No information available on the acceptance.

[source](#)

`MathOptInterface.HeuristicSolution` – Type.

| `HeuristicSolution(callback_data)`

Heuristically obtained feasible solution. The solution is submitted as `variables`, `values` where `values[i]` gives the value of `variables[i]`, similarly to `set`. The `submit` call returns a `HeuristicSolutionStatus` indicating whether the provided solution was accepted or rejected.

This can be submitted only from the `HeuristicCallback`. The field `callback_data` is a solver-specific callback type that is passed as the argument to the heuristic callback.

Some solvers require a complete solution, others only partial solutions.

[source](#)

Chapter 25

Errors

When an MOI call fails on a model, precise errors should be thrown when possible instead of simply calling error with a message. The docstrings for the respective methods describe the errors that the implementation should throw in certain situations. This error-reporting system allows code to distinguish between internal errors (that should be shown to the user) and unsupported operations which may have automatic workarounds.

When an invalid index is used in an MOI call, an [InvalidIndex](#) is thrown:

[MathOptInterface.InvalidIndex](#) – Type.

```
struct InvalidIndex{IndexType<:Index} <: Exception
    index::IndexType
end
```

An error indicating that the index `index` is invalid.

[source](#)

When an invalid result index is used to retrieve an attribute, a [ResultIndexBoundsError](#) is thrown:

[MathOptInterface.ResultIndexBoundsError](#) – Type.

```
struct ResultIndexBoundsError{AttrType} <: Exception
    attr::AttrType
    result_count::Int
end
```

An error indicating that the requested attribute `attr` could not be retrieved, because the solver returned too few results compared to what was requested. For instance, the user tries to retrieve `VariablePrimal(2)` when only one solution is available, or when the model is infeasible and has no solution.

See also: [check_result_index_bounds](#).

[source](#)

[MathOptInterface.check_result_index_bounds](#) – Function.

```
check_result_index_bounds(model::ModelLike, attr)
```

This function checks whether enough results are available in the model for the requested `attr`, using its `result_index` field. If the model does not have sufficient results to answer the query, it throws a [ResultIndexBoundsError](#).

[source](#)

As discussed in [JuMP mapping](#), for scalar constraint with a nonzero function constant, a [ScalarFunctionConstantNotZero](#) exception may be thrown:

[MathOptInterface.ScalarFunctionConstantNotZero](#) – Type.

```
struct ScalarFunctionConstantNotZero{T, F, S} <: Exception
    constant::T
end
```

An error indicating that the constant part of the function in the constraint F-in-S is nonzero.

[source](#)

Some [VariableIndex](#) constraints cannot be combined on the same variable:

[MathOptInterface.LowerBoundAlreadySet](#) – Type.

```
| LowerBoundAlreadySet{S1, S2}
```

Error thrown when setting a [VariableIndex](#)-in-S2 when a [VariableIndex](#)-in-S1 has already been added and the sets S1, S2 both set a lower bound, i.e. they are [EqualTo](#), [GreaterThan](#), [Interval](#), [Semicontinuous](#) or [Semiinteger](#).

[source](#)

[MathOptInterface.UpperBoundAlreadySet](#) – Type.

```
| UpperBoundAlreadySet{S1, S2}
```

Error thrown when setting a [VariableIndex](#)-in-S2 when a [VariableIndex](#)-in-S1 has already been added and the sets S1, S2 both set an upper bound, i.e. they are [EqualTo](#), [LessThan](#), [Interval](#), [Semicontinuous](#) or [Semiinteger](#).

[source](#)

As discussed in [AbstractCallback](#), trying to [get](#) attributes inside a callback may throw:

[MathOptInterface.OptimizeInProgress](#) – Type.

```
struct OptimizeInProgress{AttrType<:AnyAttribute} <: Exception
    attr::AttrType
end
```

Error thrown from optimizer when `MOI.get(optimizer, attr)` is called inside an [AbstractCallback](#) while it is only defined once [optimize!](#) has completed. This can only happen when `is_set_by_optimize(attr)` is true.

[source](#)

Trying to submit the wrong type of [AbstractSubmittable](#) inside an [AbstractCallback](#) (e.g., a [UserCut](#) inside a [LazyConstraintCallback](#)) will throw:

[MathOptInterface.InvalidCallbackUsage](#) – Type.

```
struct InvalidCallbackUsage{C, S} <: Exception
    callback::C
    submittable::S
end
```

An error indicating that submittable cannot be submitted inside callback.

For example, `UserCut` cannot be submitted inside `LazyConstraintCallback`.

source

The rest of the errors defined in MOI fall in two categories represented by the following two abstract types:

`MathOptInterface.UnsupportedError` – Type.

```
| UnsupportedError <: Exception
```

Abstract type for error thrown when an element is not supported by the model.

source

`MathOptInterface.NotAllowedError` – Type.

```
| NotAllowedError <: Exception
```

Abstract type for error thrown when an operation is supported but cannot be applied in the current state of the model.

source

The different `UnsupportedError` and `NotAllowedError` are the following errors:

`MathOptInterface.UnsupportedAttribute` – Type.

```
| struct UnsupportedAttribute{AttrType} <: UnsupportedError
|     attr::AttrType
|     message::String
| end
```

An error indicating that the attribute `attr` is not supported by the model, i.e. that `supports` returns false.

source

`MathOptInterface.SetAttributeNotAllowed` – Type.

```
| struct SetAttributeNotAllowed{AttrType} <: NotAllowedError
|     attr::AttrType
|     message::String # Human-friendly explanation why the attribute cannot be set
| end
```

An error indicating that the attribute `attr` is supported (see `supports`) but cannot be set for some reason (see the error string).

source

`MathOptInterface.AddVariableNotAllowed` – Type.

```
| struct AddVariableNotAllowed <: NotAllowedError
|     message::String # Human-friendly explanation why the attribute cannot be set
| end
```

An error indicating that variables cannot be added to the model.

source

`MathOptInterface.UnsupportedConstraint` – Type.

```

struct UnsupportedConstraint{F<:AbstractFunction, S<:AbstractSet} <: UnsupportedError
  message::String # Human-friendly explanation why the attribute cannot be set
end

```

An error indicating that constraints of type F-in-S are not supported by the model, i.e. that `supports_constraint` returns false.

[source](#)

`MathOptInterface.AddConstraintNotAllowed` – Type.

```

struct AddConstraintNotAllowed{F<:AbstractFunction, S<:AbstractSet} <: NotAllowedError
  message::String # Human-friendly explanation why the attribute cannot be set
end

```

An error indicating that constraints of type F-in-S are supported (see `supports_constraint`) but cannot be added.

[source](#)

`MathOptInterface.ModifyConstraintNotAllowed` – Type.

```

struct ModifyConstraintNotAllowed{F<:AbstractFunction, S<:AbstractSet,
                                C<:AbstractFunctionModification} <: NotAllowedError
  constraint_index::ConstraintIndex{F, S}
  change::C
  message::String
end

```

An error indicating that the constraint modification change cannot be applied to the constraint of index ci.

[source](#)

`MathOptInterface.ModifyObjectiveNotAllowed` – Type.

```

struct ModifyObjectiveNotAllowed{C<:AbstractFunctionModification} <: NotAllowedError
  change::C
  message::String
end

```

An error indicating that the objective modification change cannot be applied to the objective.

[source](#)

`MathOptInterface.DeleteNotAllowed` – Type.

```

struct DeleteNotAllowed{IndexType <: Index} <: NotAllowedError
  index::IndexType
  message::String
end

```

An error indicating that the index index cannot be deleted.

[source](#)

`MathOptInterface.UnsupportedSubmittable` – Type.

```

struct UnsupportedSubmittable{SubmitType} <: UnsupportedError
  sub::SubmitType
  message::String
end

```

An error indicating that the submittable `sub` is not supported by the model, i.e. that `supports` returns `false`.

source

`MathOptInterface.SubmitNotAllowed` – Type.

```
struct SubmitNotAllowed{SubmitType<:AbstractSubmittable} <: NotAllowedError
  sub::SubmitType
  message::String # Human-friendly explanation why the attribute cannot be set
end
```

An error indicating that the submittable `sub` is supported (see `supports`) but cannot be added for some reason (see the error string).

source

Note that setting the `ConstraintFunction` of a `VariableIndex` constraint is not allowed:

`MathOptInterface.SettingVariableIndexNotAllowed` – Type.

```
SettingVariableIndexNotAllowed()
```

Error type that should be thrown when the user calls `set` to change the `ConstraintFunction` of a `VariableIndex` constraint.

source

Part VI

Submodules

Chapter 26

Benchmarks

26.1 Overview

The Benchmarks submodule

To aid the development of efficient solver wrappers, MathOptInterface provides benchmarking functionality. Benchmarking a wrapper follows a two-step process.

First, prior to making changes, run and save the benchmark results on a given benchmark suite as follows:

```
using SolverPackage # Replace with your choice of solver.

using MathOptInterface
const MOI = MathOptInterface

suite = MOI.Benchmarks.suite() do
    SolverPackage.Optimizer()
end

MOI.Benchmarks.create_baseline(
    suite, "current"; directory = "/tmp", verbose = true
)
```

Use the `exclude` argument to `Benchmarks.suite` to exclude benchmarks that the solver doesn't support.

Second, after making changes to the package, re-run the benchmark suite and compare to the prior saved results:

```
using SolverPackage, MathOptInterface

const MOI = MathOptInterface

suite = MOI.Benchmarks.suite() do
    SolverPackage.Optimizer()
end

MOI.Benchmarks.compare_against_baseline(
    suite, "current"; directory = "/tmp", verbose = true
)
```

This comparison will create a report detailing improvements and regressions.

26.2 API Reference

Benchmarks

Functions to help benchmark the performance of solver wrappers. See [The Benchmarks submodule](#) for more details.

[MathOptInterface.Benchmarks.suite](#) – Function.

```
suite(
  new_model::Function;
  exclude::Vector{Regex} = Regex[]
)
```

Create a suite of benchmarks. `new_model` should be a function that takes no arguments, and returns a new instance of the optimizer you wish to benchmark.

Use `exclude` to exclude a subset of benchmarks.

Examples

```
suite() do
  GLPK.Optimizer()
end
suite(exclude = [r"delete"]) do
  Gurobi.Optimizer(OutputFlag=0)
end
```

[source](#)

[MathOptInterface.Benchmarks.create_baseline](#) – Function.

```
create_baseline(suite, name::String; directory::String = ""; kwargs...)
```

Run all benchmarks in `suite` and save to files called `name` in `directory`.

Extra `kwargs` are based to `BenchmarkTools.run`.

Examples

```
my_suite = suite() -> GLPK.Optimizer()
create_baseline(my_suite, "glpk_master"; directory = "/tmp", verbose = true)
```

[source](#)

[MathOptInterface.Benchmarks.compare_against_baseline](#) – Function.

```
compare_against_baseline(
  suite, name::String; directory::String = "",
  report_filename::String = "report.txt"
)
```

Run all benchmarks in `suite` and compare against files called `name` in `directory` that were created by a call to `create_baseline`.

A report summarizing the comparison is written to `report_filename` in `directory`.

Extra `kwargs` are based to `BenchmarkTools.run`.

Examples

```
my_suite = suite() -> GLPK.Optimizer()  
compare_against_baseline(  
  my_suite, "glpk_master"; directory = "/tmp", verbose = true  
)
```

[source](#)

Chapter 27

Bridges

27.1 Overview

The Bridges submodule

The Bridges module simplifies the process of converting models between equivalent formulations.

Tip

[Read our paper](#) for more details on how bridges are implemented.

Why bridges?

A constraint can often be written in a number of equivalent formulations. For example, the constraint $l \leq a^\top x \leq u$ ([ScalarAffineFunction-in-Interval](#)) could be re-formulated as two constraints: $a^\top x \geq l$ ([ScalarAffineFunction-in-GreaterThan](#)) and $a^\top x \leq u$ ([ScalarAffineFunction-in-LessThan](#)). An alternative re-formulation is to add a dummy variable y with the constraints $l \leq y \leq u$ ([VariableIndex-in-Interval](#)) and $a^\top x - y = 0$ ([ScalarAffineFunction-in-EqualTo](#)).

To avoid each solver having to code these transformations manually, MathOptInterface provides bridges.

A bridge is a small transformation from one constraint type to another (potentially collection of) constraint type.

Because these bridges are included in MathOptInterface, they can be re-used by any optimizer. Some bridges also implement constraint modifications and constraint primal and dual translations.

Several bridges can be used in combination to transform a single constraint into a form that the solver may understand. Choosing the bridges to use takes the form of finding a shortest path in the hypergraph of bridges. The methodology is detailed in [the MOI paper](#).

The three types of bridges

There are three types of bridges in MathOptInterface:

1. Constraint bridges
2. Variable bridges
3. Objective bridges

Constraint bridges Constraint bridges convert constraints formulated by the user into an equivalent form supported by the solver. Constraint bridges are subtypes of `Bridges.Constraint.AbstractBridge`.

The equivalent formulation may add constraints (and possibly also variables) in the underlying model.

In particular, constraint bridges can focus on rewriting the function of a constraint, and do not change the set. Function bridges are subtypes of `Bridges.Constraint.AbstractFunctionConversionBridge`.

Read the [list of implemented constraint bridges](#) for more details on the types of transformations that are available. Function bridges are `Bridges.Constraint.ScalarFunctionizeBridge` and `Bridges.Constraint.VectorFunctionizeBridge`.

Variable bridges Variable bridges convert variables added by the user, either free with `add_variable/add_variables`, or constrained with `add_constrained_variable/add_constrained_variables`, into an equivalent form supported by the solver. Variable bridges are subtypes of `Bridges.Variable.AbstractBridge`.

The equivalent formulation may add constraints (and possibly also variables) in the underlying model.

Read the [list of implemented variable bridges](#) for more details on the types of transformations that are available.

Objective bridges Objective bridges convert the `ObjectiveFunction` set by the user into an equivalent form supported by the solver. Objective bridges are subtypes of `Bridges.Objective.AbstractBridge`.

The equivalent formulation may add constraints (and possibly also variables) in the underlying model.

Read the [list of implemented objective bridges](#) for more details on the types of transformations that are available.

Bridges.full_bridge_optimizer

Tip

Unless you have an advanced use-case, this is probably the only function you need to care about.

To enable the full power of `MathOptInterface`'s bridges, wrap an optimizer in a `Bridges.full_bridge_optimizer`.

```
julia> inner_optimizer = MOI.Utilities.Model{Float64}()
MOIU.Model{Float64}

julia> optimizer = MOI.Bridges.full_bridge_optimizer(inner_optimizer, Float64)
MOIB.LazyBridgeOptimizer{MOIU.Model{Float64}}
with 0 variable bridges
with 0 constraint bridges
with 0 objective bridges
with inner model MOIU.Model{Float64}
```

That's all you have to do! Use `optimizer` as normal, and bridging will happen lazily behind the scenes. By lazily, we mean that bridging will only happen if the constraint is not supported by the `inner_optimizer`.

Info

Most bridges are added by default in `Bridges.full_bridge_optimizer`. However, for technical reasons, some bridges are not added by default. Three examples include `Bridges.Constraint.SOCtoPSDBridge`, `Bridges.Constraint.SOCtoNonConvexQuadBridge` and `Bridges.Constraint.RSOCtoNonConvexQuadBridge`. See the docs of those bridges for more information.

Add a single bridge

If you don't want to use `Bridges.full_bridge_optimizer`, you can wrap an optimizer in a single bridge.

However, this will force the constraint to be bridged, even if the `inner_optimizer` supports it.

```
julia> inner_optimizer = MOI.Utilities.Model{Float64}()
MOIU.Model{Float64}

julia> optimizer = MOI.Bridges.Constraint.SplitInterval{Float64}(inner_optimizer)
MOIB.Constraint.SingleBridgeOptimizer{MOIB.Constraint.SplitIntervalBridge{Float64, F, S, LS, US}
↪ where {F<:MOI.AbstractFunction, S<:MOI.AbstractSet, LS<:MOI.AbstractSet, US<:MOI.AbstractSet},
↪ MOIU.Model{Float64}}
with 0 constraint bridges
with inner model MOIU.Model{Float64}

julia> x = MOI.add_variable(optimizer)
MOI.VariableIndex(1)

julia> MOI.add_constraint(optimizer, x, MOI.Interval(0.0, 1.0))
MathOptInterface.ConstraintIndex{MathOptInterface.VariableIndex,
↪ MathOptInterface.Interval{Float64}}(1)

julia> MOI.get(optimizer, MOI.ListOfConstraintTypesPresent())
1-element Vector{Tuple{Type, Type}}:
 (MathOptInterface.VariableIndex, MathOptInterface.Interval{Float64})

julia> MOI.get(inner_optimizer, MOI.ListOfConstraintTypesPresent())
2-element Vector{Tuple{Type, Type}}:
 (MathOptInterface.VariableIndex, MathOptInterface.GreaterThan{Float64})
 (MathOptInterface.VariableIndex, MathOptInterface.LessThan{Float64})
```

Bridges.LazyBridgeOptimizer

If you don't want to use `Bridges.full_bridge_optimizer`, but you need more than a single bridge (or you want the bridging to happen lazily), you can manually construct a `Bridges.LazyBridgeOptimizer`.

First, wrap an inner optimizer:

```
julia> inner_optimizer = MOI.Utilities.Model{Float64}()
MOIU.Model{Float64}

julia> optimizer = MOI.Bridges.LazyBridgeOptimizer(inner_optimizer)
MOIB.LazyBridgeOptimizer{MOIU.Model{Float64}}
with 0 variable bridges
with 0 constraint bridges
with 0 objective bridges
with inner model MOIU.Model{Float64}
```

Then use `Bridges.add_bridge` to add individual bridges:

```
julia> MOI.Bridges.add_bridge(optimizer, MOI.Bridges.Constraint.SplitIntervalBridge{Float64})

julia> MOI.Bridges.add_bridge(optimizer, MOI.Bridges.Objective.FunctionizeBridge{Float64})
```

Now the constraints will be bridged only if needed:

```

julia> x = MOI.add_variable(optimizer)
MOI.VariableIndex(1)

julia> MOI.add_constraint(optimizer, x, MOI.Interval(0.0, 1.0))
MathOptInterface.ConstraintIndex{MathOptInterface.VariableIndex,
↳ MathOptInterface.Interval{Float64}}(1)

julia> MOI.get(optimizer, MOI.ListOfConstraintTypesPresent())
1-element Vector{Tuple{Type, Type}}:
 (MathOptInterface.VariableIndex, MathOptInterface.Interval{Float64})

julia> MOI.get(inner_optimizer, MOI.ListOfConstraintTypesPresent())
1-element Vector{Tuple{Type, Type}}:
 (MathOptInterface.VariableIndex, MathOptInterface.Interval{Float64})

```

27.2 Implementation

Bridge interface

To be usable by a bridge optimizer, a bridge must implement the following functions:

[MathOptInterface.Bridges.added_constrained_variable_types](#) – Function.

```

added_constrained_variable_types(
    BT::Type{<:Variable.AbstractBridge},
)::Vector{Tuple{Type}}

```

Return a list of the types of constrained variables that bridges of concrete type BT add. This is used by the [LazyBridgeOptimizer](#).

[source](#)

[MathOptInterface.Bridges.added_constraint_types](#) – Function.

```

added_constraint_types(
    BT::Type{<:Constraint.AbstractBridge},
)::Vector{Tuple{Type, Type}}

```

Return a list of the types of constraints that bridges of concrete type BT add. This is used by the [LazyBridgeOptimizer](#).

[source](#)

Additionally, variable bridges must implement:

[MathOptInterface.Bridges.Variable.supports_constrained_variable](#) – Function.

```

supports_constrained_variable(
    ::Type{<:AbstractBridge},
    ::Type{<:MOI.AbstractSet},
)::Bool

```

Return a Bool indicating whether the bridges of type BT support bridging constrained variables in S.

[source](#)

[MathOptInterface.Bridges.Variable.concrete_bridge_type](#) – Function.


```
concrete_bridge_type(
  BT::Type{<:AbstractBridge},
  S::Type{<:MOI.AbstractSet},
)::Type
```

Return the concrete type of the bridge supporting variables in S constraints. This function can only be called if `MOI.supports_constrained_variable(BT, S)` is true.

Examples

As a variable in `MathOptInterface.GreaterThan` is bridged into variables in `MathOptInterface.Nonnegatives` by the `VectorizeBridge`:

```
MOI.Bridges.Variable.concrete_bridge_type(
  MOI.Bridges.Variable.VectorizeBridge{Float64},
  MOI.GreaterThan{Float64},
)

# output

MathOptInterface.Bridges.Variable.VectorizeBridge{Float64, MathOptInterface.Nonnegatives}
```

[source](#)

`MathOptInterface.Bridges.Variable.bridge_constrained_variable` – Function.

```
bridge_constrained_variable(
  BT::Type{<:AbstractBridge},
  model::MOI.ModelLike,
  set::MOI.AbstractSet,
)
```

Bridge the constrained variable in set using bridge BT to model and returns a bridge object of type BT. The bridge type BT should be a concrete type, that is, all the type parameters of the bridge should be set. Use `concrete_bridge_type` to obtain a concrete type for given set types.

[source](#)

constraint bridges must implement:

`MathOptInterface.supports_constraint` – Method.

```
MOI.supports_constraint(
  BT::Type{<:AbstractBridge},
  F::Type{<:MOI.AbstractFunction},
  S::Type{<:MOI.AbstractSet},
)::Bool
```

Return a Bool indicating whether the bridges of type BT support bridging F-in-S constraints.

[source](#)

`MathOptInterface.Bridges.Constraint.concrete_bridge_type` – Function.

```
concrete_bridge_type(
  BT::Type{<:AbstractBridge},
  F::Type{<:MOI.AbstractFunction},
  S::Type{<:MOI.AbstractSet},
)::Type
```

Return the concrete type of the bridge supporting F-in-S constraints. This function can only be called if `MOI.supports_constraint(BT, F, S)` is true.

Examples

As a `MathOptInterface.VariableIndex-in-MathOptInterface.Interval` constraint is bridged into a `MathOptInterface.VariableIndex-in-MathOptInterface.GreaterThan` and a `MathOptInterface.VariableIndex-in-MathOptInterface.LessThan` by the `SplitIntervalBridge`:

```
MOI.Bridges.Constraint.concrete_bridge_type(
    MOI.Bridges.Constraint.SplitIntervalBridge{Float64},
    MOI.VariableIndex,
    MOI.Interval{Float64},
)

# output

MathOptInterface.Bridges.Constraint.SplitIntervalBridge{Float64, MathOptInterface.VariableIndex,
↳ MathOptInterface.Interval{Float64}, MathOptInterface.GreaterThan{Float64},
↳ MathOptInterface.LessThan{Float64}}
```

[source](#)

`MathOptInterface.Bridges.Constraint.bridge_constraint` – Function.

```
bridge_constraint(
    BT::Type{<:AbstractBridge},
    model::MOI.ModelLike,
    func::AbstractFunction,
    set::MOI.AbstractSet,
)
```

Bridge the constraint `func-in-set` using bridge `BT` to `model` and returns a bridge object of type `BT`. The bridge type `BT` should be a concrete type, that is, all the type parameters of the bridge should be set. Use `concrete_bridge_type` to obtain a concrete type for given function and set types.

[source](#)

and objective bridges must implement:

`MathOptInterface.Bridges.set_objective_function_type` – Function.

```
set_objective_function_type(
    BT::Type{<:Objective.AbstractBridge},
)::Type{<:MOI.AbstractScalarFunction}
```

Return the type of objective function that bridges of concrete type `BT` set. This is used by the `LazyBridgeOptimizer`.

[source](#)

`MathOptInterface.Bridges.Objective.concrete_bridge_type` – Function.

```
concrete_bridge_type(
    BT::Type{<:MOI.Bridges.Objective.AbstractBridge},
    F::Type{<:MOI.AbstractScalarFunction},
)::Type
```

Return the concrete type of the bridge supporting objective functions of type `F`. This function can only be called if `MOI.supports_objective_function(BT, F)` is true.

[source](#)

[MathOptInterface.Bridges.Objective.bridge_objective](#) – Function.

```
bridge_objective(
  BT::Type{<:MOI.Bridges.Objective.AbstractBridge},
  model::MOI.ModelLike,
  func::MOI.AbstractScalarFunction,
)
```

Bridge the objective function `func` using bridge `BT` to `model` and returns a bridge object of type `BT`. The bridge type `BT` should be a concrete type, that is, all the type parameters of the bridge should be set. Use [concrete_bridge_type](#) to obtain a concrete type for a given function type.

[source](#)

When querying the [NumberOfVariables](#), [NumberOfConstraints](#), [ListOfVariableIndices](#), and [ListOfConstraintIndices](#), the variables and constraints created by the bridges in the underlying model are hidden by the bridge optimizer. For this purpose, the bridge must provide access to the variables and constraints it has created by implementing the following methods of [get](#):

[MathOptInterface.get](#) – Method.

```
MOI.get(b::AbstractBridge, ::MOI.NumberOfVariables)
```

The number of variables created by the bridge `b` in the model.

[source](#)

[MathOptInterface.get](#) – Method.

```
MOI.get(b::AbstractBridge, ::MOI.ListOfVariableIndices)
```

The list of variables created by the bridge `b` in the model.

[source](#)

[MathOptInterface.get](#) – Method.

```
MOI.get(b::AbstractBridge, ::MOI.NumberOfConstraints{F, S}) where {F, S}
```

The number of constraints of the type `F`-in-`S` created by the bridge `b` in the model.

[source](#)

[MathOptInterface.get](#) – Method.

```
MOI.get(b::AbstractBridge, ::MOI.ListOfConstraintIndices{F, S}) where {F, S}
```

A `Vector{ConstraintIndex{F,S}}` with indices of all constraints of type `F`-in-`S` created by the bridge `b` in the model (i.e., of length equal to the value of `NumberOfConstraints{F,S}()`).

[source](#)

SetMap bridges

Implementing a constraint bridge relying on linear transformation between two sets is easier thanks to the [SetMap interface](#). The bridge simply needs to be a subtype of `[Bridges.Variable.SetMapBridge]` for a variable bridge and `[Bridges.Constraint.SetMapBridge]` for a constraint bridge and the linear transformation is represented with [Bridges.map_set](#), [Bridges.map_function](#), [Bridges.inverse_map_set](#), [Bridges.inverse_map_function](#),

`Bridges.adjoint_map_function` and `Bridges.inverse_adjoint_map_function`. Note that the implementing last 4 methods is optional in the sense that if they are not implemented, bridging constraint would still work but some features would be missing as described in the docstrings. See [L20, Section 2.1.2] for more details including [L20, Example 2.1.1] that illustrates the idea for `Bridges.Variable.SOCtoRSOCBridge`, `Bridges.Variable.RSOCtoSOCBridge`, `Bridges.Constraint.SOCtoRSOCBridge` and `Bridges.Constraint.RSOCtoSOCBridge`.

[L20] Legat, Benoît. Set Programming: Theory and Computation. PhD thesis. 2020.

27.3 API Reference

Bridges

`MathOptInterface.Bridges.AbstractBridge` – Type.

```
| AbstractBridge
```

Represents a bridged constraint or variable in a `MathOptInterface.Bridges.AbstractBridgeOptimizer`. It contains the indices of the variables and constraints that it has created in the model. These can be obtained using `MathOptInterface.NumberOfVariables`, `MathOptInterface.ListOfVariableIndices`, `MathOptInterface.NumberOfConstraints` and `MathOptInterface.ListOfConstraintIndices` using `MathOptInterface.get` with the bridge in place of the `MathOptInterface.ModelLike`. Attributes of the bridged model such as `MathOptInterface.ConstraintDual` and `MathOptInterface.ConstraintPrimal`, can be obtained using `MathOptInterface.get` with the bridge in place of the constraint index. These calls are used by the `MathOptInterface.Bridges.AbstractBridgeOptimizer` to communicate with the bridge so they should be implemented by the bridge.

[source](#)

`MathOptInterface.Bridges.AbstractBridgeOptimizer` – Type.

```
| AbstractBridgeOptimizer
```

A bridge optimizer applies given constraint bridges to a given optimizer thus extending the types of supported constraints. The attributes of the inner optimizer are automatically transformed to make the bridges transparent, e.g. the variables and constraints created by the bridges are hidden.

By convention, the inner optimizer should be stored in a `model` field and the dictionary mapping constraint indices to bridges should be stored in a `bridges` field. If a bridge optimizer deviates from these conventions, it should implement the functions `MOI.optimize!` and `bridge` respectively.

[source](#)

`MathOptInterface.Bridges.LazyBridgeOptimizer` – Type.

```
| LazyBridgeOptimizer{OT<:MOI.ModelLike} <: AbstractBridgeOptimizer
```

The `LazyBridgeOptimizer` combines several bridges, which are added using the `add_bridge` function.

Whenever a constraint is added, it only attempts to bridge it if it is not supported by the internal model (hence its name `Lazy`).

When bridging a constraint, it selects the minimal number of bridges needed.

For example, if a constraint F-in-S can be bridged into a constraint F1-in-S1 (supported by the internal model) using bridge 1 or bridged into a constraint F2-in-S2 (unsupported by the internal model) using bridge 2 which can then be bridged into a constraint F3-in-S3 (supported by the internal model) using bridge 3, it will choose bridge 1 as it allows to bridge F-in-S using only one bridge instead of two if it uses bridge 2 and 3.

[source](#)

[MathOptInterface.Bridges.add_bridge](#) – Function.

```
| add_bridge(b::LazyBridgeOptimizer, BT::Type{<:AbstractBridge})
```

Enable the use of the bridges of type BT by b.

[source](#)

[MathOptInterface.Bridges.remove_bridge](#) – Function.

```
| remove_bridge(b::LazyBridgeOptimizer, BT::Type{<:AbstractBridge})
```

Disable the use of the bridges of type BT by b.

[source](#)

[MathOptInterface.Bridges.has_bridge](#) – Function.

```
| has_bridge(b::LazyBridgeOptimizer, BT::Type{<:AbstractBridge})
```

Return a Bool indicating whether the bridges of type BT are used by b.

[source](#)

[MathOptInterface.Bridges.full_bridge_optimizer](#) – Function.

```
| full_bridge_optimizer(model::MOI.ModelLike, ::Type{T}) where {T}
```

Returns a [LazyBridgeOptimizer](#) bridging model for every bridge defined in this package (see below for the few exceptions) and for the coefficient type T in addition to the bridges in the list returned by `MOI.get(model, MOI.Bridges.ListOfNonstandardBridges{T}())`.

See also [ListOfNonstandardBridges](#).

Note

The following bridges are not added by `full_bridge_optimizer` except if they are in the list returned by `MOI.get(model, MOI.Bridges.ListOfNonstandardBridges{T}())` (see the doc-strings of the corresponding bridge for the reason they are not added):

- [Constraint.SOCtoNonConvexQuadBridge](#), [Constraint.RSOCtoNonConvexQuadBridge](#) and [Constraint.SOCtoPSDBridge](#).
- The subtypes of [Constraint.AbstractToIntervalBridge](#) (i.e. [Constraint.GreaterToIntervalBridge](#) and [Constraint.LessToIntervalBridge](#)) if T is not a subtype of `AbstractFloat`.

[source](#)

[MathOptInterface.Bridges.ListOfNonstandardBridges](#) – Type.

```
| ListOfNonstandardBridges{T}() <: MOI.AbstractOptimizerAttribute
```

Any optimizer can be wrapped in a [LazyBridgeOptimizer](#) using `full_bridge_optimizer`. However, by default [LazyBridgeOptimizer](#) uses a limited set of bridges that are:

1. implemented in `MOI.Bridges`
2. generally applicable for all optimizers.

For some optimizers however, it is useful to add additional bridges, such as those that are implemented in external packages (e.g., within the solver package itself) or only apply in certain circumstances (e.g., [Constraint.SOCtoNonConvexQuadBridge](#)).

Such optimizers should implement the `ListOfNonstandardBridges` attribute to return a vector of bridge types that are added by [full_bridge_optimizer](#) in addition to the list of default bridges.

Note that optimizers implementing `ListOfNonstandardBridges` may require package-specific functions or sets to be used if the non-standard bridges are not added. Therefore, you are recommended to use `model = MOI.instantiate(Package.Optimizer; with_bridge_type = T)` instead of `model = MOI.instantiate(Package.Optimizer)`. See [MathOptInterface.instantiate](#).

Examples

An optimizer using a non-default bridge in MOI.Bridges

Solvers supporting [MOI.ScalarQuadraticFunction](#) can support [MOI.SecondOrderCone](#) and [MOI.RotatedSecondOrderCone](#) by defining:

```
function MOI.get(::MyQuadraticOptimizer, ::ListofNonstandardBridges{Float64})
    return Type[
        MOI.Bridges.Constraint.SOCtoNonConvexQuadBridge{Float64},
        MOI.Bridges.Constraint.RSOCtoNonConvexQuadBridge{Float64},
    ]
end
```

An optimizer defining an internal bridge

Suppose an optimizer can exploit specific structure of a constraint, e.g., it can exploit the structure of the matrix A in the linear system of equations $A * x = b$.

The optimizer can define the function:

```
struct MatrixAffineFunction{T} <: MOI.AbstractVectorFunction
    A::SomeStructuredMatrixType{T}
    b::Vector{T}
end
```

and then a bridge

```
struct MatrixAffineFunctionBridge{T} <: MOI.Constraint.AbstractBridge
    # ...
end
# ...
```

from `VectorAffineFunction{T}` to the `MatrixAffineFunction`. Finally, it defines:

```
function MOI.get(::Optimizer{T}, ::ListofNonstandardBridges{T}) where {T}
    return Type[MatrixAffineFunctionBridge{T}]
end
```

[source](#)

[MathOptInterface.Bridges.debug_supports_constraint](#) – Function.

```
debug_supports_constraint(
  b::LazyBridgeOptimizer,
  F::Type{<:MOI.AbstractFunction},
  S::Type{<:MOI.AbstractSet};
  io::IO = Base.stdout,
)
```

Prints to `io` explanations for the value of `MOI.supports_constraint` with the same arguments.

[source](#)

`MathOptInterface.Bridges.debug_supports` – Function.

```
debug_supports(
  b::LazyBridgeOptimizer,
  ::MOI.ObjectiveFunction{F};
  io::IO = Base.stdout,
) where F
```

Prints to `io` explanations for the value of `MOI.supports` with the same arguments.

[source](#)

`MathOptInterface.Bridges.bridged_variable_function` – Function.

```
bridged_variable_function(
  b::AbstractBridgeOptimizer,
  vi::MOI.VariableIndex,
)
```

Return a `MOI.AbstractScalarFunction` of variables of `b.model` that equals `vi`. That is, if the variable `vi` is bridged, it returns its expression in terms of the variables of `b.model`. Otherwise, it returns `vi`.

[source](#)

`MathOptInterface.Bridges.unbridged_variable_function` – Function.

```
unbridged_variable_function(
  b::AbstractBridgeOptimizer,
  vi::MOI.VariableIndex,
)
```

Return a `MOI.AbstractScalarFunction` of variables of `b` that equals `vi`. That is, if the variable `vi` is an internal variable of `b.model` created by a bridge but not visible to the user, it returns its expression in terms of the variables of bridged variables. Otherwise, it returns `vi`.

[source](#)

`MathOptInterface.Bridges.bridged_function` – Function.

```
bridged_function(b::AbstractBridgeOptimizer, value)::typeof(value)
```

Substitute any bridged `MOI.VariableIndex` in `value` by an equivalent expression in terms of variables of `b.model`.

[source](#)

`MathOptInterface.Bridges.Variable.unbridged_map` – Function.

```
unbridged_map( bridge::MOI.Bridges.Variable.AbstractBridge, vi::MOI.VariableIndex, )
```

For a bridged variable in a scalar set, return a tuple of pairs mapping the variables created by the bridge to an affine expression in terms of the bridged variable `vi`.

```
unbridged_map(
    bridge::MOI.Bridges.Variable.AbstractBridge,
    vis::Vector{MOI.VariableIndex},
)
```

For a bridged variable in a vector set, return a tuple of pairs mapping the variables created by the bridge to an affine expression in terms of the bridged variable `vis`. If this method is not implemented, it falls back to calling the following method for every variable of `vis`.

```
unbridged_map(
    bridge::MOI.Bridges.Variable.AbstractBridge,
    vi::MOI.VariableIndex,
    i::MOI.IndexInVector,
)
```

For a bridged variable in a vector set, return a tuple of pairs mapping the variables created by the bridge to an affine expression in terms of the bridged variable `vi` corresponding to the `i`th variable of the vector.

If there is no way to recover the expression in terms of the bridged variable(s) `vi(s)`, return nothing. See [ZerosBridge](#) for an example of bridge returning nothing.

[source](#)

Constraint bridges

[MathOptInterface.Bridges.Constraint.AbstractBridge](#) – Type.

```
| AbstractBridge
```

Subtype of [MathOptInterface.Bridges.AbstractBridge](#) for constraint bridges.

[source](#)

[MathOptInterface.Bridges.Constraint.AbstractFunctionConversionBridge](#) – Type.

```
| abstract type AbstractFunctionConversionBridge{F, S} <: AbstractBridge end
```

Bridge a constraint `G-in-S` into a constraint `F-in-S` where `F` and `G` are equivalent representations of the same function. By convention, the transformed function is stored in the `constraint` field.

[source](#)

[MathOptInterface.Bridges.Constraint.SingleBridgeOptimizer](#) – Type.

```
| SingleBridgeOptimizer{BT<:AbstractBridge, OT<:MOI.ModelLike} <:
| AbstractBridgeOptimizer
```

The `SingleBridgeOptimizer` bridges any constraint supported by the bridge `BT`. This is in contrast with the [MathOptInterface.Bridges.LazyBridgeOptimizer](#) which only bridges the constraints that are unsupported by the internal model, even if they are supported by one of its bridges.

[source](#)

[MathOptInterface.Bridges.Constraint.add_all_bridges](#) – Function.


```
| add_all_bridges(bridged_model, ::Type{T}) where {T}
```

Add all bridges defined in the `Bridges.Constraint` submodule to `bridged_model`. The coefficient type used is `T`.

[source](#)

SetMap bridges `MathOptInterface.Bridges.Variable.SetMapBridge` – Type.

```
| abstract type SetMapBridge{T,S1,S2} <: AbstractBridge end
```

Consider two type of sets `S1`, `S2` and a linear mapping `A` that the image of a set of type `S1` under `A` is a set of type `S2`. A `SetMapBridge{T,S1,S2}` is a bridge that substitutes constrained variables in `S2` into the image through `A` of constrained variables in `S1`.

The linear map `A` is described by `MathOptInterface.Bridges.map_set`, `MathOptInterface.Bridges.map_function`. Implementing a method for these two functions is sufficient to bridge constrained variables. In order for the getters and setters of dual solutions, starting values, etc... to work as well a method for the following functions should be implemented as well: `MathOptInterface.Bridges.inverse_map_set`, `MathOptInterface.Bridges.inverse_map_function`, `MathOptInterface.Bridges.adjoint_map_function` and `MathOptInterface.Bridges.inverse_adjoint_map_function`. See the docstrings of the function to see which feature would be missing if it was not implemented for a given bridge.

[source](#)

`MathOptInterface.Bridges.Constraint.SetMapBridge` – Type.

```
| abstract type SetMapBridge{T,S2,S1,F,G} <: AbstractBridge end
```

Consider two type of sets `S1`, `S2` and a linear mapping `A` that the image of a set of type `S1` under `A` is a set of type `S2`. A `SetMapBridge{T,S2,S1,F,G}` is a bridge that maps `G`-in-`S2` constraints into `F`-in-`S1` by mapping the function through `A`.

The linear map `A` is described by `MathOptInterface.Bridges.map_set`, `MathOptInterface.Bridges.map_function`. Implementing a method for these two functions is sufficient to bridge constraints. In order for the getters and setters of dual solutions, starting values, etc... to work as well a method for the following functions should be implemented as well: `MathOptInterface.Bridges.inverse_map_set`, `MathOptInterface.Bridges.inverse_map_function`, `MathOptInterface.Bridges.adjoint_map_function` and `MathOptInterface.Bridges.inverse_adjoint_map_function`. See the docstrings of the function to see which feature would be missing if it was not implemented for a given bridge.

[source](#)

`MathOptInterface.Bridges.map_set` – Function.

```
| map_set(::Type{BT}, set) where {BT}
```

Return the image of set through the linear map `A` defined in `Variable.SetMapBridge` and `Constraint.SetMapBridge`. This is used for bridging the constraint and setting the `MathOptInterface.ConstraintSet`.

[source](#)

`MathOptInterface.Bridges.inverse_map_set` – Function.

```
| inverse_map_set(::Type{BT}, set) where {BT}
```

Return the preimage of set through the linear map A defined in [Variable.SetMapBridge](#) and [Constraint.SetMapBridge](#). This is used for getting the [MathOptInterface.ConstraintSet](#).

source

[MathOptInterface.Bridges.map_function](#) – Function.

```
| map_function(::Type{BT}, func) where {BT}
```

Return the image of func through the linear map A defined in [Variable.SetMapBridge](#) and [Constraint.SetMapBridge](#). This is used for getting the [MathOptInterface.ConstraintPrimal](#) of variable bridges. For constraint bridges, this is used for bridging the constraint, setting the [MathOptInterface.ConstraintFunction](#) and [MathOptInterface.ConstraintPrimalStart](#) and modifying the function with [MathOptInterface.modify](#).

```
| map_function(::Type{BT}, func, i::IndexInVector) where {BT}
```

Return the scalar function at the *i*th index of the vector function that would be returned by [map_function](#)(BT, func) except that it may compute the *i*th element. This is used by [bridged_function](#) and for getting the [MathOptInterface.VariablePrimal](#) and [MathOptInterface.VariablePrimalStart](#) of variable bridges.

source

[MathOptInterface.Bridges.inverse_map_function](#) – Function.

```
| inverse_map_function(::Type{BT}, func) where {BT}
```

Return the image of func through the inverse of the linear map A defined in [Variable.SetMapBridge](#) and [Constraint.SetMapBridge](#). This is used by [Variable.unbridged_map](#) and for setting the [MathOptInterface.VariablePrimal](#) of variable bridges and for getting the [MathOptInterface.ConstraintFunction](#), the [MathOptInterface.ConstraintPrimal](#) and the [MathOptInterface.ConstraintPrimalStart](#) of constraint bridges.

source

[MathOptInterface.Bridges.adjoint_map_function](#) – Function.

```
| adjoint_map_function(::Type{BT}, func) where {BT}
```

Return the image of func through the adjoint of the linear map A defined in [Variable.SetMapBridge](#) and [Constraint.SetMapBridge](#). This is used for getting the [MathOptInterface.ConstraintDual](#) and [MathOptInterface.ConstraintDualStart](#) of constraint bridges.

source

[MathOptInterface.Bridges.inverse_adjoint_map_function](#) – Function.

```
| inverse_adjoint_map_function(::Type{BT}, func) where {BT}
```

Return the image of func through the inverse of the adjoint of the linear map A defined in [Variable.SetMapBridge](#) and [Constraint.SetMapBridge](#). This is used for getting the [MathOptInterface.ConstraintDual](#) of variable bridges and setting the [MathOptInterface.ConstraintDualStart](#) of constraint bridges.

source

Bridges implemented [MathOptInterface.Bridges.Constraint.FlipSignBridge](#) – Type.

```
| FlipSignBridge{T, S1, S2, F, G}
```

Bridge a G-in-S1 constraint into an F-in-S2 constraint by multiplying the function by -1 and taking the point reflection of the set across the origin. The flipped F-in-S constraint is stored in the constraint field by convention.

[source](#)

[MathOptInterface.Bridges.Constraint.AbstractToIntervalBridge](#) - Type.

```
| AbstractToIntervalBridge{T, S1, F}
```

Bridge a F-in-Interval constraint into an F-in-Interval{T} constraint where we have either:

- S1 = MOI.GreaterThan{T}
- S1 = MOI.LessThan{T}

The F-in-Interval{T} constraint is stored in the constraint field by convention.

Warning

It is required that T be a AbstractFloat type because otherwise typemin and typemax would either be not implemented (e.g. BigInt) or would not give infinite value (e.g. Int). For this reason, this bridge is only added to [MathOptInterface.Bridges.full_bridge_optimizer](#). when T is a subtype of AbstractFloat.

[source](#)

[MathOptInterface.Bridges.Constraint.GreaterToIntervalBridge](#) - Type.

```
| GreaterToIntervalBridge{T, F<:MOI.AbstractScalarFunction} <:
|   AbstractToIntervalBridge{T, MOI.GreaterThan{T}, F}
```

Transforms a F-in-GreaterThan{T} constraint into an F-in-Interval{T} constraint.

[source](#)

[MathOptInterface.Bridges.Constraint.LessToIntervalBridge](#) - Type.

```
| LessToIntervalBridge{T, F<:MOI.AbstractScalarFunction} <:
|   AbstractToIntervalBridge{T, MOI.LessThan{T}, F}
```

Transforms a F-in-LessThan{T} constraint into an F-in-Interval{T} constraint.

[source](#)

[MathOptInterface.Bridges.Constraint.GreaterToLessBridge](#) - Type.

```
| GreaterToLessBridge{
|   T,
|   F<:MOI.AbstractScalarFunction,
|   G<:MOI.AbstractScalarFunction
| } <: FlipSignBridge{T, MOI.GreaterThan{T}, MOI.LessThan{T}, F, G}
```

Transforms a G-in-GreaterThan{T} constraint into an F-in-LessThan{T} constraint.

[source](#)

[MathOptInterface.Bridges.Constraint.LessToGreaterBridge](#) - Type.

```
LessToGreaterBridge{
  T,
  F<:MOI.AbstractScalarFunction,
  G<:MOI.AbstractScalarFunction
} <: FlipSignBridge{T, MOI.LessThan{T}, MOI.GreaterThan{T}, F, G}
```

Transforms a G-in-LessThan{T} constraint into an F-in-GreaterThan{T} constraint.

[source](#)

[MathOptInterface.Bridges.Constraint.NonnegToNonposBridge](#) – Type.

```
NonnegToNonposBridge{
  T,
  F<:MOI.AbstractVectorFunction,
  G<:MOI.AbstractVectorFunction
} <: FlipSignBridge{T, MOI.Nonnegatives, MOI.Nonpositives, F, G}
```

Transforms a G-in-Nonnegatives constraint into a F-in-Nonpositives constraint.

[source](#)

[MathOptInterface.Bridges.Constraint.NonposToNonnegBridge](#) – Type.

```
NonposToNonnegBridge{
  T,
  F<:MOI.AbstractVectorFunction,
  G<:MOI.AbstractVectorFunction,
} <: FlipSignBridge{T, MOI.Nonpositives, MOI.Nonnegatives, F, G}
```

Transforms a G-in-Nonpositives constraint into a F-in-Nonnegatives constraint.

[source](#)

[MathOptInterface.Bridges.Constraint.VectorizeBridge](#) – Type.

```
VectorizeBridge{T,F,S,G}
```

Transforms a constraint G-in-scalar_set_type(S, T) where S <: VectorLinearSet to F-in-S.

Examples

The constraint VariableIndex-in-LessThan{Float64} becomes VectorAffineFunction{Float64}-in-Nonpositives, where T = Float64, F = VectorAffineFunction{Float64}, S = Nonpositives, and G = VariableIndex.

[source](#)

[MathOptInterface.Bridges.Constraint.ScalarizeBridge](#) – Type.

```
ScalarizeBridge{T, F, S}
```

Transforms a constraint AbstractVectorFunction-in-vector_set_type(S) where S <: LPCone{T} to F-in-S.

[source](#)

[MathOptInterface.Bridges.Constraint.ScalarSlackBridge](#) – Type.

```
ScalarSlackBridge{T, F, S}
```

The `ScalarSlackBridge` converts a constraint G -in- S where G is a function different from `VariableIndex` into the constraints F -in-`EqualTo{T}` and `VariableIndex`-in- S .

F is the result of subtracting a `VariableIndex` from G . Typically G is the same as F , but that is not mandatory.

[source](#)

`MathOptInterface.Bridges.Constraint.VectorSlackBridge` – Type.

```
| VectorSlackBridge{T, F, S}
```

The `VectorSlackBridge` converts a constraint G -in- S where G is a function different from `VectorOfVariables` into the constraints F -in-`Zeros` and `VectorOfVariables`-in- S .

F is the result of subtracting a `VectorOfVariables` from G . Typically G is the same as F , but that is not mandatory.

[source](#)

`MathOptInterface.Bridges.Constraint.ScalarFunctionizeBridge` – Type.

```
| ScalarFunctionizeBridge{T, S}
```

The `ScalarFunctionizeBridge` converts a constraint `VariableIndex`-in- S into the constraint `ScalarAffineFunction{T}`-in- S .

[source](#)

`MathOptInterface.Bridges.Constraint.VectorFunctionizeBridge` – Type.

```
| VectorFunctionizeBridge{T, S}
```

The `VectorFunctionizeBridge` converts a constraint `VectorOfVariables`-in- S into the constraint `VectorAffineFunction{T}`-in- S .

[source](#)

`MathOptInterface.Bridges.Constraint.SplitIntervalBridge` – Type.

```
| SplitIntervalBridge{T, F, S, LS, US}
```

The `SplitIntervalBridge` splits a F -in- S constraint into a F -in- LS and a F -in- US constraint where we have either:

- $S = \text{MOI.Interval}\{T\}$, $LS = \text{MOI.GreaterThan}\{T\}$ and $US = \text{MOI.LessThan}\{T\}$,
- $S = \text{MOI.EqualTo}\{T\}$, $LS = \text{MOI.GreaterThan}\{T\}$ and $US = \text{MOI.LessThan}\{T\}$, or
- $S = \text{MOI.Zeros}$, $LS = \text{MOI.Nonnegatives}$ and $US = \text{MOI.Nonpositives}$.

For instance, if F is `MOI.ScalarAffineFunction` and S is `MOI.Interval`, it transforms the constraint $la, x + u$ into the constraints $a, x + l$ and $a, x + u$.

Note

If $T <: \text{AbstractFloat}$ and S is `MOI.Interval{T}` then no lower (resp. upper) bound constraint is created if the lower (resp. upper) bound is `typemin(T)` (resp. `typemax(T)`). Similarly, when `MathOptInterface.ConstraintSet` is set, a lower or upper bound constraint may be deleted or created accordingly.

source

`MathOptInterface.Bridges.Constraint.SOCtoRSOCBridge` – Type.

```
| SOCtoRSOCBridge{T, F, G}
```

We simply do the inverse transformation of `RSOCtoSOCBridge`. In fact, as the transformation is an involution, we do the same transformation.

source

`MathOptInterface.Bridges.Constraint.RSOCtoSOCBridge` – Type.

```
| RSOCtoSOCBridge{T, F, G}
```

The `RotatedSecondOrderCone` is `SecondOrderCone` representable; see [BN01, p. 104]. Indeed, we have $2tu = (t/\sqrt{2} + u/\sqrt{2})^2 - (t/\sqrt{2} - u/\sqrt{2})^2$ hence

$$2tu \geq \|x\|_2^2$$

is equivalent to

$$(t/\sqrt{2} + u/\sqrt{2})^2 \geq \|x\|_2^2 + (t/\sqrt{2} - u/\sqrt{2})^2.$$

We can therefore use the transformation $(t, u, x) \mapsto (t/\sqrt{2} + u/\sqrt{2}, t/\sqrt{2} - u/\sqrt{2}, x)$. Note that the linear transformation is a symmetric involution (i.e. it is its own transpose and its own inverse). That means in particular that the norm of constraint primal and dual values are preserved by the transformation.

[BN01] Ben-Tal, Aharon, and Nemirovski, Arkadi. Lectures on modern convex optimization: analysis, algorithms, and engineering applications. Society for Industrial and Applied Mathematics, 2001.

source

`MathOptInterface.Bridges.Constraint.SOCtoNonConvexQuadBridge` – Type.

```
| SOCtoNonConvexQuadBridge{T}
```

Constraints of the form `VectorOfVariables-in-SecondOrderCone` can be transformed into a `ScalarQuadraticFunction-in-LessThan` and a `ScalarAffineFunction-in-GreaterThan`. Indeed, the definition of the second-order cone

$$t \geq \|x\|_2 \quad (1)$$

is equivalent to

$$\sum x_i^2 \leq t^2 \quad (2)$$

with $t \geq 0$. (3)

Warning

This transformation starts from a convex constraint (1) and creates a non-convex constraint (2), because the Q matrix associated with the constraint (2) has one negative eigenvalue. This might be wrongly interpreted by a solver. Some solvers can look at (2) and understand that it is a second order cone, but this is not a general rule. For these reasons this bridge is not automatically added by `MOI.Bridges.full_bridge_optimizer`. Care is recommended when adding this bridge to an optimizer.

[source](#)

`MathOptInterface.Bridges.Constraint.RSOCtoNonConvexQuadBridge` – Type.

`| RSOCtoNonConvexQuadBridge{T}`

Constraints of the form `VectorOfVariables-in-SecondOrderCone` can be transformed into a `ScalarQuadraticFunction-in-LessThan` and a `ScalarAffineFunction-in-GreaterThan`. Indeed, the definition of the second-order cone

$$2tu \geq \|x\|_2^2, t, u \geq 0 \quad (1)$$

is equivalent to

$$\sum x_i^2 \leq 2tu \quad (2)$$

with $t, u \geq 0$. (3)

WARNING This transformation starts from a convex constraint (1) and creates a non-convex constraint (2), because the Q matrix associated with the constraint 2 has two negative eigenvalues. This might be wrongly interpreted by a solver. Some solvers can look at (2) and understand that it is a rotated second order cone, but this is not a general rule. For these reasons, this bridge is not automatically added by `MOI.Bridges.full_bridge_optimizer`. Care is recommended when adding this bridge to an optimizer.

[source](#)

`MathOptInterface.Bridges.Constraint.QuadtoSOCBridge` – Type.

`| QuadtoSOCBridge{T}`

The set of points x satisfying the constraint

$$\frac{1}{2}x^T Q x + a^T x + b \leq 0$$

is a convex set if Q is positive semidefinite and is the union of two convex cones if a and b are zero (i.e. homogeneous case) and Q has only one negative eigenvalue. Currently, only the non-homogeneous transformation is implemented, see the Note section below for more details.

Non-homogeneous case

If Q is positive semidefinite, there exists U such that $Q = U^T U$, the inequality can then be rewritten as

$$\|Ux\|_2^2 \leq 2(-a^T x - b)$$

which is equivalent to the membership of $(1, -a^T x - b, Ux)$ to the rotated second-order cone.

Homogeneous case

If Q has only one negative eigenvalue, the set of x such that $x^T Q x \leq 0$ is the union of a convex cone and its opposite. We can choose which one to model by checking the existence of bounds on variables as shown below.

Second-order cone

If Q is diagonal and has eigenvalues $(1, 1, -1)$, the inequality $x^2 + x^2 \leq z^2$ combined with $z \geq 0$ defines the Lorenz cone (i.e. the second-order cone) but when combined with $z \leq 0$, it gives the opposite of the second order cone. Therefore, we need to check if the variable z has a lower bound 0 or an upper bound 0 in order to determine which cone is

Rotated second-order cone

The matrix Q corresponding to the inequality $x^2 \leq 2yz$ has one eigenvalue 1 with eigenvectors $(1, 0, 0)$ and $(0, 1, -1)$ and one eigenvalue -1 corresponding to the eigenvector $(0, 1, 1)$. Hence if we intersect this union of two convex cone with the halfspace $x + y \geq 0$, we get the rotated second-order cone and if we intersect it with the halfspace $x + y \leq 0$ we get the opposite of the rotated second-order cone. Note that y and z have the same sign since yz is nonnegative hence $x + y \geq 0$ is equivalent to $x \geq 0$ and $y \geq 0$.

Note

The check for existence of bound can be implemented (but inefficiently) with the current interface but if bound is removed or transformed (e.g. ≤ 0 transformed into ≥ 0) then the bridge is no longer valid. For this reason the homogeneous version of the bridge is not implemented yet.

[source](#)

`MathOptInterface.Bridges.Constraint.SOCtoPSDBridge` - Type.

The `SOCtoPSDBridge` transforms the second order cone constraint $\|x\| \leq t$ into the semidefinite cone constraints

$$\begin{pmatrix} t & x^T \\ x & tI \end{pmatrix} \succeq 0$$

Indeed by the Schur Complement, it is positive definite iff

$$\begin{aligned} tI &\succ 0 \\ t - x^T (tI)^{-1} x &\succ 0 \end{aligned}$$

which is equivalent to

$$\begin{aligned} t &> 0 \\ t^2 &> x^T x \end{aligned}$$

Warning

This bridge is not added by default by `MOI.Bridges.full_bridge_optimizer` as bridging second order cone constraints to semidefinite constraints can be achieved by the `SOCtoRSOCBridge` followed by the `RSOCtoPSDBridge` while creating a smaller semidefinite constraint.

source

`MathOptInterface.Bridges.Constraint.RS0CtoPSDBridge` - Type.

The `RS0CtoPSDBridge` transforms the second order cone constraint $\|x\| \leq 2tu$ with $u \geq 0$ into the semidefinite cone constraints

$$\begin{pmatrix} t & x^\top \\ x & 2uI \end{pmatrix} \succeq 0$$

Indeed by the Schur Complement, it is positive definite iff

$$\begin{aligned} uI &\succ 0 \\ t - x^\top (2uI)^{-1} x &\succ 0 \end{aligned}$$

which is equivalent to

$$\begin{aligned} u &> 0 \\ 2tu &> x^\top x \end{aligned}$$

source

`MathOptInterface.Bridges.Constraint.NormInfinityBridge` - Type.

`| NormInfinityBridge{T}`

The `NormInfinityCone` is representable with LP constraints, since $t \geq \max_i |x_i|$ if and only if $t \geq x_i$ and $t \geq -x_i$ for all i .

source

`MathOptInterface.Bridges.Constraint.NormOneBridge` - Type.

`| NormOneBridge{T}`

The `NormOneCone` is representable with LP constraints, since $t \geq \sum_i |x_i|$ if and only if there exists a vector y such that $t \geq \sum_i y_i$ and $y_i \geq x_i, y_i \geq -x_i$ for all i .

source

`MathOptInterface.Bridges.Constraint.GeoMeantoRelEntrBridge` - Type.

`| GeoMeantoRelEntrBridge{T}`

The geometric mean cone is representable with a relative entropy constraint and a nonnegative auxiliary variable.

This is because $u \leq \prod_{i=1}^n w_i^{1/n}$ is equivalent to $y \geq 0$ and $0 \leq u + y \leq \prod_{i=1}^n w_i^{1/n}$, and the latter inequality is equivalent to $1 \leq \prod_{i=1}^n (\frac{w_i}{u+y})^{1/n}$, which is equivalent to $0 \leq \sum_{i=1}^n \log(\frac{w_i}{u+y})^{1/n}$, which is equivalent to $0 \geq \sum_{i=1}^n (u+y) \log(\frac{u+y}{w_i})$.

Thus $(u, w) \in \text{GeometricMeanCone}(1+n)$ is representable as $y \geq 0, (0, w, (u+y)e) \in \text{RelativeEntropyCone}(1+2n)$, where e is a vector of ones.

source

`MathOptInterface.Bridges.Constraint.GeoMeanBridge` - Type.

`| GeoMeanBridge{T, F, G, H}`

The `GeometricMeanCone` is `SecondOrderCone` representable; see [1, p. 105].

The reformulation is best described in an example.

Consider the cone of dimension 4:

$$t \leq \sqrt[3]{x_1 x_2 x_3}$$

This can be rewritten as $\exists x_{21} \geq 0$ such that:

$$\begin{aligned} t &\leq x_{21}, \\ x_{21}^4 &\leq x_1 x_2 x_3 x_{21}. \end{aligned}$$

Note that we need to create x_{21} and not use t^4 directly as t is allowed to be negative. Now, this is equivalent to:

$$\begin{aligned} t &\leq x_{21}/\sqrt{4}, \\ x_{21}^2 &\leq 2x_1 x_{12}, \\ x_{11}^2 &\leq 2x_1 x_2, & x_{12}^2 &\leq 2x_3(x_{21}/\sqrt{4}). \end{aligned}$$

[1] Ben-Tal, Aharon, and Arkadi Nemirovski. Lectures on modern convex optimization: analysis, algorithms, and engineering applications. Society for Industrial and Applied Mathematics, 2001.

[source](#)

`MathOptInterface.Bridges.Constraint.RelativeEntropyBridge` - Type.

`| RelativeEntropyBridge{T}`

The `RelativeEntropyCone` is representable with exponential cone and LP constraints, since $u \geq \sum_{i=1}^n w_i \log(\frac{w_i}{v_i})$ if and only if there exists a vector y such that $u \geq \sum_i y_i$ and $y_i \geq w_i \log(\frac{w_i}{v_i})$ or equivalently $v_i \geq w_i \exp(\frac{-y_i}{w_i})$ or equivalently $(-y_i, w_i, v_i) \in \text{ExponentialCone}$, for all i .

[source](#)

`MathOptInterface.Bridges.Constraint.NormSpectralBridge` - Type.

`| NormSpectralBridge{T}`

The `NormSpectralCone` is representable with a PSD constraint, since $t \geq \sigma_1(X)$ if and only if $[tIX^\top; XtI] \succ 0$.

[source](#)

`MathOptInterface.Bridges.Constraint.NormNuclearBridge` - Type.

`| NormNuclearBridge{T}`

The `NormNuclearCone` is representable with an SDP constraint and extra variables, since $t \geq \sum_i \sigma_i(X)$ if and only if there exists symmetric matrices U, V such that $[UX^\top; XV] \succ 0$ and $t \geq (tr(U) + tr(V))/2$.

[source](#)

`MathOptInterface.Bridges.Constraint.SquareBridge` – Type.

```

SquareBridge{T, F<:MOI.AbstractVectorFunction,
             G<:MOI.AbstractScalarFunction,
             TT<:MOI.AbstractSymmetricMatrixSetTriangle,
             ST<:MOI.AbstractSymmetricMatrixSetSquare} <: AbstractBridge

```

The `SquareBridge` reformulates the constraint of a square matrix to be in `ST` to a list of equality constraints for pair or off-diagonal entries with different expressions and a `TT` constraint the upper triangular part of the matrix.

For instance, the constraint for the matrix

$$\begin{pmatrix} 1 & 1+x & 2-3x \\ 1+x & 2+x & 3-x \\ 2-3x & 2+x & 2x \end{pmatrix}$$

to be PSD can be broken down to the constraint of the symmetric matrix

$$\begin{pmatrix} 1 & 1+x & 2-3x \\ \cdot & 2+x & 3-x \\ \cdot & \cdot & 2x \end{pmatrix}$$

and the equality constraint between the off-diagonal entries (2, 3) and (3, 2) $2x == 1$. Note that now symmetrization constraint need to be added between the off-diagonal entries (1, 2) and (2, 1) or between (1, 3) and (3, 1) since the expressions are the same.

[source](#)

`MathOptInterface.Bridges.Constraint.RootDetBridge` – Type.

```

RootDetBridge{T,F,G,H}

```

The `RootDetConeTriangle` is representable by a `PositiveSemidefiniteConeTriangle` and an `GeometricMeanCone` constraints; see [1, p. 149].

Indeed, $t \leq \det(X)^{1/n}$ if and only if there exists a lower triangular matrix such that:

$$\begin{pmatrix} X \\ \top & \text{Diag}() \end{pmatrix} \succeq 0$$

$$t \leq (x_{11}x_{22} \cdots x_{nn})^{1/n}$$

[1] Ben-Tal, Aharon, and Arkadi Nemirovski. Lectures on modern convex optimization: analysis, algorithms, and engineering applications. Society for Industrial and Applied Mathematics, 2001.

[source](#)

`MathOptInterface.Bridges.Constraint.LogDetBridge` – Type.

```
| LogDetBridge{T,F,G,H,I}
```

The `LogDetConeTriangle` is representable by a `PositiveSemidefiniteConeTriangle` and `ExponentialCone` constraints.

Indeed, $\log \det(X) = \log(\delta_1) + \dots + \log(\delta_n)$ where $\delta_1, \dots, \delta_n$ are the eigenvalues of X .

Adapting the method from [1, p. 149], we see that $t \leq u \log(\det(X/u))$ for $u > 0$ if and only if there exists a lower triangular matrix such that

$$\begin{pmatrix} X & \\ & \text{Diag}() \end{pmatrix} \succeq 0$$

$$t \leq u \log(11/u) + u \log(22/u) + \dots + u \log(nn/u)$$

[1] Ben-Tal, Aharon, and Arkadi Nemirovski. Lectures on modern convex optimization: analysis, algorithms, and engineering applications. Society for Industrial and Applied Mathematics, 2001. ""

[source](#)

`MathOptInterface.Bridges.Constraint.IndicatorActiveOnFalseBridge` – Type.

```
| IndicatorActiveOnFalseBridge{T}
```

The `IndicatorActiveOnFalseBridge` replaces an indicator constraint activated on 0 with a variable z_0 with the constraint activated on 1, with a variable z_1 . It stores the added variable and added constraints:

- $z_1 \in \mathbb{B}$ in `zero_one_cons`
- $z_0 + z_1 == 1$ in `'indisjunction_cons'`
- The added `ACTIVATE_ON_ONE` indicator constraint in `indicator_cons_index`.

[source](#)

`MathOptInterface.Bridges.Constraint.IndicatorSOS1Bridge` – Type.

```
| IndicatorSOS1Bridge{T,S<:MOI.AbstractScalarSet}
```

The `IndicatorSOS1Bridge` replaces an indicator constraint of the following form: $z \in \mathbb{B}, z == 1 \implies f(x) \in S$ with a `SOS1` constraint: $z \in \mathbb{B}, \text{slack free}, f(x) + \text{slack} \in S, \text{SOS1}(\text{slack}, z)$.

[source](#)

`MathOptInterface.Bridges.Constraint.SemiToBinaryBridge` – Type.

```
| SemiToBinaryBridge{T, S <: MOI.AbstractScalarSet}
```

The `SemiToBinaryBridge` replaces a `Semicontinuous` constraint: $x \in \text{Semicontinuous}(l, u)$ is replaced by: $z \in \{0, 1\}, x \leq z \cdot u, x \geq z \cdot l$.

The `SemiToBinaryBridge` replaces a `Semiinteger` constraint: $x \in \text{Semiinteger}(l, u)$ is replaced by: $z \in \{0, 1\}, x \in \mathbb{Z}, x \leq z \cdot u, x \geq z \cdot l$.

[source](#)

`MathOptInterface.Bridges.Constraint.ZeroOneBridge` – Type.

```
| ZeroOneBridge{T}
```

The `ZeroOneBridge` splits a `MOI.VariableIndex`-in-`MOI.ZeroOne` constraint into a `MOI.VariableIndex`-in-`MOI.Integer` constraint and a `MOI.VariableIndex`-in-`MOI.Interval(0, 1)` constraint.

[source](#)

Variable bridges

`MathOptInterface.Bridges.Variable.AbstractBridge` – Type.

```
| AbstractBridge
```

Subtype of `MathOptInterface.Bridges.AbstractBridge` for variable bridges.

[source](#)

`MathOptInterface.Bridges.Variable.SingleBridgeOptimizer` – Type.

```
| SingleBridgeOptimizer{BT<:AbstractBridge, OT<:MOI.ModelLike} <:  
| AbstractBridgeOptimizer
```

The `SingleBridgeOptimizer` bridges any constrained variables supported by the bridge `BT`. This is in contrast with the `MathOptInterface.Bridges.LazyBridgeOptimizer` which only bridges the constrained variables that are unsupported by the internal model, even if they are supported by one of its bridges.

Note

Two bridge optimizers using variable bridges cannot be used together as both of them assume that the underlying model only returns variable indices with nonnegative values.

[source](#)

`MathOptInterface.Bridges.Variable.add_all_bridges` – Function.

```
| add_all_bridges(bridged_model, ::Type{T}) where {T}
```

Add all bridges defined in the `Bridges.Variable` submodule to `bridged_model`. The coefficient type used is `T`.

[source](#)

Bridges implemented `MathOptInterface.Bridges.Variable.FlipSignBridge` – Type.

```
| FlipSignBridge{T, S1, S2}
```

Bridge constrained variables in `S1` into constrained variables in `S2` by multiplying the variables by `-1` and taking the point reflection of the set across the origin. The flipped `MOI.VectorOfVariables-in-S` constraint is stored in the `flipped_constraint` field by convention.

[source](#)

`MathOptInterface.Bridges.Variable.ZerosBridge` – Type.

```
| ZerosBridge{T} <: Bridges.Variable.AbstractBridge
```

Transforms constrained variables in `MathOptInterface.Zeros` to zeros, which ends up creating no variables in the underlying model.

The bridged variables are therefore similar to parameters with zero values. Parameters with non-zero value can be created with constrained variables in `MOI.EqualTo` by combining a `VectorizeBridge` and this bridge. The functions cannot be unbridged, given a function, we cannot determine, if the bridged variables were used.

The dual values cannot be determined by the bridge but they can be determined by the bridged optimizer using `MathOptInterface.Utilities.get_fallback` if a `CachingOptimizer` is used (since `ConstraintFunction` cannot be got as functions cannot be unbridged).

[source](#)

[MathOptInterface.Bridges.Variable.FreeBridge](#) - Type.

```
| FreeBridge{T} <: Bridges.Variable.AbstractBridge
```

Transforms constrained variables in [MOI.Reals](#) to the difference of constrained variables in [MOI.Nonnegatives](#).

[source](#)

[MathOptInterface.Bridges.Variable.NonposToNonnegBridge](#) - Type.

```
| NonposToNonnegBridge{T} <:
|   FlipSignBridge{T, MOI.Nonpositives, MOI.Nonnegatives}
```

Transforms constrained variables in Nonpositives into constrained variables in Nonnegatives.

[source](#)

[MathOptInterface.Bridges.Variable.VectorizeBridge](#) - Type.

```
| VectorizeBridge{T, S}
```

Transforms a constrained variable in `scalar_set_type(S, T)` where `S <: VectorLinearSet` into a constrained vector of one variable in `S`. For instance, `VectorizeBridge{Float64, MOI.Nonnegatives}` transforms a constrained variable in `MOI.GreaterThan{Float64}` into a constrained vector of one variable in `MOI.Nonnegatives`.

[source](#)

[MathOptInterface.Bridges.Variable.SOCToRSOCBridge](#) - Type.

```
| SOCToRSOCBridge{T} <:
|   ↪ Bridges.Variable.SetMapBridge{T, MOI.RotatedSecondOrderCone, MOI.SecondOrderCone}
```

Same transformation as [MOI.Bridges.Constraint.SOCToRSOCBridge](#).

[source](#)

[MathOptInterface.Bridges.Variable.RSOCtoSOCBridge](#) - Type.

```
| RSOCtoSOCBridge{T} <:
|   ↪ Bridges.Variable.SetMapBridge{T, MOI.SecondOrderCone, MOI.RotatedSecondOrderCone}
```

Same transformation as [MOI.Bridges.Constraint.RSOCtoSOCBridge](#).

[source](#)

[MathOptInterface.Bridges.Variable.RSOCtoPSDBridge](#) - Type.

```
| RSOCtoPSDBridge{T} <: Bridges.Variable.AbstractBridge
```

Transforms constrained variables in [MathOptInterface.RotatedSecondOrderCone](#) to constrained variables in [MathOptInterface.PositiveSemidefiniteConeTriangle](#).

[source](#)

Objective bridges

`MathOptInterface.Bridges.Objective.AbstractBridge` – Type.

```
| AbstractBridge
```

Subtype of `MathOptInterface.Bridges.AbstractBridge` for objective bridges.

[source](#)

`MathOptInterface.Bridges.Objective.SingleBridgeOptimizer` – Type.

```
| SingleBridgeOptimizer{BT<:AbstractBridge, OT<:MOI.ModelLike} <: AbstractBridgeOptimizer
```

The `SingleBridgeOptimizer` bridges any objective functions supported by the bridge `BT`. This is in contrast with the `MathOptInterface.Bridges.LazyBridgeOptimizer` which only bridges the objective functions that are unsupported by the internal model, even if they are supported by one of its bridges.

[source](#)

`MathOptInterface.Bridges.Objective.add_all_bridges` – Function.

```
| add_all_bridges(bridged_model, ::Type{T}) where {T}
```

Add all bridges defined in the `Bridges.Objective` submodule to `bridged_model`. The coefficient type used is `T`.

[source](#)

Bridges implemented `MathOptInterface.Bridges.Objective.SlackBridge` – Type.

```
| SlackBridge{T, F, G}
```

The `SlackBridge` converts an objective function of type `G` into a `MOI.VariableIndex` objective by creating a slack variable and a `F`-in-`MOI.LessThan` constraint for minimization or `F`-in-`MOI.LessThan` constraint for maximization where `F` is `MOI.Utilities.promote_operation(-, T, G, MOI.VariableIndex)`. Note that when using this bridge, changing the optimization sense is not supported. Set the sense to `MOI.FEASIBILITY_SENSE` first to delete the bridge in order to change the sense, then re-add the objective.

[source](#)

`MathOptInterface.Bridges.Objective.FunctionizeBridge` – Type.

```
| FunctionizeBridge{T}
```

The `FunctionizeBridge` converts a `VariableIndex` objective into a `ScalarAffineFunction{T}` objective.

[source](#)

Chapter 28

FileFormats

28.1 Overview

The FileFormats submodule

The FileFormats module provides functionality for reading and writing MOI models using [write_to_file](#) and [read_from_file](#).

Supported file types

You must read and write files to a [FileFormats.Model](#) object. Specify the file-type by passing a [FileFormats.FileFormat](#) enum. For example:

The Conic Benchmark Format

```
| julia> model = MOI.FileFormats.Model(format = MOI.FileFormats.FORMAT_CBF)
| A Conic Benchmark Format (CBF) model
```

The LP file format

```
| julia> model = MOI.FileFormats.Model(format = MOI.FileFormats.FORMAT_LP)
| A .LP-file model
```

The MathOptFormat file format

```
| julia> model = MOI.FileFormats.Model(format = MOI.FileFormats.FORMAT_MOF)
| A MathOptFormat Model
```

The MPS file format

```
| julia> model = MOI.FileFormats.Model(format = MOI.FileFormats.FORMAT_MPS)
| A Mathematical Programming System (MPS) model
```

The NL file format

```
| julia> model = MOI.FileFormats.Model(format = MOI.FileFormats.FORMAT_NL)
| An AMPL (.nl) model
```


The REW file format

```
julia> model = MOI.FileFormats.Model(format = MOI.FileFormats.FORMAT_REW)
A Mathematical Programming System (MPS) model
```

Note that the **REW format** is identical to the MPS file format, except that all names are replaced with generic identifiers.

The SDPA file format

```
julia> model = MOI.FileFormats.Model(format = MOI.FileFormats.FORMAT_SDPA)
A SemiDefinite Programming Algorithm Format (SDPA) model
```

Write to file

To write a model src to a **MathOptFormat file**, use:

```
julia> src = MOI.Utilities.Model{Float64}()
MOIU.Model{Float64}

julia> MOI.add_variable(src)
MathOptInterface.VariableIndex(1)

julia> dest = MOI.FileFormats.Model(format = MOI.FileFormats.FORMAT_MOF)
A MathOptFormat Model

julia> MOI.copy_to(dest, src)
MathOptInterface.Utilities.IndexMap with 1 entry:
  VariableIndex(1) => VariableIndex(1)

julia> MOI.write_to_file(dest, "file.mof.json")

julia> print(read("file.mof.json", String))
{
  "name": "MathOptFormat Model",
  "version": {
    "major": 1,
    "minor": 0
  },
  "variables": [
    {
      "name": "x1"
    }
  ],
  "objective": {
    "sense": "feasibility"
  },
  "constraints": []
}
```

Read from file

To read a MathOptFormat file, use:

```
julia> dest = MOI.FileFormats.Model(format = MOI.FileFormats.FORMAT_MOF)
A MathOptFormat Model

julia> MOI.read_from_file(dest, "file.mof.json")

julia> MOI.get(dest, MOI.ListOfVariableIndices())
1-element Vector{MathOptInterface.VariableIndex}:
 MathOptInterface.VariableIndex(1)

julia> rm("file.mof.json") # Clean up after ourselves.
```

Detecting the filetype automatically

Instead of the `format` keyword, you can also use the `filename` keyword argument to `FileFormats.Model`. This will attempt to automatically guess the format from the file extension. For example:

```
julia> src = MOI.Utilities.Model{Float64}()
MOIU.Model{Float64}

julia> dest = MOI.FileFormats.Model(filename = "file.cbf.gz")
A Conic Benchmark Format (CBF) model

julia> MOI.copy_to(dest, src)
MathOptInterface.Utilities.IndexMap()

julia> MOI.write_to_file(dest, "file.cbf.gz")

julia> src_2 = MOI.FileFormats.Model(filename = "file.cbf.gz")
A Conic Benchmark Format (CBF) model

julia> src = MOI.Utilities.Model{Float64}()
MOIU.Model{Float64}

julia> dest = MOI.FileFormats.Model(filename = "file.cbf.gz")
A Conic Benchmark Format (CBF) model

julia> MOI.copy_to(dest, src)
MathOptInterface.Utilities.IndexMap()

julia> MOI.write_to_file(dest, "file.cbf.gz")

julia> src_2 = MOI.FileFormats.Model(filename = "file.cbf.gz")
A Conic Benchmark Format (CBF) model

julia> MOI.read_from_file(src_2, "file.cbf.gz")

julia> rm("file.cbf.gz") # Clean up after ourselves.
```

Note how the compression format (GZip) is also automatically detected from the filename.

Unsupported constraints

In some cases `src` may contain constraints that are not supported by the file format (e.g., the CBF format supports integer variables but not binary). If so, copy `src` to a bridged model using `Bridges.full_bridge_optimizer`:

```
src = MOI.Utilities.Model{Float64}()
x = MOI.add_variable(model)
MOI.add_constraint(model, x, MOI.ZeroOne())
dest = MOI.FileFormats.Model(format = MOI.FileFormats.FORMAT_CBF)
bridged = MOI.Bridges.full_bridge_optimizer(dest, Float64)
MOI.copy_to(bridged, src)
MOI.write_to_file(dest, "my_model.cbf")
```

Note

Even after bridging, it may still not be possible to write the model to file because of unsupported constraints (e.g., PSD variables in the LP file format).

Read and write to io

In addition to `write_to_file` and `read_from_file`, you can read and write directly from IO streams using `Base.write` and `Base.read!`:

```
julia> src = MOI.Utilities.Model{Float64}()
MOIU.Model{Float64}

julia> dest = MOI.FileFormats.Model(format = MOI.FileFormats.FORMAT_MPS)
A Mathematical Programming System (MPS) model

julia> MOI.copy_to(dest, src)
MathOptInterface.Utilities.IndexMap()

julia> io = IOBuffer();

julia> write(io, dest)

julia> seekstart(io);

julia> src_2 = MOI.FileFormats.Model(format = MOI.FileFormats.FORMAT_MPS)
A Mathematical Programming System (MPS) model

julia> read!(io, src_2);
```

Validating MOF files

MathOptFormat files are governed by a schema. Use [JSONSchema.jl](#) to check if a `.mof.json` file satisfies the schema.

First, construct the schema object as follows:

```
julia> import JSON, JSONSchema

julia> schema = JSONSchema.Schema(JSON.parsefile(MOI.FileFormats.MOF.SCHEMA_PATH))
A JSONSchema
```

Then, check if a model file is valid using `isvalid`:

```
julia> good_model = JSON.parse("""
{
```

```

        "version": {
            "major": 1,
            "minor": 0
        },
        "variables": [{"name": "x"}],
        "objective": {"sense": "feasibility"},
        "constraints": []
    }
    """
)

julia> isvalid(schema, good_model)
true

```

If we construct an invalid file, for example by mis-typing name as NaMe, the validation fails:

```

julia> bad_model = JSON.parse("""
{
    "version": {
        "major": 1,
        "minor": 0
    },
    "variables": [{"NaMe": "x"}],
    "objective": {"sense": "feasibility"},
    "constraints": []
}
""");

julia> isvalid(schema, bad_model)
false

```

Use JSONSchema.validate to obtain more insight into why the validation failed:

```

julia> JSONSchema.validate(schema, bad_model)
Validation failed:
path:      [variables][1]
instance:  Dict{String, Any}{"NaMe" => "x"}
schema key: required
schema value: Any["name"]

```

28.2 API Reference

File Formats

Functions to help read and write MOI models to/from various file formats. See [The FileFormats submodule](#) for more details.

[MathOptInterface.FileFormats.Model](#) – Function.

```

Model(
    ;
    format::FileFormat = FORMAT_AUTOMATIC,
    filename::Union{Nothing, String} = nothing,
    kwargs...
)

```

Return model corresponding to the FileFormat format, or, if format == FORMAT_AUTOMATIC, guess the format from filename.

The filename argument is only needed if format == FORMAT_AUTOMATIC.

kwargs are passed to the underlying model constructor.

[source](#)

`MathOptInterface.FileFormats.FileFormat` – Type.

| FileFormat

List of accepted export formats.

- FORMAT_AUTOMATIC: try to detect the file format based on the file name
- FORMAT_CBF: the Conic Benchmark format
- FORMAT_LP: the LP file format
- FORMAT_MOF: the MathOptFormat file format
- FORMAT_MPS: the MPS file format
- FORMAT_NL: the AMPL .nl file format
- FORMAT_REW: the .rew file format, which is MPS with generic names
- FORMAT_SDPA: the SemiDefinite Programming Algorithm format

[source](#)

`MathOptInterface.FileFormats.CBF.Model` – Type.

| Model()

Create an empty instance of FileFormats.CBF.Model.

[source](#)

`MathOptInterface.FileFormats.LP.Model` – Type.

| Model(; kwargs...)

Create an empty instance of FileFormats.LP.Model.

Keyword arguments are:

- maximum_length::Int=255: the maximum length for the name of a variable. Ip_solve 5.0 allows only 16 characters, while CPLEX 12.5+ allow 255.
- warn::Bool=false: print a warning when variables or constraints are renamed.

[source](#)

`MathOptInterface.FileFormats.MOF.Model` – Type.

| Model(; kwargs...)

Create an empty instance of FileFormats.MOF.Model.

Keyword arguments are:

- `print_compact::Bool=false`: print the JSON file in a compact format without spaces or newlines.
- `warn::Bool=false`: print a warning when variables or constraints are renamed

source

`MathOptInterface.FileFormats.MPS.Model` – Type.

```
| Model(; kwargs...)
```

Create an empty instance of `FileFormats.MPS.Model`.

Keyword arguments are:

- `warn::Bool=false`: print a warning when variables or constraints are renamed.
- `print_objsense::Bool=false`: print the OBJSENSE section when writing
- `generic_names::Bool=false`: strip all names in the model and replace them with the generic names `C$i` and `R$i` for the i 'th column and row respectively.

source

`MathOptInterface.FileFormats.NL.Model` – Type.

```
| Model()
```

Create a new Optimizer object.

source

`MathOptInterface.FileFormats.SDPA.Model` – Type.

```
| Model(; number_type::Type = Float64)
```

Create an empty instance of `FileFormats.SDPA.Model{number_type}`.

It is important to be aware that the SDPA file format is interpreted in geometric form and not standard conic form. The standard conic form and geometric conic form are two dual standard forms for semidefinite programs (SDPs). The geometric conic form of an SDP is as follows:

$$\min \quad b^T y \quad (28.1)$$

$$\text{s.t.} \quad \sum_{i=1}^m A_i y_i - C \in \mathbb{K} \quad (28.2)$$

where \mathcal{K} is a cartesian product of nonnegative orthant and positive semidefinite matrices that align with a block diagonal structure shared with the matrices A_i and C .

In other words, the geometric conic form contains free variables and affine constraints in either the nonnegative orthant or the positive semidefinite cone. That is, in the `MathOptInterface`'s terminology, `MathOptInterface.VectorAff` in `MathOptInterface.Nonnegatives` and `MathOptInterface.VectorAffineFunction` in `MathOptInterface.PositiveSemi` constraints.

The corresponding standard conic form of the dual SDP is as follows:

$$\max \quad \text{tr}(CX) \quad (28.3)$$

$$\text{s.t.} \quad \text{tr}(A_i X) = b_i \quad i = 1, \dots, m. \quad (28.4)$$

In other words, the standard conic form contains nonnegative and positive semidefinite variables with equality constraints. That is, in the `MathOptInterface`'s terminology, `MathOptInterface.VectorOfVariables`-in-`MathOptInterface.Nonnegatives`, `MathOptInterface.VectorOfVariables`-in-`MathOptInterface.PositiveSemidefinite` and `MathOptInterface.ScalarAffineFunction`-in-`MathOptInterface.EqualTo` constraints.

If a model is in standard conic form, use `Dualization.jl` to transform it into the geometric conic form before writing it. Otherwise, the nonnegative (resp. positive semidefinite) variables will be bridged into free variables with affine constraints constraining them to belong to the nonnegative orthant (resp. positive semidefinite cone) by the `MathOptInterface.Bridges.Constraint.VectorFunctionizeBridge`. Moreover, equality constraints will be bridged into pairs of affine constraints in the nonnegative orthant by the `MathOptInterface.Bridges.Constraint.SplitIntervalBridge` and then the `MathOptInterface.Bridges.Constraint.VectorFunctionizeBridge`.

If a solver is in standard conic form, use `Dualization.jl` to transform the model read into standard conic form before copying it to the solver. Otherwise, the free variables will be bridged into pairs of variables in the nonnegative orthant by the `MathOptInterface.Bridges.Variable.FreeBridge` and affine constraints will be bridged into equality constraints by creating a slack variable by the `MathOptInterface.Bridges.Constraint.VectorFunctionizeBridge`.

source

Chapter 29

Nonlinear

29.1 Overview

Nonlinear

Warning

The Nonlinear submodule is experimental. Until this message is removed, breaking changes may be introduced in any minor or patch release of MathOptInterface.

The Nonlinear submodule contains data structures and functions for working with a nonlinear optimization problem in the form of an expression graph. This page explains the API and describes the rationale behind its design.

Standard form

Nonlinear programs (NLPs) are a class of optimization problems in which some of the constraints or the objective function are nonlinear:

$$\min f_0(x) \tag{29.1}$$

$$\text{s.t. } l_j \leq f_j(x) \leq u_j \quad j = 1 \dots m \tag{29.2}$$

There may be additional constraints, as well as things like variable bounds and integrality restrictions, but we do not consider them here because they are best dealt with by other components of MathOptInterface.

API overview

The core element of the Nonlinear submodule is [Nonlinear.Model](#):

```
julia> const Nonlinear = MathOptInterface.Nonlinear;

julia> model = Nonlinear.Model()
A Nonlinear.Model with:
  0 objectives
  0 parameters
  0 expressions
  0 constraints
```

[Nonlinear.Model](#) is a mutable struct that stores all of the nonlinear information added to the model.

Decision variables Decision variables are represented by `VariableIndexes`. The user is responsible for creating these using `MOI.VariableIndex(i)`, where `i` is the column associated with the variable.

Expressions The input data structure is a Julia Expr. The input expressions can incorporate `VariableIndexes`, but these must be interpolated into the expression with `$`:

```
julia> x = MOI.VariableIndex(1)
MathOptInterface.VariableIndex{1}

julia> input = :(1 + sin($x)^2)
:(1 + sin(MathOptInterface.VariableIndex{1}) ^ 2)
```

There are a number of restrictions on the input Expr:

- It cannot contain macros
- It cannot contain broadcasting
- It cannot contain splatting (except in limited situations)
- It cannot contain linear algebra, such as matrix-vector products
- It cannot contain generator expressions, including `sum(i for i in S)`

Given an input expression, add an expression using `Nonlinear.add_expression`:

```
julia> expr = Nonlinear.add_expression(model, input)
MathOptInterface.Nonlinear.ExpressionIndex{1}
```

The return value, `expr`, is a `Nonlinear.ExpressionIndex` that can then be interpolated into other input expressions.

Looking again at `model`, we see:

```
julia> model
A Nonlinear.Model with:
 0 objectives
 0 parameters
 1 expression
 0 constraints
```

Parameters In addition to constant literals like 1 or 1.23, you can create parameters. Parameters are placeholders whose values can change before passing the expression to the solver. Create a parameter using `Nonlinear.add_parameter`, which accepts a default value:

```
julia> p = Nonlinear.add_parameter(model, 1.23)
MathOptInterface.Nonlinear.ParameterIndex{1}
```

The return value, `p`, is a `Nonlinear.ParameterIndex` that can then be interpolated into other input expressions.

Looking again at `model`, we see:

```
julia> model
A Nonlinear.Model with:
 0 objectives
 1 parameter
 1 expression
 0 constraints
```

Update a parameter as follows:

```
julia> model[p]
1.23

julia> model[p] = 4.56
4.56

julia> model[p]
4.56
```

Objectives Set a nonlinear objective using `Nonlinear.set_objective`:

```
julia> Nonlinear.set_objective(model, :($p + $expr + $x))

julia> model
A Nonlinear.Model with:
 1 objective
 1 parameter
 1 expression
 0 constraints
```

Clear a nonlinear objective by passing nothing:

```
julia> Nonlinear.set_objective(model, nothing)

julia> model
A Nonlinear.Model with:
 0 objectives
 1 parameter
 1 expression
 0 constraints
```

But we'll re-add the objective for later:

```
julia> Nonlinear.set_objective(model, :($p + $expr + $x));
```

Constraints Add a constraint using `Nonlinear.add_constraint`:

```
julia> c = Nonlinear.add_constraint(model, :(1 + sqrt($x)), MOI.LessThan(2.0))
MathOptInterface.Nonlinear.ConstraintIndex(1)

julia> model
A Nonlinear.Model with:
 1 objective
```

```

1 parameter
1 expression
1 constraint

```

The return value, `c`, is a `Nonlinear.ConstraintIndex` that is a unique identifier for the constraint. Interval constraints are also supported:

```

julia> c2 = Nonlinear.add_constraint(model, :(1 + sqrt($x)), MOI.Interval(-1.0, 2.0))
MathOptInterface.Nonlinear.ConstraintIndex{2}

julia> model
A Nonlinear.Model with:
 1 objective
 1 parameter
 1 expression
 2 constraints

```

Delete a constraint using `Nonlinear.delete`:

```

julia> Nonlinear.delete(model, c2)

julia> model
A Nonlinear.Model with:
 1 objective
 1 parameter
 1 expression
 1 constraint

```

User-defined operators By default, `Nonlinear` supports a wide range of univariate and multivariate operators. However, you can also define your own operators by registering them.

Univariate operators Register a univariate user-defined operator using `Nonlinear.register_operator`:

```

julia> f(x) = 1 + sin(x)^2
f (generic function with 1 method)

julia> Nonlinear.register_operator(model, :my_f, 1, f)

```

Now, you can use `:my_f` in expressions:

```

julia> new_expr = Nonlinear.add_expression(model, :(my_f($x + 1)))
MathOptInterface.Nonlinear.ExpressionIndex{2}

```

By default, `Nonlinear` will compute first- and second-derivatives of the registered operator using `ForwardDiff.jl`. Override this by passing functions which compute the respective derivative:

```

julia> f'(x) = 2 * sin(x) * cos(x)
f' (generic function with 1 method)

julia> Nonlinear.register_operator(model, :my_f2, 1, f, f')

```

or

```
julia> f''(x) = 2 * (cos(x)^2 - sin(x)^2)
f'' (generic function with 1 method)

julia> Nonlinear.register_operator(model, :my_f3, 1, f, f', f'')
```

Multivariate operators Register a multivariate user-defined operator using `Nonlinear.register_operator`:

```
julia> g(x...) = x[1]^2 + x[1] * x[2] + x[2]^2
g (generic function with 1 method)

julia> Nonlinear.register_operator(model, :my_g, 2, g)
```

Now, you can use `:my_g` in expressions:

```
julia> new_expr = Nonlinear.add_expression(model, :(my_g($x + 1, $x)))
MathOptInterface.Nonlinear.ExpressionIndex(3)
```

By default, `Nonlinear` will compute the gradient of the registered operator using `ForwardDiff.jl`. (Hessian information is not supported.) Override this by passing a function to compute the gradient:

```
julia> function ∇g(ret, x...)
    ret[1] = 2 * x[1] + x[2]
    ret[2] = x[1] + 2 * x[2]
    return
end
∇g (generic function with 1 method)

julia> Nonlinear.register_operator(model, :my_g2, 2, g, ∇g)
```

MathOptInterface `MathOptInterface` communicates the nonlinear portion of an optimization problem to solvers using concrete subtypes of `AbstractNLP evaluator`, which implement the `Nonlinear programming API`.

Create an `AbstractNLP evaluator` from `Nonlinear.Model` using `Nonlinear.Evaluator`.

`Nonlinear.Evaluator` requires an `Nonlinear.AbstractAutomaticDifferentiation` backend and an ordered list of the variables that are included in the model.

The following backends are available to choose from within MOI, although other packages may add more options by sub-typing `Nonlinear.AbstractAutomaticDifferentiation`:

- `Nonlinear.ExprGraphOnly`
- `Nonlinear.SparseReverseMode`.

```
julia> evaluator = Nonlinear.Evaluator(model, Nonlinear.ExprGraphOnly(), [x])
Nonlinear.Evaluator with available features:
* :ExprGraph
```

The functions of the `Nonlinear programming` API implemented by `Nonlinear.Evaluator` depends upon the chosen `Nonlinear.AbstractAutomaticDifferentiation` backend.

The `:ExprGraph` feature means we can call `objective_expr` and `constraint_expr` to retrieve the expression graph of the problem. However, we cannot call gradient terms such as `eval_objective_gradient` because `Nonlinear.ExprGraphOnly` does not have the capability to differentiate a nonlinear expression.

If, instead, we pass `Nonlinear.SparseReverseMode`, then we get access to `:Grad`, the gradient of the objective function, `:Jac`, the Jacobian matrix of the constraints, `:JacVec`, the ability to compute Jacobian-vector products, and `:ExprGraph`.

```
julia> evaluator = Nonlinear.Evaluator(
    model,
    Nonlinear.SparseReverseMode(),
    [x],
)
Nonlinear.Evaluator with available features:
* :Grad
* :Jac
* :JacVec
* :ExprGraph
```

However, before using the evaluator, we need to call `initialize`:

```
julia> MOI.initialize(evaluator, [:Grad, :Jac, :JacVec, :ExprGraph])
```

Now we can call methods like `eval_objective`:

```
julia> x = [1.0]
1-element Vector{Float64}:
 1.0

julia> MOI.eval_objective(evaluator, x)
7.268073418273571
```

and `eval_objective_gradient`:

```
julia> grad = [0.0]
1-element Vector{Float64}:
 0.0

julia> MOI.eval_objective_gradient(evaluator, grad, x)

julia> grad
1-element Vector{Float64}:
 1.909297426825682
```

Instead of passing `Nonlinear.Evaluator` directly to solvers, solvers query the `NLPBlock` attribute, which returns an `NLPBlockData`. This object wraps an `Nonlinear.Evaluator` and includes other information such as constraint bounds and whether the evaluator has a nonlinear objective. Create and set `NLPBlockData` as follows:

```
julia> block = MOI.NLPBlockData(evaluator);

julia> model = MOI.Utilities.UniversalFallback(MOI.Utilities.Model{Float64}());

julia> MOI.set(model, MOI.NLPBlock(), block);
```

Warning

Only call `NLPBlockData` once you have finished modifying the problem in `model`.

Putting everything together, you can create a nonlinear optimization problem in `MathOptInterface` as follows:

```
import MathOptInterface
const MOI = MathOptInterface

function build_model(
    model::MOI.ModelLike;
    backend::MOI.Nonlinear.AbstractAutomaticDifferentiation,
)
    x = MOI.add_variable(model)
    y = MOI.add_variable(model)
    MOI.set(model, MOI.ObjectiveSense(), MOI.MIN_SENSE)
    nl_model = MOI.Nonlinear.Model()
    MOI.Nonlinear.set_objective(nl_model, :($x^2 + $y^2))
    evaluator = MOI.Nonlinear.Evaluator(nl_model, backend, [x, y])
    MOI.set(model, MOI.NLPBlock(), MOI.NLPBlockData(evaluator))
    return
end

# Replace `model` and `backend` with your optimizer and backend of choice.
model = MOI.Utilities.UniversalFallback(MOI.Utilities.Model{Float64}())
build_model(model; backend = MOI.Nonlinear.SparseReverseMode())
```

Expression-graph representation

`Nonlinear.Model` stores nonlinear expressions in `Nonlinear.Expressions`. This section explains the design of the expression graph data structure in `Nonlinear.Expression`.

Given a nonlinear function like $f(x) = \sin(x)^2 + x$, a conceptual aid for thinking about the graph representation of the expression is to convert it into **Polish prefix notation**:

```
f(x, y) = (+ (^ (sin x) 2) x)
```

This format identifies each operator (function), as well as a list of arguments. Operators can be univariate, like `sin`, or multivariate, like `+`.

A common way of representing Polish prefix notation in code is as follows:

```
julia> x = MOI.VariableIndex(1);

julia> struct ExprNode
    op::Symbol
    children::Vector{Union{ExprNode, Float64, MOI.VariableIndex}}
end

julia> expr = ExprNode(:+, [ExprNode(:^, [ExprNode(:sin, [x]), 2.0]), x]);
```

This data structure follows our Polish prefix notation very closely, and we can easily identify the arguments to an operator. However, it has a significant draw-back: each node in the graph requires a `Vector`, which is heap-allocated and tracked by Julia's garbage collector (GC). For large models, we can expect to have millions of nodes in the expression graph, so this overhead quickly becomes prohibitive for computation.

An alternative is to record the expression as a linear tape:

```
julia> expr = Any[:, 2, :^, 2, :sin, 1, x, 2.0, x]
9-element Vector{Any}:
 :+
 2
 :^
 2
 :sin
 1
 MathOptInterface.VariableIndex(1)
 2.0
 MathOptInterface.VariableIndex(1)
```

The `Int` after each operator `Symbol` specifies the number of arguments.

This data-structure is a single vector, which resolves our problem with the GC, but each element is the abstract type, `Any`, and so any operations on it will lead to slower dynamic dispatch. It's also hard to identify the children of each operation without reading the entire tape.

To summarize, representing expression graphs in Julia has the following challenges:

- Nodes in the expression graph should not contain a heap-allocated object
- All data-structures should be concretely typed
- It should be easy to identify the children of a node

Sketch of the design in `Nonlinear` `Nonlinear` overcomes these problems by decomposing the data structure into a number of different concrete-typed vectors.

First, we create vectors of the supported uni- and multivariate operators.

```
julia> const UNIVARIATE_OPERATORS = [:sin];
julia> const MULTIVARIATE_OPERATORS =[:, :^];
```

In practice, there are many more supported operations than the ones listed here.

Second, we create an enum to represent the different types of nodes present in the expression graph:

```
julia> @enum(
    NodeType,
    NODE_CALL_MULTIVARIATE,
    NODE_CALL_UNIVARIATE,
    NODE_VARIABLE,
    NODE_VALUE,
)
```

In practice, there are node types other than the ones listed here.

Third, we create two concretely-typed structs as follows:

```
julia> struct Node
    type::NodeType
    parent::Int
    index::Int
end

julia> struct Expression
    nodes::Vector{Node}
    values::Vector{Float64}
end
```

For each node `node` in the `.nodes` field, if `node.type` is:

- `NODE_CALL_MULTIVARIATE`, we look up `MULTIVARIATE_OPERATORS[node.index]` to retrieve the operator
- `NODE_CALL_UNIVARIATE`, we look up `UNIVARIATE_OPERATORS[node.index]` to retrieve the operator
- `NODE_VARIABLE`, we create `MOI.VariableIndex(node.index)`
- `NODE_VALUE`, we look up `values[node.index]`

The `.parent` field of each node is the integer index of the parent node in `.nodes`. For the first node, the parent is -1 by convention.

Therefore, we can represent our function as:

```
julia> expr = Expression(
    [
        Node(NODE_CALL_MULTIVARIATE, -1, 1),
        Node(NODE_CALL_MULTIVARIATE, 1, 2),
        Node(NODE_CALL_UNIVARIATE, 2, 1),
        Node(NODE_VARIABLE, 3, 1),
        Node(NODE_VALUE, 2, 1),
        Node(NODE_VARIABLE, 1, 1),
    ],
    [2.0],
);
```

This is less readable than the other options, but does this data structure meet our design goals?

Instead of a heap-allocated object for each node, we only have two Vectors for each expression, `nodes` and `values`, as well as two constant vectors for the `OPERATORS`. In addition, all fields are concretely typed, and there are no Union or Any types.

For our third goal, it is not easy to identify the children of a node, but it is easy to identify the parent of any node. Therefore, we can use `Nonlinear.adjacency_matrix` to compute a sparse matrix that maps parents to their children.

The tape is also ordered topologically, so that a reverse pass of the nodes evaluates all children nodes before their parent.

The design in practice In practice, `Node` and `Expression` are exactly `Nonlinear.Node` and `Nonlinear.Expression`. However, `Nonlinear.NodeType` has more fields to account for comparison operators such as `:>=` and `:<=`, logic operators such as `:&&` and `:||`, nonlinear parameters, and nested subexpressions.

Moreover, instead of storing the operators as global constants, they are stored in `Nonlinear.OperatorRegistry`, and it also stores a vector of logic operators and a vector of comparison operators. In addition to `Nonlinear.DEFAULT_UNIVARIATE_OPERATORS` and `Nonlinear.DEFAULT_MULTIVARIATE_OPERATORS`, you can register user-defined functions using `Nonlinear.register_operator`.

`Nonlinear.Model` is a struct that stores the `Nonlinear.OperatorRegistry`, as well as a list of parameters and subexpressions in the model.

ReverseAD

`Nonlinear.ReverseAD` is a submodule for computing derivatives of a nonlinear optimization problem using sparse reverse-mode automatic differentiation (AD).

This section does not attempt to explain how sparse reverse-mode AD works, but instead explains why MOI contains its own implementation, and highlights notable differences from similar packages.

Warning

Don't use the API in `ReverseAD` to compute derivatives. Instead, create a `Nonlinear.Evaluator` object with `Nonlinear.SparseReverseMode` as the backend, and then query the MOI API methods.

Design goals The JuliaDiff organization maintains a [list of packages](#) for doing AD in Julia. At last count, there were at least ten packages--not including `ReverseAD`--for reverse-mode AD in Julia. `ReverseAD` exists because it has a different set of design goals.

- **Goal: handle scale and sparsity** The types of nonlinear optimization problems that MOI represents can be large scale (10^5 or more functions across 10^5 or more variables) with very sparse derivatives. The ability to compute a sparse Hessian matrix is essential. To the best of our knowledge, `ReverseAD` is the only reverse-mode AD system in Julia that handles sparsity by default.
- **Goal: limit the scope to improve robustness** Most other AD packages accept arbitrary Julia functions as input and then trace an expression graph using operator overloading. This means they must deal (or detect and ignore) with control flow, I/O, and other vagaries of Julia. In contrast, `ReverseAD` only accepts functions in the form of `Nonlinear.Expression`, which greatly limits the range of syntax that it must deal with. By reducing the scope of what we accept as input to functions relevant for mathematical optimization, we can provide a simpler implementation with various performance optimizations.
- **Goal: provide outputs which match what solvers expect** Other AD packages focus on differentiating individual Julia functions. In contrast, `ReverseAD` has a very specific use-case: to generate outputs needed by the MOI nonlinear API. This means it needs to efficiently compute sparse Hessians, and it needs subexpression handling to avoid recomputing subexpressions that are shared between functions.

History `ReverseAD` started life as `ReverseDiffSparse.jl`, development of which began in early 2014(!). This was well before the other packages started development. Because we had a well-tested, working AD in JuMP, there was less motivation to contribute to and explore other AD packages. The lack of historical interaction also meant that other packages were not optimized for the types of problems that JuMP is built for (i.e., large-scale sparse problems). When we first created `MathOptInterface`, we kept the AD in JuMP to simplify the transition, and post-poned the development of a first-class nonlinear interface in `MathOptInterface`.

Prior to the introduction of `Nonlinear`, JuMP's nonlinear implementation was a confusing mix of functions and types spread across the code base and in the private `_Derivatives` submodule. This made it hard to swap

the AD system for another. The main motivation for refactoring JuMP to create the Nonlinear submodule in MathOptInterface was to abstract the interface between JuMP and the AD system, allowing us to swap-in and test new AD systems in the future.

29.2 API Reference

Nonlinear Modeling

More information can be found in the [Nonlinear](#) section of the manual.

[MathOptInterface.Nonlinear](#) – Module.

| Nonlinear

Warning

The Nonlinear submodule is experimental. Until this message is removed, breaking changes may be introduced in any minor or patch release of MathOptInterface.

[source](#)

[MathOptInterface.Nonlinear.Model](#) – Type.

| Model()

The core datastructure for representing a nonlinear optimization problem.

It has the following fields:

- `objective::Union{Nothing,Expression}` : holds the nonlinear objective function, if one exists, otherwise nothing.
- `expressions::Vector{Expression}` : a vector of expressions in the model.
- `constraints::OrderedDict{ConstraintIndex,Constraint}` : a map from [ConstraintIndex](#) to the corresponding [Constraint](#). An `OrderedDict` is used instead of a `Vector` to support constraint deletion.
- `parameters::Vector{Float64}` : holds the current values of the parameters.
- `operators::OperatorRegistry` : stores the operators used in the model.

[source](#)

Expressions

[MathOptInterface.Nonlinear.ExpressionIndex](#) – Type.

| ExpressionIndex

An index to a nonlinear expression that is returned by [add_expression](#).

Given `data::Model` and `ex::ExpressionIndex`, use `data[ex]` to retrieve the corresponding [Expression](#).

[source](#)

[MathOptInterface.Nonlinear.add_expression](#) – Function.

| `add_expression(model::Model, expr)::ExpressionIndex`

Parse `expr` into a [Expression](#) and add to `model`. Returns an [ExpressionIndex](#) that can be interpolated into other input expressions.

`expr` must be a type that is supported by [parse_expression](#).

Examples

```
model = Model()
x = MOI.VariableIndex(1)
ex = add_expression(model, :($x^2 + 1))
set_objective(model, :(sqrt($ex)))
```

[source](#)

Parameters

[MathOptInterface.Nonlinear.ParameterIndex](#) – Type.

```
ParameterIndex
```

An index to a nonlinear parameter that is returned by [add_parameter](#). Given `data::Model` and `p::ParameterIndex`, use `data[p]` to retrieve the current value of the parameter and `data[p] = value` to set a new value.

[source](#)

[MathOptInterface.Nonlinear.add_parameter](#) – Function.

```
add_parameter(model::Model, value::Float64)::ParameterIndex
```

Add a new parameter to `model` with the default value `value`. Returns a [ParameterIndex](#) that can be interpolated into other input expressions and used to modify the value of the parameter.

Examples

```
model = Model()
x = MOI.VariableIndex(1)
p = add_parameter(model, 1.2)
c = add_constraint(model, :($x^2 - $p), MOI.LessThan(0.0))
```

[source](#)

Objectives

[MathOptInterface.Nonlinear.set_objective](#) – Function.

```
set_objective(model::Model, obj)::Nothing
```

Parse `obj` into a [Expression](#) and set as the objective function of `model`.

`obj` must be a type that is supported by [parse_expression](#).

To remove the objective, pass `nothing`.

Examples

```
model = Model()
x = MOI.VariableIndex(1)
set_objective(model, :($x^2 + 1))
set_objective(model, x)
set_objective(model, nothing)
```

[source](#)

Constraints

`MathOptInterface.Nonlinear.ConstraintIndex` – Type.

```
| ConstraintIndex
```

An index to a nonlinear constraint that is returned by `add_constraint`.

Given `data::Model` and `c::ConstraintIndex`, use `data[c]` to retrieve the corresponding `Constraint`.

[source](#)

`MathOptInterface.Nonlinear.add_constraint` – Function.

```
| add_constraint(  
|     model::Model,  
|     func,  
|     set::Union{  
|         MOI.GreaterThan{Float64},  
|         MOI.LessThan{Float64},  
|         MOI.Interval{Float64},  
|         MOI.EqualTo{Float64},  
|     },  
| )
```

Parse `func` and `set` into a `Constraint` and add to `model`. Returns a `ConstraintIndex` that can be used to delete the constraint or query solution information.

Examples

```
| model = Model()  
| x = MOI.VariableIndex(1)  
| c = add_constraint(model, :($x^2), MOI.LessThan(1.0))
```

[source](#)

`MathOptInterface.Nonlinear.delete` – Function.

```
| delete(model::Model, c::ConstraintIndex)::Nothing
```

Delete the constraint index `c` from `model`.

Examples

```
| model = Model()  
| x = MOI.VariableIndex(1)  
| c = add_constraint(model, :($x^2), MOI.LessThan(1.0))  
| delete(model, c)
```

[source](#)

User-defined operators

`MathOptInterface.Nonlinear.OperatorRegistry` – Type.

```
| OperatorRegistry()
```

Create a new `OperatorRegistry` to store and evaluate univariate and multivariate operators.

[source](#)

[MathOptInterface.Nonlinear.DEFAULT_UNIVARIATE_OPERATORS](#) – Constant.

```
| DEFAULT_UNIVARIATE_OPERATORS
```

The list of univariate operators that are supported by default.

[source](#)

[MathOptInterface.Nonlinear.DEFAULT_MULTIVARIATE_OPERATORS](#) – Constant.

```
| DEFAULT_MULTIVARIATE_OPERATORS
```

The list of multivariate operators that are supported by default.

[source](#)

[MathOptInterface.Nonlinear.register_operator](#) – Function.

```
register_operator(
    model::Model,
    op::Symbol,
    nargs::Int,
    f::Function,
    [∇f::Function],
    [∇²f::Function],
)
```

Register the user-defined operator `op` with `nargs` input arguments in `model`.

Univariate functions

- $f(x::T)::T$ must be a function that takes a single input argument x and returns the function evaluated at x . If ∇f and $\nabla^2 f$ are not provided, f must support any `Real` input type T .
- $\nabla f(x::T)::T$ is a function that takes a single input argument x and returns the first derivative of f with respect to x . If $\nabla^2 f$ is not provided, ∇f must support any `Real` input type T .
- $\nabla^2 f(x::T)::T$ is a function that takes a single input argument x and returns the second derivative of f with respect to x .

Multivariate functions

- $f(x::T...)::T$ must be a function that takes a `nargs` input arguments x and returns the function evaluated at x . If ∇f and $\nabla^2 f$ are not provided, f must support any `Real` input type T .
- $\nabla f(g::AbstractVector{T}, x::T...)::T$ is a function that takes a cache vector g of length `length(x)`, and fills each element $g[i]$ with the partial derivative of f with respect to $x[i]$.
- $\nabla^2 f(H::AbstractMatrix, x::T...)::T$ is a function that takes a matrix H and fills the lower-triangular components $H[i, j]$ with the Hessian of f with respect to $x[i]$ and $x[j]$ for $i \geq j$.

Notes for multivariate Hessians

- H has `size(H) == (length(x), length(x))`, but you must not access elements $H[i, j]$ for $i > j$.
- H is dense, but you do not need to fill structural zeros.

[source](#)

[MathOptInterface.Nonlinear.register_operator_if_needed](#) – Function.

```

register_operator_if_needed(
  registry::OperatorRegistry,
  op::Symbol,
  nargs::Int,
  f::Function;
)

```

Similar to `register_operator`, but this function warns if the function is not registered, and skips silently if it already is.

[source](#)

`MathOptInterface.Nonlinear.assert_registered` – Function.

```

assert_registered(registry::OperatorRegistry, op::Symbol, nargs::Int)

```

Throw an error if `op` is not registered in `registry` with `nargs` arguments.

[source](#)

`MathOptInterface.Nonlinear.check_return_type` – Function.

```

check_return_type(::Type{T}, ret::S) where {T,S}

```

Overload this method for new types `S` to throw an informative error if a user-defined function returns the type `S` instead of `T`.

[source](#)

`MathOptInterface.Nonlinear.eval_univariate_function` – Function.

```

eval_univariate_function(
  registry::OperatorRegistry,
  op::Symbol,
  x::T,
) where {T}

```

Evaluate the operator `op(x)::T`, where `op` is a univariate function in `registry`.

[source](#)

`MathOptInterface.Nonlinear.eval_univariate_gradient` – Function.

```

eval_univariate_gradient(
  registry::OperatorRegistry,
  op::Symbol,
  x::T,
) where {T}

```

Evaluate the first-derivative of the operator `op(x)::T`, where `op` is a univariate function in `registry`.

[source](#)

`MathOptInterface.Nonlinear.eval_univariate_hessian` – Function.

```

eval_univariate_hessian(
  registry::OperatorRegistry,
  op::Symbol,
  x::T,
) where {T}

```

Evaluate the second-derivative of the operator $op(x) :: T$, where op is a univariate function in registry.

source

`MathOptInterface.Nonlinear.eval_multivariate_function` – Function.

```
eval_multivariate_function(
  registry::OperatorRegistry,
  op::Symbol,
  x::AbstractVector{T},
) where {T}
```

Evaluate the operator $op(x) :: T$, where op is a multivariate function in registry.

source

`MathOptInterface.Nonlinear.eval_multivariate_gradient` – Function.

```
eval_multivariate_gradient(
  registry::OperatorRegistry,
  op::Symbol,
  g::AbstractVector{T},
  x::AbstractVector{T},
) where {T}
```

Evaluate the gradient of operator $g = \nabla op(x)$, where op is a multivariate function in registry.

source

`MathOptInterface.Nonlinear.eval_multivariate_hessian` – Function.

```
eval_multivariate_hessian(
  registry::OperatorRegistry,
  op::Symbol,
  H::AbstractMatrix,
  x::AbstractVector{T},
) where {T}
```

Evaluate the Hessian of operator $\nabla^2 op(x)$, where op is a multivariate function in registry.

The Hessian is stored in the lower-triangular part of the matrix H .

Note

Implementations of the Hessian operators will not fill structural zeros. Therefore, before calling this function you should pre-populate the matrix H with \emptyset .

source

`MathOptInterface.Nonlinear.eval_logic_function` – Function.

```
eval_logic_function(
  registry::OperatorRegistry,
  op::Symbol,
  lhs::T,
  rhs::T,
)::Bool where {T}
```

Evaluate $(lhs \ op \ rhs) :: Bool$, where op is a logic operator in registry.

source

[MathOptInterface.Nonlinear.eval_comparison_function](#) – Function.

```
eval_comparison_function(
    registry::OperatorRegistry,
    op::Symbol,
    lhs::T,
    rhs::T,
)::Bool where {T}
```

Evaluate $(lhs \text{ op } rhs) :: \text{Bool}$, where op is a comparison operator in `registry`.

[source](#)

Automatic-differentiation backends

[MathOptInterface.Nonlinear.Evaluator](#) – Type.

```
Evaluator(
    model::Model,
    backend::AbstractAutomaticDifferentiation,
    ordered_variables::Vector{MOI.VariableIndex},
)
```

Create `Evaluator`, a subtype of `MOI.AbstractNLP evaluator`, from `Model`.

[source](#)

[MathOptInterface.Nonlinear.AbstractAutomaticDifferentiation](#) – Type.

```
AbstractAutomaticDifferentiation
```

An abstract type for extending `Evaluator`.

[source](#)

[MathOptInterface.Nonlinear.ExprGraphOnly](#) – Type.

```
ExprGraphOnly() <: AbstractAutomaticDifferentiation
```

The default implementation of `AbstractAutomaticDifferentiation`. The only supported feature is `:ExprGraph`.

[source](#)

[MathOptInterface.Nonlinear.SparseReverseMode](#) – Type.

```
SparseReverseMode() <: AbstractAutomaticDifferentiation
```

An implementation of `AbstractAutomaticDifferentiation` that uses sparse reverse-mode automatic differentiation to compute derivatives. Supports all features in the MOI nonlinear interface.

[source](#)

Data-structure

[MathOptInterface.Nonlinear.Node](#) – Type.

```
struct Node
    type::NodeType
    index::Int
    parent::Int
end
```


A single node in a nonlinear expression tree. Used by [Expression](#).

See the `MathOptInterface` documentation for information on how the nodes and values form an expression tree.

[source](#)

`MathOptInterface.Nonlinear.NodeType` – Type.

```
| NodeType
```

An enum describing the possible node types. Each [Node](#) has a `.index` field, which should be interpreted as follows:

- `NODE_CALL_MULTIVARIATE`: the index into `operators.multivariate_operators`
- `NODE_CALL_UNIVARIATE`: the index into `operators.univariate_operators`
- `NODE_LOGIC`: the index into `operators.logic_operators`
- `NODE_COMPARISON`: the index into `operators.comparison_operators`
- `NODE_MOI_VARIABLE`: the value of `MOI.VariableIndex(index)` in the user's space of the model.
- `NODE_VARIABLE`: the 1-based index of the internal vector
- `NODE_VALUE`: the index into the `.values` field of [Expression](#)
- `NODE_PARAMETER`: the index into `data.parameters`
- `NODE_SUBEXPRESSION`: the index into `data.expressions`

[source](#)

`MathOptInterface.Nonlinear.Expression` – Type.

```
| struct Expression
|     nodes::Vector{Node}
|     values::Vector{Float64}
| end
```

The core type that represents a nonlinear expression. See the `MathOptInterface` documentation for information on how the nodes and values form an expression tree.

[source](#)

`MathOptInterface.Nonlinear.Constraint` – Type.

```
| struct Constraint
|     expression::Expression
|     set::Union{
|         MOI.LessThan{Float64},
|         MOI.GreaterThan{Float64},
|         MOI.EqualTo{Float64},
|         MOI.Interval{Float64},
|     }
| end
```

A type to hold information relating to the nonlinear constraint $f(x)$ in S , where $f(x)$ is defined by `.expression`, and S is `.set`.

[source](#)

`MathOptInterface.Nonlinear.adjacency_matrix` – Function.

```
| adjacency_matrix(nodes::Vector{Node})
```

Compute the sparse adjacency matrix describing the parent-child relationships in nodes.

The element (i, j) is true if there is an edge from node $[j]$ to node $[i]$. Since we get a column-oriented matrix, this gives us a fast way to look up the edges leaving any node (i.e., the children).

source

`MathOptInterface.Nonlinear.parse_expression` – Function.

```
| parse_expression(data::Model, input)::Expression
```

Parse input into a `Expression`.

source

```
| parse_expression(
    data::Model,
    expr::Expression,
    input::Any,
    parent_index::Int,
)::Expression
```

Parse input into a `Expression`, and add it to `expr` as a child of `expr.nodes[parent_index]`. Existing subexpressions and parameters are stored in `data`.

You can extend parsing support to new types of objects by overloading this method with a different type on `input::Any`.

source

`MathOptInterface.Nonlinear.convert_to_expr` – Function.

```
| convert_to_expr(data::Model, expr::Expression)
```

Convert the `Expression` `expr` into a Julia `Expr`.

- subexpressions are represented by a `ExpressionIndex` object.
- parameters are represented by a `ParameterIndex` object.
- variables are represented by an `MOI.VariableIndex` object.

source

```
| convert_to_expr(
    evaluator::Evaluator,
    expr::Expression;
    moi_output_format::Bool,
)
```

Convert the `Expression` `expr` into a Julia `Expr`.

If `moi_output_format = true`:

- subexpressions will be converted to Julia `Expr` and substituted into the output expression.
- the current value of each parameter will be interpolated into the expression

- variables will be represented in the form `x[MOI.VariableIndex(i)]`

If `moi_output_format = false`:

- subexpressions will be represented by a `ExpressionIndex` object.
- parameters will be represented by a `ParameterIndex` object.
- variables will be represented by an `MOI.VariableIndex` object.

Warning

To use `moi_output_format = true`, you must have first called `MOI.initialize` with `:ExprGraph` as a requested feature.

[source](#)

`MathOptInterface.Nonlinear.ordinal_index` – Function.

```
| ordinal_index(evaluator::Evaluator, c::ConstraintIndex)::Int
```

Return the 1-indexed value of the constraint index `c` in evaluator.

Examples

```
| model = Model()
| x = MOI.VariableIndex(1)
| c1 = add_constraint(model, :($x^2), MOI.LessThan(1.0))
| c2 = add_constraint(model, :($x^2), MOI.LessThan(1.0))
| evaluator = Evaluator(model)
| MOI.initialize(evaluator, Symbol[])
| ordinal_index(evaluator, c2) # Returns 2
| delete(model, c1)
| evaluator = Evaluator(model)
| MOI.initialize(evaluator, Symbol[])
| ordinal_index(model, c2) # Returns 1
```

[source](#)

Chapter 30

Utilities

30.1 Overview

The Utilities submodule

The Utilities submodule provides a variety of functionality for managing `MOI.ModelLike` objects.

Utilities.Model

`Utilities.Model` provides an implementation of a `ModelLike` that efficiently supports all functions and sets defined within MOI. However, given the extensibility of MOI, this might not cover all use cases.

Create a model as follows:

```
julia> model = MOI.Utilities.Model{Float64}()
MOIU.Model{Float64}
```

Utilities.UniversalFallback

`Utilities.UniversalFallback` is a layer that sits on top of any `ModelLike` and provides non-specialized (slower) fallbacks for constraints and attributes that the underlying `ModelLike` does not support.

For example, `Utilities.Model` doesn't support some variable attributes like `VariablePrimalStart`, so JuMP uses a combination of Universal fallback and `Utilities.Model` as a generic problem cache:

```
julia> model = MOI.Utilities.UniversalFallback(MOI.Utilities.Model{Float64}())
MOIU.UniversalFallback{MOIU.Model{Float64}}
fallback for MOIU.Model{Float64}
```

Warning

Adding a `UniversalFallback` means that your model will now support all constraints, even if the inner-model does not! This can lead to unexpected behavior.

Utilities.@model

For advanced use cases that need efficient support for functions and sets defined outside of MOI (but still known at compile time), we provide the `Utilities.@model` macro.

The `@model` macro takes a name (for a new type, which must not exist yet), eight tuples specifying the types of constraints that are supported, and then a `Bool` indicating the type is a subtype of `MOI.AbstractOptimizer` (if `true`) or `MOI.ModelLike` (if `false`).

The eight tuples are in the following order:

1. Un-typed scalar sets, e.g., `Integer`
2. Typed scalar sets, e.g., `LessThan`
3. Un-typed vector sets, e.g., `Nonnegatives`
4. Typed vector sets, e.g., `PowerCone`
5. Un-typed scalar functions, e.g., `VariableIndex`
6. Typed scalar functions, e.g., `ScalarAffineFunction`
7. Un-typed vector functions, e.g., `VectorOfVariables`
8. Typed vector functions, e.g., `VectorAffineFunction`

The tuples can contain more than one element. Typed-sets must be specified without their type parameter, i.e., `MOI.LessThan`, not `MOI.LessThan{Float64}`.

Here is an example:

```
julia> MOI.Utilities.@model(
    MyNewModel,
    (MOI.Integer,),           # Un-typed scalar sets
    (MOI.GreaterThan,),      # Typed scalar sets
    (MOI.Nonnegatives,),     # Un-typed vector sets
    (MOI.PowerCone,),        # Typed vector sets
    (MOI.VariableIndex,),    # Un-typed scalar functions
    (MOI.ScalarAffineFunction,), # Typed scalar functions
    (MOI.VectorOfVariables,), # Un-typed vector functions
    (MOI.VectorAffineFunction,), # Typed vector functions
    true,                    # <:MOI.AbstractOptimizer?
)
MathOptInterface.Utilities.GenericOptimizer{T, MathOptInterface.Utilities.ObjectiveContainer{T},
↳ MathOptInterface.Utilities.VariablesContainer{T}, MyNewModelFunctionConstraints{T}} where T

julia> model = MyNewModel{Float64}()
MOIU.GenericOptimizer{Float64, MOIU.ObjectiveContainer{Float64}, MOIU.VariablesContainer{Float64},
↳ MyNewModelFunctionConstraints{Float64}}
```

Warning

`MyNewModel` supports every `VariableIndex`-in-Set constraint, as well as `VariableIndex`, `ScalarAffineFunction`, and `ScalarQuadraticFunction` objective functions. Implement `MOI.supports` as needed to forbid constraint and objective function combinations.

As another example, `PATHSolver`, which only supports `VectorAffineFunction`-in-`Complements` defines its optimizer as:

```
julia> MOI.Utilities.@model(
    PathOptimizer,
    (), # Scalar sets
    (), # Typed scalar sets
    (MOI.Complements,), # Vector sets
```

```

        (), # Typed vector sets
        (), # Scalar functions
        (), # Typed scalar functions
        (), # Vector functions
        (MOI.VectorAffineFunction,), # Typed vector functions
        true, # is_optimizer
    )
MathOptInterface.Utilities.GenericOptimizer{T, MathOptInterface.Utilities.ObjectiveContainer{T},
↪ MathOptInterface.Utilities.VariablesContainer{T},
↪ MathOptInterface.Utilities.VectorOfConstraints{MathOptInterface.VectorAffineFunction{T},
↪ MathOptInterface.Complements}} where T

```

However, `PathOptimizer` does not support some `VariableIndex-in-Set` constraints, so we must explicitly define:

```

julia> function MOI.supports_constraint(
    ::PathOptimizer,
    ::Type{MOI.VariableIndex},
    ::Type{Union{<:MOI.Semiinteger, MOI.Semicontinuous, MOI.ZeroOne, MOI.Integer}}
)
    return false
end

```

Finally, `PATH` doesn't support an objective function, so we need to add:

```

julia> MOI.supports(::PathOptimizer, ::MOI.ObjectiveFunction) = false

```

Warning

This macro creates a new type, so it must be called from the top-level of a module, e.g., it cannot be called from inside a function.

Utilities.CachingOptimizer

A `[Utilities.CachingOptimizer]` is an `MOI` layer that abstracts the difference between solvers that support incremental modification (e.g., they support adding variables one-by-one), and solvers that require the entire problem in a single API call (e.g., they only accept the `A`, `b` and `c` matrices of a linear program).

It has two parts:

1. A cache, where the model can be built and modified incrementally
2. An optimizer, which is used to solve the problem

```

julia> model = MOI.Utilities.CachingOptimizer(
    MOI.Utilities.Model{Float64}(),
    PathOptimizer{Float64}(),
)
MOIU.CachingOptimizer{MOIU.GenericOptimizer{Float64, MOIU.ObjectiveContainer{Float64},
↪ MOIU.VariablesContainer{Float64}, MOIU.VectorOfConstraints{MOI.VectorAffineFunction{Float64},
↪ MOI.Complements}}, MOIU.Model{Float64}}
in state EMPTY_OPTIMIZER
in mode AUTOMATIC

```

```

with model cache MOIU.Model{Float64}
with optimizer MOIU.GenericOptimizer{Float64, MOIU.ObjectiveContainer{Float64},
↪ MOIU.VariablesContainer{Float64}, MOIU.VectorOfConstraints{MOI.VectorAffineFunction{Float64}},
↪ MOI.Complements}}

```

A `Utilities.CachingOptimizer` may be in one of three possible states:

- `NO_OPTIMIZER`: The `CachingOptimizer` does not have any optimizer.
- `EMPTY_OPTIMIZER`: The `CachingOptimizer` has an empty optimizer, and it is not synchronized with the cached model. Modifications are forwarded to the cache, but not to the optimizer.
- `ATTACHED_OPTIMIZER`: The `CachingOptimizer` has an optimizer, and it is synchronized with the cached model. Modifications are forwarded to the optimizer. If the optimizer does not support modifications, and error will be thrown.

Use `Utilities.attach_optimizer` to go from `EMPTY_OPTIMIZER` to `ATTACHED_OPTIMIZER`:

```

julia> MOI.Utilities.attach_optimizer(model)

julia> model
MOIU.CachingOptimizer{MOIU.GenericOptimizer{Float64, MOIU.ObjectiveContainer{Float64},
↪ MOIU.VariablesContainer{Float64}, MOIU.VectorOfConstraints{MOI.VectorAffineFunction{Float64}},
↪ MOI.Complements}}, MOIU.Model{Float64}}
in state ATTACHED_OPTIMIZER
in mode AUTOMATIC
with model cache MOIU.Model{Float64}
with optimizer MOIU.GenericOptimizer{Float64, MOIU.ObjectiveContainer{Float64},
↪ MOIU.VariablesContainer{Float64}, MOIU.VectorOfConstraints{MOI.VectorAffineFunction{Float64}},
↪ MOI.Complements}}

```

Info

You must be in `ATTACHED_OPTIMIZER` to use `optimize!`.

Use `Utilities.reset_optimizer` to go from `ATTACHED_OPTIMIZER` to `EMPTY_OPTIMIZER`:

```

julia> MOI.Utilities.reset_optimizer(model)

julia> model
MOIU.CachingOptimizer{MOIU.GenericOptimizer{Float64, MOIU.ObjectiveContainer{Float64},
↪ MOIU.VariablesContainer{Float64}, MOIU.VectorOfConstraints{MOI.VectorAffineFunction{Float64}},
↪ MOI.Complements}}, MOIU.Model{Float64}}
in state EMPTY_OPTIMIZER
in mode AUTOMATIC
with model cache MOIU.Model{Float64}
with optimizer MOIU.GenericOptimizer{Float64, MOIU.ObjectiveContainer{Float64},
↪ MOIU.VariablesContainer{Float64}, MOIU.VectorOfConstraints{MOI.VectorAffineFunction{Float64}},
↪ MOI.Complements}}

```

Info

Calling `MOI.empty!(model)` also resets the state to `EMPTY_OPTIMIZER`. So after emptying a model, the modification will only be applied to the cache.

Use `Utilities.drop_optimizer` to go from any state to `NO_OPTIMIZER`:

```
julia> MOI.Utilities.drop_optimizer(model)

julia> model
MOIU.CachingOptimizer{MOIU.GenericOptimizer{Float64, MOIU.ObjectiveContainer{Float64},
↪ MOIU.VariablesContainer{Float64}, MOIU.VectorOfConstraints{MOI.VectorAffineFunction{Float64},
↪ MOI.Complements}}, MOIU.Model{Float64}}
in state NO_OPTIMIZER
in mode AUTOMATIC
with model cache MOIU.Model{Float64}
with optimizer nothing
```

Pass an empty optimizer to `Utilities.reset_optimizer` to go from `NO_OPTIMIZER` to `EMPTY_OPTIMIZER`:

```
julia> MOI.Utilities.reset_optimizer(model, PathOptimizer{Float64}())

julia> model
MOIU.CachingOptimizer{MOIU.GenericOptimizer{Float64, MOIU.ObjectiveContainer{Float64},
↪ MOIU.VariablesContainer{Float64}, MOIU.VectorOfConstraints{MOI.VectorAffineFunction{Float64},
↪ MOI.Complements}}, MOIU.Model{Float64}}
in state EMPTY_OPTIMIZER
in mode AUTOMATIC
with model cache MOIU.Model{Float64}
with optimizer MOIU.GenericOptimizer{Float64, MOIU.ObjectiveContainer{Float64},
↪ MOIU.VariablesContainer{Float64}, MOIU.VectorOfConstraints{MOI.VectorAffineFunction{Float64},
↪ MOI.Complements}}
```

Deciding when to attach and reset the optimizer is tedious, and you will often write code like this:

```
try
    # modification
catch
    MOI.Utilities.reset_optimizer(model)
    # Re-try modification
end
```

To make this easier, `Utilities.CachingOptimizer` has two modes of operation:

- **AUTOMATIC:** The `CachingOptimizer` changes its state when necessary. Attempting to add a constraint or perform a modification not supported by the optimizer results in a drop to `EMPTY_OPTIMIZER` mode.
- **MANUAL:** The user must change the state of the `CachingOptimizer`. Attempting to perform an operation in the incorrect state results in an error.

By default, `AUTOMATIC` mode is chosen. However, you can create a `CachingOptimizer` in `MANUAL` mode as follows:

```
julia> model = MOI.Utilities.CachingOptimizer(
    MOI.Utilities.Model{Float64}(),
    MOI.Utilities.MANUAL,
)
```



```

MOIU.CachingOptimizer{MOI.AbstractOptimizer, MOIU.Model{Float64}}
in state NO_OPTIMIZER
in mode MANUAL
with model cache MOIU.Model{Float64}
with optimizer nothing

julia> MOI.Utilities.reset_optimizer(model, PathOptimizer{Float64}())

julia> model
MOIU.CachingOptimizer{MOI.AbstractOptimizer, MOIU.Model{Float64}}
in state EMPTY_OPTIMIZER
in mode MANUAL
with model cache MOIU.Model{Float64}
with optimizer MOIU.GenericOptimizer{Float64, MOIU.ObjectiveContainer{Float64},
↪ MOIU.VariablesContainer{Float64}, MOIU.VectorOfConstraints{MOI.VectorAffineFunction{Float64},
↪ MOI.Complements}}

```

Printing

Use `print` to print the formulation of the model.

```

julia> model = MOI.Utilities.Model{Float64}();

julia> x = MOI.add_variable(model)
MathOptInterface.VariableIndex(1)

julia> MOI.set(model, MOI.VariableName(), x, "x_var")

julia> MOI.add_constraint(model, x, MOI.ZeroOne())
MathOptInterface.ConstraintIndex{MathOptInterface.VariableIndex, MathOptInterface.ZeroOne}(1)

julia> MOI.set(model, MOI.ObjectiveFunction{typeof(x)}(), x)

julia> MOI.set(model, MOI.ObjectiveSense(), MOI.MAX_SENSE)

julia> print(model)
Maximize VariableIndex:
  x_var

Subject to:

VariableIndex-in-ZeroOne
  x_var ∈ {0, 1}

```

Use `Utilities.latex_formulation` to display the model in LaTeX form:

```

julia> MOI.Utilities.latex_formulation(model)
$$ \begin{aligned}
& \max \quad x_{\text{var}} \\
& \text{Subject to} \\
& \quad \text{VariableIndex-in-ZeroOne} \\
& \quad x_{\text{var}} \in \{0, 1\}
\end{aligned}
$$

```

Tip

In Julia, calling `print` or ending a cell with `Utilities.latex_formulation` will render the model in LaTeX.

Utilities.MatrixOfConstraints

The constraints of `Utilities.Model` are stored as a vector of tuples of function and set in a `Utilities.VectorOfConstraints`. Other representations can be used by parametrizing the type `Utilities.GenericModel` (resp. `Utilities.GenericOptimizer`). For instance, if all non-VariableIndex constraints are affine, the coefficients of all the constraints can be stored in a single sparse matrix using `Utilities.MatrixOfConstraints`. The constraints storage can even be customized up to a point where it exactly matches the storage of the solver of interest, in which case `copy_to` can be implemented for the solver by calling `copy_to` to this custom model.

For instance, `Clp` defines the following model

```
MOI.Utilities.@product_of_scalar_sets(LP, MOI.EqualTo{T}, MOI.LessThan{T}, MOI.GreaterThan{T})
const Model = MOI.Utilities.GenericModel{
    Float64,
    MOI.Utilities.MatrixOfConstraints{
        Float64,
        MOI.Utilities.MutableSparseMatrixCSC{Float64,Cint,MOI.Utilities.ZeroBasedIndexing},
        MOI.Utilities.Hyperrectangle{Float64},
        LP{Float64},
    },
}
```

The `copy_to` operation can now be implemented as follows (assuming that the `Model` definition above is in the `Clp` module so that it can be referred to as `Model`, to be distinguished with `Utilities.Model`):

```
function _copy_to(dest::Optimizer, src::Model)
    @assert MOI.is_empty(dest)
    A = src.constraints.coefficients
    row_bounds = src.constraints.constants
    Clp_loadProblem(
        dest,
        A.n,
        A.m,
        A.colptr,
        A.rowval,
        A.nzval,
        src.lower_bound,
        src.upper_bound,
        # (...) objective vector (omitted),
        row_bounds.lower,
        row_bounds.upper,
    )
    # Set objective sense and constant (omitted)
    return
end

function MOI.copy_to(dest::Optimizer, src::Model)
    _copy_to(dest, src)
    return MOI.Utilities.identity_index_map(src)
end

function MOI.copy_to(
```

```

    dest::Optimizer,
    src::MOI.Utilities.UniversalFallback{Model},
  )
  # Copy attributes from `src` to `dest` and error in case any unsupported
  # constraints or attributes are set in `UniversalFallback`.
  return MOI.copy_to(dest, src.model)
end

function MOI.copy_to(
  dest::Optimizer,
  src::MOI.ModelLike,
)
  model = Model()
  index_map = MOI.copy_to(model, src)
  _copy_to(dest, model)
  return index_map
end

```

ModelFilter

Utilities provides [Utilities.ModelFilter](#) as a useful tool to copy a subset of a model. For example, given an infeasible model, we can copy the irreducible infeasible subsystem (for models implementing [ConstraintConflictStatus](#)) as follows:

```

my_filter(::Any) = true
function my_filter(ci::MOI.ConstraintIndex)
  status = MOI.get(dest, MOI.ConstraintConflictStatus(), ci)
  return status != MOI.NOT_IN_CONFLICT
end
filtered_src = MOI.Utilities.ModelFilter(my_filter, src)
index_map = MOI.copy_to(dest, filtered_src)

```

Fallbacks

The value of some attributes can be inferred from the value of other attributes.

For example, the value of [ObjectiveValue](#) can be computed using [ObjectiveFunction](#) and [VariablePrimal](#).

When a solver gives direct access to an attribute, it is better to return this value. However, if this is not the case, [Utilities.get_fallback](#) can be used instead. For example:

```

function MOI.get(model::Optimizer, attr::MOI.ObjectiveFunction)
  return MOI.Utilities.get_fallback(model, attr)
end

```

DoubleDicts

When writing MOI interfaces, we often need to handle situations in which we map [ConstraintIndex](#)s to different values. For example, to a string for [ConstraintName](#).

One option is to use a dictionary like `Dict{MOI.ConstraintIndex, String}`. However, this incurs a performance cost because the key is not a concrete type.

The `DoubleDicts` submodule helps this situation by providing two types main types [Utilities.DoubleDicts.DoubleDict](#) and [Utilities.DoubleDicts.IndexDoubleDict](#). These types act like normal dictionaries, but internally they use more efficient dictionaries specialized to the type of the function-set pair.

The most common usage of a `DoubleDict` is in the `index_map` returned by `copy_to`. Performance can be improved, by using a function barrier. That is, instead of code like:

```
index_map = MOI.copy_to(dest, src)
for (F, S) in MOI.get(src, MOI.ListOfConstraintTypesPresent())
    for ci in MOI.get(src, MOI.ListOfConstraintIndices{F,S}())
        dest_ci = index_map[ci]
        # ...
    end
end
```

use instead:

```
function function_barrier(
    dest,
    src,
    index_map::MOI.Utilities.DoubleDicts.IndexDoubleDictInner{F,S},
) where {F,S}
    for ci in MOI.get(src, MOI.ListOfConstraintIndices{F,S}())
        dest_ci = index_map[ci]
        # ...
    end
    return
end

index_map = MOI.copy_to(dest, src)
for (F, S) in MOI.get(src, MOI.ListOfConstraintTypesPresent())
    function_barrier(dest, src, index_map[F, S])
end
```

30.2 API Reference

Utilities.Model

`MathOptInterface.Utilities.Model` – Type.

An implementation of `ModelLike` that supports all functions and sets defined in `MOI`. It is parameterized by the coefficient type.

Examples

```
model = Model{Float64}()
x = add_variable(model)
```

[source](#)

Utilities.UniversalFallback

`MathOptInterface.Utilities.UniversalFallback` – Type.

```
UniversalFallback
```

The `UniversalFallback` can be applied on a `MathOptInterface.ModelLike` model to create the model `UniversalFallback(model)` supporting any constraint and attribute. This allows to have a specialized

implementation in `model` for performance critical constraints and attributes while still supporting other attributes with a small performance penalty. Note that `model` is unaware of constraints and attributes stored by `UniversalFallback` so this is not appropriate if `model` is an optimizer (for this reason, `MathOptInterface.optimize!` has not been implemented). In that case, optimizer bridges should be used instead.

source

Utilities.@model

`MathOptInterface.Utilities.@model` – Macro.

```
macro model(
    model_name,
    scalar_sets,
    typed_scalar_sets,
    vector_sets,
    typed_vector_sets,
    scalar_functions,
    typed_scalar_functions,
    vector_functions,
    typed_vector_functions,
    is_optimizer = false
)
```

Creates a type `model_name` implementing the MOI model interface and containing `scalar_sets` scalar sets, `typed_scalar_sets` typed scalar sets, `vector_sets` vector sets, `typed_vector_sets` typed vector sets, `scalar_functions` scalar functions, `typed_scalar_functions` typed scalar functions, `vector_functions` vector functions and `typed_vector_functions` typed vector functions. To give no set/function, write `()`, to give one set `S`, write `(S,)`.

The function `MathOptInterface.VariableIndex` should not be given in `scalar_functions`. The model supports `MathOptInterface.VariableIndex`-in-`S` constraints where `S` is `MathOptInterface.EqualTo`, `MathOptInterface.GreaterThan`, `MathOptInterface.LessThan`, `MathOptInterface.Interval`, `MathOptInterface.Integer`, `MathOptInterface.ZeroOne`, `MathOptInterface.Semicontinuous` or `MathOptInterface.Semiinteger`. The sets supported with the `MathOptInterface.VariableIndex` cannot be controlled from the macro, use the `UniversalFallback` to support more sets.

This macro creates a model specialized for specific types of constraint, by defining specialized structures and methods. To create a model that, in addition to be optimized for specific constraints, also support arbitrary constraints and attributes, use `UniversalFallback`.

If `is_optimizer = true`, the resulting struct is a of `GenericOptimizer`, which is a subtype of `MathOptInterface.AbstractOptimizer`, otherwise, it is a `GenericModel`, which is a subtype of `MathOptInterface.ModelLike`.

Examples

The model describing an linear program would be:

```
@model(LPModel,
    (),
    (MOI.EqualTo, MOI.GreaterThan, MOI.LessThan, MOI.Interval),
    (MOI.Zeros, MOI.Nonnegatives, MOI.Nonpositives),
    (),
    (),
    (MOI.ScalarAffineFunction,),
    (MOI.VectorOfVariables,),
    (MOI.VectorAffineFunction,),
    # Name of model
    # untyped scalar sets
    # typed scalar sets
    # untyped vector sets
    # typed vector sets
    # untyped scalar functions
    # typed scalar functions
    # untyped vector functions
    # typed vector functions
```

```

    false
)

```

Let `MOI` denote `MathOptInterface`, `MOIU` denote `MOI.Utilities`. The macro would create the following types with `struct_of_constraint_code`:

```

struct LPModelScalarConstraints{T, C1, C2, C3, C4} <: MOIU.StructOfConstraints
    moi_equalto::C1
    moi_greaterthan::C2
    moi_lessthan::C3
    moi_interval::C4
end
struct LPModelVectorConstraints{T, C1, C2, C3} <: MOIU.StructOfConstraints
    moi_zeros::C1
    moi_nonnegatives::C2
    moi_nonpositives::C3
end
struct LPModelFunctionConstraints{T} <: MOIU.StructOfConstraints
    moi_scalaraffinefunction::LPModelScalarConstraints{
        T,
        MOIU.VectorOfConstraints{MOI.ScalarAffineFunction{T}, MOI.EqualTo{T}},
        MOIU.VectorOfConstraints{MOI.ScalarAffineFunction{T}, MOI.GreaterThan{T}},
        MOIU.VectorOfConstraints{MOI.ScalarAffineFunction{T}, MOI.LessThan{T}},
        MOIU.VectorOfConstraints{MOI.ScalarAffineFunction{T}, MOI.Interval{T}}
    }
    moi_vectorofvariables::LPModelVectorConstraints{
        T,
        MOIU.VectorOfConstraints{MOI.VectorOfVariables, MOI.Zeros},
        MOIU.VectorOfConstraints{MOI.VectorOfVariables, MOI.Nonnegatives},
        MOIU.VectorOfConstraints{MOI.VectorOfVariables, MOI.Nonpositives}
    }
    moi_vectoraffinefunction::LPModelVectorConstraints{
        T,
        MOIU.VectorOfConstraints{MOI.VectorAffineFunction{T}, MOI.Zeros},
        MOIU.VectorOfConstraints{MOI.VectorAffineFunction{T}, MOI.Nonnegatives},
        MOIU.VectorOfConstraints{MOI.VectorAffineFunction{T}, MOI.Nonpositives}
    }
end
const LPModel{T} =
↳ MOIU.GenericModel{T, MOIU.ObjectiveContainer{T}, MOIU.VariablesContainer{T}, LPModelFunctionConstraints{T}}

```

The type `LPModel` implements the `MathOptInterface` API except methods specific to optimizers like `optimize!` or `get` with `VariablePrimal`.

[source](#)

`MathOptInterface.Utilities.GenericModel` – Type.

```
mutable struct GenericModel{T,O,V,C} <: AbstractModelLike{T}
```

Implements a model supporting coefficients of type `T` and:

- An objective function stored in `.objective::O`
- Variables and `VariableIndex` constraints stored in `.variable_bounds::V`
- F-in-S constraints (excluding `VariableIndex` constraints) stored in `.constraints::C`

All interactions should take place via the MOI interface, so the types `O`, `V`, and `C` should implement the API as needed for their functionality.

source

`MathOptInterface.Utilities.GenericOptimizer` – Type.

```
mutable struct GenericOptimizer{T,O,V,C} <: AbstractOptimizer{T}
```

Implements a model supporting coefficients of type `T` and:

- An objective function stored in `.objective::O`
- Variables and `VariableIndex` constraints stored in `.variable_bounds::V`
- F-in-S constraints (excluding `VariableIndex` constraints) stored in `.constraints::C`

All interactions should take place via the MOI interface, so the types `O`, `V`, and `C` should implement the API as needed for their functionality.

source

`.objective` `MathOptInterface.Utilities.ObjectiveContainer` – Type.

```
ObjectiveContainer{T}
```

A helper struct to simplify the handling of objective functions in `Utilities.Model`.

source

`.variables` `MathOptInterface.Utilities.VariablesContainer` – Type.

```
struct VariablesContainer{T} <: AbstractVectorBounds
    set_mask::Vector{UInt16}
    lower::Vector{T}
    upper::Vector{T}
end
```

A struct for storing variables and `VariableIndex`-related constraints. Used in `MOI.Utilities.Model` by default.

source

`MathOptInterface.Utilities.FreeVariables` – Type.

```
mutable struct FreeVariables <: MOI.ModelLike
    n::Int64
    FreeVariables() = new(0)
end
```

A struct for storing free variables that can be used as the `variables` field of `GenericModel` or `GenericModel`. It represents a model that does not support any constraint nor objective function.

Example

The following model type represents a conic model in geometric form. As opposed to `VariablesContainer`, `FreeVariables` does not support constraint bounds so they are bridged into an affine constraint in the `MathOptInterface.Nonnegatives` cone as expected for the geometric conic form.

```

julia> MOI.Utilities.@product_of_sets(
    Cones,
    MOI.Zeros,
    MOI.Nonnegatives,
    MOI.SecondOrderCone,
    MOI.PositiveSemidefiniteConeTriangle,
);

julia> const ConicModel{T} = MOI.Utilities.GenericOptimizer{
    T,
    MOI.Utilities.ObjectiveContainer{T},
    MOI.Utilities.FreeVariables,
    MOI.Utilities.MatrixOfConstraints{
        T,
        MOI.Utilities.MutableSparseMatrixCSC{
            T,
            Int,
            MOI.Utilities.OneBasedIndexing,
        },
        Vector{T},
        Cones{T},
    },
};

julia> model = MOI.instantiate(ConicModel{Float64}, with_bridge_type=Float64);

julia> x = MOI.add_variable(model)
MathOptInterface.VariableIndex(1)

julia> c = MOI.add_constraint(model, x, MOI.GreaterThan(1.0))
MathOptInterface.ConstraintIndex{MathOptInterface.VariableIndex, MathOptInterface.GreaterThan{
    Float64}}(1)

julia> MOI.Bridges.is_bridged(model, c)
true

julia> bridge = MOI.Bridges.bridge(model, c)
MathOptInterface.Bridges.Constraint.VectorizeBridge{Float64, MathOptInterface.
    VectorAffineFunction{Float64}, MathOptInterface.Nonnegatives, MathOptInterface.VariableIndex
}(MathOptInterface.ConstraintIndex{MathOptInterface.VectorAffineFunction{Float64},
    MathOptInterface.Nonnegatives}(1), 1.0)

julia> bridge.vector_constraint
MathOptInterface.ConstraintIndex{MathOptInterface.VectorAffineFunction{Float64}, MathOptInterface.
    .Nonnegatives}(1)

julia> MOI.Bridges.is_bridged(model, bridge.vector_constraint)
false

```

[source](#)

.constraints [MathOptInterface.Utilities.VectorOfConstraints](#) - Type.

```

mutable struct VectorOfConstraints{
    F<:MOI.AbstractFunction,
    S<:MOI.AbstractSet,

```



```

} <: MOI.ModelLike
  constraints::CleverDicts.CleverDict{
    MOI.ConstraintIndex{F,S},
    Tuple{F,S},
    typeof(CleverDicts.key_to_index),
    typeof(CleverDicts.index_to_key),
  }
end

```

A struct storing F-in-S constraints as a mapping between the constraint indices to the corresponding tuple of function and set.

[source](#)

[MathOptInterface.Utilities.StructOfConstraints](#) – Type.

```
| abstract type StructOfConstraints <: MOI.ModelLike end
```

A struct storing a subfields other structs storing constraints of different types.

See [Utilities.@struct_of_constraints_by_function_types](#) and [Utilities.@struct_of_constraints_by_set_types](#).

[source](#)

[MathOptInterface.Utilities.@struct_of_constraints_by_function_types](#) – Macro.

```
| Utilities.@struct_of_constraints_by_function_types(name, func_types...)
```

Given a vector of n function types (F_1, F_2, \dots, F_n) in `func_types`, defines a subtype of `StructOfConstraints` of name `name` and which type parameters $\{T, C_1, C_2, \dots, C_n\}$. It contains n field where the i th field has type C_i and stores the constraints of function type F_i .

The expression F_i can also be a union in which case any constraint for which the function type is in the union is stored in the field with type C_i .

[source](#)

[MathOptInterface.Utilities.@struct_of_constraints_by_set_types](#) – Macro.

```
| Utilities.@struct_of_constraints_by_set_types(name, func_types...)
```

Given a vector of n set types (S_1, S_2, \dots, S_n) in `func_types`, defines a subtype of `StructOfConstraints` of name `name` and which type parameters $\{T, C_1, C_2, \dots, C_n\}$. It contains n field where the i th field has type C_i and stores the constraints of set type S_i . The expression S_i can also be a union in which case any constraint for which the set type is in the union is stored in the field with type C_i . This can be useful if C_i is a [MatrixOfConstraints](#) in order to concatenate the coefficients of constraints of several different set types in the same matrix.

[source](#)

[MathOptInterface.Utilities.struct_of_constraint_code](#) – Function.

```
| struct_of_constraint_code(struct_name, types, field_types = nothing)
```

Given a vector of n `Union{SymbolFun, UnionSymbolFS{SymbolFun}}` or `Union{SymbolSet, UnionSymbolFS{SymbolSet}}` in `types`, defines a subtype of `StructOfConstraints` of name `name` and which type parameters $\{T, F_1, F_2, \dots, F_n\}$ if `field_types` is `nothing` and a $\{T\}$ otherwise. It contains n field where the i th field has type C_i if `field_types` is `nothing` and type `field_types[i]` otherwise. If `types` is vector of

`Union{SymbolFun, _UnionSymbolFS{SymbolFun}}` (resp. `Union{SymbolSet, _UnionSymbolFS{SymbolSet}}`) then the constraints of that function (resp. set) type are stored in the corresponding field.

This function is used by the macros `@model`, `@struct_of_constraints_by_function_types` and `@struct_of_constraints_b`

source

Caching optimizer

`MathOptInterface.Utilities.CachingOptimizer` – Type.

| `CachingOptimizer`

`CachingOptimizer` is an intermediate layer that stores a cache of the model and links it with an optimizer. It supports incremental model construction and modification even when the optimizer doesn't.

Constructors

| `CachingOptimizer(cache::MOI.ModelLike, optimizer::AbstractOptimizer)`

Creates a `CachingOptimizer` in `AUTOMATIC` mode, with the optimizer `optimizer`.

The type of the optimizer returned is `CachingOptimizer{typeof(optimizer), typeof(cache)}` so it does not support the function `reset_optimizer(::CachingOptimizer, new_optimizer)` if the type of `new_optimizer` is different from the type of `optimizer`.

| `CachingOptimizer(cache::MOI.ModelLike, mode::CachingOptimizerMode)`

Creates a `CachingOptimizer` in the `NO_OPTIMIZER` state and mode `mode`.

The type of the optimizer returned is `CachingOptimizer{MOI.AbstractOptimizer, typeof(cache)}` so it does support the function `reset_optimizer(::CachingOptimizer, new_optimizer)` if the type of `new_optimizer` is different from the type of `optimizer`.

About the type

States

A `CachingOptimizer` may be in one of three possible states (`CachingOptimizerState`):

- `NO_OPTIMIZER`: The `CachingOptimizer` does not have any optimizer.
- `EMPTY_OPTIMIZER`: The `CachingOptimizer` has an empty optimizer. The optimizer is not synchronized with the cached model.
- `ATTACHED_OPTIMIZER`: The `CachingOptimizer` has an optimizer, and it is synchronized with the cached model.

Modes

A `CachingOptimizer` has two modes of operation (`CachingOptimizerMode`):

- `MANUAL`: The only methods that change the state of the `CachingOptimizer` are `Utilities.reset_optimizer`, `Utilities.drop_optimizer`, and `Utilities.attach_optimizer`. Attempting to perform an operation in the incorrect state results in an error.
- `AUTOMATIC`: The `CachingOptimizer` changes its state when necessary. For example, `optimize!` will automatically call `attach_optimizer` (an optimizer must have been previously set). Attempting to add a constraint or perform a modification not supported by the optimizer results in a drop to `EMPTY_OPTIMIZER` mode.

source

[MathOptInterface.Utilities.attach_optimizer](#) – Function.

```
| attach_optimizer(model::CachingOptimizer)
```

Attaches the optimizer to `model`, copying all model data into it. Can be called only from the `EMPTY_OPTIMIZER` state. If the copy succeeds, the `CachingOptimizer` will be in state `ATTACHED_OPTIMIZER` after the call, otherwise an error is thrown; see [MathOptInterface.copy_to](#) for more details on which errors can be thrown.

source

[MathOptInterface.Utilities.reset_optimizer](#) – Function.

```
| reset_optimizer(m::CachingOptimizer, optimizer::MOI.AbstractOptimizer)
```

Sets or resets `m` to have the given empty optimizer `optimizer`.

Can be called from any state. An assertion error will be thrown if `optimizer` is not empty.

The `CachingOptimizer` `m` will be in state `EMPTY_OPTIMIZER` after the call.

source

```
| reset_optimizer(m::CachingOptimizer)
```

Detaches and empties the current optimizer. Can be called from `ATTACHED_OPTIMIZER` or `EMPTY_OPTIMIZER` state. The `CachingOptimizer` will be in state `EMPTY_OPTIMIZER` after the call.

source

[MathOptInterface.Utilities.drop_optimizer](#) – Function.

```
| drop_optimizer(m::CachingOptimizer)
```

Drops the optimizer, if one is present. Can be called from any state. The `CachingOptimizer` will be in state `NO_OPTIMIZER` after the call.

source

[MathOptInterface.Utilities.state](#) – Function.

```
| state(m::CachingOptimizer)::CachingOptimizerState
```

Returns the state of the `CachingOptimizer` `m`. See [Utilities.CachingOptimizer](#).

source

[MathOptInterface.Utilities.mode](#) – Function.

```
| mode(m::CachingOptimizer)::CachingOptimizerMode
```

Returns the operating mode of the `CachingOptimizer` `m`. See [Utilities.CachingOptimizer](#).

source

Mock optimizer

[MathOptInterface.Utilities.MockOptimizer](#) – Type.

```
| MockOptimizer
```

`MockOptimizer` is a fake optimizer especially useful for testing. Its main feature is that it can store the values that should be returned for each attribute.

source

Printing

`MathOptInterface.Utilities.latex_formulation` – Function.

```
| latex_formulation(model::MOI.ModelLike; kwargs...)
```

Wrap model in a type so that it can be pretty-printed as text/latex in a notebook like IJulia, or in Documenter.

To render the model, end the cell with `latex_formulation(model)`, or call `display(latex_formulation(model))` in to force the display of the model from inside a function.

Possible keyword arguments are:

- `simplify_coefficients` : Simplify coefficients if possible by omitting them or removing trailing zeros.
- `default_name` : The name given to variables with an empty name.
- `print_types` : Print the MOI type of each function and set for clarity.

[source](#)

Copy utilities

`MathOptInterface.Utilities.default_copy_to` – Function.

```
| default_copy_to(dest::MOI.ModelLike, src::MOI.ModelLike)
```

A default implementation of `MOI.copy_to(dest, src)` for models that implement the incremental interface, i.e., `MOI.supports_incremental_interface` returns true.

[source](#)

`MathOptInterface.Utilities.IndexMap` – Type.

```
| IndexMap()
```

The dictionary-like object returned by `MathOptInterface.copy_to`.

[source](#)

`MathOptInterface.Utilities.identity_index_map` – Function.

```
| identity_index_map(model::MOI.ModelLike)
```

Return an `IndexMap` that maps all variable and constraint indices of model to themselves.

[source](#)

`MathOptInterface.Utilities.ModelFilter` – Type.

```
| ModelFilter(filter::Function, model::MOI.ModelLike)
```

A layer to filter out various components of model.

The filter function takes a single argument, which is each element from the list returned by the attributes below. It returns true if the element should be visible in the filtered model and false otherwise.

The components that are filtered are:

- Entire constraint types via:
 - `MOI.ListOfConstraintTypesPresent`
- Individual constraints via:
 - `MOI.ListOfConstraintIndices{F,S}`
- Specific attributes via:
 - `MOI.ListOfModelAttributeSet`
 - `MOI.ListOfConstraintAttributesSet`
 - `MOI.ListOfVariableAttributesSet`

Warning

The list of attributes filtered may change in a future release. You should write functions that are generic and not limited to the five types listed above. Thus, you should probably define a fallback filter(`::Any`) = `true`.

See below for examples of how this works.

Note

This layer has a limited scope. It is intended to be used in conjunction with `MOI.copy_to`.

Example: copy model excluding integer constraints

Use the `do` syntax to provide a single function.

```
filtered_src = MOI.Utilities.ModelFilter(src) do item
    return item != (MOI.VariableIndex, MOI.Integer)
end
MOI.copy_to(dest, filtered_src)
```

Example: copy model excluding names

Use type dispatch to simplify the implementation:

```
my_filter(::Any) = true # Note the generic fallback!
my_filter(::MOI.VariableName) = false
my_filter(::MOI.ConstraintName) = false
filtered_src = MOI.Utilities.ModelFilter(my_filter, src)
MOI.copy_to(dest, filtered_src)
```

Example: copy irreducible infeasible subsystem

```
my_filter(::Any) = true # Note the generic fallback!
function my_filter(ci::MOI.ConstraintIndex)
    status = MOI.get(dest, MOI.ConstraintConflictStatus(), ci)
    return status != MOI.NOT_IN_CONFLICT
end
filtered_src = MOI.Utilities.ModelFilter(my_filter, src)
MOI.copy_to(dest, filtered_src)
```

[source](#)

MatrixOfConstraints

`MathOptInterface.Utilities.MatrixOfConstraints` – Type.

```
mutable struct MatrixOfConstraints{T,AT,BT,ST} <: MOI.ModelLike
    coefficients::AT
    constants::BT
    sets::ST
    caches::Vector{Any}
    are_indices_mapped::Vector{BitSet}
    final_touch::Bool
end
```

Represent `ScalarAffineFunction` and `VectorAffineFunction` constraints in a matrix form where the linear coefficients of the functions are stored in the `coefficients` field, the constants of the functions or sets are stored in the `constants` field. Additional information about the sets are stored in the `sets` field.

This model can only be used as the constraints field of a `MOI.Utilities.AbstractModel`.

When the constraints are added, they are stored in the `caches` field. They are only loaded in the `coefficients` and `constants` fields once `MOI.Utilities.final_touch` is called. For this reason, `MatrixOfConstraints` should not be used by an incremental interface. Use `MOI.copy_to` instead.

The constraints can be added in two different ways:

1. With `add_constraint`, in which case a canonicalized copy of the function is stored in `caches`.
2. With `pass_nonvariable_constraints`, in which case the functions and sets are stored themselves in `caches` without mapping the variable indices. The corresponding index in `caches` is added in `are_indices_mapped`. This avoids doing a copy of the function in case the getter of `CanonicalConstraintFunction` does not make a copy for the source model, e.g., this is the case of `VectorOfConstraints`.

We illustrate this with an example. Suppose a model is copied from a `src::MOI.Utilities.Model` to a bridged model with a `MatrixOfConstraints`. For all the types that are not bridged, the constraints will be copied with `pass_nonvariable_constraints`. Hence the functions stored in `caches` are exactly the same as the ones stored in `src`. This is ok since this is only during the `copy_to` operation during which `src` cannot be modified. On the other hand, for the types that are bridged, the functions added may contain duplicates even if the functions did not contain duplicates in `src` so duplicates are removed with `MOI.Utilities.canonical`.

Interface

The `.coefficients::AT` type must implement:

- `AT()`
- `MOI.empty(::AT)!`
- `MOI.Utilities.add_column`
- `MOI.Utilities.set_number_of_rows`
- `MOI.Utilities.allocate_terms`
- `MOI.Utilities.load_terms`
- `MOI.Utilities.final_touch`

The `.constants::BT` type must implement:

- `BT()`

- `Base.empty! (::BT)`
- `Base.resize (::BT)`
- `MOI.Utilities.load_constants`
- `MOI.Utilities.function_constants`
- `MOI.Utilities.set_from_constants`

The `.sets :: ST` type must implement:

- `ST()`
- `MOI.is_empty (::ST)`
- `MOI.empty (::ST)`
- `MOI.dimension (::ST)`
- `MOI.is_valid (::ST, ::MOI.ConstraintIndex)`
- `MOI.get (::ST, ::MOI.ListOfConstraintTypesPresent)`
- `MOI.get (::ST, ::MOI.NumberOfConstraints)`
- `MOI.get (::ST, ::MOI.ListOfConstraintIndices)`
- `MOI.Utilities.set_types`
- `MOI.Utilities.set_index`
- `MOI.Utilities.add_set`
- `MOI.Utilities.rows`
- `MOI.Utilities.final_touch`

[source](#)

`.coefficients` `MathOptInterface.Utilities.add_column` – Function.

```
| add_column(coefficients) :: Nothing
```

Tell coefficients to pre-allocate datastructures as needed to store one column.

[source](#)

`MathOptInterface.Utilities.allocate_terms` – Function.

```
| allocate_terms(coefficients, index_map, func) :: Nothing
```

Tell coefficients that the terms of the function `func` where the variable indices are mapped with `index_map` will be loaded with `load_terms`.

The function `func` must be canonicalized before calling `allocate_terms`. See [is_canonical](#).

[source](#)

`MathOptInterface.Utilities.set_number_of_rows` – Function.

```
| set_number_of_rows(coefficients, n) :: Nothing
```

Tell coefficients to pre-allocate datastructures as needed to store `n` rows.

[source](#)

[MathOptInterface.Utilities.load_terms](#) – Function.

```
| load_terms(coefficients, index_map, func, offset)::Nothing
```

Loads the terms of func to coefficients, mapping the variable indices with index_map.

The ith dimension of func is loaded at the (offset + i)th row of coefficients.

The function must be allocated first with [allocate_terms](#).

The function func must be canonicalized, see [is_canonical](#).

[source](#)

[MathOptInterface.Utilities.final_touch](#) – Function.

```
| final_touch(coefficients)::Nothing
```

Informs the coefficients that all functions have been added with load_terms. No more modification is allowed unless MOI.empty! is called.

```
| final_touch(sets)::Nothing
```

Informs the sets that all functions have been added with add_set. No more modification is allowed unless MOI.empty! is called.

[source](#)

[MathOptInterface.Utilities.extract_function](#) – Function.

```
| extract_function(coefficients, row::Integer, constant::T) where {T}
```

Return the MOI.ScalarAffineFunction{T} function corresponding to row row in coefficients.

```
| extract_function(
    coefficients,
    rows::UnitRange,
    constants::Vector{T},
) where {T}
```

Return the MOI.VectorAffineFunction{T} function corresponding to rows rows in coefficients.

[source](#)

[MathOptInterface.Utilities.MutableSparseMatrixCSC](#) – Type.

```
| mutable struct MutableSparseMatrixCSC{Tv,Ti<:Integer,I<:AbstractIndexing}
    indexing::I
    m::Int
    n::Int
    colptr::Vector{Ti}
    rowval::Vector{Ti}
    nzval::Vector{Tv}
    nz_added::Vector{Ti}
end
```

Matrix type loading sparse matrices in the Compressed Sparse Column format. The indexing used is indexing, see [AbstractIndexing](#). The other fields have the same meaning than for SparseArrays.SparseMatrixCSC except that the indexing is different unless indexing is OneBasedIndexing. In addition, nz_added is used to cache the number of non-zero terms that have been added to each column due to the incremental nature of load_terms.

The matrix is loaded in 5 steps:

1. `MOI.empty!` is called.
2. `MOI.Utilities.add_column` and `MOI.Utilities.allocate_terms` are called in any order.
3. `MOI.Utilities.set_number_of_rows` is called.
4. `MOI.Utilities.load_terms` is called for each affine function.
5. `MOI.Utilities.final_touch` is called.

[source](#)

`MathOptInterface.Utilities.AbstractIndexing` - Type.

```
| abstract type AbstractIndexing end
```

Indexing to be used for storing the row and column indices of `MutableSparseMatrixCSC`. See [ZeroBasedIndexing](#) and [OneBasedIndexing](#).

[source](#)

`MathOptInterface.Utilities.ZeroBasedIndexing` - Type.

```
| struct ZeroBasedIndexing <: AbstractIndexing end
```

Zero-based indexing: the i th row or column has index $i - 1$. This is useful when the vectors of row and column indices need to be communicated to a library using zero-based indexing such as C libraries.

[source](#)

`MathOptInterface.Utilities.OneBasedIndexing` - Type.

```
| struct ZeroBasedIndexing <: AbstractIndexing end
```

One-based indexing: the i th row or column has index i . This enables an allocation-free conversion of `MutableSparseMatrixCSC` to `SparseArrays.SparseMatrixCSC`.

[source](#)

.constants `MathOptInterface.Utilities.load_constants` - Function.

```
| load_constants(constants, offset, func_or_set)::Nothing
```

This function loads the constants of `func_or_set` in `constants` at an offset of `offset`. Where `offset` is the sum of the dimensions of the constraints already loaded. The storage should be preallocated with `resize!` before calling this function.

This function should be implemented to be usable as storage of constants for [MatrixOfConstraints](#).

The constants are loaded in three steps:

1. `Base.empty!` is called.
2. `Base.resize!` is called with the sum of the dimensions of all constraints.
3. `MOI.Utilities.load_constants` is called for each function for vector constraint or set for scalar constraint.

[source](#)

`MathOptInterface.Utilities.function_constants` - Function.

```
| function_constants(constants, rows)
```

This function returns the function constants that were loaded with `load_constants` at the rows `rows`.

This function should be implemented to be usable as storage of constants for `MatrixOfConstraints`.

[source](#)

`MathOptInterface.Utilities.set_from_constants` – Function.

```
| set_from_constants(constants, S::Type, rows)::S
```

This function returns an instance of the set `S` for which the constants were loaded with `load_constants` at the rows `rows`.

This function should be implemented to be usable as storage of constants for `MatrixOfConstraints`.

[source](#)

`MathOptInterface.Utilities.Hyperrectangle` – Type.

```
| struct Hyperrectangle{T} <: AbstractVectorBounds
|   lower::Vector{T}
|   upper::Vector{T}
| end
```

A struct for the `.constants` field in `MatrixOfConstraints`.

[source](#)

`.sets MathOptInterface.Utilities.set_index` – Function.

```
| set_index(sets, ::Type{S})::Union{Int,Nothing} where {S<:MOI.AbstractSet}
```

Return an integer corresponding to the index of the set type in the list given by `set_types`.

If `S` is not part of the list, return nothing.

[source](#)

`MathOptInterface.Utilities.set_types` – Function.

```
| set_types(sets)::Vector{Type}
```

Return the list of the types of the sets allowed in `sets`.

[source](#)

`MathOptInterface.Utilities.add_set` – Function.

```
| add_set(sets, i)::Int64
```

Add a scalar set of type index `i`.

```
| add_set(sets, i, dim)::Int64
```

Add a vector set of type index `i` and dimension `dim`.

Both methods return a unique `Int64` of the set that can be used to reference this set.

[source](#)

`MathOptInterface.Utilities.rows` – Function.

```
| rows(sets, ci::MOI.ConstraintIndex)::Union{Int,UnitRange{Int}}
```

Return the rows in `1:MOI.dimension(sets)` corresponding to the set of id `ci.value`.

For scalar sets, this returns an `Int`. For vector sets, this returns an `UnitRange{Int}`.

[source](#)

`MathOptInterface.Utilities.num_rows` – Function.

```
| num_rows(sets::OrderedProductOfSets, ::Type{S}) where {S}
```

Return the number of rows corresponding to a set of type `S`. That is, it is the sum of the dimensions of the sets of type `S`.

[source](#)

`MathOptInterface.Utilities.set_with_dimension` – Function.

```
| set_with_dimension(::Type{S}, dim) where {S<:MOI.AbstractVectorSet}
```

Returns the instance of `S` of `MathOptInterface.dimension` `dim`. This needs to be implemented for sets of type `S` to be useable with `MatrixOfConstraints`.

[source](#)

`MathOptInterface.Utilities.ProductOfSets` – Type.

```
| abstract type ProductOfSets{T} end
```

Represents a cartesian product of sets of given types.

[source](#)

`MathOptInterface.Utilities.MixOfScalarSets` – Type.

```
| abstract type MixOfScalarSets{T} <: ProductOfSets{T} end
```

Product of scalar sets in the order the constraints are added, mixing the constraints of different types.

Use `@mix_of_scalar_sets` to generate a new subtype.

[source](#)

`MathOptInterface.Utilities.@mix_of_scalar_sets` – Macro.

```
| @mix_of_scalar_sets(name, set_types...)
```

Generate a new `MixOfScalarSets` subtype.

Example

```
| @mix_of_scalar_sets(
|     MixedIntegerLinearProgramSets,
|     MOI.GreaterThan{T},
|     MOI.LessThan{T},
|     MOI.EqualTo{T},
|     MOI.Integer,
| )
```

source

MathOptInterface.Utilities.OrderedProductOfSets - Type.

```
abstract type OrderedProductOfSets{T} <: ProductOfSets{T} end
```

Product of sets in the order the constraints are added, grouping the constraints of the same types contiguously.

Use `@product_of_sets` to generate new subtypes.

source

MathOptInterface.Utilities.@product_of_sets - Macro.

```
@product_of_sets(name, set_types...)
```

Generate a new `OrderedProductOfSets` subtype.

Example

```
@product_of_sets(
    LinearOrthants,
    MOI.Zeros,
    MOI.Nonnegatives,
    MOI.Nonpositives,
    MOI.ZeroOne,
)
```

source

Fallbacks

[MathOptInterface.Utilities.get_fallback](#) - Function.

```
get_fallback(model::MOI.ModelLike, ::MOI.ObjectiveValue)
```

Compute the objective function value using the `VariablePrimal` results and the `ObjectiveFunction` value.

source

```
get_fallback(model::MOI.ModelLike, ::MOI.DualObjectiveValue, T::Type)::T
```

Compute the dual objective value of type T using the ConstraintDual results and the ConstraintFunction and ConstraintSet values. Note that the nonlinear part of the model is ignored.

source

```
get_fallback(model::MOI.ModelLike, ::MOI.ConstraintPrimal,  
            constraint_index::MOI.ConstraintIndex)
```

Compute the value of the function of the constraint of index `constraint_index` using the `VariablePrimal` results and the `ConstraintFunction` values.

source

[illegible]

Compute the dual of the constraint of index `ci` using the `ConstraintDual` of other constraints and the `ConstraintFunction` values. Throws an error if some constraints are quadratic or if there is one another `MOI.VariableIndex-in-S` or `MOI.VectorOfVariables-in-S` constraint with one of the variables in the function of the constraint `ci`.

[source](#)

Function utilities

The following utilities are available for functions:

`MathOptInterface.Utilities.eval_variables` – Function.

```
| eval_variables(varval::Function, f::AbstractFunction)
```

Returns the value of function `f` if each variable index `vi` is evaluated as `varval(vi)`. Note that `varval` should return a number, see [substitute_variables](#) for a similar function where `varval` returns a function.

[source](#)

`MathOptInterface.Utilities.map_indices` – Function.

```
| map_indices(index_map::Function, attr::MOI.AnyAttribute, x::X)::X where {X}
```

Substitute any `MOI.VariableIndex` (resp. `MOI.ConstraintIndex`) in `x` by the `MOI.VariableIndex` (resp. `MOI.ConstraintIndex`) of the same type given by `index_map(x)`.

When to implement this method for new types `X`

This function is used by implementations of `MOI.copy_to` on constraint functions, attribute values and submittable values. If you define a new attribute whose values `x::X` contain variable or constraint indices, you must also implement this function.

[source](#)

```
| map_indices(
|     variable_map::AbstractDict{T,T},
|     x::X,
| )::X where {T<:MOI.Index,X}
```

Shortcut for `map_indices(vi -> variable_map[vi], x)`.

[source](#)

`MathOptInterface.Utilities.substitute_variables` – Function.

```
| substitute_variables(variable_map::Function, x)
```

Substitute any `MOI.VariableIndex` in `x` by `variable_map(x)`. The `variable_map` function returns either `MOI.VariableIndex` or `MOI.ScalarAffineFunction`, see [eval_variables](#) for a similar function where `variable_map` returns a number.

This function is used by bridge optimizers on constraint functions, attribute values and submittable values when at least one variable bridge is used hence it needs to be implemented for custom types that are meant to be used as attribute or submittable value.

WARNING: Don't use `substitute_variables(::Function, ...)` because Julia will not specialize on this. Use instead `substitute_variables(::F, ...)` where `{F<:Function}`.

[source](#)

`MathOptInterface.Utilities.filter_variables` – Function.

```
| filter_variables(keep::Function, f::AbstractFunction)
```

Return a new function `f` with the variable `vi` such that `!keep(vi)` removed.

WARNING: Don't define `filter_variables(::Function, ...)` because Julia will not specialize on this. Define instead `filter_variables(::F, ...)` where `{F<:Function}`.

[source](#)

`MathOptInterface.Utilities.remove_variable` – Function.

```
| remove_variable(f::AbstractFunction, vi::VariableIndex)
```

Return a new function `f` with the variable `vi` removed.

[source](#)

```
| remove_variable(f::MOI.AbstractFunction, s::MOI.AbstractSet, vi::MOI.VariableIndex)
```

Return a tuple `(g, t)` representing the constraint `f`-in-`s` with the variable `vi` removed. That is, the terms containing the variable `vi` in the function `f` are removed and the dimension of the set `s` is updated if needed (e.g. when `f` is a `VectorOfVariables` with `vi` being one of the variables).

[source](#)

`MathOptInterface.Utilities.all_coefficients` – Function.

```
| all_coefficients(p::Function, f::MOI.AbstractFunction)
```

Determine whether predicate `p` returns true for all coefficients of `f`, returning false as soon as the first coefficient of `f` for which `p` returns false is encountered (short-circuiting). Similar to `all`.

[source](#)

`MathOptInterface.Utilities.unsafe_add` – Function.

```
| unsafe_add(t1::MOI.ScalarAffineTerm, t2::MOI.ScalarAffineTerm)
```

Sums the coefficients of `t1` and `t2` and returns an output `MOI.ScalarAffineTerm`. It is unsafe because it uses the variable of `t1` as the variable of the output without checking that it is equal to that of `t2`.

[source](#)

```
| unsafe_add(t1::MOI.ScalarQuadraticTerm, t2::MOI.ScalarQuadraticTerm)
```

Sums the coefficients of `t1` and `t2` and returns an output `MOI.ScalarQuadraticTerm`. It is unsafe because it uses the variable's of `t1` as the variable's of the output without checking that they are the same (up to permutation) to those of `t2`.

[source](#)

```
| unsafe_add(t1::MOI.VectorAffineTerm, t2::MOI.VectorAffineTerm)
```

Sums the coefficients of `t1` and `t2` and returns an output `MOI.VectorAffineTerm`. It is unsafe because it uses the `output_index` and `variable` of `t1` as the `output_index` and `variable` of the output term without checking that they are equal to those of `t2`.

[source](#)

`MathOptInterface.Utilities.isapprox_zero` – Function.

```
| isapprox_zero(f::MOI.AbstractFunction, tol)
```

Return a Bool indicating whether the function `f` is approximately zero using `tol` as a tolerance.

Important note

This function assumes that `f` does not contain any duplicate terms, you might want to first call `canonical` if that is not guaranteed. For instance, given

```
| f = MOI.ScalarAffineFunction(MOI.ScalarAffineTerm{Float64}([1, -1], [x, x]), 0)`.
```

then `isapprox_zero(f)` is false but `isapprox_zero(MOIU.canonical(f))` is true.

[source](#)

`MathOptInterface.Utilities.modify_function` – Function.

```
| modify_function(f::AbstractFunction, change::AbstractFunctionModification)
```

Return a new function `f` modified according to `change`.

[source](#)

`MathOptInterface.Utilities.zero_with_output_dimension` – Function.

```
| zero_with_output_dimension(::Type{T}, output_dimension::Integer) where {T}
```

Create an instance of type `T` with the output dimension `output_dimension`.

This is mostly useful in Bridges, when code needs to be agnostic to the type of vector-valued function that is passed in.

[source](#)

The following functions can be used to canonicalize a function:

`MathOptInterface.Utilities.is_canonical` – Function.

```
| is_canonical(f::Union{ScalarAffineFunction, VectorAffineFunction})
```

Returns a Bool indicating whether the function is in canonical form. See `canonical`.

[source](#)

```
| is_canonical(f::Union{ScalarQuadraticFunction, VectorQuadraticFunction})
```

Returns a Bool indicating whether the function is in canonical form. See `canonical`.

[source](#)

`MathOptInterface.Utilities.canonical` – Function.

```
| canonical(
    f::Union{
        ScalarAffineFunction,
        VectorAffineFunction,
        ScalarQuadraticFunction,
        VectorQuadraticFunction,
    },
)
```

Returns the function in a canonical form, i.e.

- A term appear only once.
- The coefficients are nonzero.
- The terms appear in increasing order of variable where there the order of the variables is the order of their value.
- For a `AbstractVectorFunction`, the terms are sorted in ascending order of output index.

The output of `canonical` can be assumed to be a copy of `f`, even for `VectorOfVariables`.

Examples

If `x` (resp. `y`, `z`) is `VariableIndex(1)` (resp. `2`, `3`). The canonical representation of `ScalarAffineFunction([y, x, z, x, z], [2, 1, 3, -2, -3], 5)` is `ScalarAffineFunction([x, y], [-1, 2], 5)`.

[source](#)

`MathOptInterface.Utilities.canonicalize!` – Function.

```
| canonicalize!(f::Union{ScalarAffineFunction, VectorAffineFunction})
```

Convert a function to canonical form in-place, without allocating a copy to hold the result. See [canonical](#).

[source](#)

```
| canonicalize!(f::Union{ScalarQuadraticFunction, VectorQuadraticFunction})
```

Convert a function to canonical form in-place, without allocating a copy to hold the result. See [canonical](#).

[source](#)

The following functions can be used to manipulate functions with basic algebra:

`MathOptInterface.Utilities.scalar_type` – Function.

```
| scalar_type(F::Type{<:MOI.AbstractVectorFunction})
```

Type of functions obtained by indexing objects obtained by calling `eachscalar` on functions of type `F`.

[source](#)

`MathOptInterface.Utilities.scalarize` – Function.

```
| scalarize(func::MOI.VectorOfVariables, ignore_constants::Bool = false)
```

Returns a vector of scalar functions making up the vector function in the form of a `Vector{MOI.SingleVariable}`.

See also [eachscalar](#).

[source](#)

```
| scalarize(func::MOI.VectorAffineFunction{T}, ignore_constants::Bool = false)
```

Returns a vector of scalar functions making up the vector function in the form of a `Vector{MOI.ScalarAffineFunction{T}}`.

See also [eachscalar](#).

[source](#)

```
| scalarize(func::MOI.VectorQuadraticFunction{T}, ignore_constants::Bool = false)
```


Returns a vector of scalar functions making up the vector function in the form of a `Vector{MOI.ScalarQuadraticFunction{T}}`

See also [eachscalar](#).

[source](#)

[MathOptInterface.Utilities.eachscalar](#) – Function.

```
| eachscalar(f::MOI.AbstractVectorFunction)
```

Returns an iterator for the scalar components of the vector function.

See also [scalarize](#).

[source](#)

```
| eachscalar(f::MOI.AbstractVector)
```

Returns an iterator for the scalar components of the vector.

[source](#)

[MathOptInterface.Utilities.promote_operation](#) – Function.

```
| promote_operation(
|     op::Function,
|     ::Type{T},
|     ArgsTypes::Type{<:Union{T, MOI.AbstractFunction}}...,
| ) where {T}
```

Returns the type of the `MOI.AbstractFunction` returned to the call `operate(op, T, args...)` where the types of the arguments `args` are `ArgsTypes`.

[source](#)

[MathOptInterface.Utilities.operate](#) – Function.

```
| operate(
|     op::Function,
|     ::Type{T},
|     args::Union{T, MOI.AbstractFunction}...,
| )::MOI.AbstractFunction where {T}
```

Returns an `MOI.AbstractFunction` representing the function resulting from the operation `op(args...)` on functions of coefficient type `T`. No argument can be modified.

[source](#)

[MathOptInterface.Utilities.operate!](#) – Function.

```
| operate!(
|     op::Function,
|     ::Type{T},
|     args::Union{T, MOI.AbstractFunction}...,
| )::MOI.AbstractFunction where {T}
```

Returns an `MOI.AbstractFunction` representing the function resulting from the operation `op(args...)` on functions of coefficient type `T`. The first argument can be modified. The return type is the same than the method `operate(op, T, args...)` without `!`.

[source](#)

[MathOptInterface.Utilities.operate_output_index!](#) – Function.

```
operate_output_index!(
  op::Function,
  ::Type{T},
  output_index::Integer,
  func::MOI.AbstractVectorFunction
  args::Union{T, MOI.AbstractScalarFunction}...
)::MOI.AbstractFunction where {T}
```

Returns an `MOI.AbstractVectorFunction` where the function at `output_index` is the result of the operation `op` applied to the function at `output_index` of `func` and `args`. The functions at output index different to `output_index` are the same as the functions at the same output index in `func`. The first argument can be modified.

[source](#)

[MathOptInterface.Utilities.vectorize](#) – Function.

```
vectorize(x::AbstractVector{MOI.VariableIndex})
```

Returns the vector of scalar affine functions in the form of a `MOI.VectorAffineFunction{T}`.

[source](#)

```
vectorize(funcs::AbstractVector{MOI.ScalarAffineFunction{T}}) where T
```

Returns the vector of scalar affine functions in the form of a `MOI.VectorAffineFunction{T}`.

[source](#)

```
vectorize(funcs::AbstractVector{MOI.ScalarQuadraticFunction{T}}) where T
```

Returns the vector of scalar quadratic functions in the form of a `MOI.VectorQuadraticFunction{T}`.

[source](#)

Constraint utilities

The following utilities are available for moving the function constant to the set for scalar constraints:

[MathOptInterface.Utilities.shift_constant](#) – Function.

```
shift_constant(set::MOI.AbstractScalarSet, offset)
```

Returns a new scalar set `new_set` such that `func-in-set` is equivalent to `func + offset-in-new_set`.

Only define this function if it makes sense to!

Use [supports_shift_constant](#) to check if the set supports shifting:

```
if supports_shift_constant(typeof(old_set))
  new_set = shift_constant(old_set, offset)
  f.constant = 0
  add_constraint(model, f, new_set)
else
  add_constraint(model, f, old_set)
end
```

See also [supports_shift_constant](#).

Examples

The call `shift_constant(MOI.Interval(-2, 3), 1)` is equal to `MOI.Interval(-1, 4)`.

[source](#)

[MathOptInterface.Utilities.supports_shift_constant](#) – Function.

```
| supports_shift_constant(::Type{S}) where {S<:MOI.AbstractSet}
```

Return true if [shift_constant](#) is defined for set S.

See also [shift_constant](#).

[source](#)

[MathOptInterface.Utilities.normalize_and_add_constraint](#) – Function.

```
| normalize_and_add_constraint(  
|     model::MOI.ModelLike,  
|     func::MOI.AbstractScalarFunction,  
|     set::MOI.AbstractScalarSet;  
|     allow_modify_function::Bool = false,  
| )
```

Adds the scalar constraint obtained by moving the constant term in `func` to the set in `model`. If `allow_modify_function` is true then the function `func` can be modified.

[source](#)

[MathOptInterface.Utilities.normalize_constant](#) – Function.

```
| normalize_constant(  
|     func::MOI.AbstractScalarFunction,  
|     set::MOI.AbstractScalarSet;  
|     allow_modify_function::Bool = false,  
| )
```

Return the func-in-set constraint in normalized form. That is, if `func` is [MOI.ScalarQuadraticFunction](#) or [MOI.ScalarAffineFunction](#), the constant is moved to the set. If `allow_modify_function` is true then the function `func` can be modified.

[source](#)

The following utility identifies those constraints imposing bounds on a given variable, and returns those bound values:

[MathOptInterface.Utilities.get_bounds](#) – Function.

```
| get_bounds(model::MOI.ModelLike, ::Type{T}, x::MOI.VariableIndex)
```

Return a tuple (lb, ub) of type `Tuple{T, T}`, where lb and ub are lower and upper bounds, respectively, imposed on x in `model`.

[source](#)

The following utilities are useful when working with symmetric matrix cones.

[MathOptInterface.Utilities.is_diagonal_vectorized_index](#) – Function.

```
| is_diagonal_vectorized_index(index::Base.Integer)
```

Return whether index is the index of a diagonal element in a [MOI.AbstractSymmetricMatrixSetTriangle](#) set.

[source](#)

[MathOptInterface.Utilities.side_dimension_for_vectorized_dimension](#) – Function.

```
| side_dimension_for_vectorized_dimension(n::Integer)
```

Return the dimension d such that `MOI.dimension(MOI.PositiveSemidefiniteConeTriangle(d))` is n.

[source](#)

DoubleDicts

[MathOptInterface.Utilities.DoubleDicts.DoubleDict](#) – Type.

```
| DoubleDict{V}
```

An optimized dictionary to map `MOI.ConstraintIndex` to values of type `V`.

Works as a `AbstractDict{MOI.ConstraintIndex, V}` with minimal differences.

If `V` is also a `MOI.ConstraintIndex`, use [IndexDoubleDict](#).

Note that `MOI.ConstraintIndex` is not a concrete type, opposed to `MOI.ConstraintIndex{MOI.VariableIndex, MOI.Integers}`, which is a concrete type.

When looping through multiple keys of the same Function-in-Set type, use

```
| inner = dict[F, S]
```

to return a type-stable [DoubleDictInner](#).

[source](#)

[MathOptInterface.Utilities.DoubleDicts.DoubleDictInner](#) – Type.

```
| DoubleDictInner{F,S,V}
```

A type stable inner dictionary of [DoubleDict](#).

[source](#)

[MathOptInterface.Utilities.DoubleDicts.IndexDoubleDict](#) – Type.

```
| IndexDoubleDict
```

A specialized version of `[DoubleDict]` in which the values are of type `MOI.ConstraintIndex`

When looping through multiple keys of the same Function-in-Set type, use

```
| inner = dict[F, S]
```

to return a type-stable [IndexDoubleDictInner](#).

[source](#)

[MathOptInterface.Utilities.DoubleDicts.IndexDoubleDictInner](#) – Type.

| `IndexDoubleDictInner{F,S}`

A type stable inner dictionary of [IndexDoubleDict](#).

[source](#)

Chapter 31

Test

31.1 Overview

The Test submodule

The Test submodule provides tools to help solvers implement unit tests in order to ensure they implement the MathOptInterface API correctly, and to check for solver-correctness.

We use a centralized repository of tests, so that if we find a bug in one solver, instead of adding a test to that particular repository, we add it here so that all solvers can benefit.

How to test a solver

The skeleton below can be used for the wrapper test file of a solver named FooBar.

```
# ===== /test/MOI_wrapper.jl =====
module TestFooBar

import FooBar
using MathOptInterface
using Test

const MOI = MathOptInterface

const OPTIMIZER = MOI.instantiate(
    MOI.OptimizerWithAttributes(FooBar.Optimizer, MOI.Silent() => true),
)

const BRIDGED = MOI.instantiate(
    MOI.OptimizerWithAttributes(FooBar.Optimizer, MOI.Silent() => true),
    with_bridge_type = Float64,
)

# See the docstring of MOI.Test.Config for other arguments.
const CONFIG = MOI.Test.Config(
    # Modify tolerances as necessary.
    atol = 1e-6,
    rtol = 1e-6,
    # Use MOI.LOCALLY_SOLVED for local solvers.
    optimal_status = MOI.OPTIMAL,
    # Pass attributes or MOI functions to `exclude` to skip tests that
```

```

    # rely on this functionality.
    exclude = Any[MOI.VariableName, MOI.delete],
)

"""
    runtests()

This function runs all functions in the this Module starting with `test_`.
"""
function runtests()
    for name in names(@__MODULE__; all = true)
        if startswith("$name", "test_")
            @testset "$name" begin
                getfield(@__MODULE__, name)()
            end
        end
    end
end

"""
    test_runtests()

This function runs all the tests in MathOptInterface.Test.

Pass arguments to `exclude` to skip tests for functionality that is not
implemented or that your solver doesn't support.
"""
function test_runtests()
    MOI.Test.runtests(
        BRIDGED,
        CONFIG,
        exclude = [
            "test_attribute_NumberOfThreads",
            "test_quadratic_",
        ],
        # This argument is useful to prevent tests from failing on future
        # releases of MOI that add new tests. Don't let this number get too far
        # behind the current MOI release though! You should periodically check
        # for new tests in order to fix bugs and implement new features.
        exclude_tests_after = v"0.10.5",
    )
    return
end

"""
    test_SolverName()

You can also write new tests for solver-specific functionality. Write each new
test as a function with a name beginning with `test_`.
"""
function test_SolverName()
    @test MOI.get(FooBar.Optimizer(), MOI.SolverName()) == "FooBar"
    return
end

end # module TestFooBar

```

```
# This line at the end of the file runs all the tests!
TestFooBar.runtests()
```

Then modify your `runtests.jl` file to include the `MOI_wrapper.jl` file:

```
# ===== /test/runtests.jl =====

using Test

@testset "MOI" begin
    include("test/MOI_wrapper.jl")
end
```

Info

The optimizer BRIDGED constructed with `instantiate` automatically bridges constraints that are not supported by OPTIMIZER using the bridges listed in [Bridges](#). It is recommended for an implementation of MOI to only support constraints that are natively supported by the solver and let bridges transform the constraint to the appropriate form. For this reason it is expected that tests may not pass if OPTIMIZER is used instead of BRIDGED.

How to debug a failing test

When writing a solver, it's likely that you will initially fail many tests! Some failures will be bugs, but other failures you may choose to exclude.

There are two ways to exclude tests:

- Exclude tests whose names contain a string using:

```
MOI.Test.runtests(
    model,
    config;
    exclude = String["test_to_exclude", "test_conic_"],
)
```

This will exclude tests whose name contains either of the two strings provided.

- Exclude tests which rely on specific functionality using:

```
MOI.Test.Config(exclude = Any[MOI.VariableName, MOI.optimize!])
```

This will exclude tests which use the `MOI.VariableName` attribute, or which call `MOI.optimize!`.

Each test that fails can be independently called as:

```
model = FooBar.Optimizer()
config = MOI.Test.Config()
MOI.empty!(model)
MOI.Test.test_category_name_that_failed(model, config)
```

You can look-up the source code of the test that failed by searching for it in the `src/Test/test_category.jl` file.

Tip

Each test function also has a docstring that explains what the test is for. Use `? MOI.Test.test_category_name_that_fails` from the REPL to read it.

Periodically, you should re-run excluded tests to see if they now pass. The easiest way to do this is to swap the `exclude` keyword argument of `runtests` to `include`. For example:

```
MOI.Test.runtests(
    model,
    config;
    exclude = String["test_to_exclude", "test_conic_"],
)
```

becomes

```
MOI.Test.runtests(
    model,
    config;
    include = String["test_to_exclude", "test_conic_"],
)
```

How to add a test

To detect bugs in solvers, we add new tests to `MOI.Test`.

As an example, `ECOS` errored calling `optimize!` twice in a row. (See [ECOS.jl PR #72](#).) We could add a test to `ECOS.jl`, but that would only stop us from re-introducing the bug to `ECOS.jl` in the future, but it would not catch other solvers in the ecosystem with the same bug! Instead, if we add a test to `MOI.Test`, then all solvers will also check that they handle a double `optimize!` call!

For this test, we care about correctness, rather than performance. therefore, we don't expect solvers to efficiently decide that they have already solved the problem, only that calling `optimize!` twice doesn't throw an error or give the wrong answer.

Step 1

Install the `MathOptInterface` julia package in dev mode ([ref](#)):

```
julia> ]
(@v1.6) pkg> dev MathOptInterface
```

Step 2

From here on, proceed with making the following changes in the `~/.julia/dev/MathOptInterface` folder (or equivalent dev path on your machine).

Step 3

Since the double-optimize error involves solving an optimization problem, add a new test to `src/Test/UnitTests/solve.jl`:

```
"""
    test_unit_optimize!_twice(model::MOI.ModelLike, config::Config)
```

Test that calling `MOI.optimize!` twice does not error.

This problem was first detected in ECOS.jl PR#72:

<https://github.com/jump-dev/ECOS.jl/pull/72>

"""

```
function test_unit_optimize!_twice(
    model::MOI.ModelLike,
    config::Config{T},
) where {T}
    # Use the `@requires` macro to check conditions that the test function
    # requires in order to run. Models failing this `@requires` check will
    # silently skip the test.
    @requires MOI.supports_constraint(
        model,
        MOI.VariableIndex,
        MOI.GreaterThan{Float64},
    )
    @requires _supports(config, MOI.optimize!)
    # If needed, you can test that the model is empty at the start of the test.
    # You can assume that this will be the case for tests run via `runtests`.
    # User's calling tests individually need to call `MOI.empty!` themselves.
    @test MOI.is_empty(model)
    # Create a simple model. Try to make this as simple as possible so that the
    # majority of solvers can run the test.
    x = MOI.add_variable(model)
    MOI.add_constraint(model, x, MOI.GreaterThan(one(T)))
    MOI.set(model, MOI.ObjectiveSense(), MOI.MIN_SENSE)
    MOI.set(
        model,
        MOI.ObjectiveFunction{MOI.VariableIndex}(),
        x,
    )
    # The main component of the test: does calling `optimize!` twice error?
    MOI.optimize!(model)
    MOI.optimize!(model)
    # Check we have a solution.
    @test MOI.get(model, MOI.TerminationStatus()) == MOI.OPTIMAL
    # There is a three-argument version of `Base.isapprox` for checking
    # approximate equality based on the tolerances defined in `config`:
    @test isapprox(MOI.get(model, MOI.VariablePrimal(), x), one(T), config)
    # For code-style, these tests should always `return` `nothing`.
    return
end
```

Info

Make sure the function is agnostic to the number type T! Don't assume it is a Float64 capable solver!

We also need to write a test for the test. Place this function immediately below the test you just wrote in the same file:

```
function setup_test(
    ::typeof(test_unit_optimize!_twice),
    model::MOI.Utilities.MockOptimizer,
```

```

        ::Config,
    )
    MOI.Utilities.set_mock_optimize!(
        model,
        (mock::MOI.Utilities.MockOptimizer) -> MOIU.mock_optimize!(
            mock,
            MOI.OPTIMAL,
            (MOI.FEASIBLE_POINT, [1.0]),
        ),
    )
    return
end

```

Finally, you also need to implement `Test.version_added`. If we added this test when the latest released version of MOI was v0.10.5, define:

```
version_added(::typeof(test_unit_optimize!_twice)) = v"0.10.6"
```

Step 6

Commit the changes to git from `~/julia/dev/MathOptInterface` and submit the PR for review.

Tip

If you need help writing a test, [open an issue on GitHub](#), or ask the [Developer Chatroom](#)

31.2 API Reference

The Test submodule

Functions to help test implementations of MOI. See [The Test submodule](#) for more details.

[MathOptInterface.Test.Config](#) – Type.

```

Config(
    ::Type{T} = Float64;
    atol::Real = Base.rtoldefault(T),
    rtol::Real = Base.rtoldefault(T),
    optimal_status::MOI.TerminationStatusCode = MOI.OPTIMAL,
    infeasible_status::MOI.TerminationStatusCode = MOI.INFEASIBLE,
    exclude::Vector{Any} = Any[],
) where {T}

```

Return an object that is used to configure various tests.

Configuration arguments

- `atol::Real = Base.rtoldefault(T)`: Control the absolute tolerance used when comparing solutions.
- `rtol::Real = Base.rtoldefault(T)`: Control the relative tolerance used when comparing solutions.
- `optimal_status = MOI.OPTIMAL`: Set to `MOI.LOCALLY_SOLVED` if the solver cannot prove global optimality.

- `infeasible_status = MOI.INFEASIBLE`: Set to `MOI.LOCALLY_INFEASIBLE` if the solver cannot prove global infeasibility.
- `exclude = Vector{Any}`: Pass attributes or functions to exclude to skip parts of tests that require certain functionality. Common arguments include:
 - `MOI.delete` to skip deletion-related tests
 - `MOI.optimize!` to skip optimize-related tests
 - `MOI.ConstraintDual` to skip dual-related tests
 - `MOI.VariableName` to skip setting variable names
 - `MOI.ConstraintName` to skip setting constraint names

Examples

For a nonlinear solver that finds local optima and does not support finding dual variables or constraint names:

```
Config(
    Float64;
    optimal_status = MOI.LOCALLY_SOLVED,
    exclude = Any[
        MOI.ConstraintDual,
        MOI.VariableName,
        MOI.ConstraintName,
        MOI.delete,
    ],
)
```

[source](#)

`MathOptInterface.Test.runtests` – Function.

```
runtests(
    model::MOI.ModelLike,
    config::Config;
    include::Vector{String} = String[],
    exclude::Vector{String} = String[],
    warn_unsupported::Bool = false,
    exclude_tests_after::VersionNumber = v"999.0.0",
)
```

Run all tests in `MathOptInterface.Test` on `model`.

Configuration arguments

- `config` is a `Test.Config` object that can be used to modify the behavior of tests.
- If `include` is not empty, only run tests that contain an element from `include` in their name.
- If `exclude` is not empty, skip tests that contain an element from `exclude` in their name.
- `exclude` takes priority over `include`.
- If `warn_unsupported` is `false`, `runtests` will silently skip tests that fail with `UnsupportedConstraint` or `UnsupportedAttribute`. When `warn_unsupported` is `true`, a warning will be printed. For most cases the default behavior (`false`) is what you want, since these tests likely test functionality that is not supported by `model`. However, it can be useful to run `warn_unsupported = true` to check you are not skipping tests due to a missing `supports_constraint` method or equivalent.

- `exclude_tests_after` is a version number that excludes any tests to MOI added after that version number. This is useful for solvers who can declare a fixed set of tests, and not cause their tests to break if a new patch of MOI is released with a new test.

See also: [setup_test](#).

Example

```
config = MathOptInterface.Test.Config()
MathOptInterface.Test.runtests(
    model,
    config;
    include = ["test_linear_"],
    exclude = ["VariablePrimalStart"],
    warn_unsupported = true,
    exclude_tests_after = v"0.10.5",
)
```

[source](#)

[MathOptInterface.Test.setup_test](#) - Function.

```
setup_test(::typeof(f), model::MOI.ModelLike, config::Config)
```

Overload this method to modify `model` before running the test function `f` on `model` with `config`. You can also modify the fields in `config` (e.g., to loosen the default tolerances).

This function should either return nothing, or return a function which, when called with zero arguments, undoes the setup to return the model to its previous state. You do not need to undo any modifications to `config`.

This function is most useful when writing new tests of the tests for MOI, but it can also be used to set test-specific tolerances, etc.

See also: [runtests](#)

Example

```
function MOI.Test.setup_test(
    ::typeof(MOI.Test.test_linear_VariablePrimalStart_partial),
    mock::MOIU.MockOptimizer,
    ::MOI.Test.Config,
)
    MOIU.set_mock_optimize!(
        mock,
        (mock::MOIU.MockOptimizer) -> MOIU.mock_optimize!(mock, [1.0, 0.0]),
    )
    mock.eval_variable_constraint_dual = false

    function reset_function()
        mock.eval_variable_constraint_dual = true
        return
    end
    return reset_function
end
```

[source](#)

[MathOptInterface.Test.version_added](#) - Function.

```
| version_added(::typeof(function_name))
```

Returns the version of MOI in which the test `function_name` was added.

This method should be implemented for all new tests.

See the `exclude_tests_after` keyword of [runtests](#) for more details.

[source](#)

[MathOptInterface.Test.@requires](#) - Macro.

```
| @requires(x)
```

Check that the condition `x` is true. Otherwise, throw an [RequirementUnmet](#) error to indicate that the model does not support something required by the test function.

Examples

```
| @requires MOI.supports(model, MOI.Silent())  
| @test MOI.get(model, MOI.Silent())
```

[source](#)

[MathOptInterface.Test.RequirementUnmet](#) - Type.

```
| RequirementUnmet(msg::String) <: Exception
```

An error for throwing in tests to indicate that the model does not support some requirement expected by the test function.

[source](#)